

# Binary Graph Neural Networks

Mehdi Bahri<sup>1</sup>

Gaétan Bahl<sup>2,3</sup>

Stefanos Zafeiriou<sup>1</sup>

<sup>1</sup>Imperial College London, UK

<sup>2</sup>Université Côte d’Azur - Inria

<sup>3</sup>IRT Saint-Exupéry

{m.bahri, s.zafeiriou}@imperial.ac.uk, gaetan.bahl@inria.fr

## Abstract

Graph Neural Networks (GNNs) have emerged as a powerful and flexible framework for representation learning on irregular data. As they generalize the operations of classical CNNs on grids to arbitrary topologies, GNNs also bring much of the implementation challenges of their Euclidean counterparts. Model size, memory footprint, and energy consumption are common concerns for many real-world applications. Network binarization allocates a single bit to network parameters and activations, thus dramatically reducing the memory requirements (up to 32x compared to single-precision floating-point parameters) and maximizing the benefits of fast SIMD instructions of modern hardware for measurable speedups. However, in spite of the large body of work on binarization for classical CNNs, this area remains largely unexplored in geometric deep learning. In this paper, we present and evaluate different strategies for the binarization of graph neural networks. We show that through careful design of the models, and control of the training process, binary graph neural networks can be trained at only a moderate cost in accuracy on challenging benchmarks. In particular, we present the first dynamic graph neural network in Hamming space, able to leverage efficient  $k$ -NN search on binary vectors to speed-up the construction of the dynamic graph. We further verify that the binary models offer significant savings on embedded devices.

## 1. Introduction

Standard CNNs assume their input to be regularly-sampled points on a grid, and are therefore suitable for data that can be well-represented in an Euclidean space, such as images, sound, or videos. However, many increasingly relevant types of data do not fit this framework [4]. Graph theory offers a broad mathematical formalism for modeling interactions, and is therefore commonly used in fields such as network sciences [11], bioinformatics [24, 38], and recommender systems [35], as well as for studying discretisations of continuous mathematical structures such as in computer graphics. This motivates the development of machine learn-

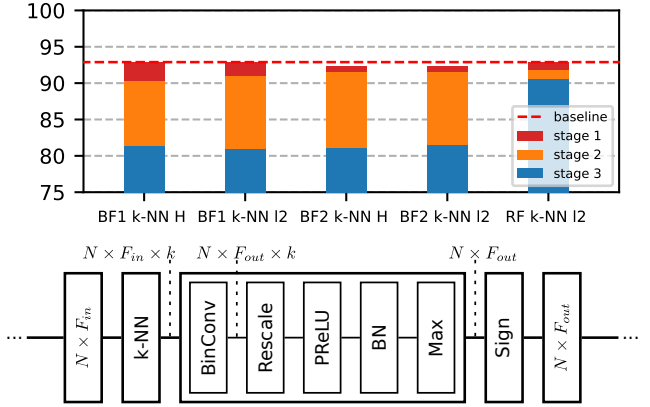


Figure 1. **Top:** Test accuracy of different binarization schemes at all stages of our cascaded distillation protocol (baseline: 92.89%). **Bottom:** The “BF2” variant of our XorEdgeConv operator.

ing methods able to effectively deal with graph-supported data. Among them, Graph Neural Networks (GNNs) generalize the operations of CNNs to arbitrary topologies by extending the basic building blocks of CNNs such as convolutions and pooling to graphs. Similarly to CNNs, GNNs learn deep representations of graphs or graph elements, and have emerged as the best performing models for learning on graphs as well as on 3D data with the development of advanced and increasingly deep architectures [32, 17].

As the computational complexity of the networks and the scale of graph datasets increase, so does the need for faster and smaller models. The motivations for resource-efficient deep learning are numerous and also apply to deep learning on graphs and 3D shapes. Computer vision models are routinely deployed on embedded devices, such as mobile phones or satellites [2, 31], where energy and storage constraints are important. The development of smart devices and IoT may bring about the need for power-efficient graph learning models [27, 57, 7]. Finally, models that require GPUs for inference can be expensive to serve, whereas CPUs are typically more affordable. This latter point is especially relevant to the applications of GNNs on large relational datasets, such as popular social networks, large-scale data mining, or large bioinformatics datasets [59].

While recent work has proposed algorithmic changes to make graph neural networks more scalable, such as the use of sampling [20, 56] or architectural improvements [14, 9] and simplifications [51], our approach is orthogonal to these advances and focuses on compressing existing graph neural network architectures while preserving model performance. Model compression is a well-researched area for Euclidean neural networks, but has seen very little application in geometric deep learning. In this paper, we study different strategies for binarizing graph neural networks.

Our contributions are as follows:

- We present a binarization strategy inspired by the latest developments in binary neural networks for images [6, 34] and knowledge distillation for graph networks
- We develop an efficient dynamic graph neural network model that constructs the dynamic graph in Hamming space, thus paving the way for significant speedups at inference time, with negligible loss of accuracy when using real-valued weights
- We conduct a thorough ablation study of the hyperparameters and techniques used in our approach
- We demonstrate real-world acceleration of our models on a budget ARM device

**Notations** Throughout the paper, matrices and vectors are denoted by upper and lowercase bold letters (*e.g.*,  $\mathbf{X}$  and  $\mathbf{x}$ ), respectively.  $\mathbf{I}$  denotes the identity matrix of compatible dimensions. The  $i^{th}$  column of  $\mathbf{X}$  is denoted as  $\mathbf{x}_i$ . The sets of real numbers is denoted by  $\mathbb{R}$ . A *graph*  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$  consists of *vertices*  $\mathcal{V} = \{1, \dots, n\}$  and *edges*  $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$ . The graph structure can be encoded in the *adjacency matrix*  $\mathbf{A}$ , where  $a_{ij} = 1$  if  $(i, j) \in \mathcal{E}$  (in which case  $i$  and  $j$  are said to be *adjacent*) and zero otherwise. The *neighborhood* of vertex  $i$ , denoted by  $\mathcal{N}(i) = \{j : (i, j) \in \mathcal{E}\}$ , is the set of vertices adjacent to  $i$ . We write  $\parallel$  the concatenation operator.

## 2. Related Work

**Knowledge distillation** is a popular approach for model compression, where a pretrained “teacher” network is used to supervise and inform the training of a smaller “student” network by matching the features of both networks at specific knowledge transfer points through specialized losses. In logit matching [22], a cross-entropy loss is used to regularize the output logits of the student by matching them with a blurred version of the teacher’s logits computed using a softmax with an additional temperature hyperparameter. More recent works also focus on matching internal activations of both networks, such as attention volumes in [55].

**Quantized and Binary Neural Networks** Network quantization [18, 58] refers to the practice of lowering the numerical precision of a model in a bid to reduce its size and speed-up inference. Binary Neural Networks (BNNs) [25] push it to the extreme and use a single bit for weights and activations. The seminal work of XNOR-Net [43] showed that re-introducing a small number of floating point operations in BNNs can drastically improve the performance compared to using pure binary operations by reducing the quantization error. In XNOR-Net, a dot product  $\star$  between real tensors  $\mathbf{A}$  and  $\mathbf{B}$  of compatible dimensions is approximated by  $\mathbf{A} \star \mathbf{B} \approx (\text{sign}(\mathbf{A}) \otimes \text{sign}(\mathbf{B}))\alpha\beta$ , where  $\beta = \frac{1}{n}\|\mathbf{A}\|_1$  and  $\alpha = \frac{1}{n}\|\mathbf{B}\|_1$  are rescaling constants. XNOR-Net++ [6] proposed to instead learn a rescaling tensor  $\Gamma$ , with shared factors to limit the number of trainable parameters and avoid overfitting. Finally, in Real to Binary networks [34], the authors compile state of the art techniques and improve the performance of binary models with knowledge distillation.

**Graph Neural Networks** Graph Neural Networks were initially proposed in [19, 44] as a form of recursive neural network. Later formulations relied on Fourier analysis on graphs using the eigen decomposition of the graph Laplacian [5] and approximations of such [10], but suffered from the connectivity-specific nature of the Laplacian. Attention-based models [36, 13, 46, 45] are purely spatial approaches that compute a vertex’s features as a dynamic weighting of its neighbours’. Spatial and spectral approaches have been unified [29] and shown to derive from the more general neural message passing [16] framework. We refer to recent reviews on GNNs, such as [52] for a comprehensive overview, and focus only on the operators we binarize in this paper.

The message-passing framework offers a general formulation of graph neural networks:

$$\mathbf{x}_i^{(l)} = \gamma^{(l)} \left( \mathbf{x}_i^{(l-1)}, \bigoplus_{j \in \mathcal{N}(i)} \phi^{(l)} \left( \mathbf{x}_i^{(l-1)}, \mathbf{x}_j^{(l-1)}, \mathbf{e}_{ij}^{(l-1)} \right) \right), \quad (1)$$

where  $\bigoplus$  denotes a differentiable symmetric (permutation-invariant) function, (*e.g.* max or  $\sum$ ),  $\phi$  a differentiable kernel function,  $\gamma$  is an MLP, and  $\mathbf{x}_i$  and  $\mathbf{e}_{ij}$  are features associated with vertex  $i$  and edge  $(i, j)$ , respectively.

The EdgeConv operator is a special case introduced as part of the Dynamic Graph CNN (DGCNN) model [49] and defines an edge message as a function of  $\mathbf{x}_j - \mathbf{x}_i$ :

$$\mathbf{e}_{ij}^{(l)} = \text{ReLU} \left( \boldsymbol{\theta}^{(l)} (\mathbf{x}_j^{(l-1)} - \mathbf{x}_i^{(l-1)}) + \boldsymbol{\phi}^{(l)} \mathbf{x}_i^{(l-1)} \right) \quad (2)$$

$$= \text{ReLU} \left( \boldsymbol{\Theta}^{(l)} \tilde{\mathbf{X}}^{(l-1)} \right) \quad (3)$$

where  $\tilde{\mathbf{X}}^{(l-1)} = \left[ \mathbf{x}_i^{(l-1)} \parallel \mathbf{x}_j^{(l-1)} - \mathbf{x}_i^{(l-1)} \right]$ .

And the output of the convolution as the max aggregation ( $\square = \max$ ) as the new node feature:

$$\mathbf{x}_i^{(l)} = \max_{j \in \mathcal{N}(i)} \mathbf{e}_{ij}^{(l)} \quad (4)$$

While the EdgeConv operator is applicable to graph inputs, the main use case presented in [49] is for point clouds, where the neighbours are found by  $k$ -Nearest Neighbours ( $k$ -NN) search in feature space before each convolutional layer. DGCNN is the first example of a dynamic graph architecture, with follow-up work in [26].

The GraphSAGE [20] operator introduced inductive learning on large graphs with sampling and can also be phrased as a message passing operator:

$$\mathbf{x}_i^{(l)} = \text{Norm} \left( \text{ReLU} \left( \mathbf{W}^{(l)} \left[ \mathbf{x}_i^{(l-1)} \parallel \text{Aggr}_{j \in \mathcal{N}(i)} \mathbf{x}_j^{(l-1)} \right] \right) \right) \quad (5)$$

Where Aggr is a symmetric aggregation function such as max, sum or mean, and Norm denotes the  $\ell_2$  normalization.

**Model Compression in Geometric Deep Learning** In [47], the authors propose to binarize the Graph Attention (GAT) operator [45], and evaluate their method on small-scale datasets such as Cora [37] and Pubmed [29]. In [48], the authors apply the XNOR-Net method to the GCN model [29] with success, but also on small-scale datasets. Finally, [42] propose to binarize PointNet with tailored aggregation and scaling functions. At the time of writing, the Local Structure Preserving (LSP) module of [54] is the only knowledge distillation method specifically designed for GNNs. LSP defines local structure vectors  $LS_i$  for each node in the graph:

$$LS_{ij} = \frac{\exp(\text{SIM}(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k \in \mathcal{N}(i)} \exp(\text{SIM}(\mathbf{x}_i, \mathbf{x}_k))} \quad (6)$$

where SIM denotes a similarity measure, *e.g.*,  $\|\cdot\|_2^2$  or a kernel function such as a Gaussian RBF kernel. The total local structure preserving loss between a student network  $s$  and a teacher  $t$  is then defined as

$$L_{LSP} = \frac{1}{|\mathcal{V}|} \sum_{i \in \mathcal{V}} \sum_{j \in \mathcal{N}^u(i)} LS_{ij}^s \log \frac{LS_{ij}^s}{LS_{ij}^t}, \quad (7)$$

$\mathcal{N}^u(i) = \mathcal{N}^s(i) \cup \mathcal{N}^t(i)$  to enable dynamic graph models.

### 3. Method

Eq. 1 is more general than the vanilla Euclidean convolution, which boils down to a single matrix product to quantize. We must therefore choose which elements of Eq. 1 to binarize and how: the node features  $\mathbf{x}_i$ , the edge messages  $\mathbf{e}_{ij}$ , and the functions  $\square$ ,  $\gamma$  and  $\phi$  may all need to be adapted.

**Quantization** We follow the literature and adopt the sign operator as the binarization function

$$\text{sign}(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ -1 & \text{if } x < 0 \end{cases} \quad (8)$$

The gradient of sign is zero almost everywhere. We therefore follow the literature and employ the straight-through estimator [3] to provide a valid gradient. We use this method for both network weights and activations. Furthermore, we mean center and clip the real latent network weights after their update in the backpropagation step.

**Learnable rescaling** Assuming a dot product operation (*e.g.* a fully-connected or convolutional layer)  $\mathbf{A} \star \mathbf{B} \in \mathbb{R}^{o \times h \times w}$ , we approximate it as in [6]:

$$\mathbf{A} \star \mathbf{B} \approx (\text{sign}(\mathbf{A}) \otimes \text{sign}(\mathbf{B})) \odot \Gamma, \quad (9)$$

with  $\Gamma$  a learned rescaling tensor. We use two constructions of  $\Gamma$  depending on the model. Channel-wise:

$$\Gamma = \alpha \in \mathbb{R}^{o \times 1 \times 1} \quad (10)$$

and one rank-1 factor per mode:

$$\Gamma = \alpha \otimes \beta \otimes \gamma, \quad \alpha \in \mathbb{R}^o, \beta \in \mathbb{R}^h, \gamma \in \mathbb{R}^w \quad (11)$$

**Activation functions** Recent work [34] has shown using non-linear activations in XNOR-Net - type blocks can improve the performance of binary neural networks, with PReLU bringing the most improvement.

**Knowledge Distillation** Inspired by [34], we investigate the applicability of knowledge distillation for the binarization of graph neural networks. For classification tasks, we use a logit matching loss [22] as the base distillation method. We also implemented the LSP module of [54].

**Multi-stage training** We employ a cascaded distillation scheme [34], an overview of which is shown in Figure 2.

*Stage 1:* We first build a real-valued and real-weighted network with the same architecture as the desired binary network by replacing the quantization function with tanh. We distillate the original (base) network into this first student network. We employ weight decay with weight  $\lambda = 1e - 5$ , logit matching, and LSP. We use the same initial learning rate and learning rate schedule as for the base network.

*Stage 2:* The model of stage 1 becomes the teacher, the student is a binary network with real-valued weights but binary activations. We initialize the student with the weights of the teacher. We employ weight decay with  $\lambda = 1e - 5$ , logit matching, and LSP. We use a smaller learning rate (*e.g.* half) than for stage 1 and the same learning rate schedule.

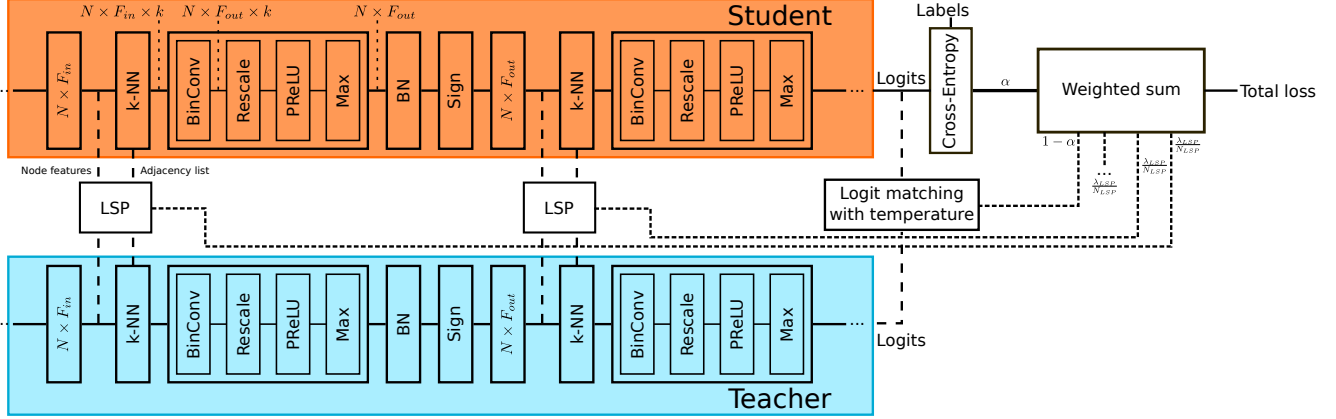


Figure 2. **Distillation with the "BF1" variant of XorEdgeConv**: the student model is more heavily quantized than the teacher. Knowledge transfer points equipped with LSP modules encourage similar dynamic graph feature distributions after each  $k$ -NN graph computation (except for the first, performed on the fixed 3D coordinates). Logit matching is used to further inform the training of the student.

**Stage 3:** The model of stage 2 becomes the teacher, the student is a binary network with binary weights and binary activations. We initialize the student with the weights of the teacher. We use logit matching and LSP but no weight decay. We further decrease the initial learning rate (*e.g.* one tenth of that of stage 1) and decrease it when the loss plateaus.

**Batch Normalization** We investigate the importance of the order of the dot product and batch normalization operations for discretizing dot product operations within graph convolution operators. However, our base approach is to follow the XNOR-Net block structure [43] with learnable rescaling (*i.e.* XNOR-Net++ block). In particular, all fully-connected layers of MLPs that follow graph feature extraction layers are binarized using the XNOR-Net++ block.

## 4. Models

We choose the Dynamic Graph CNN model, built around the EdgeConv operator of Eq. 3 as our main case study. DGCNN has several characteristics that make it an interesting candidate for binarization. First, the EdgeConv operator is widely applicable to graphs and point clouds. Second, the operator relies on both node features and edge messages, contrary to other operators previously studied in GNN binarization such as GCN. Third, the time complexity of DGCNN is strongly impacted by the  $k$ -NN search in feature space.  $k$ -NN search can be made extremely efficient in Hamming space, and fast algorithms could theoretically be implemented for the construction of the dynamic graph at inference provided that the graph features used in the search are binary, which requires a different binarization strategy than merely approximating the dense layer in EdgeConv.

For completeness, we derive a binary SAGE operator.

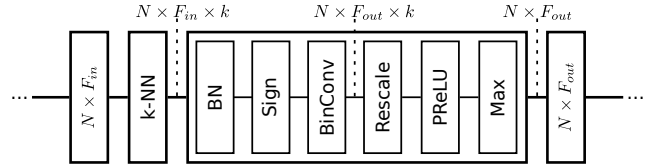


Figure 3. **The BinEdgeConv operator** ("RF" model in the experiments) can be deployed as a drop-in replacement for EdgeConv and follows the XNOR-Net++ approach to binarization.

### 4.1. Direct binarization

Our first approach binarizes the network weights and the graph features at the input of the graph convolution layers, but keeps the output real-valued. The network, therefore, produces real-valued node features. We replace the EdgeConv operator by a block similar to XNOR-Net++, using learnable rescaling and batch normalization pre-quantization:

We define the BinEdgeConv operator as:

$$\mathbf{e}_{ij}^{(l)} = \sigma \left( \text{sign}(\Theta^{(l)}) \otimes \text{sign} \left( \text{BN} \left( \tilde{\mathbf{X}}^{(l-1)} \right) \right) \odot \Gamma^{(l)} \right) \quad (12)$$

$$\mathbf{x}_i^{(l)} = \max_{j \in \mathcal{N}(i)} \mathbf{e}_{ij}^{(l)} \quad (13)$$

with  $\sigma$  the PReLU activation,  $\Theta^{(l)}$  a set of learnable real parameters and  $\Gamma^{(l)}$  a real rescaling tensor implemented as per Eq. 11. BinEdgeConv is visualized in Figure 3.

We use the same structure to approximate the MLP classifier. Similarly, we binarize Eq. 5 to get:

$$\mathbf{h}^{(l)} = \text{sign} \left( \text{BN} \left( \left[ \mathbf{x}_i^{(l-1)} \parallel \text{Aggr}_{j \in \mathcal{N}(i)} \mathbf{x}_j^{(l-1)} \right] \right) \right) \quad (14)$$

$$\mathbf{x}_i^{(l)} = \text{Norm} \left( \sigma \left( \left( \text{sign}(\mathbf{W}^{(l)}) \otimes \mathbf{h}^{(l)} \right) \odot \Gamma^{(l)} \right) \right) \quad (15)$$

with  $\sigma$  the PReLU activation and  $\Gamma^{(l)}$  following Eq. 10.



## 4.2. Dynamic Graph in Hamming Space

As mentioned, one advantage of binary node features is to enable fast computation of the  $k$ -Nearest Neighbours graph at inference time by replacing the  $\ell_2$  norm with the Hamming distance. We detail our approach to enable quantization-aware training with  $k$ -NN search on binary vectors.

**Edge feature** The central learnable operation of EdgeConv is  $\Theta [\mathbf{x}_i \parallel \mathbf{x}_j - \mathbf{x}_i]$  as per Eq. 3, where the edge feature is  $\mathbf{x}_j - \mathbf{x}_i$ . Assuming binary node features, the standard subtraction operation becomes meaningless. Formally, for  $\mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n$  with  $\mathbb{R}^n$  the  $n$ -dimensional Euclidean vector space over the field of real numbers

$$\mathbf{x}_1 - \mathbf{x}_2 := \mathbf{x}_1 + (-\mathbf{x}_2) \quad (16)$$

by definition, with  $(-\mathbf{x}_2)$  the additive inverse of  $\mathbf{x}_2$ . Seeing binary vectors as elements of vector spaces over the finite field  $\mathbb{F}_2$ , we can adapt Eq. 16 with the operations of boolean algebra. The addition therefore becomes the boolean exclusive or (XOR)  $\oplus$ , and the additive inverse of  $(-x)_{\mathbb{F}_2}$  is  $x$  itself ( $x \oplus x = 0$ ). With our choice of quantizer (Eq. 8),  $\mathbf{x}_i, \mathbf{x}_j \in \{-1, 1\}^n$  and we observe that  $\mathbf{x}_i \oplus \mathbf{x}_j = -\mathbf{x}_i \odot \mathbf{x}_j$ . We therefore base our binary EdgeConv operator for binary node features, XorEdgeConv, on the following steps:

$$\mathbf{e}_{ij}^{(l)} = \sigma \left( \text{sign}(\Theta^{(l)}) \otimes \tilde{\mathbf{X}}_b^{(l-1)} \odot \Gamma^{(l)} \right) \quad (17)$$

$$\mathbf{x}_i^{(l)} = \text{sign} \left( \max_{j \in \mathcal{N}(i)} \mathbf{e}_{ij}^{(l)} \right) \quad (18)$$

with  $\tilde{\mathbf{X}}_b^{(l-1)} = [\mathbf{x}_i^{(l-1)} \parallel -\mathbf{x}_j^{(l-1)} \odot \mathbf{x}_i^{(l-1)}]$ ,  $\Theta^{(l)}$  a set of learnable real parameters and  $\Gamma^{(l)}$  a real rescaling tensor implemented as per Eq. 11. We further investigate the practical importance of the placement of the batch normalization operation, either before or after the aggregation function, by proposing two variants:

$$\mathbf{x}_i^{(l)} = \text{sign} \left( \text{BN} \left( \max_{j \in \mathcal{N}(i)} \mathbf{e}_{ij}^{(l)} \right) \right) \quad (19)$$

shown as part of Figure 2 and

$$\mathbf{x}_i^{(l)} = \text{sign} \left( \max_{j \in \mathcal{N}(i)} \text{BN} \left( \mathbf{e}_{ij}^{(l)} \right) \right) \quad (20)$$

drawn in Figure 1. Here, the main difference lies in the distribution of the features pre-quantization.

**Nearest Neighbours Search** The Hamming distance between two binary vectors  $\mathbf{x}, \mathbf{y}$  is  $d_H(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} \oplus \mathbf{y}\|_H$  where  $\|\cdot\|_H$  is the number of non-zero bits, and can be efficiently implemented as  $\text{popcount}(\mathbf{x} \text{ xor } \mathbf{y})$ . We note that

this relates our approach to previous work on efficient hashing [39, 41, 30] and metric learning [40], especially given the dynamic nature of the graph. Unfortunately, like the sign function, the hamming distance has an ill-defined gradient, which hinders its use as-is for training. We therefore investigate two continuous relaxations. (1) we use the standard  $\ell_2$  norm for training, since all norms are equivalent in finite dimensions. (2) we observe that the matrix of pairwise Hamming distances between  $d$ -dimensional vectors  $\mathbf{x}_i$  valued in  $\{-1, 1\}$  can be computed in a single matrix-matrix product up to a factor 2 as (see Eq. 5 of [33]):

$$\mathbf{D} = -(\mathbf{X}\mathbf{X}^T - d\mathbf{I}_d) \quad (21)$$

with  $\mathbf{X}$  the matrix of the  $\mathbf{x}_i$  stacked row-wise and  $\mathbf{I}_d$  the identity. We investigate both options.

**Local structure** With binary node features, we now have to choose how to define the local structure similarity measure of Eq. 6. One option is to use the standard Gaussian RBF as in the real-valued case. Another option is to define the similarity in Hamming space, like for the  $k$ -NN search. We therefore investigate the following similarity metric:

$$\text{SIM}(\mathbf{x}_i, \mathbf{x}_j) = e^{-\|\mathbf{x}_i \oplus \mathbf{x}_j\|_H} \quad (22)$$

For vectors  $\mathbf{x}, \mathbf{y} \in \{-1, 1\}^n$ , we note that  $\|\mathbf{x} \oplus \mathbf{y}\|_H = \frac{1}{2} \sum_{k=1}^n (-x_k y_k + 1)$ .

## 5. Experimental Evaluation

We perform a thorough ablation study of our method on Dynamic Graph CNN. The model binarized according to the method of Section 4.1 and using the BinEdgeConv operator of Eq. 12 is referred to as **RF** for "Real graph Features". The model binarized according to Section 4.2 and using the XorEdgeConv operator is referred to as **BF1** if following Eq. 19 or **BF2** if following Eq. 20. We evaluate DGCNN on the ModelNet40 classification benchmark, as in [49].

**Balance functions** Recent work [42] has uncovered possible limitations in binary graph and point cloud learning models when quantizing the output of max-pooling aggregation of batch-normalized high-dimensional features. Similarly, [47] claim that a balance function is necessary to avoid large values in the outputs of the dot product operations when most pre-quantization inputs are positive. We evaluate two strategies for re-centering the input of sign after max aggregation, namely mean-centering, and median-centering (thus ensuring a perfectly balanced distribution of positive and negative values pre-quantization). We test these techniques for the max aggregation of edge messages ("edge balance", e.g. between the framed block and the sign operation in Figure 1) and for the max and average pooling operations before the MLP classifier ("global balance").

#	Model	Distillation	Stage	KNN	LSP	$\lambda_{LSP}$	Activation	Global balance	Edge balance	Acc (%)
1	BF1	None	-	H	-	-	PreLU	-	-	61.26
2	BF1	Direct	-	H	-	-	PreLU	-	-	62.84
3	BF1	Cascade	3	H	-	-	PreLU	-	-	<b>81.04</b>
4	BF1	Cascade	3	H	-	-	None	-	-	74.80
5	BF1	Cascade	2	H	$\ell_2$	100	PreLU	-	-	90.32
6	BF1	Cascade	3	H	$\ell_2$	100	PreLU	-	-	<b>81.32</b>
7	BF1	Cascade	3	H	$\ell_2$	100	ReLU	-	-	76.46
8	BF1	None	-	$\ell_2$	-	-	PreLU	-	-	62.97
9	BF1	Direct	-	$\ell_2$	-	-	PreLU	-	-	61.95
10	BF1	Cascade	3	$\ell_2$	-	-	PreLU	-	-	<b>81.00</b>
11	BF1	Cascade	2	$\ell_2$	$\ell_2$	100	PreLU	-	-	91.00
12	BF1	Cascade	3	$\ell_2$	$\ell_2$	100	PreLU	-	-	<b>81.00</b>
13	BF2	None	-	H	-	-	PreLU	-	-	59.72
14	BF2	Direct	-	H	-	-	PreLU	-	-	59.00
15	BF2	Cascade	3	H	-	-	PreLU	-	-	<b>79.58</b>
16	BF2	Cascade	3	H	-	-	PreLU	-	Mean	51.05
17	BF2	Cascade	3	H	-	-	PreLU	-	Median	75.93
18	BF2	Cascade	3	H	-	-	None	-	Median	71.96
19	BF2	Cascade	2	H	$\ell_2$	100	PreLU	-	-	91.57
20	BF2	Cascade	3	H	$\ell_2$	100	PreLU	-	-	<b>81.08</b>
21	BF2	Cascade	3	H	$\ell_2$	100	None	-	-	76.09
22	BF2	Cascade	3	H	$\ell_2$	100	ReLU	-	-	76.22
23	BF2	Cascade	3	H	$\ell_2$	100	PreLU	Mean	-	67.87
24	BF2	Cascade	3	H	$\ell_2$	100	PreLU	Median	-	60.82
25	BF2	None	-	$\ell_2$	-	-	PreLU	-	-	57.90
26	BF2	Direct	-	$\ell_2$	-	-	PreLU	-	-	59.12
27	BF2	Cascade	3	$\ell_2$	-	-	PreLU	-	-	80.11
28	BF2	Cascade	2	$\ell_2$	$\ell_2$	100	PreLU	-	-	91.53
29	BF2	Cascade	3	$\ell_2$	$\ell_2$	100	PreLU	-	-	<b>81.52</b>
30	RF	None	-	$\ell_2$	-	-	PreLU	-	-	79.30
31	RF	Direct	-	$\ell_2$	-	-	PreLU	-	-	72.69
32	RF	Cascade	3	$\ell_2$	-	-	PreLU	-	-	91.05
33	RF	Cascade	3	$\ell_2$	$\ell_2$	100	PreLU	-	-	90.52
34	RF	Cascade	3	$\ell_2$	$\ell_2$	100	None	-	-	89.71
35	RF	Cascade	3	$\ell_2$	$\ell_2$	100	ReLU	-	-	89.59
36	Baseline	-	-	-	-	ReLU	-	-	-	92.89

Table 1. Different variants and ablations of our models.

We can see in Table 1 that in all cases, the addition of balance functions actually lowered the performance of the models. This suggests that using batch normalization prior to quantization, as is common in the binary CNN literature, is sufficient at the message aggregation level and for producing graph embedding vectors.

**Effect of activation** Since the sign operation can be seen as acting as an activation applied on the output and to the weights of the XorEdgeConv operator, we first compare the models with binary node features with PReLU, ReLU, or no additional activation in Table 1. We can see the PReLU non-linearity offers significant improvements over the mod-

els trained with ReLU or without non-linearity in the edge messages at the cost of a single additional `fp32` parameter - the largest improvement being observed for the models that apply either median centering or batch normalization before the quantization operation.

**Binary node features and  $k$ -NN** We now study the final performance of our models depending on whether we use BinEdgeConv (real node features) or XorEdgeConv. Looking at the final models (stage 3) in Table 1, the model with real-valued node features that performs  $k$ -NN search with the  $\ell_2$  norm performs comparably with the full floating-point model. On the other hand, we saw a greater reduction in accu-

racy with the binary node features for the full binary models, and comparable accuracy whether we use the  $\ell_2$  norm or the relaxed Hamming distance. However, as reported in Table 1, using real weights (stage 2) with binary node features and  $k$ -NN search performed in Hamming space matched the performance of the original floating point model. This suggests that, although more research is necessary to maximize the performance of the full binary networks in Hamming space, dynamic graph networks that learn binary codes and construct the dynamic graph in Hamming space can be trained with minimal reduction in performance.

**Impact of LSP** The node features of the teacher and of the students are always real-valued at stage 1. Stage 2 was carried out using either the Gaussian RBF similarity or Eq. 22 for the student (which may have binary node features) and the Gaussian RBF for the teacher. Stage 3 uses either similarity measure for both the teacher and student. We also report the results of distilling the baseline DGCNN (full floating-point) model into a BF1 or BF2 full-binary model using the similarity in Hamming space for the student.

We saw inconsistent improvements when using LSP with the Gaussian RBF ( $\ell_2$ ), as seen in Table 1. This suggests the usefulness of the additional structure preserving knowledge is situational, as it can both increase or decrease model performance. Contrary to the models trained using  $k$ -NN search performed in Hamming space, the models trained with distillation using the similarity measure of Eq. 22 did not match the performance of their Gaussian  $\ell_2$  counterparts, as shown in Table 2, which we conjecture to be due to poor gradients.

Model	Stage	KNN	LSP	$\lambda_{LSP}$	Acc. (%)
BF1	2	H	H	100	38.21
BF1	2	$\ell_2$	H	100	38.94
BF2	2	H	H	100	63.25
BF2	2	$\ell_2$	H	100	64.71
BF1	3	H	H	100	16.29
BF1	3	$\ell_2$	H	100	20.34
BF2	3	$\ell_2$	H	100	9.40
BF2	3	H	H	100	11.47
BF1	Direct	$\ell_2$	H	100	23.34
BF2	Direct	H	H	100	30.23
BF2	Direct	$\ell_2$	H	100	32.17
BF1	Direct	H	H	100	36.47

Table 2. Performance of models trained with LSP using the Hamming-based similarity of Eq. 22 (H) at different stages and for direct distillation. Compared to the models trained using the Gaussian RBF ( $\ell_2$ ) similarity, low performance was observed.

**Cascaded distillation** Table 1 compares distilling the full floating point networks directly into a binary network, training from scratch, and the three stage distillation method. We

observe consistently higher performance with the progressive distillation, confirming its effectiveness.

**Large-scale inductive learning with GraphSAGE** We benchmark our binarized GraphSAGE on the OGB-Product and OGB-Protein node property prediction datasets [23], which are recent and challenging (2,449,029 nodes, 61,859,140 edges for OGB-product) benchmarks with standardized evaluation procedures, compared to the more commonly used ones, such as Cora [37] (2708 nodes, 5429 edges) used in [47] or Reddit [21] (232,965 nodes, 114,615,892 edges) used in [48]. Notably, the Protein dataset is challenging due to the high average node degree and small graph diameter which may exacerbates limitations of GNNs [1].

Model	OGBN-Products		OGBN-Proteins	
	Mean acc.	Std. acc.	Mean acc.	Std. acc.
SAGE fp32	0.7862	0.0043	0.7734	0.0041
SAGE bin f.s.	0.7300	0.0156	0.7497	0.0047
SAGE bin l.m.	0.7260	0.0153	-	-
GCN fp32	0.7564	0.0021	0.7254	0.0044

Table 3. Final test accuracy on the OGB-Product and OGB-Protein node property prediction benchmarks, averaged over 10 runs.

We implemented BinSAGE according to Eq. 15, details of the architecture can be found in Section 5.2. For OGB-product, we use logit matching only for distillation and no PReLU activation. For OGB-protein, we use PReLU activations and no distillation due to the very large number of edges in the dataset. We use channel-wise rescaling only for both to maximize scalability. On OGB-Product, we did not observe a statistically significant different between training the model from scratch and three-stage distillation with logit matching: in both cases, the full binary model came within 5-6% of the full-precision model. On OGB-protein, the simple binary network trained from scratch is within 3% of the accuracy of the full-precision network and outperforms the full-precision GCN. This suggests other strategies to improve model scalability, in this case sampling, can be successfully combined with our binarisation method.

### 5.1. Speed on embedded hardware

In order to measure the speed improvements yielded by our binary conversion scheme, we chose to benchmark it on a Raspberry Pi 4B board with 4GB of RAM and a Broadcom BCM2711 Quad core Cortex-A72 (ARM v8) 64-bit SoC @ 1.5GHz, running Manjaro 64-bit. The Pi is a popular, cheap, and readily available ARM-based platform, and is thus a good fit for our experiments.

We benchmark five DGCNN models, in order to measure the speedup for each successive optimization. The specificities of each model are given in Table 4. The input size is set to 1024 points with a batch size of 8, 40 output classes, and 20 nearest neighbors.

Model	Bin Weights	Bin Feats	Hamm Dist	Optim
DGCNN				
BDGCNN RF	✓			
BDGCNN BF	✓	✓		
BDGCNN BF H	✓	✓	✓	
BDGCNN BF H O	✓	✓	✓	✓

Table 4. Features of benchmarked models. Hamm Dist = Pairwise Hamming distance instead of L2. Optim = Hamming distance implemented in ARM NEON operations on bit-packed features (simulated).

We implement and convert our models to Tensorflow Lite using LARQ [15], an open-source library for binarized neural networks, and benchmark them using the LARQ Compute Engine (LCE) tool. Once converted, the model file size is only 340KB down from 7.2MB, for a 20x reduction.

Figure 4 shows the benchmark results. Our optimized binary model halves the run-time, thus achieving a substantial speedup. Peak memory usage is also significantly reduced, from 575MB to 346MB.

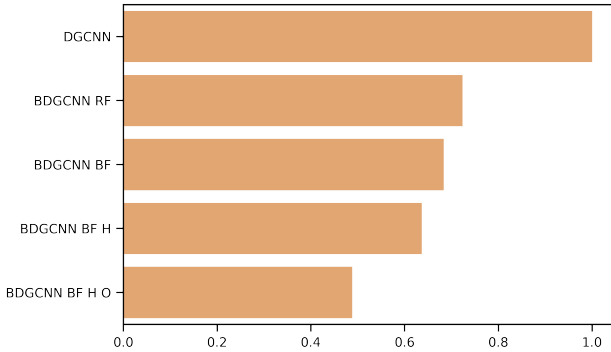


Figure 4. Relative run time on Raspberry Pi 4 compared to the base DGCNN model. A 2x speedup is achieved by our final optimized model. Run times computed with LCE benchmark tool over 50 runs.

Unfortunately, we did not have an optimized version of the hamming distance in the LARQ runtime at the time of writing. Thus, the final result is simulated by profiling the run-time of the un-optimized Hamming distance, and dividing it by 32, which is the speedup we would obtain by implementing this operation using ARM NEON instructions.

## 5.2. Implementation details

For DGCNN, we follow the architecture of [49]. For GraphSAGE, we use the baseline architecture of the OGB benchmarks [23]; that is, three layers with 256 hidden features and mean aggregation. We use three knowledge transfer points for LSP on DGCNN, one after each EdgeConv layer except for the first layer (the  $k$ -NN and graph features are computed on the 3D coordinates of the point cloud and do not change). All binary models assume binary inputs. For DGCNN, the final layer of the MLP classifier is kept real-weighted, as is customary in the binary neural network literature due to the small number of parameters, but the input features are binarized. For GraphSAGE, all three layers

have binary weights.

Our models are implemented in Pytorch. We use the reference implementation of DGCNN by the authors as a starting point, and Pytorch Geometric [12] for GraphSAGE and OGB [23]. We use the Adam optimizer [28]. We train the DGCNN models for 350 epochs on 4 Nvidia 2080 Ti GPUs, the initial learning rate of stage 1 is set to  $1e-3$  and for stage 2 to  $5e-4$ , with learning rate decay of 0.5 at 50% and 75% of the total number of epochs. For stage 3, we set the learning rate to  $1e-4$  and decay by a factor of 0.5 every 50 epochs. We trained GraphSAGE according to the OGB benchmark methodology, using the provided training, validation, and test sets. We trained all models for 20 epochs and averaged the performance over 10 runs. We used  $\ell_2$  regularization on the learnable scaling factors only for GraphSAGE with a weight  $\lambda = 1e-4$ . For logit matching, we set  $T = 3$  and  $\alpha = 1e-1$ . For LSP, we set  $\lambda_{LSP} = 1e2$ .

## 6. Conclusion

In this work, we introduce a binarization scheme for graph neural networks based on the XNOR-Net++ methodology and knowledge distillation. We study the impact of various schemes and design decisions on the final performance of binarized graph neural networks, and show that our approach allows us to closely match or equal the performance of floating-point models on difficult benchmarks with significant savings in terms of memory consumption and inference time. We further demonstrate that dynamic graph neural networks can be trained to high accuracy with binary node features, enabling fast construction of the dynamic graph at inference time through efficient Hamming-based algorithms, and further relating dynamic graph models to metric learning and fast hashing. Our DGCNN in Hamming space equals the performance of the full floating point model when trained with floating point weights, and offers large speed and memory savings when trained with binary weights. Future work will investigate further improving the accuracy of full binary dynamic graph models, and inference with fast  $k$ -NN search in Hamming space.

**Acknowledgements** M. B. is supported by a Department of Computing scholarship from Imperial College London, and a Qualcomm Innovation Fellowship. S.Z. was partially funded by the EPSRC Fellowship DEFORM: Large Scale Shape Analysis of Deformable Models of Humans (EP/S010203/1) and an Amazon AWS Machine Learning Research Award. The authors are grateful to Inria Sophia Antipolis - Méditerranée "Nef" computation cluster for providing resources and support.



## A. DGCNN and ModelNet40

In this appendix, we provide details of the DGCNN model and of the ModelNet40 dataset omitted from the main text for brevity.

**ModelNet40 classification** The ModelNet40 dataset [53] contains 12311 shapes representing 3D CAD models of man-made objects pertaining to 40 categories. We follow the experimental setting of [50] and [8]. We keep 9843 shapes for training and 2468 for testing. We uniformly sample 1024 points on mesh faces weighted by surface area and normalize the resulting point clouds in the unit sphere. The original meshes are discarded. Only the 3D cartesian coordinates  $(x, y, z)$  of the points are used as input. We use the same data augmentation techniques (random scaling and perturbations) as [50] and base our implementation on the author’s public code<sup>1</sup>. We report the overall accuracy as the model score.

**Model architecture** All DGCNN models use 4 EdgeConv (or BinEdgeConv or XorEdgeConv) layers with 64, 64, 128, and 256 output channels and no spatial transformer networks. According to the architecture of [50], the output of the four graph convolution layers are concatenated and transformed to node embeddings of dimension 1024. We use both global average pooling and global max pooling to obtain graph embeddings from all node embeddings; the resulting features are concatenated and fed to a three layer MLP classifier with output dimensions 512, 256, and 40 (the number of classes in the dataset). We use dropout with probability  $p = 0.5$ .

## B. Low-level implementation

This appendix provides further details on the low-level implementation and memory cost of our models.

### B.1. Parameter counts

We report the counts of binary and floating-point parameters for the baseline DGCNN and our binary models (stage 3) in Table 5.

Model	FP32 Param.	Bin. param.	Total param.
Baseline	1,812,648	0	1,812,648
BF1	11,064	1,804,672	1,815,736
BF2	11,064	1,804,672	1,815,736
RF	15,243	1,804,672	1,819,915

Table 5. Number of parameters given by torchsummaryX. Separated into FP and binary operations. 99.39% of the parameters are binary for **BF1** and **BF2**, 99.16% of the parameters are binary for **RF**.

As can be seen in Table 5, our binarization procedure introduces a few extra parameters, but over 99% of the network parameters are binary.

<sup>1</sup><https://github.com/WangYueFt/dgcnn/tree/master/pytorch>

### B.2. Profiling and optimization of DGCNN

In order to obtain the data from Section 5.1 of the main paper, we convert our models with the LARQ converter and benchmark them using the LCE benchmark utility.

The pairwise Hamming distance is naively implemented as a matrix multiplication operation (Eq. 21 of the main text), and we obtain the profiler data in Table 6, where we have highlighted the nodes used by that operation. However, not all nodes of these types belong to the three pairwise distances calculations. We thus provide in Table 7 the complete profiler output for only one distance calculation in binary space, of which there are three in the DGCNN models.

Node Type	Avg. ms	Avg %	Times called
TOPK_V2	488.007	22.18%	4
CONCATENATION	384.707	17.485%	6
FULLY_CONNECTED	171.175	7.77994%	32
PRELU	143.086	6.50329%	7
TILE	136.443	6.20137%	4
LceBconv2d	127.371	5.78904%	6
MAX_POOL_2D	122.743	5.5787%	5
MUL	105.993	4.81741%	11
SUB	92.382	4.19878%	4
LceQuantize	91.168	4.14361%	10
NEG	78.453	3.56571%	4
PACK	56.301	2.55889%	4
GATHER	55.989	2.54471%	4
CONV_2D	39.096	1.77692%	2
RESHAPE	35.091	1.59489%	82
ADD	28.557	1.29792%	6
TRANSPPOSE	23.829	1.08303%	36
AVERAGE_POOL_2D	8.071	0.366829%	1
SLICE	5.278	0.239886%	64
LceDequantize	5.174	0.235159%	4
SUM	1.132	0.0514497%	1
SQUARE	0.153	0.00695389%	1
SOFTMAX	0.01	0.000454502%	1

Table 6. LCE Profiler data for "BDGCNN BF H", summary by node types. In red: nodes that appear in Matmul op which can be rewritten as NEON operations for Hamming distance.

These operations account for 24% of the network’s run time. Thus, a speed-up of 32x of these operations would reduce them to around 1% of the network’s run time, which is negligible.

While we did not have an optimized version integrated with the LARQ runtime at the time of writing, optimizing the pairwise Hamming distance computation in binary space with ARM NEON (SIMD) operations is quite simple, since it can be implemented as `popcount(x XOR y)`. On bit-packed 64-bit data (conversion handled by LCE), with feature vectors of dimension 64, this can be written as:

```

1 #include "arm_neon.h"
2
3 // input data in feats
4 int8_t n_outs = npoints*(npoints-1)/2
5 int8_t* out = malloc(n_outs*sizeof(int8_t));
6 for(int i = 0; i < npoints; i++) {
7
8     // load first feature
9     uint32x2_t a = vld1_u32(feats + 8*i);

```

Node type	Avg. time	Avg? %	Operation name
TRANPOSE	5.91618	0.268874%	[bin_dgcnn_b.f1_h/MatMul_315]:208
SLICE	0.45752	0.020793%	[bin_dgcnn_b.f1_h/MatMul_316]:209
RESHAPE	0.43792	0.0199023%	[bin_dgcnn_b.f1_h/MatMul_317]:210
SLICE	0.14872	0.00675892%	[bin_dgcnn_b.f1_h/MatMul_318]:211
RESHAPE	0.22018	0.0100066%	[bin_dgcnn_b.f1_h/MatMul_319]:212
SLICE	0.16006	0.0072743%	[bin_dgcnn_b.f1_h/MatMul_320]:213
RESHAPE	0.22242	0.0101084%	[bin_dgcnn_b.f1_h/MatMul_321]:214
SLICE	0.16268	0.00739337%	[bin_dgcnn_b.f1_h/MatMul_322]:215
RESHAPE	0.21126	0.0096012%	[bin_dgcnn_b.f1_h/MatMul_323]:216
SLICE	0.15406	0.00700161%	[bin_dgcnn_b.f1_h/MatMul_324]:217
RESHAPE	0.20686	0.00940123%	[bin_dgcnn_b.f1_h/MatMul_325]:218
SLICE	0.1494	0.00678983%	[bin_dgcnn_b.f1_h/MatMul_326]:219
RESHAPE	0.21348	0.00970209%	[bin_dgcnn_b.f1_h/MatMul_327]:220
SLICE	0.15514	0.00705069%	[bin_dgcnn_b.f1_h/MatMul_328]:221
RESHAPE	0.2117	0.00962119%	[bin_dgcnn_b.f1_h/MatMul_329]:222
SLICE	0.15154	0.00688708%	[bin_dgcnn_b.f1_h/MatMul_330]:223
RESHAPE	0.17318	0.00787056%	[bin_dgcnn_b.f1_h/MatMul_331]:224
SLICE	0.15244	0.00692799%	[bin_dgcnn_b.f1_h/MatMul_332]:225
RESHAPE	0.16864	0.00766423%	[bin_dgcnn_b.f1_h/MatMul_333]:226
SLICE	0.15614	0.00709614%	[bin_dgcnn_b.f1_h/MatMul_334]:227
RESHAPE	0.1736	0.00788965%	[bin_dgcnn_b.f1_h/MatMul_335]:228
SLICE	0.15118	0.00687072%	[bin_dgcnn_b.f1_h/MatMul_336]:229
RESHAPE	0.17086	0.00776513%	[bin_dgcnn_b.f1_h/MatMul_337]:230
SLICE	0.149	0.00677165%	[bin_dgcnn_b.f1_h/MatMul_338]:231
RESHAPE	0.16924	0.0076915%	[bin_dgcnn_b.f1_h/MatMul_339]:232
SLICE	0.1505	0.00683982%	[bin_dgcnn_b.f1_h/MatMul_340]:233
RESHAPE	0.16894	0.00767787%	[bin_dgcnn_b.f1_h/MatMul_341]:234
SLICE	0.14994	0.00681437%	[bin_dgcnn_b.f1_h/MatMul_342]:235
RESHAPE	0.17058	0.0077524%	[bin_dgcnn_b.f1_h/MatMul_343]:236
SLICE	0.15018	0.00682528%	[bin_dgcnn_b.f1_h/MatMul_344]:237
RESHAPE	0.16996	0.00772422%	[bin_dgcnn_b.f1_h/MatMul_345]:238
SLICE	0.1496	0.00679892%	[bin_dgcnn_b.f1_h/MatMul_346]:239
RESHAPE	0.17112	0.00777694%	[bin_dgcnn_b.f1_h/MatMul_347]:240
TRANPOSE	0.41792	0.0189933%	[bin_dgcnn_b.f1_h/MatMul_348]:241
FULLY.CONNECTED	8.78396	0.399207%	[bin_dgcnn_b.f1_h/MatMul_349]:242
TRANPOSE	0.72016	0.0327293%	[bin_dgcnn_b.f1_h/MatMul_350]:243
FULLY.CONNECTED	8.64452	0.39287%	[bin_dgcnn_b.f1_h/MatMul_351]:244
TRANPOSE	0.71804	0.032633%	[bin_dgcnn_b.f1_h/MatMul_352]:245
FULLY.CONNECTED	8.63224	0.392312%	[bin_dgcnn_b.f1_h/MatMul_353]:246
TRANPOSE	0.72162	0.0327957%	[bin_dgcnn_b.f1_h/MatMul_354]:247
FULLY.CONNECTED	8.62624	0.392039%	[bin_dgcnn_b.f1_h/MatMul_355]:248
TRANPOSE	0.68654	0.0312014%	[bin_dgcnn_b.f1_h/MatMul_356]:249
FULLY.CONNECTED	8.6722	0.394128%	[bin_dgcnn_b.f1_h/MatMul_357]:250
TRANPOSE	0.69886	0.0317613%	[bin_dgcnn_b.f1_h/MatMul_358]:251
FULLY.CONNECTED	8.6892	0.394901%	[bin_dgcnn_b.f1_h/MatMul_359]:252
TRANPOSE	0.71076	0.0323021%	[bin_dgcnn_b.f1_h/MatMul_360]:253
FULLY.CONNECTED	8.70248	0.395504%	[bin_dgcnn_b.f1_h/MatMul_361]:254
TRANPOSE	0.71256	0.0323839%	[bin_dgcnn_b.f1_h/MatMul_362]:255
FULLY.CONNECTED	8.76456	0.398326%	[bin_dgcnn_b.f1_h/MatMul_363]:256
PACK	13.822	0.628173%	[bin_dgcnn_b.f1_h/MatMul_364]:257
SUB	29.9335	1.3604%	[bin_dgcnn_b.f1_h/sub_3;bin_dgcnn_b.f1_h/MatMul_3;b]:258

Table 7. LCE Profiler data for a single Hamming distance computation as a matrix multiplication, in "BDGCNN BF H".

```

10 for(int j = i; j < npoints; j++) {
11
12     //load second feature
13     uint32x2_t b = vld1_u32(feats + 8*j);
14
15     b = veor_u32(a, b); // XOR op
16
17
18     // popcount op
19     int8x8_t c = vreinterpret_u32_s8(b);
20     c = vcnt_s8(c);
21
22     // reduce to single number
23     // by adding as a tree

```

```

24     int64x1_t res;
25     res = vpaddd_s32(vpaddd_s16(vpaddd_s8(c)));
26
27     //store the output (last 8 bits)
28     int8x8_t res8 = vreinterpret_s64_s8(res);
29     out[j + npoints*j] = vget_lane_s8(res8, 7);
30 }
31

```

Listing 1. Implementation of pairwise Hamming distance in ARM NEON intrinsics (for readability). Note that this code actually treats 64 features at a time and could thus provide a 64x speedup (or more by grouping loads and writes with vld4). We use 32x

as a conservative estimate since we couldn’t account for LCE’s bit-packed conversion.

”TopK” operations account for 22% of the runtime and we view them as incompressible in our simulation (Table 6). It is possible that they could be written in NEON as well, however, this optimization is not as trivial as the Hamming distance one. Remaining operations, such as ”Concatenation”, cannot be optimized further.

Contrary to simpler GNNs such as GCN, DGCNN is quite computationally intensive and involves a variety of operations on top of simple dot products, which makes it an interesting challenge for binarization, and illustrate that for complex graph neural networks more efforts are required, such as redefining suitable edge messages for binary graph features, or speeding-up pairwise distances computations, as done in this work. The inherent complexity also limits the attainable speedups from binarization, as shown by the large portion of the runtime taken by memory operations (concatenation) and top-k.

## C. Details regarding GraphSAGE

In all experiments, the architecture used is identical to that used as a baseline by the OGB team. We report the accuracy following verbatim the experimental procedure of the OGB benchmark, using the suitable provided evaluators and dataset splits. Due to the very large number of edges in the dataset, we were unable to implement LSP in a sufficiently scalable manner (although the forward pass of the similarity computation can be implemented efficiently, the gradient of the similarity with respect to the node features is a tensor of size  $|\mathcal{E}| \times |\mathcal{V}| \times D$  where  $|\mathcal{E}|$  is the number of edges in the graph,  $|\mathcal{V}|$  the number of nodes, and  $D$  the dimension of the features. Although the tensor is sparse, Pytorch currently did not have sufficient support of sparse tensors for gradients. We therefore chose not to include the results in the main text. We report the results of our binary GraphSAGE models, against two floating-point baselines: GraphSAGE and GCN.

## D. Balance functions

For completeness, we also report the results at stage 2 of the multi-stage distillation scheme in Table 8. It is apparent that the additional operations degraded the performance not only for the full-binary models of stage 3, but also for the models for which all inputs are binary but weights are real.

## E. Table of mathematical operators

## References

[1] Uri Alon and Eran Yahav. On the Bottleneck of Graph Neural Networks and its Practical Implications. jun 2020. 7

[2] Gaétan Bahl, Lionel Daniel, Matthieu Moretti, and Florent Lafarge. Low-power neural networks for semantic segmentation of satellite images. In *Proceedings of the IEEE International Conference on Computer Vision Workshops*, page 0, 2019. 1

[3] Yoshua Bengio, Nicholas Léonard, and Aaron Courville. Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv preprint arXiv:1308.3432*, 2013. 3

[4] Michael M. Bronstein, Joan Bruna, Yann Lecun, Arthur Szlam, and Pierre Vandergheynst. Geometric Deep Learning: Going beyond Euclidean data, 2017. 1

[5] Joan Bruna, Wojciech Zaremba, Arthur Szlam, and Yann Lecun. Spectral networks and deep locally connected networks on graphs. *2nd International Conference on Learning Representations, ICLR 2014 - Conference Track Proceedings*, pages 1–14, 2014. 2

[6] Adrian Bulat and Georgios Tzimiropoulos. XNOR-Net++: Improved binary neural networks. In *30th British Machine Vision Conference 2019, {BMVC} 2019, Cardiff, UK, September 9-12, 2019*, page 62. {BMVA} Press, 2019. 2, 3

[7] Sergio Casas, Cole Gulino, Renjie Liao, and R. Urtasun. Spaggn: Spatially-aware graph neural networks for relational behavior forecasting from sensor data. *2020 IEEE International Conference on Robotics and Automation (ICRA)*, pages 9491–9497, 2020. 1

[8] R Qi Charles, Hao Su, Mo Kaichun, and Leonidas J Guibas. PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation. In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 77–85. IEEE, jul 2017. 9

[9] Wei Lin Chiang, Yang Li, Xuanqing Liu, Samy Bengio, Si Si, and Cho Jui Hsieh. Cluster-GCN: An efficient algorithm for training deep and large graph convolutional networks. In *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2019. 2

[10] Michaël Defferrard, Xavier Bresson, and Pierre Vandergheynst. Convolutional neural networks on graphs with fast localized spectral filtering. In D. D. Lee, M. Sugiyama, U. V. Luxburg, I. Guyon, and R. Garnett, editors, *Advances in Neural Information Processing Systems 29*, pages 3844–3852. Curran Associates, Inc., 2016. 2

[11] Sybil Derrible and Christopher Kennedy. Applications of graph theory and network science to transit network design. *Transport reviews*, 31(4):495–519, 2011. 1

[12] Matthias Fey and Jan E Lenssen. Fast Graph Representation Learning with PyTorch Geometric. In *ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019. 8

[13] Matthias Fey, Jan Eric Lenssen, Frank Weichert, and Heinrich Muller. SplineCNN: Fast Geometric Deep Learning with Continuous B-Spline Kernels. *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, pages 869–877, 2018. 2

[14] Fabrizio Frasca, Emanuele Rossi, Davide Eynard, Ben Chamberlain, Michael Bronstein, and Federico Monti. SIGN: Scalable Inception Graph Neural Networks. apr 2020. 2

[15] Lukas Geiger and Plumerai Team. Larc: An Open-Source Library for Training Binarized Neural Networks. *Journal of Open Source Software*, 5(45):1746, 2020. 8

- [16] Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In *34th International Conference on Machine Learning, ICML 2017*, 2017. 2
- [17] Shunwang Gong, Mehdi Bahri, Michael M. Bronstein, and Stefanos Zafeiriou. Geometrically Principled Connections in Graph Neural Networks. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11412–11421. IEEE, jun 2020. 1
- [18] Yunchao Gong, Liu Liu, Ming Yang, and Lubomir Bourdev. Compressing Deep Convolutional Networks using Vector Quantization, 2014. 2
- [19] Marco Gori, Gabriele Monfardini, and Franco Scarselli. A new model for learning in Graph domains. *Proceedings of the International Joint Conference on Neural Networks*, 2(May 2014):729–734, 2005. 2
- [20] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*, volume 2017-Decem, pages 1025–1035, 2017. 2, 3
- [21] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs, 2018. 7
- [22] Geoffrey Hinton, Oriol Vinyals, and Jeffrey Dean. Distilling the Knowledge in a Neural Network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015. 2, 3
- [23] Weihua Hu, Matthias Fey, Marinka Zitnik, Yuxiao Dong, Hongyu Ren, Bowen Liu, Michele Catasta, and Jure Leskovec. Open Graph Benchmark: Datasets for Machine Learning on Graphs, 2020. 7, 8
- [24] Kexin Huang, Cao Xiao, Lucas Glass, Marinka Zitnik, and Jimeng Sun. Skipgnn: Predicting molecular interactions with skip-graph networks. *arXiv preprint arXiv:2004.14949*, 2020. 1
- [25] Itay Hubara, Matthieu Courbariaux, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio. Binarized Neural Networks. In D Lee, M Sugiyama, U Luxburg, I Guyon, and R Garnett, editors, *Advances in Neural Information Processing Systems*, volume 29, pages 4107–4115. Curran Associates, Inc., 2016. 2
- [26] Anees Kazi, Luca Cosmo, Nassir Navab, and Michael Bronstein. Differentiable graph module (DGM) for graph convolutional networks, 2020. 3
- [27] Abdullah Khanfor, Amal Nammouchi, Hakim Ghazzai, Ye Yang, Mohammad R. Haider, and Yehia Massoud. Graph neural networks-based clustering for social internet of things, 2020. 1
- [28] Diederik P Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. *Pattern Recognition Letters*, 94:172–179, dec 2014. 8
- [29] Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *5th International Conference on Learning Representations, ICLR 2017 - Conference Track Proceedings*, 2017. 2, 3
- [30] H Lai, Y Pan, S Liu, Z Weng, and J Yin. Improved Search in Hamming Space Using Deep Multi-Index Hashing. *IEEE Transactions on Circuits and Systems for Video Technology*, 29(9):2844–2855, 2019. 5
- [31] Edgar Lemaire, Matthieu Moretti, Lionel Daniel, Benoît Miramond, Philippe Millet, Frederic Feresin, and Sébastien Bilavarn. An FPGA-based Hybrid Neural Network accelerator for embedded satellite image classification. In *2020 IEEE International Symposium on Circuits and Systems (ISCAS)*, pages 1–5. IEEE, 2020. 1
- [32] Guohao Li, Matthias Muller, Ali Thabet, and Bernard Ghanem. DeepGCNs: Can GCNs go as deep as CNNs? In *Proceedings of the IEEE International Conference on Computer Vision*, 2019. 1
- [33] Wei Liu, Jun Wang, Rongrong Ji, Yu Gang Jiang, and Shih Fu Chang. Supervised hashing with kernels. In *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 2012. 5
- [34] Brais Martinez, Jing Yang, Adrian Bulat, and Georgios Tzimiropoulos. Training binary neural networks with real-to-binary convolutions. In *International Conference on Learning Representations*, 2020. 2, 3
- [35] Batul J Mirza, Benjamin J Keller, and Naren Ramakrishnan. Studying recommendation algorithms by graph analysis. *Journal of intelligent information systems*, 20(2):131–160, 2003. 1
- [36] Federico Monti, Davide Boscaini, Jonathan Masci, Emanuele Rodolá, Jan Svoboda, and Michael M Bronstein. Geometric deep learning on graphs and manifolds using mixture model CNNs. *Proceedings - 30th IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2017*, 2017-Janua:5425–5434, 2017. 2
- [37] Jan Motl and Oliver Schulte. The ctu prague relational learning repository, 2015. 3, 7
- [38] Walter Nelson, Marinka Zitnik, Bo Wang, Jure Leskovec, Anna Goldenberg, and Roded Sharan. To Embed or Not: Network Embedding as a Paradigm in Computational Biology. *Frontiers in Genetics*, 10:381, 2019. 1
- [39] Mohammad Norouzi and David J. Fleet. Minimal loss hashing for compact binary codes. In *Proceedings of the 28th International Conference on Machine Learning, ICML 2011*, 2011. 5
- [40] Mohammad Norouzi, David J. Fleet, and Ruslan Salakhutdinov. Hamming distance metric learning. In *Advances in Neural Information Processing Systems*, 2012. 5
- [41] Mohammad Norouzi, Ali Punjani, and David J. Fleet. Fast exact search in hamming space with multi-index hashing. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2014. 5
- [42] Haotong Qin, Zhongang Cai, Mingyuan Zhang, Yifu Ding, Haiyu Zhao, Shuai Yi, Xianglong Liu, and Hao Su. BiPointNet: Binary Neural Network for Point Clouds. oct 2020. 3, 5
- [43] Mohammad Rastegari, Vicente Ordonez, Joseph Redmon, and Ali Farhadi. XNOR-Net: ImageNet Classification Using Binary. *Eccv2016*, 2016. 2, 4, 14
- [44] Franco Scarselli, Marco Gori, Ah Chung Tsoi, Markus Hagenbuchner, and Gabriele Monfardini. The graph neural network model. *IEEE transactions on neural networks / a publication of the IEEE Neural Networks Council*, 20(1):61–80, 2009. 2



- [45] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lió, and Yoshua Bengio. Graph Attention Networks. *ICLR*, oct 2018. 2, 3
- [46] Nitika Verma, Edmond Boyer, and Jakob Verbeek. FeaStNet: Feature-Steered Graph Convolutions for 3D Shape Analysis. In *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, jun 2018. 2
- [47] Hanchen Wang, Defu Lian, Ying Zhang, Lu Qin, Xiangjian He, Yiguang Lin, and Xuemin Lin. Binarized Graph Neural Network. apr 2020. 3, 5, 7
- [48] Junfu Wang, Yunhong Wang, Zhen Yang, Liang Yang, and Yuanfang Guo. Bi-GCN: Binary Graph Convolutional Network. oct 2020. 3, 7
- [49] Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E. Sarma, Michael M. Bronstein, and Justin M. Solomon. Dynamic graph Cnn for learning on point clouds. *ACM Transactions on Graphics*, 2019. 2, 3, 5, 8
- [50] Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E Sarma, Michael M Bronstein, and Justin M Solomon. Dynamic Graph CNN for Learning on Point Clouds. *ACM Trans. Graph.*, 38(5):146:1—146:12, oct 2019. 9
- [51] Felix Wu, Tianyi Zhang, Amauri Holanda de Souza, Christopher Fifty, Tao Yu, and Kilian Q. Weinberger. Simplifying graph convolutional networks. In *36th International Conference on Machine Learning, ICML 2019*, 2019. 2
- [52] Zonghan Wu, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and Philip S. Yu. A Comprehensive Survey on Graph Neural Networks. *IEEE Transactions on Neural Networks and Learning Systems*, 2020. 2
- [53] Zhirong Wu, Shuran Song, Aditya Khosla, Fisher Yu, Linguang Zhang, Xiaoou Tang, and Jianxiong Xiao. 3D ShapeNets: A deep representation for volumetric shapes. In *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 2015. 9
- [54] Yiding Yang, Jiayan Qiu, Mingli Song, Dacheng Tao, and Xinchao Wang. Distilling Knowledge From Graph Convolutional Networks. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, jun 2020. 3
- [55] Sergey Zagoruyko and Nikos Komodakis. Paying More Attention to Attention: Improving the Performance of Convolutional Neural Networks via Attention Transfer. In *ICLR*, 2017. 2
- [56] Hanqing Zeng, Hongkuan Zhou, Ajitesh Srivastava, Rajgopal Kannan, and Viktor Prasanna. GraphSAINT: Graph Sampling Based Inductive Learning Method. In *International Conference on Learning Representations*, 2020. 2
- [57] Weishan Zhang, Yafei Zhang, Liang Xu, Jiehan Zhou, Yan Liu, Mu Gu, Xin Liu, and Su Yang. Modeling iot equipment with graph neural networks. *IEEE Access*, PP:1–1, 03 2019. 1
- [58] Shuchang Zhou, Yuxin Wu, Zekun Ni, Xinyu Zhou, He Wen, and Yuheng Zou. DoReFa-Net: Training Low Bitwidth Convolutional Neural Networks with Low Bitwidth Gradients. *CoRR*, abs/1606.06160, 2016. 2
- [59] Marinka Zitnik, SM Rok Sasic, and Jure Leskovec. Biosnap datasets: Stanford biomedical network dataset collection. Note: <http://snap.stanford.edu/biodata> Cited by, 5(1), 2018. 1

Model	Stage	KNN	LSP	Global balance	Edge balance	Acc
BF2	2	H	-	-	Median	90.07
BF2	2	H	-	-	Mean	83.87
BF2	2	H	$\ell_2$	Median		87.60
BF2	2	H	$\ell_2$	Mean		89.47
Baseline BF2	2	H	$\ell_2$	None		<b>91.57</b>

Table 8. Effect of additional balance functions on models with binary activations but floating-point weights. The performance of the baseline model suffers with the introduction of either mean or median centering prior to quantization.

Symbol	Name	Description
$\ \cdot\ _H$	Hamming norm	Number of non-zero (or not -1) bits in a binary vector
$d(\cdot, \cdot)_H$	Hamming distance	Number of bits that differ between two binary vectors, equivalent to <code>popcount(xor())</code>
$\oplus$	Exclusive OR (XOR)	$1 \oplus 1 = -1 \oplus -1 = -1$ , $-1 \oplus 1 = 1 \oplus -1 = 1$
$\odot$	Hadamard product	Element-wise product between tensors
$\otimes$	Binary-real or Binary-binary dot product or convolution	Equivalent to <code>popcount(xnor())</code> ( <i>i.e.</i> no multiplications) for binary tensors
$\otimes$	Outer product	
$\star$	Dot product or convolution	Denoted by $*$ in [43]
$ \mathcal{X} $	Cardinal of a set $\mathcal{X}$	Number of elements in the set
$\cdot  \cdot$	Concatenation	
$:=$	Definition	

Table 9. Table of the mathematical operators used in the manuscript.