

Machine Learning 10-315

Maria Florina Balcan
Machine Learning Department
Carnegie Mellon University

03/29/2019

Today:


- Artificial neural networks
- Backpropagation

Reading:

- Mitchell: Chapter 4
- Bishop: Chapter 5

Artificial Neural Network (ANN)

- **Biological systems** built of very complex webs of interconnected neurons.
- Highly connected to other neurons, and performs computations by combining signals from other neurons.
- Outputs of these computations may be transmitted to one or more other neurons.



- **Artificial Neural Networks** built out of a densely interconnected set of simple units (e.g., sigmoid units).
- Each unit takes real-valued inputs (possibly the outputs of other units) and produces a real-valued output (which may become input to many other units).

Connectionist Models

Consider humans:



- Neuron switching time $\sim .001$ second
- Number of neurons $\sim 10^{10}$
- Connections per neuron $\sim 10^{4-5}$
- Scene recognition time $\sim .1$ second
- 100 inference steps doesn't seem like enough

→ much parallel computation

Properties of artificial neural nets (ANN's):

- Many neuron-like threshold switching units
- Many weighted interconnections among units
- Highly parallel, distributed process

Multilayer Networks of Sigmoid Units





input: two features from spectral analysis of a spoken sound

output: vowel sound occurring in the context “h__d”


ALVINN

[Pomerleau 1993]



Artificial Neural Networks to learn $f: X \rightarrow Y$

- f_w typically a non-linear function, $f_w: X \rightarrow Y$
- X feature space: (vector of) continuous and/or discrete vars
- Y output space: (vector of) continuous and/or discrete vars
- f_w network of basic units




Learning algorithm: given $(x_d, t_d)_{d \in D}$, train weights w of all units to minimize sum of squared errors of predicted network outputs.

Find parameters w to minimize
$$\sum_{d \in D} (f_w(x_d) - t_d)^2$$

Use gradient descent!


What type of units should we use?

- Classifier is a multilayer *network of units*.
- Each *unit* takes some inputs and produces one output. Output of one unit can be the input of another.



Multilayer network of Linear units?

- Advantage: we know how to do gradient descent on linear units



Problem: linear of linear is just linear.

$$z = w_{3,1}(\mathbf{w}_1 \cdot \mathbf{x}) + w_{3,2}(\mathbf{w}_2 \cdot \mathbf{x}) = (\mathbf{w}_{3,1}\mathbf{w}_1 + \mathbf{w}_{3,2}\mathbf{w}_2) \cdot \mathbf{x} = \text{linear}$$


Multilayer network of Perceptron units?

- Advantage: Can produce highly non-linear decision boundaries!

Input layer

Hidden layer

Output layer



Threshold function: $\int x = 1$ if x is positive, 0 if x is negative.

Problem: discontinuous threshold is not differentiable. Can't do gradient descent.


Multilayer network of sigmoid units

- Advantage: Can produce highly non-linear decision boundaries!
- Sigmoid is differentiable, so can use gradient descent

Input layer

Hidden layer


Output layer



$$\sigma(x) = \frac{1}{1 + e^{-x}} = \text{graph of sigmoid function}$$

Very useful in practice!

The Sigmoid Unit



σ is the sigmoid function; $\sigma(x) = \frac{1}{1+e^{-x}}$

Nice property: $\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$

We can derive gradient descent rules to train

- One sigmoid unit
- *Multilayer networks* of sigmoid units \rightarrow Backpropagation


Gradient Descent to Minimize Squared Error

Goal: Given $(x_d, t_d)_{d \in D}$ find w to minimize $E_D[w] = \frac{1}{2} \sum_{d \in D} (f_w(x_d) - t_d)^2$

Batch mode Gradient Descent:

Do until satisfied

1. Compute the gradient $\nabla E_D[w]$
2. $w \leftarrow w - \eta \nabla E_D[w]$



$$\nabla E[w] \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

Incremental (stochastic) Gradient Descent:

Do until satisfied

- For each training example d in D
 1. Compute the gradient $\nabla E_d[w]$
 2. $w \leftarrow w - \eta \nabla E_d[w]$

$$E_d[w] \equiv \frac{1}{2} (t_d - o_d)^2$$

Note: Incremental Gradient Descent can approximate Batch Gradient Descent arbitrarily closely if η made small enough

Gradient descent in weight space

Goal: Given $(x_d, t_d)_{d \in D}$ find w to minimize $E_D[w] = \frac{1}{2} \sum_{d \in D} (f_w(x_d) - t_d)^2$

This error measure defines a surface over the hypothesis (i.e. weight) space





figure from Cho & Chow, *Neurocomputing* 1999

Gradient descent in weight space

Gradient descent is an iterative process aimed at finding a minimum in the error surface.

on each iteration

- current weights define a point in this space
- find direction in which error surface descends most steeply
- take a step (i.e. update weights) in that direction




Gradient descent in weight space

Calculate the gradient of E : $\nabla E(\mathbf{w}) = \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$

Take a step in the opposite direction

$$D\mathbf{w} = -h \nabla E(\mathbf{w})$$

$$Dw_i = -h \frac{\partial E}{\partial w_i}$$



Taking derivative: chain rule


Recall the chain rule from calculus

$$y = f(u)$$


$$u = g(x)$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$

Gradient Descent for the Sigmoid Unit†



Given $(x_d, t_d)_{d \in D}$ find \mathbf{w} to minimize $\sum_{d \in D} (o_d - t_d)^2$



o_d = observed unit output for x_d

$$o_d = \sigma(\text{net}_d); \quad \text{net}_d = \sum_i w_i x_{i,d}$$


$$\begin{aligned} \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 = \frac{1}{2} \sum_{d \in D} \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_{d \in D} 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) = \sum_{d \in D} (t_d - o_d) \left(-\frac{\partial o_d}{\partial w_i} \right) \\ &= - \sum_{d \in D} (t_d - o_d) \frac{\partial o_d}{\partial \text{net}_d} \frac{\partial \text{net}_d}{\partial w_i} \end{aligned}$$

But we know: $\frac{\partial o_d}{\partial \text{net}_d} = \frac{\partial \sigma(\text{net}_d)}{\partial \text{net}_d} = o_d(1 - o_d)$ and $\frac{\partial \text{net}_d}{\partial w_i} = \frac{\partial (\mathbf{w} \cdot \mathbf{x}_d)}{\partial w_i} = x_{i,d}$

So: $\frac{\partial E}{\partial w_i} = - \sum_{d \in D} (t_d - o_d) o_d (1 - o_d) x_{i,d}$

Gradient Descent for the Sigmoid Unit

Given $(x_d, t_d)_{d \in D}$ find \mathbf{w} to minimize $\sum_{d \in D} (o_d - t_d)^2$



o_d = observed unit output for x_d

$$o_d = \sigma(\text{net}_d); \quad \text{net}_d = \sum_i w_i x_{i,d}$$

$$\frac{\partial E}{\partial w_i} = - \sum_{d \in D} (t_d - o_d) o_d (1 - o_d) x_{i,d}$$


δ_d error term $t_d - o_d$ multiplied by $o_d(1 - o_d)$ that comes from the derivative of the sigmoid function

$$\frac{\partial E}{\partial w_i} = - \sum_{d \in D} \delta_d x_{i,d}$$

Update rule: $w \leftarrow w - \eta \nabla E[w]$

Gradient Descent for Multilayer Networks

Given $(x_d, t_d)_{d \in D}$ find \mathbf{w} to minimize $\frac{1}{2} \sum_{d \in D} \sum_{k \in \text{Outputs}} (o_{k,d} - t_{kd})^2$



Backpropagation Algorithm

Incremental/stochastic gradient descent

Initialize all weights to small random numbers.

Until satisfied, Do:

- **For** each training example (x, t) **do**:
 1. Input the training example to the network and compute the network outputs
 2. For each output unit k :

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$


3. For each hidden unit h :

$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{h,k} \delta_k$$

4. Update each network weight $w_{i,j}$

$$w_{i,j} \leftarrow w_{i,j} + \Delta w_{i,j}$$

where $\Delta w_{i,j} = \eta \delta_j x_{i,j}$



o = observed unit output

t = target output

x = input

$x_{i,j}$ = i th input to j th unit

w_{ij} = wt from i to j

More on Backpropagation



- Gradient descent over entire *network* weight vector
- Easily generalized to arbitrary directed graphs
- Will find a local, not necessarily global error minimum
 - In practice, often works well (can run multiple times)

- Often include weight *momentum* α

$$\Delta w_{i,j}(n) = \eta \delta_j x_{i,j} + \alpha \Delta w_{i,j}(n-1)$$


- Minimizes error over *training* examples
 - Will it generalize well to subsequent examples?
- Training can take thousands of iterations \rightarrow slow!
- Using network after training is very fast

Overfitting in ANNs



- Validation/generalization error first decreases, then increases.
- Weights tuned to fit the idiosyncrasies of the training examples that are not representative of the general distribution.
- Stop when lowest error over validation set.
- Not always obvious when lowest error over validation set has been reached.

Dealing with Overfitting



Our learning algorithm involves a parameter

n = number of gradient descent iterations

How do we choose n to optimize future error?

- Separate available data into training and validation set
- Use training to perform gradient descent
- $n \leftarrow$ number of iterations that optimizes validation set error

Dealing with Overfitting

- Regularization techniques
 - norm constraint
 - dropout
 - batch normalization
 - data augmentation
 - early stopping
 - ...

Convergence of Backpropagation

Gradient descent to some local minimum

- Perhaps not global minimum...
- Add momentum
- Stochastic gradient descent
- Train multiple nets with different initial weights

Nature of convergence

- Initialize weights near zero
- Therefore, initial networks near-linear
- Increasingly non-linear functions possible as training progresses

Expressive Capabilities of ANNs

Boolean functions:

- Every Boolean function can be represented by a network with a single hidden layer
- But might require exponential (in number of inputs) hidden units

Continuous functions:


- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer [Cybenko 1989; Hornik et al. 1989]
- Any function can be approximated to arbitrarily accuracy by a network with two hidden layers [Cybenko 1988]

Representing Simple Boolean Functions

Inputs $x_i \in \{0,1\}$

Or function

$$x_{i_1} \vee x_{i_2} \vee \cdots \vee x_{i_k}$$




$w_i = 1$ if i is an i_j

$w_i = 0$ otherwise

And function

$$x_{i_1} \wedge x_{i_2} \wedge \cdots \wedge x_{i_k}$$




$w_i = 1$ if i is an i_j

$w_i = 0$ otherwise

And with negations

$$x_{i_1} \wedge \bar{x}_{i_2} \wedge \cdots \wedge x_{i_k}$$



$w_i = 1$ if i is i_j not negated

$w_i = -1$ if i is i_j negated

$w_i = 0$ otherwise

$t = \#$ not negated

General Boolean functions

Every Boolean function can be represented by a network with a single hidden layer; might require exponential # of hidden units

Can write any Boolean function as a truth table:


000		+
001		—
010		—
011		+
100		—
101		—
110		+
111		+

View as OR of ANDs, with one AND for each positive entry.

$$\bar{x}_1\bar{x}_2\bar{x}_3 \vee \bar{x}_1x_2x_3 \vee x_1x_2\bar{x}_3 \vee x_1x_2x_3$$

Then combine AND and OR networks into a 2-layer network.

Learning Hidden Layer Representations




A target function:

Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned??

Learning Hidden Layer Representations

A network:



Learned hidden layer representation:

Input		Hidden Values		Output
10000000	→	.89 .04 .08	→	10000000
01000000	→	.01 .11 .88	→	01000000
00100000	→	.01 .97 .27	→	00100000
00010000	→	.99 .97 .71	→	00010000
00001000	→	.03 .05 .02	→	00001000
00000100	→	.22 .99 .99	→	00000100
00000010	→	.80 .01 .98	→	00000010
00000001	→	.60 .94 .01	→	00000001

Artificial Neural Networks: Summary

- Highly non-linear regression/classification
- Vector valued inputs and outputs
- Potentially millions of parameters to estimate
- Actively used to model distributed computation in the brain
- Hidden layers learn intermediate representations
- Stochastic gradient descent, local minima problems
- Overfitting and how to deal with it.

Other Activation Functions

- Problem with sigmoid: saturation




Figure borrowed from *Pattern Recognition and Machine Learning*, Bishop

Other Activation Functions

- Activation function ReLU (rectified linear unit)





Figure from *Deep learning*, by Goodfellow, Bengio, Courville.

Other Activation Functions

- Activation function ReLU (rectified linear unit)

Gradient 0

The Rectified Linear Activation Function



Gradient 1

Other Activation Functions

- Generalizations of ReLU $\text{gReLU}(z) = \max\{z, 0\} + \alpha \min\{z, 0\}$
 - Leaky-ReLU(z) = $\max\{z, 0\} + 0.01 \min\{z, 0\}$
 - Parametric-ReLU(z): α learnable

