


10701: Introduction to Machine Learning

Ensemble methods: Boosting from Weak Learners

Eric Xing




Lecture 12, October 14, 2020

Reading: see class homepage



Weak Learners: Fighting the bias-variance tradeoff

- **Simple (a.k.a. weak) learners** e.g., naïve Bayes, logistic regression, decision stumps (or shallow decision trees)



Are good 😊 - Low variance, don't usually overfit

Are bad ☹️ - High bias, can't solve hard learning problems

- **Can we make weak learners always good???**
 - **No!!!** But often yes...



Why boost weak learners?

Goal: Automatically categorize type of call requested
(Collect, Calling card, Person-to-person, etc.)



- yes I'd like to place a collect call long distance please (Collect)
- operator I need to make a call but I need to bill it to my office (ThirdNumber)
- yes I'd like to place a call on my master card please (CallingCard)

- ❑ Easy to find “rules of thumb” that are “often” correct.
E.g. If ‘card’ occurs in utterance, then predict ‘calling card’
- ❑ Hard to find single highly accurate prediction rule.



Voting (Ensemble Methods)

- Instead of learning a single (weak) classifier, learn **many weak classifiers** that are **good at different parts of the input space**
- **Output class:** (Weighted) vote of each classifier
 - Classifiers that are most “sure” will vote with more conviction
 - Classifiers will be most “sure” about a particular part of the space
 - On average, do better than single classifier!



$$H: X \rightarrow Y (-1,1)$$

$$H(X) = h_1(X) + h_2(X)$$

$$H(X) = \text{sign}\left(\sum_t \alpha_t h_t(X)\right)$$

\downarrow
weights



Voting (Ensemble Methods)

- ❑ Instead of learning a single (weak) classifier, learn **many weak classifiers** that are **good at different parts of the input space**
- ❑ **Output class:** (Weighted) vote of each classifier
 - ❑ Classifiers that are most “sure” will vote with more conviction
 - ❑ Classifiers will be most “sure” about a particular part of the space
 - ❑ On average, do better than single classifier!
- ❑ **But how do you ???**
 - ❑ force classifiers h_t to learn about different parts of the input space?
 - ❑ weigh the votes of different classifiers? α_t



Bagging

- ❑ Recall decision trees (lecture 3)
 - ❑ Pros: interpretable, can handle discrete and continuous features, robust to outliers, **low bias**, etc.
 - ❑ Cons: **high variance**
- ❑ Trees are perfect candidates for ensembles
 - ❑ Consider averaging many (nearly) unbiased tree estimators
 - ❑ **Bias** remains similar, but **variance** is reduced
- ❑ This is called **bagging** (bootstrap aggregating) (Breiman, 1996)
 - ❑ Train many trees on bootstrapped data, then take average

$$f(x) = \frac{1}{B} \sum_{b=1}^B f_b(x)$$

- ❑ Bootstrap: statistical term for “roll n-face dice n times”



Random Forest

- Reduce correlation between trees, by introducing randomness
 1. For $b = 1, \dots, B$,
 1. Draw a bootstrap dataset Z^*
 2. Learn a tree $f_b(\cdot)$ on Z^* , in particular **select m features randomly** out of p features as candidates before splitting
 2. Output:
 - Regression: $f(x) = \frac{1}{B} \sum_{b=1}^B f_b(x)$
 - Classification: majority vote
- Typically take $m \leq \sqrt{p}$



Boosting [Schapire'89]

- ❑ **Idea:** given a weak learner, run it multiple times on (reweighted) training data, then let learned classifiers vote
- ❑ On each iteration t :
 - ❑ weight each training example by how incorrectly it was classified
 - ❑ Learn a weak hypothesis – h_t
 - ❑ A strength for this hypothesis – α_t
- ❑ Final classifier:
$$H(X) = \text{sign}(\sum \alpha_t h_t(X))$$
- ❑ **Practically useful, and theoretically interesting**
- ❑ Important issues:
 - ❑ what is the criterion that we are optimizing? (measure of loss)
 - ❑ we would like to estimate each new component classifier in the same manner (modularity)



AdaBoost

- At the k th iteration we find (any) classifier $h(\mathbf{x}; \theta_k^*)$ for which the weighted classification error:

$$\varepsilon_k = \sum_{i=1}^n W_i^{k-1} I(y_i \neq h(\mathbf{x}_i; \theta_k^*)) / \sum_{i=1}^n W_i^{k-1}$$

is better than chance.

- This is meant to be "easy" --- weak classifier
- Determine how many "votes" to assign to the new component classifier:



$$\alpha_k = 0.5 \log((1 - \varepsilon_k) / \varepsilon_k)$$

- stronger classifier gets more votes
- Update the weights on the training examples:



$$W_i^k = W_i^{k-1} \exp\{-y_i \alpha_k h(\mathbf{x}_i; \theta_k)\}$$





Boosting Example (Decision Stumps)



h_1
 $\epsilon_1=0.30$
 $\alpha_1=0.42$




$\epsilon_2=0.21$
 $\alpha_2=0.65$
 h_2



$\epsilon_3=0.14$
 $\alpha_3=0.92$



Boosting Example (Decision Stumps)



- What is the criterion that we are optimizing? (measure of loss)



Measurement of error

- Loss function:

$$\lambda(y, h(\mathbf{x})) \quad (\text{e.g. } I(y \neq h(\mathbf{x})))$$

- Generalization error:

$$L(h) = E[\lambda(y, h(\mathbf{x}))]$$

- Objective: find h with minimum *generalization* error
- Main boosting idea: minimize the *empirical* error:

$$\hat{L}(h) = \frac{1}{N} \sum_{i=1}^N \lambda(y_i, h(\mathbf{x}_i))$$



Exponential Loss

- Empirical loss:

$$\hat{L}(h) = \frac{1}{N} \sum_{i=1}^N \lambda(y_i, \hat{h}_m(\mathbf{x}_i))$$

- Another possible measure of empirical loss is

$$\hat{L}(h) = \sum_{i=1}^n \exp\{-y_i \hat{h}_m(\mathbf{x}_i)\}$$



Exponential Loss

- One possible measure of empirical loss is

$$\hat{L}(h) = \sum_{i=1}^n \exp\{-y_i \hat{h}_m(\mathbf{x}_i)\}$$

Recall that:

$$\hat{h}_m(\mathbf{x}) = \alpha_1 h(\mathbf{x}; \theta_1) + \dots + \alpha_m h(\mathbf{x}; \theta_m)$$

$$= \sum_{i=1}^n \exp\{-y_i \hat{h}_{m-1}(\mathbf{x}_i) - y_i a_m h(\mathbf{x}_i; \theta_m)\}$$

$$= \sum_{i=1}^n \exp\{-y_i \hat{h}_{m-1}(\mathbf{x}_i)\} \exp\{-y_i a_m h(\mathbf{x}_i; \theta_m)\}$$

$$= \sum_{i=1}^n W_i^{m-1} \exp\{-y_i a_m h(\mathbf{x}_i; \theta_m)\}$$

$$W_i^{m-1} = \exp\{-y_i \hat{h}_{m-1}(\mathbf{x}_i)\}$$

- The combined classifier based on $m - 1$ iterations defines a weighted loss criterion for **the next simple classifier to add**
- each training sample is weighted by its "classifiability" (or difficulty) seen by the classifier we have built so far



Linearization of loss function

- We can simplify a bit the estimation criterion for the new component classifiers (assuming α is small)

$$\exp\{-y_i a_m h(\mathbf{x}_i; \theta_m)\} \approx 1 - y_i a_m h(\mathbf{x}_i; \theta_m)$$

- Now our empirical loss criterion reduces to

$$\sum_{i=1}^n \exp\{-y_i \hat{h}_m(\mathbf{x}_i)\}$$

$$\approx \sum_{i=1}^n W_i^{m-1} (1 - y_i a_m h(\mathbf{x}_i; \theta_m))$$

$$W_i^{m-1} = \exp\{-y_i \hat{h}_{m-1}(\mathbf{x}_i)\}$$

$$= \sum_{i=1}^n W_i^{m-1} - a_m \sum_{i=1}^n W_i^{m-1} y_i h(\mathbf{x}_i; \theta_m)$$

- We could choose a new component classifier to optimize this weighted agreement



A possible algorithm

- At stage m we find θ^* that maximize (or at least give a sufficiently high) weighted agreement:

$$\sum_{i=1}^n W_i^{m-1} y_i h(\mathbf{x}_i; \theta_m^*)$$

- each sample is weighted by its "difficulty" under the previously combined $m - 1$ classifiers,
- more "difficult" samples received heavier attention as they dominates the total loss
- Then we go back and find the "votes" α_m^* associated with the new classifier by minimizing the **original** weighted (exponential) loss

$$\hat{L}(h) = \sum_{i=1}^n W_i^{m-1} \exp\{-y_i a_m h(\mathbf{x}_i; \theta_m)\}$$

\Rightarrow

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$



The AdaBoost algorithm

$$W_i^{m-1} = \exp\{-y_i \hat{h}_{m-1}(\mathbf{x}_i)\}$$

- At the k th iteration we find (any) classifier $h(\mathbf{x}; \theta_k^*)$ for which the weighted classification error:

$$\varepsilon_k = \sum_{i=1}^n W_i^{k-1} I(y_i \neq h(\mathbf{x}_i; \theta_k^*)) / \sum_{i=1}^n W_i^{k-1}$$

is better than change.

- This is meant to be "easy" --- weak classifier
- Determine how many "votes" to assign to the new component classifier:

$$\alpha_k = 0.5 \log((1 - \varepsilon_k) / \varepsilon_k)$$

- stronger classifier gets more votes
- Update the weights on the training examples:

$$W_i^k = W_i^{k-1} \exp\{-y_i \alpha_k h(\mathbf{x}_i; \theta_k)\}$$



The AdaBoost algorithm cont'd

- The final classifier after m boosting iterations is given by the sign of

$$\hat{h}(\mathbf{x}) = \frac{\alpha_1 h(\mathbf{x}; \theta_1) + \dots + \alpha_m h(\mathbf{x}; \theta_m)}{\alpha_1 + \dots + \alpha_m}$$

- the votes here are normalized for convenience



Boosting

- We have basically derived a Boosting algorithm that sequentially adds **new component classifiers**, each trained on reweighted training examples
 - each component classifier is presented with a slightly different problem
- AdaBoost preliminaries:
 - we work with *normalized weights* W_i on the training examples, initially uniform ($W_i = 1/n$)
 - the weight reflect the "*degree of difficulty*" of each datum on the latest classifier



AdaBoost: summary

- Input:
 - N examples $\mathcal{S}_N = \{(x_1, y_1), \dots, (x_N, y_N)\}$
 - a weak base learner $h = h(x, \theta)$
- Initialize: equal example weights $w_i = 1/N$ for all $i = 1..N$
- Iterate for $t = 1..T$:
 1. train base learner according to weighted example set (w_t, x) and obtain hypothesis $h_t = h(x, \theta_t)$
 2. compute hypothesis error ε_t
 3. compute hypothesis weight α_t
 4. update example weights for next iteration w_{t+1}
- Output: final hypothesis as a linear combination of h_t




Base Learners

- ❑ Weak learners used in practice:
 - ❑ Decision stumps (axis parallel splits)
 - ❑ Decision trees (e.g. C4.5 by Quinlan 1996)
 - ❑ Multi-layer neural networks
 - ❑ Radial basis function networks
- ❑ Can base learners operate on weighted examples?
 - ❑ In many cases they can be modified to accept weights along with the examples
 - ❑ In general, we can sample the examples (with replacement) according to the distribution defined by the weights



Boosting results – Digit recognition

[Schapire, 1989]




- ❑ Boosting often, **but not always**
 - ❑ Robust to overfitting
 - ❑ Test set error decreases even after training error is zero



Generalization Error Bounds

[Freund & Schapire'95]

$$error_{true}(H) \leq error_{train}(H) + \tilde{O} \left(\sqrt{\frac{Td}{m}} \right)$$

	bias	variance	
 tradeoff	large	small	T small
	small	large	T large

- ❑ T – number of boosting rounds
- ❑ d – VC dimension of weak learner, measures complexity of classifier
- ❑ m – number of training examples



Generalization Error Bounds

[Freund & Schapire'95]

$$error_{true}(H) \leq error_{train}(H) + \tilde{O} \left(\sqrt{\frac{Td}{m}} \right) \quad \text{With high probability}$$

Boosting can overfit if T is large


Boosting often, **Contradicts experimental results**

- ▣ Robust to overfitting
- ▣ Test set error decreases even after training error is zero

Need better analysis tools – margin based bounds



Experiments



The Boosting Approach to Machine Learning, by Robert E. Schapire



Training Margins

- When a vote is taken, the **more predictors agreeing**, the **more confident** you are in your prediction.

- Margin for example:

$$\text{margin}_h(\mathbf{x}_i, y_i) = y_i \left[\frac{\alpha_1 h(\mathbf{x}_i; \theta_1) + \dots + \alpha_m h(\mathbf{x}_i; \theta_m)}{\alpha_1 + \dots + \alpha_m} \right]$$

The margin lies in $[-1, 1]$ and is negative for all misclassified examples.

- Successive boosting iterations improve the majority vote or margin for the training examples



A Margin Bound

- For any γ , the generalization error is less than:

$$\Pr(\text{margin}_h(\mathbf{x}, y) \leq \gamma) + O\left(\sqrt{\frac{d}{m\gamma^2}}\right)$$

Robert E. Schapire, Yoav Freund, Peter Bartlett and Wee Sun Lee.
Boosting the margin: A new explanation for the effectiveness of voting
methods. *The Annals of Statistics*, 26(5):1651-1686, 1998.

- It does not depend on $T!!!$



Summary

- Boosting takes a weak learner and converts it to a strong one
- Works by asymptotically minimizing the empirical error
- Effectively maximizes the margin of the combined hypothesis





Supplementary

Some additional points for fun



Boosting and Logistic Regression

- Logistic regression assumes:

$$P(Y = 1|X) = \frac{1}{1 + \exp(f(x))} \quad f(x) = w_0 + \sum_j w_j x_j$$

- And tries to maximize data likelihood:

$$P(\mathcal{D}|f) \stackrel{\text{iid}}{=} \prod_{i=1}^m \frac{1}{1 + \exp(-y_i f(x_i))}$$

- Equivalent to minimizing log loss

$$-\log P(\mathcal{D}|f) = \sum_{i=1}^m \ln(1 + \exp(-y_i f(x_i)))$$



Boosting and Logistic Regression

Logistic regression equivalent to minimizing log loss

$$\sum_{i=1}^m \ln(1 + \exp(-y_i f(x_i)))$$

$$f(x) = w_0 + \sum_j w_j x_j$$

Boosting minimizes similar loss function!!

$$\frac{1}{m} \sum_{i=1}^m \exp(-y_i f(x_i)) = \prod_t Z_t$$

$$f(x) = \sum_t \alpha_t h_t(x)$$

Weighted average of weak learners



**Both smooth approximations
of 0/1 loss!**



Boosting and Logistic Regression

Boosting:

- Minimize exp loss

$$\sum_{i=1}^m \ln(1 + \exp(-y_i f(x_i)))$$

- Define

$$f(x) = \sum_j w_j x_j$$

where $h_t(x)$ defined dynamically
to fit data
(not a linear classifier)

- Weights α_t learned per iteration t
incrementally

Logistic regression:

- Minimize log loss

$$\sum_{i=1}^m \exp(-y_i f(x_i))$$

- Define

$$f(x) = \sum_t \alpha_t h_t(x)$$

where x_j predefined features
(linear classifier)

- Jointly optimize over all weights
 $W_0, W_1, W_2 \dots$



Hard & Soft Decision

- Weighted average of weak learners

$$f(x) = \sum_t \alpha_t h_t(x)$$

- Hard Decision/Predicted label:

$$H(x) = \text{sign}(f(x))$$


- Soft Decision:
 - (based on analogy with logistic regression)

$$P(Y = 1|X) = \frac{1}{1 + \exp(f(x))}$$



Effect of Outliers

- Good 😊 : Can identify outliers since focuses on examples that are hard to categorize
- Bad 😞 : Too many outliers can degrade classification performance dramatically increase time to convergence



Gradient Boosting

- Goal: Find nonlinear predictor $\hat{h}(x) \in \mathcal{H}$ such that

$$\hat{h} = \arg \min_{h \in \mathcal{H}} \mathcal{L}(h(X), Y)$$

- Gradient boosting generalizes Adaboost (exponential loss) to any smooth loss functions $\mathcal{L}(\cdot, \cdot)$

Square loss (regression)

$$\mathcal{L}(h(X), Y) = \sum_{i=1}^n (h(\mathbf{x}_i) - y_i)^2$$

Logistic loss (classification)

$$\mathcal{L}(h(X), Y) = \sum_{i=1}^n \ln(1 + e^{-h(\mathbf{x}_i)y_i})$$

Margin loss (ranking)

$$\mathcal{L}(h(X), Y) = \sum_{(i,i'): y_{(i,i')}=1} \max(0, 1 - (h(\mathbf{x}_i) - h(\mathbf{x}_{i'})))^2$$

(prefer item i over j)

Others...



Gradient Boosting Decision Tree

- Let's use decision tree to approximate g_{k-1}
- A J-leaf node decision tree can be viewed as a partition of the input space

$$q : \mathbb{R}^d \rightarrow \{1, 2, \dots, J\}$$

- and a prediction value (weight) associated with each partition

$$w \in \mathbb{R}^J$$

- Will learn q (tree structure) first, then w

