

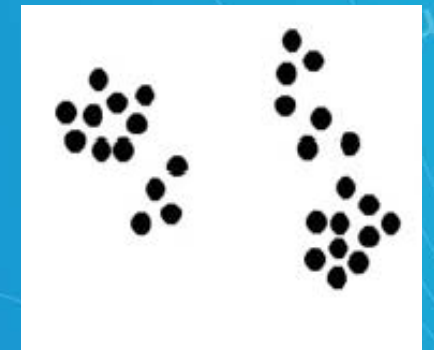
# 10701: Introduction to Machine Learning

## Mixture models and the Expectation Maximization algorithm

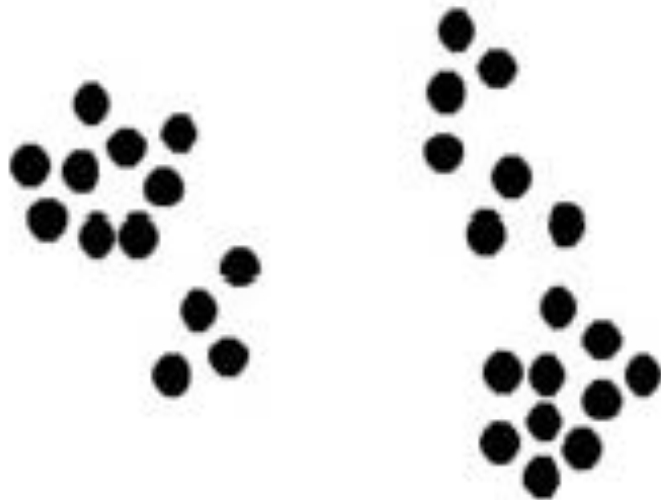
Eric Xing

Lecture 14, October 21, 2020

**Reading: Chap. 9, 13, C.B book**



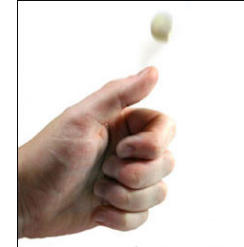
# Clustering and partially observable probabilistic models



# Recap: Indicator Variables and Discrete Distributions

- Bernoulli distribution:  $\text{Ber}(p)$

$$P(x) = \begin{cases} 1-p & \text{for } x=0 \\ p & \text{for } x=1 \end{cases} \Rightarrow P(x) = p^x (1-p)^{1-x}$$



- Multinomial distribution:  $\text{Mult}(1, \theta)$

- Multinomial (indicator) variable:

$$X = \begin{bmatrix} X_1 \\ X_2 \\ X_3 \\ X_4 \\ X_5 \\ X_6 \end{bmatrix}, \quad \text{where} \quad \begin{aligned} X_j & \in [0,1], \quad \text{and} \quad \sum_{j \in \{1, \dots, 6\}} X_j = 1 \\ X_j & = 1 \text{ w.p. } \theta_j, \quad \sum_{j \in \{1, \dots, 6\}} \theta_j = 1. \end{aligned}$$



$$\begin{aligned} P(x_i) &= P(\{x_{n,k} = 1, \text{ where } k \text{ index the die - side of the } n\text{th roll}\}) \\ &= \theta_k = \theta_1^{x_{n,1}} \times \theta_2^{x_{n,2}} \times \dots \times \theta_K^{x_{n,K}} = \prod_{k=1}^K \theta_k^{x_{n,k}} \end{aligned}$$

- It can be shown that:

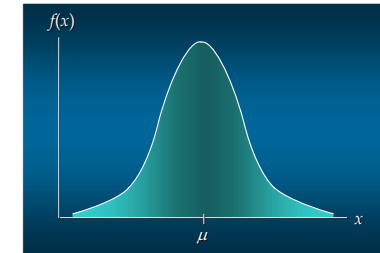
$$\hat{\theta}_{k,MLE} = \frac{n_k}{N} \quad \text{or} \quad \hat{\theta}_{k,MLE} = \frac{1}{N} \sum_n x_{n,k}$$



# Recap: Continuous Variables and Gaussian Distributions

## □ Normal (Gaussian) Probability Density Function

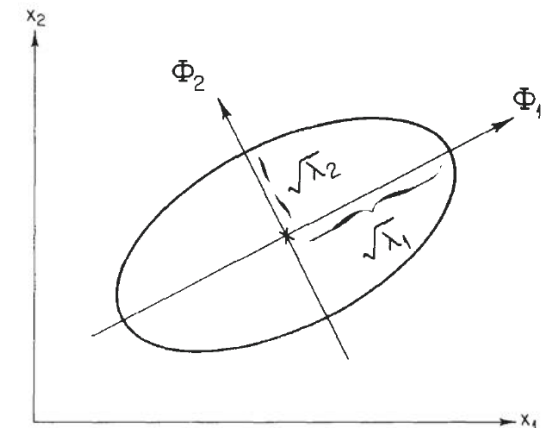
$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$



- The distribution is symmetric, and is often illustrated as a bell-shaped curve.
- Two parameters,  $\mu$  (mean) and  $\sigma$  (standard deviation), determine the location and shape of the distribution.
- The highest point on the normal curve is at the mean, which is also the median and mode.
- The mean can be any numerical value: negative, zero, or positive.

## □ Multivariate Gaussian

$$p(X; \bar{\mu}, \Sigma) = \frac{1}{(\sqrt{2\pi})^{n/2} |\Sigma|^{1/2}} \exp\left\{-\frac{1}{2}(X - \bar{\mu})^T \Sigma^{-1}(X - \bar{\mu})\right\}$$



# Recap: MLE for a multivariate-Gaussian

- It can be shown that the MLE for  $\mu$  and  $\Sigma$  is

$$\mu_{MLE} = \frac{1}{N} \sum_n (x_n)$$

$$\Sigma_{MLE} = \frac{1}{N} \sum_n (x_n - \mu_{ML})(x_n - \mu_{ML})^T = \frac{1}{N} S$$

where the scatter matrix is

$$S = \sum_n (x_n - \mu_{ML})(x_n - \mu_{ML})^T = \left( \sum_n x_n x_n^T \right) - N \mu_{ML} \mu_{ML}^T$$

$$x_n = \begin{pmatrix} x_{n,1} \\ x_{n,2} \\ \vdots \\ x_{n,K} \end{pmatrix}$$
$$X = \begin{pmatrix} ---x_1^T--- \\ ---x_2^T--- \\ \vdots \\ ---x_N^T--- \end{pmatrix}$$

- The sufficient statistics are  $\sum_n x_n$  and  $\sum_n x_n x_n^T$ .
- Note that  $X^T X = \sum_n x_n x_n^T$  may not be full rank (eg. if  $N < D$ ), in which case  $\Sigma_{ML}$  is not invertible



# Recap: Conditional Gaussian

- The data:

$$\{(x_1, y_1), (x_2, y_2), (x_3, y_3), \dots, (x_N, y_N)\}$$

- Both nodes are observed:

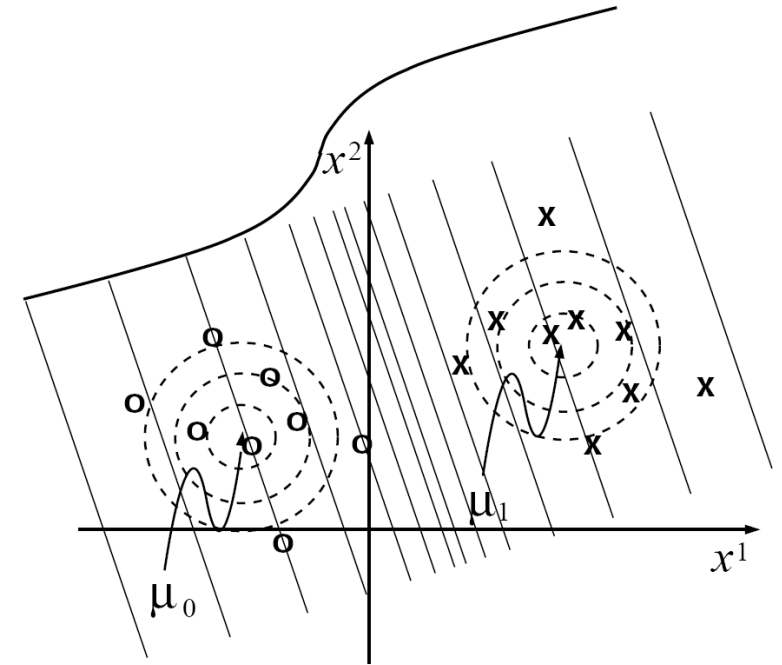
- $\mathbf{y}$  is a class indicator vector

$$p(y_n) = \text{multi}(y_n : \pi) = \prod_k \pi_k^{y_{n,k}}$$

- $\mathbf{x}$  is a conditional Gaussian variable with a class-specific mean

$$p(x_n | y_{n,k} = 1, \mu, \sigma) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp\left\{-\frac{1}{2\sigma^2} (x_n - \mu_k)^2\right\}$$

$$p(x | y, \mu, \sigma) = \prod_n \left( \prod_k N(x_n : \mu_k, \sigma)^{y_{n,k}} \right)$$



# Recap: MLE of conditional Gaussian

- Data log-likelihood

$$\ell(\boldsymbol{\theta}; D) = \log \prod_n p(x_n, y_n) = \log \prod_n p(y_n | \pi) p(x_n | y_n, \mu, \sigma)$$

- MLE

$$\hat{\pi}_{k,MLE} = \arg \max_{\pi} \ell(\boldsymbol{\theta}; D), \quad \hat{\pi}_{k,MLE} = \frac{\sum_n y_{n,k}}{N} = n_k / N$$

the fraction of  
samples of class  $m$

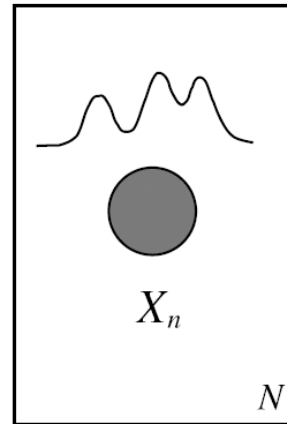
$$\hat{\mu}_{k,MLE} = \arg \max \ell(\boldsymbol{\theta}; D), \quad \hat{\mu}_{k,MLE} = \frac{\sum_n y_{n,k} x_n}{\sum_n y_{n,k}} = \frac{\sum_n y_{n,k} x_n}{n_k}$$

the average of  
samples of class  $m$

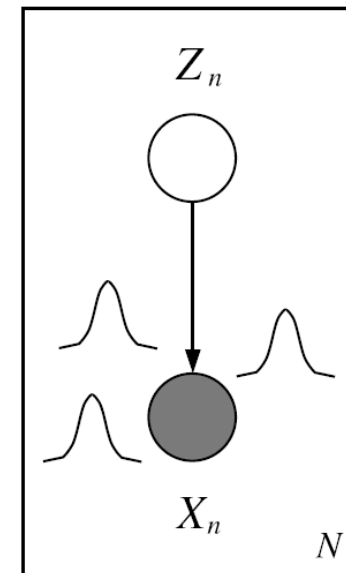


# Mixture Models

- A density model  $p(x)$  may be multi-modal.
- We may be able to model it as a mixture of uni-modal distributions (e.g., Gaussians).
- Each mode may correspond to a different sub-population (e.g., male and female).
- **Indicator variable  $Z$  is NOT observed!**



(a)



(b)





# Unobserved Variables

- A variable can be unobserved (latent) because:
  - it is an imaginary quantity meant to provide some simplified and abstractive view of the data generation process
    - e.g., speech recognition models, mixture models ...
  - it is a real-world object and/or phenomena, but difficult or impossible to measure
    - e.g., the temperature of a star, causes of a disease, evolutionary ancestors ...
  - it is a real-world object and/or phenomena, but sometimes wasn't measured, because of faulty sensors; or was measure with a noisy channel, etc.
    - e.g., traffic radio, aircraft signal on a radar screen,
- Discrete latent variables can be used to partition/cluster data into sub-groups (mixture models, forthcoming).
- Continuous latent variables (factors) can be used for dimensionality reduction (factor analysis, etc., later lectures).

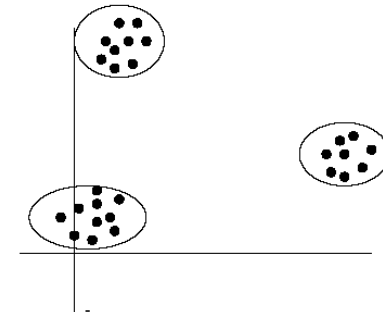
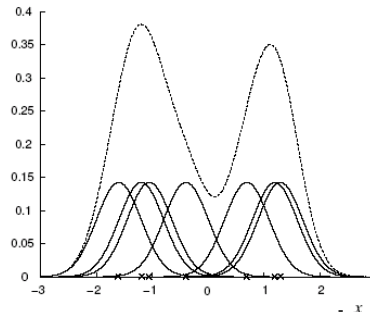


# Gaussian Mixture Models (GMMs)

- Consider a mixture of  $K$  Gaussian components:

$$p(x_n | \mu, \Sigma) = \sum_k \pi_k N(x, | \mu_k, \Sigma_k)$$

mixture proportion    mixture component



- This model can be used for unsupervised clustering.
  - This model (fit by AutoClass) has been used to discover new kinds of stars in astronomical data, etc.

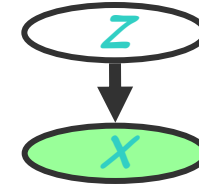


# GGM Likelihood

- Consider a mixture of  $K$  Gaussian components:

- $Z$  is a latent class indicator vector:

$$p(\mathbf{z}_n) = \text{multi}(\mathbf{z}_n : \pi) = \prod_k (\pi_k)^{z_n^k}$$



- $X$  is a conditional Gaussian variable with a class-specific mean/covariance

$$p(\mathbf{x}_n | \mathbf{z}_n^k = 1, \mu, \Sigma) = \frac{1}{(2\pi)^{m/2} |\Sigma_k|^{1/2}} \exp\left\{-\frac{1}{2} (\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k)\right\}$$

- The likelihood of a sample:

$$\begin{aligned} p(\mathbf{x}_n | \mu, \Sigma) &= \sum_k p(\mathbf{z}^k = 1 | \pi) p(\mathbf{x}_n | \mathbf{z}^k = 1, \mu, \Sigma) \\ &= \sum_{\mathbf{z}_n} \prod_k \left( (\pi_k)^{z_n^k} \mathcal{N}(\mathbf{x}_n : \mu_k, \Sigma_k)^{z_n^k} \right) = \sum_k \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k) \end{aligned}$$

mixture component

mixture proportion



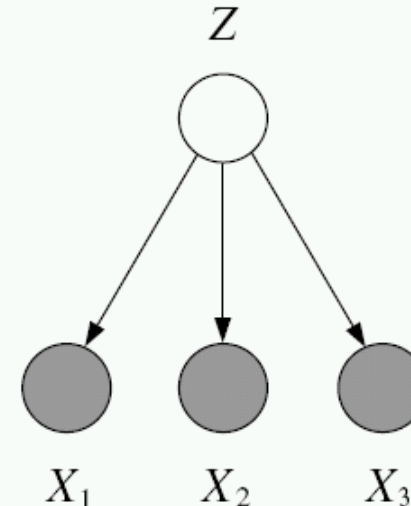
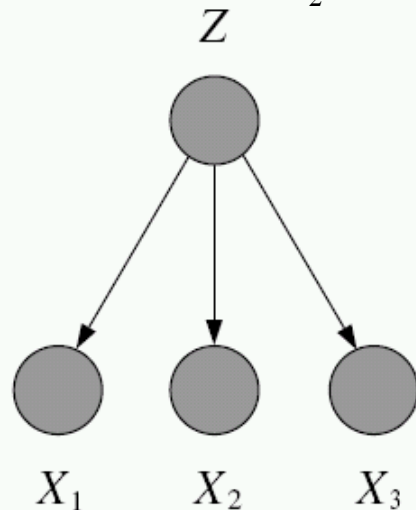
# Why is Learning Harder?

- In fully observed iid settings, the log likelihood decomposes into a sum of local terms.

$$\ell_c(\theta; D) = \log p(x, z | \theta) = \log p(z | \theta_z) + \log p(x | z, \theta_x)$$

- With latent variables, all the parameters become coupled together via *marginalization*

$$\ell_c(\theta; D) = \log \sum_z p(x, z | \theta) = \log \sum_z p(z | \theta_z) p(x | z, \theta_x)$$



# Gradient Learning for mixture models

- We can learn mixture densities using gradient descent on the log likelihood. The gradients are quite interesting:

$$\ell(\theta) = \log p(\mathbf{x} | \theta) = \log \sum_k \pi_k p_k(\mathbf{x} | \theta_k)$$

$$\begin{aligned} \frac{\partial \ell}{\partial \theta} &= \frac{1}{p(\mathbf{x} | \theta)} \sum_k \pi_k \frac{\partial p_k(\mathbf{x} | \theta_k)}{\partial \theta} \\ &= \sum_k \frac{\pi_k}{p(\mathbf{x} | \theta)} p_k(\mathbf{x} | \theta_k) \frac{\partial \log p_k(\mathbf{x} | \theta_k)}{\partial \theta} \\ &= \sum_k \pi_k \frac{p_k(\mathbf{x} | \theta_k)}{p(\mathbf{x} | \theta)} \frac{\partial \log p_k(\mathbf{x} | \theta_k)}{\partial \theta_k} = \sum_k r_k \frac{\partial \ell_k}{\partial \theta_k} \end{aligned}$$

- In other words, the gradient is the responsibility weighted sum of the individual log likelihood gradients.
- Can pass this to a conjugate gradient routine.



# Parameter Constraints

- Often we have constraints on the parameters, e.g.  $\sum_k \pi_k = 1$ ,  $\Sigma$  being symmetric positive definite (hence  $\Sigma_{ii} > 0$ ).
- We can use constrained optimization, or we can reparameterize in terms of unconstrained values.
  - For normalized weights, use the softmax transform:
  - For covariance matrices, use the Cholesky decomposition:

$$\Sigma^{-1} = \mathbf{A}^T \mathbf{A}$$

where  $\mathbf{A}$  is upper triangular with positive diagonal:

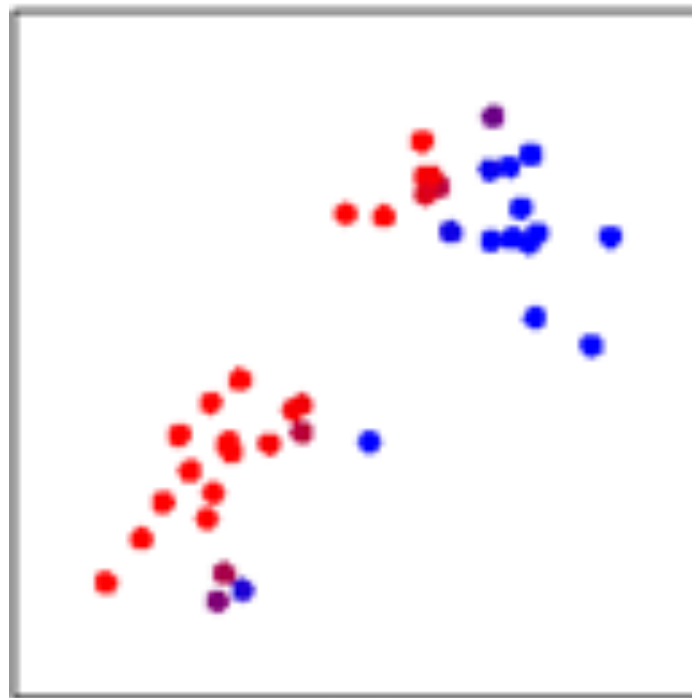
$$\mathbf{A}_{ii} = \exp(\lambda_i) > 0 \quad \mathbf{A}_{ij} = \eta_{ij} \quad (j > i) \quad \mathbf{A}_{ij} = 0 \quad (j < i)$$

the parameters  $\gamma_i, \lambda_i, \eta_{ij} \in \mathbb{R}$  are unconstrained.

- Use chain rule to compute  $\frac{\partial \ell}{\partial \pi}, \frac{\partial \ell}{\partial \mathbf{A}}$ .



# The Expectation-Maximization (EM) Algorithm



# GGM Likelihood

## -- a close look

□ E.g., mixture of K Gaussians:

□  $Z$  is a latent class indicator vector

$$p(z_n) = \text{multi}(z_n : \pi) = \prod_k (\pi_k)^{z_n^k}$$

□  $X$  is a conditional Gaussian variable with a class-specific mean/covariance

$$p(x_n | z_n^k = \mathbf{1}, \mu, \Sigma) = \frac{1}{(2\pi)^{m/2} |\Sigma_k|^{1/2}} \exp\left\{-\frac{1}{2}(x_n - \mu_k)^T \Sigma_k^{-1} (x_n - \mu_k)\right\}$$

□ The likelihood of a sample:

$$\begin{aligned} p(x_n | \mu, \Sigma) &= \sum_k p(z_n^k = \mathbf{1} | \pi) p(x_n | z_n^k = \mathbf{1}, \mu, \Sigma) \\ &= \sum_{z_n} \prod_k \left( (\pi_k)^{z_n^k} N(x_n : \mu_k, \Sigma_k)^{z_n^k} \right) = \sum_k \pi_k N(x_n | \mu_k, \Sigma_k) \end{aligned}$$

## Recap: MLE of conditional Gaussian

□ Data log-likelihood

$$\ell(\theta; D) = \log \prod_n p(x_n, y_n) = \log \prod_n p(y_n | \pi) p(x_n | y_n, \mu, \sigma)$$

□ MLE

$$\hat{\pi}_{k,MLE} = \arg \max_{\pi} \ell(\theta; D), \quad \hat{\pi}_{k,MLE} = \frac{\sum_n y_{n,k}}{N} = \frac{n_k}{N}$$

the fraction of samples of class  $m$

$$\hat{\mu}_{k,MLE} = \arg \max_{\mu} \ell(\theta; D), \quad \hat{\mu}_{k,MLE} = \frac{\sum_n y_{n,k} x_n}{\sum_n y_{n,k}} = \frac{\sum y_{n,k} x_n}{n_k}$$

the average of samples of class  $m$



❖ **If we do not know  $z_n$**

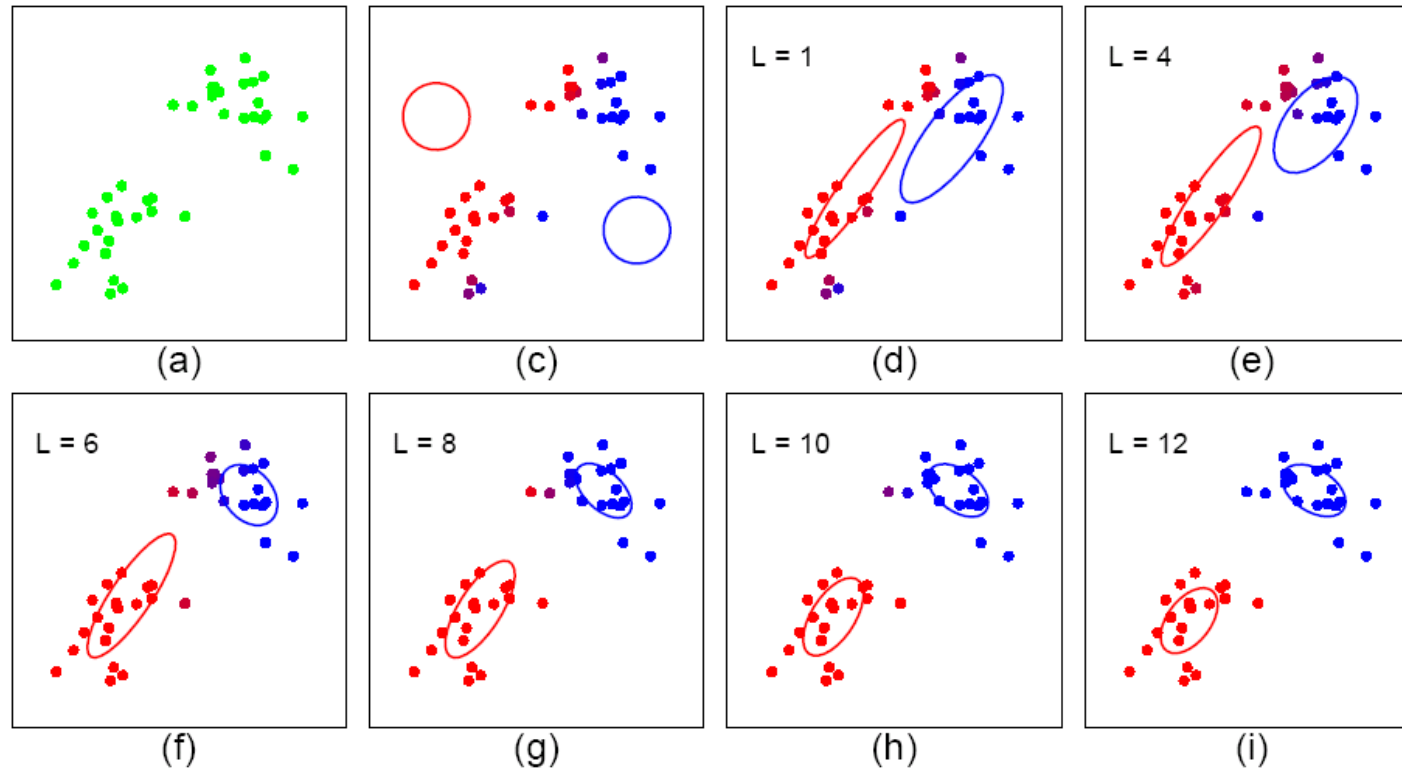
$$z_n \rightarrow p(z_n^k = \mathbf{1} | x, \mu^{(t)}, \Sigma^{(t)})$$





# An EM heuristic for GMM

- Start:
  - "Guess" the centroid  $\mu_k$  and covariance  $\Sigma_k$  of each of the K clusters
- Loop



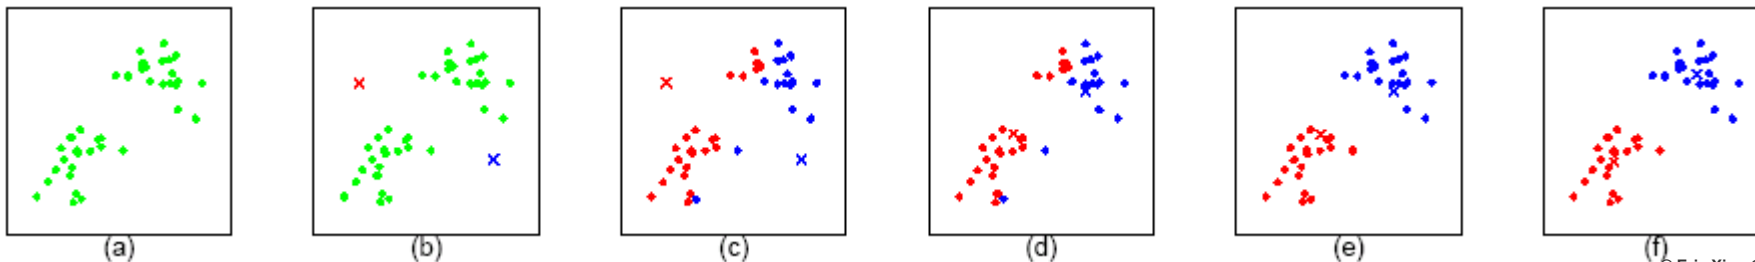
# Recall K-means

- Start:
  - "Guess" the centroid  $\mu_k$  and covariance  $\Sigma_k$  of each of the K clusters
- Loop
  - For each point  $n=1$  to  $N$ , compute its cluster label:

$$z_n^{(t)} = \arg \max_k (x_n - \mu_k^{(t)})^T \Sigma_k^{-1(t)} (x_n - \mu_k^{(t)})$$

- For each cluster  $k=1:K$

$$\mu_k^{(t+1)} = \frac{\sum_n \delta(z_n^{(t)}, k) x_n}{\sum_n \delta(z_n^{(t)}, k)} \quad \Sigma_k^{(t+1)} = \dots$$



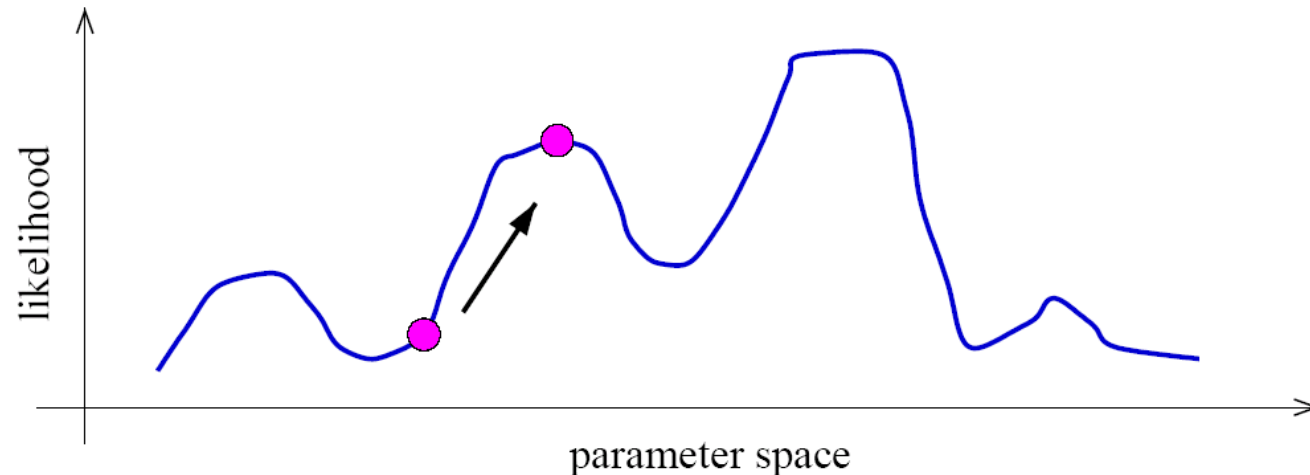
# Notes on EM

- EM is an optimization strategy for objective functions that can be interpreted as likelihoods in the presence of missing data.
- It is much simpler than gradient methods:
  - No need to choose step size.
  - Enforces constraints automatically.
  - Calls inference and fully observed learning as subroutines.
- EM is an Iterative algorithm with two linked steps:
  - E-step: fill-in hidden values using inference,  $p(\mathbf{z}|\mathbf{x}, \theta)$ .
  - M-step: update parameters  $t+1$  using standard MLE/MAP method applied to completed data
- We will prove that this procedure monotonically improves (or leaves it unchanged). Thus it always converges to a local optimum of the likelihood.



# Identifiability

- ❑ A mixture model induces a multi-modal likelihood.
- ❑ Hence gradient ascent can only find a local maximum.
- ❑ Mixture models are unidentifiable, since we can always switch the hidden labels without affecting the likelihood.
- ❑ Hence we should be careful in trying to interpret the “meaning” of latent variables.



# How is EM derived?

- A mixture of K Gaussians:

- $\mathbf{Z}$  is a latent class indicator vector

$$p(\mathbf{z}_n) = \text{multi}(\mathbf{z}_n : \boldsymbol{\pi}) = \prod_k (\pi_k)^{z_n^k}$$

- $\mathbf{X}$  is a conditional Gaussian variable with a class-specific mean/covariance

$$p(\mathbf{x}_n | \mathbf{z}_n^k = 1, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{m/2} |\boldsymbol{\Sigma}_k|^{1/2}} \exp\left\{-\frac{1}{2} (\mathbf{x}_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k)\right\}$$

- The likelihood of a sample:

$$\begin{aligned} p(x_n | \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= \sum_k p(z_n^k = 1 | \boldsymbol{\pi}) p(x_n | z_n^k = 1, \boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ &= \sum_{z_n} \prod_k \left( (\pi_k)^{z_n^k} N(x_n : \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)^{z_n^k} \right) = \sum_k \pi_k N(x_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \end{aligned}$$

- The “complete” likelihood

$$p(x_n, z_n^k = 1 | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = p(z_n^k = 1 | \boldsymbol{\pi}) p(x_n | z_n^k = 1, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \pi_k N(x_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

$$p(x_n, \mathbf{z}_n | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \prod_k [\pi_k N(x_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)]^{z_n^k}$$

**But this is itself a random variable! Not good as objective function**



# How is EM derived?

- The complete log likelihood:

$$\begin{aligned}\ell(\boldsymbol{\theta}; D) &= \log \prod_n p(z_n, x_n) = \log \prod_n p(z_n | \pi) p(x_n | z_n, \mu, \sigma) \\ &= \sum_n \log \prod_k \pi_k^{z_n^k} + \sum_n \log \prod_k N(x_n; \mu_k, \sigma)^{z_n^k} \\ &= \sum_n \sum_k z_n^k \log \pi_k - \sum_n \sum_k z_n^k \frac{1}{2\sigma^2} (x_n - \mu_k)^2 + C\end{aligned}$$

- The expected complete log likelihood

$$\begin{aligned}\langle \ell_c(\boldsymbol{\theta}; \boldsymbol{x}, \boldsymbol{z}) \rangle &= \sum_n \langle \log p(\boldsymbol{z}_n | \pi) \rangle_{p(\boldsymbol{z}|\boldsymbol{x})} + \sum_n \langle \log p(\boldsymbol{x}_n | \boldsymbol{z}_n, \mu, \Sigma) \rangle_{p(\boldsymbol{z}|\boldsymbol{x})} \\ &= \sum_n \sum_k \langle z_n^k \rangle \log \pi_k - \frac{1}{2} \sum_n \sum_k \langle z_n^k \rangle ((\boldsymbol{x}_n - \mu_k)^T \Sigma_k^{-1} (\boldsymbol{x}_n - \mu_k) + \log |\Sigma_k| + \mathcal{C})\end{aligned}$$



# E-step

- We maximize  $\langle l_c(\theta) \rangle$  iteratively using the following iterative procedure:

- **Expectation step**: computing the expected value of the sufficient statistics of the hidden variables (i.e.,  $\mathbf{z}$ ) given current est. of the parameters (i.e.,  $\pi$  and  $\mu$ ).

$$\tau_n^{k(t)} = \langle \mathbf{z}_n^k \rangle_{q^{(t)}} = p(\mathbf{z}_n^k = \mathbf{1} | \mathbf{x}, \mu^{(t)}, \Sigma^{(t)}) = \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n, | \mu_k^{(t)}, \Sigma_k^{(t)})}{\sum_i \pi_i^{(t)} \mathcal{N}(\mathbf{x}_n, | \mu_i^{(t)}, \Sigma_i^{(t)})}$$

- Here we are essentially doing **inference**



# M-step

- We maximize  $\langle l_c(\boldsymbol{\theta}) \rangle$  iteratively using the following iterative procedure:

- **Maximization step**: compute the parameters under current results of the expected value of the hidden variables

$$\pi_k^* = \arg \max \langle l_c(\boldsymbol{\theta}) \rangle, \quad \Rightarrow \frac{\partial}{\partial \pi_k} \langle l_c(\boldsymbol{\theta}) \rangle = 0, \forall k, \quad \text{s.t. } \sum_k \pi_k = 1$$

$$\Rightarrow \pi_k^* = \frac{\sum_n \langle z_n^k \rangle_{q^{(t)}}}{N} = \frac{\sum_n \tau_n^{k(t)}}{N} = \frac{\langle n_k \rangle}{N}$$

$$\mu_k^* = \arg \max \langle l(\boldsymbol{\theta}) \rangle, \quad \Rightarrow \mu_k^{(t+1)} = \frac{\sum_n \tau_n^{k(t)} \mathbf{x}_n}{\sum_n \tau_n^{k(t)}}$$

$$\Sigma_k^* = \arg \max \langle l(\boldsymbol{\theta}) \rangle, \quad \Rightarrow \Sigma_k^{(t+1)} = \frac{\sum_n \tau_n^{k(t)} (\mathbf{x}_n - \mu_k^{(t+1)})(\mathbf{x}_n - \mu_k^{(t+1)})^T}{\sum_n \tau_n^{k(t)}}$$

Fact :

$$\frac{\partial \log |\mathbf{A}^{-1}|}{\partial \mathbf{A}^{-1}} = \mathbf{A}^T$$

$$\frac{\partial \mathbf{x}^T \mathbf{A} \mathbf{x}}{\partial \mathbf{A}} = \mathbf{x} \mathbf{x}^T$$

- This is isomorphic to **MLE** except that the variables that are hidden are replaced by their expectations (in general they will be replaced by their corresponding "**sufficient statistics**")





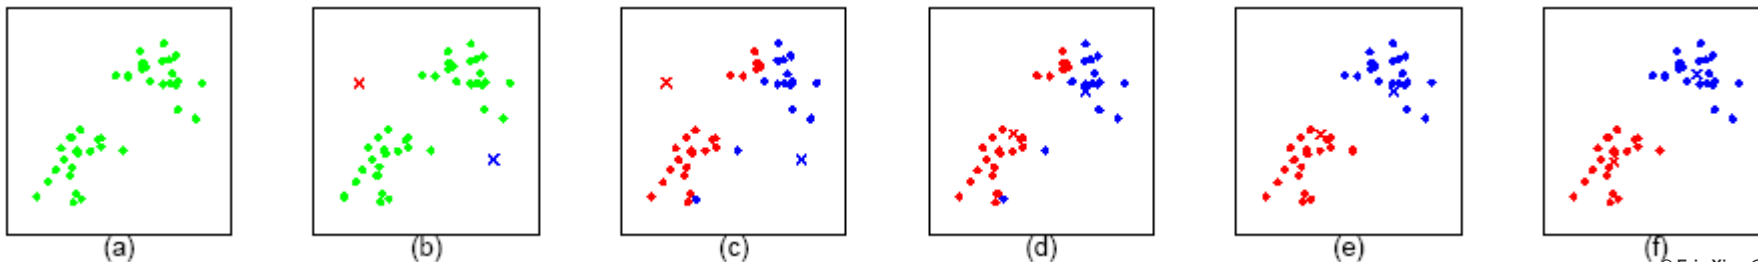
# Compare: K-means

- The EM algorithm for mixtures of Gaussians is like a "soft version" of the K-means algorithm.
- In the K-means "E-step" we do hard assignment:

$$\mathbf{z}_n^{(t)} = \arg \max_k (\mathbf{x}_n - \mu_k^{(t)})^T \Sigma_k^{-1(t)} (\mathbf{x}_n - \mu_k^{(t)})$$

- In the K-means "M-step" we update the means as the weighted sum of the data, but now the weights are 0 or 1:

$$\mu_k^{(t+1)} = \frac{\sum_n \delta(\mathbf{z}_n^{(t)}, k) \mathbf{x}_n}{\sum_n \delta(\mathbf{z}_n^{(t)}, k)}$$



# Theory underlying EM

- What are we doing?
- Recall that according to MLE, we intend to learn the model parameter that would have maximize the likelihood of the data.
- But we do not observe  $z$ , so computing

is difficult!  $\ell_c(\theta; D) = \log \sum_z p(x, z | \theta) = \log \sum_z p(z | \theta_z) p(x | z, \theta_x)$

- What shall we do?



# Complete & Incomplete Log Likelihoods

- Complete log likelihood

Let  $\mathbf{X}$  denote the observable variable(s), and  $\mathbf{Z}$  denote the latent variable(s).

If  $\mathbf{Z}$  could be observed, then

$$\ell_c(\theta; \mathbf{x}, \mathbf{z}) \stackrel{\text{def}}{=} \log p(\mathbf{x}, \mathbf{z} \mid \theta)$$

- Usually, optimizing  $\ell_c()$  given both  $\mathbf{z}$  and  $\mathbf{x}$  is straightforward (c.f. MLE for fully observed models).
- Recalled that in this case the objective for, e.g., MLE, decomposes into a sum of factors, the parameter for each factor can be estimated separately.
- But given that  $\mathbf{Z}$  is not observed,  $\ell_c()$  is a random quantity, cannot be maximized directly.

- Incomplete log likelihood

With  $\mathbf{z}$  unobserved, our objective becomes the log of a marginal probability:

$$\ell_c(\theta; \mathbf{x}) = \log p(\mathbf{x} \mid \theta) = \log \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z} \mid \theta)$$

- This objective won't decouple



# Expected Complete Log Likelihood

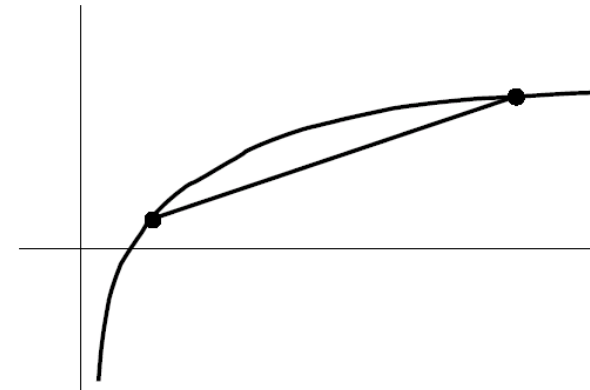
- For *any* distribution  $q(\mathbf{z})$ , define *expected complete log likelihood*:

$$\langle \ell_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_q \stackrel{\text{def}}{=} \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}, \theta) \log p(\mathbf{x}, \mathbf{z} | \theta)$$

- A deterministic function of  $\theta$
- Linear in  $\ell_c()$  --- inherit its factorizability
- Does maximizing this surrogate yield a maximizer of the likelihood?

- Jensen's inequality

$$\begin{aligned} \ell(\theta; \mathbf{x}) &= \log p(\mathbf{x} | \theta) \\ &= \log \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z} | \theta) \\ &= \log \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}) \frac{p(\mathbf{x}, \mathbf{z} | \theta)}{q(\mathbf{z} | \mathbf{x})} \\ &\geq \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}) \log \frac{p(\mathbf{x}, \mathbf{z} | \theta)}{q(\mathbf{z} | \mathbf{x})} \end{aligned}$$



$$\Rightarrow \ell(\theta; \mathbf{x}) \geq \langle \ell_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_q + H_q$$



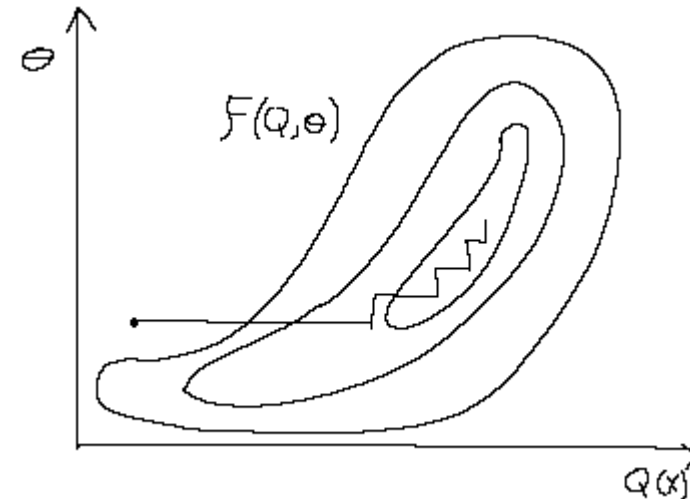
# Lower Bounds and Free Energy

- For fixed data  $x$ , define a functional called the free energy:

$$F(q, \theta) \stackrel{\text{def}}{=} \sum_z q(z | x) \log \frac{p(x, z | \theta)}{q(z | x)} \leq \ell(\theta; x)$$

- The EM algorithm is coordinate-ascent on  $F$ :

- E-step:  $q^{t+1} = \arg \max_q F(q, \theta^t)$
- M-step:  $\theta^{t+1} = \arg \max_{\theta} F(q^{t+1}, \theta)$



# E-step: maximization of expected $l_c$ w.r.t. $q$

- Claim:

$$q^{t+1} = \arg \max_q F(q, \theta^t) = p(z | x, \theta^t)$$

- This is the posterior distribution over the latent variables given the data and the parameters. Often we need this at test time anyway (e.g. to perform classification).
- Proof (easy): this setting attains the bound  $l(\theta; x) \geq F(q, \theta)$

$$\begin{aligned} F(p(z|x, \theta^t), \theta^t) &= \sum_z p(z|x, \theta^t) \log \frac{p(x, z | \theta^t)}{p(z|x, \theta^t)} \\ &= \sum_z p(z|x, \theta^t) \log p(x | \theta^t) \\ &= \log p(x | \theta^t) = \ell(\theta^t; x) \end{aligned}$$

- Can also show this result using variational calculus or the fact that

$$\ell(\theta; x) - F(q, \theta) = \text{KL}(q \| p(z | x, \theta))$$



# E-step $\equiv$ plug in posterior expectation of latent variables

- Without loss of generality: assume that  $p(\mathbf{x}, \mathbf{z} | \theta)$  is a generalized exponential family distribution:

$$p(\mathbf{x}, \mathbf{z} | \theta) = \frac{1}{Z(\theta)} h(\mathbf{x}, \mathbf{z}) \exp \left\{ \sum_i \theta_i f_i(\mathbf{x}, \mathbf{z}) \right\}$$

- Special cases: if  $p(\mathbf{x} | \mathbf{z})$  are GLIMs, then  $f_i(\mathbf{x}, \mathbf{z}) = \eta_i^T(\mathbf{z}) \xi_i(\mathbf{x})$
- The expected complete log likelihood under  $q^{t+1} = p(\mathbf{z} | \mathbf{x}, \theta^t)$  is

$$\begin{aligned} \langle \ell_c(\theta^t; \mathbf{x}, \mathbf{z}) \rangle_{q^{t+1}} &= \sum_{\mathbf{z}} q(\mathbf{z} | \mathbf{x}, \theta^t) \log p(\mathbf{x}, \mathbf{z} | \theta^t) - A(\theta) \\ &= \sum_i \theta_i^t \langle f_i(\mathbf{x}, \mathbf{z}) \rangle_{q(\mathbf{z} | \mathbf{x}, \theta^t)} - A(\theta) \\ &\stackrel{p \sim \text{GLIM}}{=} \sum_i \theta_i^t \langle \eta_i(\mathbf{z}) \rangle_{q(\mathbf{z} | \mathbf{x}, \theta^t)} \xi_i(\mathbf{x}) - A(\theta) \end{aligned}$$



# M-step: maximization of expected $l_c$ w.r.t. $\theta$

- Note that the free energy breaks into two terms:

$$\begin{aligned} F(q, \theta) &= \sum_z q(z | \mathbf{x}) \log \frac{p(\mathbf{x}, \mathbf{z} | \theta)}{q(z | \mathbf{x})} \\ &= \sum_z q(z | \mathbf{x}) \log p(\mathbf{x}, \mathbf{z} | \theta) - \sum_z q(z | \mathbf{x}) \log q(z | \mathbf{x}) \\ &= \langle \ell_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_q + H_q \end{aligned}$$

- The first term is the expected complete log likelihood (energy) and the second term, which does not depend on  $\theta$ , is the entropy.
- Thus, in the M-step, maximizing with respect to  $\theta$  for fixed  $q$  we only need to consider the first term:

$$\theta^{t+1} = \arg \max_{\theta} \langle \ell_c(\theta; \mathbf{x}, \mathbf{z}) \rangle_{q^{t+1}} = \arg \max_{\theta} \sum_z q(z | \mathbf{x}) \log p(\mathbf{x}, \mathbf{z} | \theta)$$

- Under optimal  $q^{t+1}$ , this is equivalent to solving a standard MLE of fully observed model  $p(\mathbf{x}, \mathbf{z} | \theta)$ , with the **sufficient statistics** involving  $\mathbf{z}$  replaced by their expectations w.r.t.  $p(\mathbf{z} | \mathbf{x}, \theta)$ .





# Summary: EM Algorithm

- A way of maximizing likelihood function for latent variable models. Finds MLE of parameters when the original (hard) problem can be broken up into two (easy) pieces:
  1. Estimate some “missing” or “unobserved” data from observed data and current parameters.
  2. Using this “complete” data, find the maximum likelihood parameter estimates.
- Alternate between filling in the latent variables using the best guess (posterior) and updating the parameters based on this guess:
  - E-step:  $q^{t+1} = \arg \max_q F(q, \theta^t)$
  - M-step:  $\theta^{t+1} = \arg \max_{\theta} F(q^{t+1}, \theta)$
- In the M-step we optimize a lower bound on the likelihood. In the E-step we close the gap, making bound=likelihood.



# EM Variants

- ❑ Sparse EM:

Do not re-compute exactly the posterior probability on each data point under all models, because it is almost zero. Instead keep an “active list” which you update every once in a while.

- ❑ Generalized (Incomplete) EM:

It might be hard to find the ML parameters in the M-step, even given the completed data. We can still make progress by doing an M-step that improves the likelihood a bit (e.g. gradient step). Recall the IRLS step in the mixture of experts model.



# A Report Card for EM

- ❑ Some good things about EM:
  - ❑ no learning rate (step-size) parameter
  - ❑ automatically enforces parameter constraints
  - ❑ very fast for low dimensions
  - ❑ each iteration guaranteed to improve likelihood
- ❑ Some bad things about EM:
  - ❑ can get stuck in local minima
  - ❑ can be slower than conjugate gradient (especially near convergence)
  - ❑ requires expensive inference step
  - ❑ is a maximum likelihood/MAP method





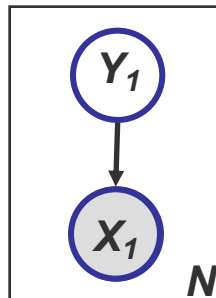
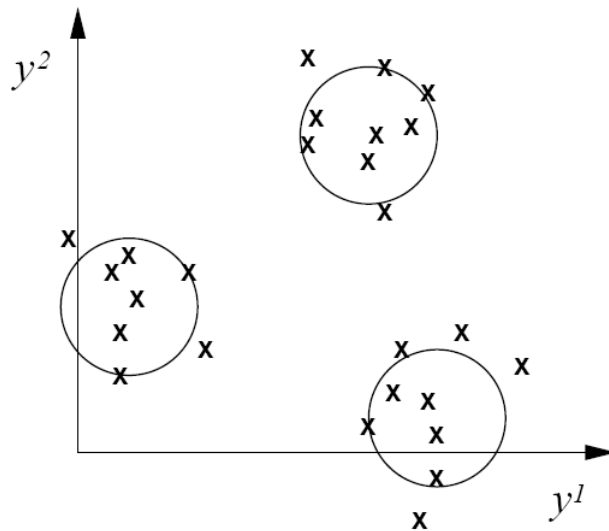
# Supplementary

## EM for Hidden Markov Models

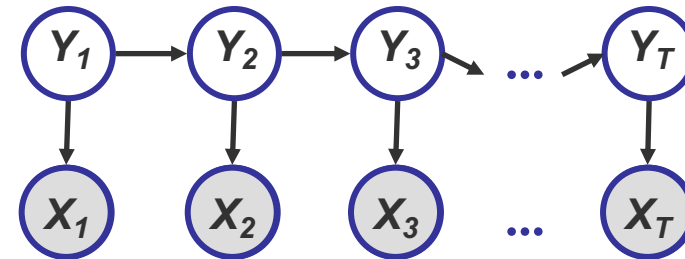
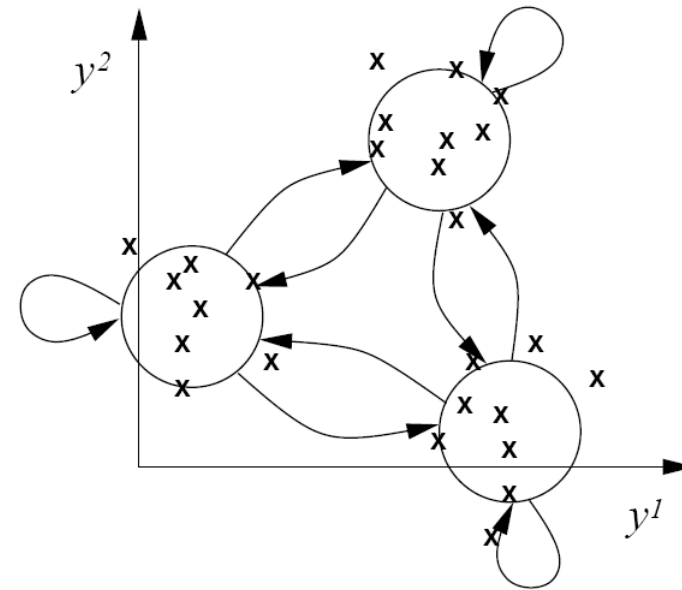


# From static to dynamic mixture models

Static mixture



Dynamic mixture



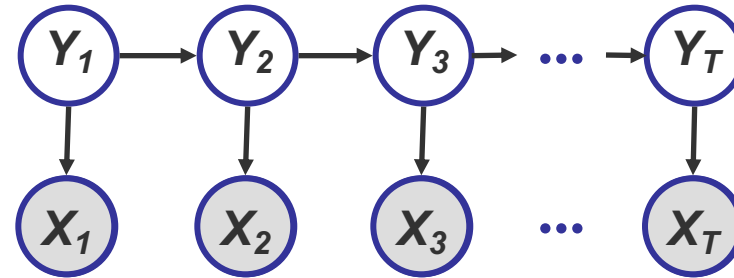
# Hidden Markov Models

**The underlying source:**

genomic entities,  
dice,

**The sequence:**

CGH signal,  
sequence of rolls,



Markov property:



# This problem is IMPORTANT!!! -☺

## An experience in a casino

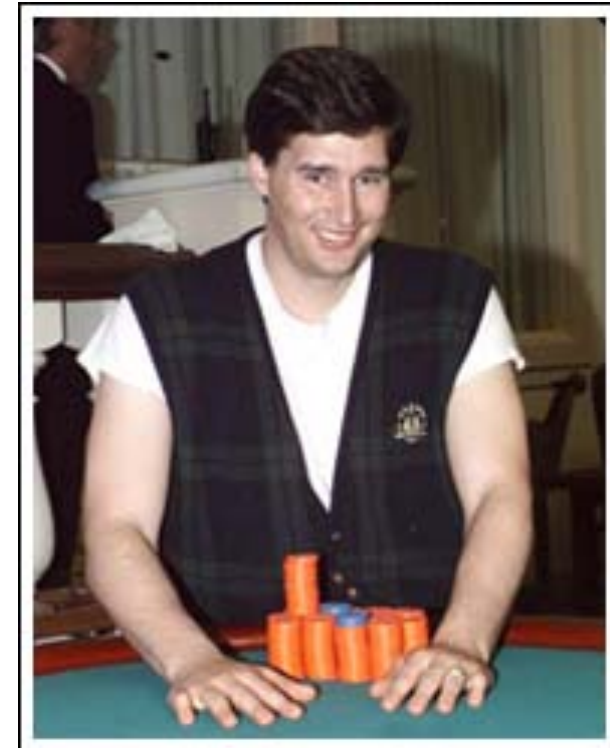
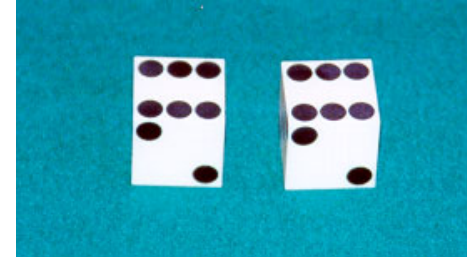
### Game:

1. You bet \$1
2. You roll (always with a fair die)
3. Casino player rolls (maybe with fair die, maybe with loaded die)
4. Highest number wins \$2

### Question:

1245526462146146136136661664661636  
616366163616515615115146123562344

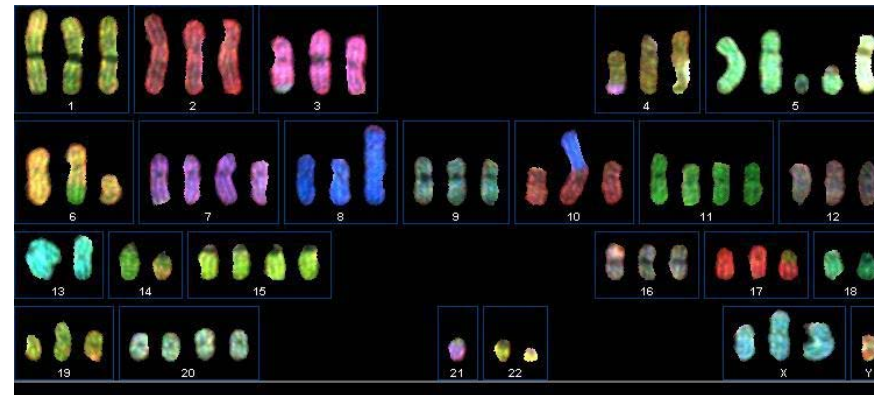
Which die is being used in each play?



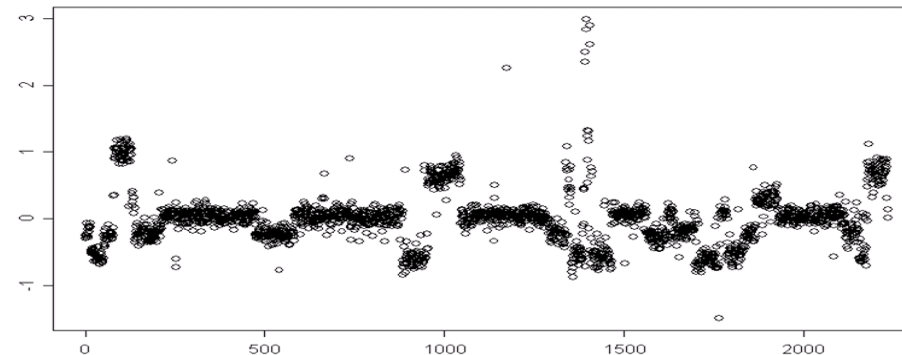
# A more serious question ...



- Naturally, data points arrive one at a time
  - Does the ordering index carry (additional) clustering information besides the data value itself ?
- Example:  
Chromosomes of tumor cell:

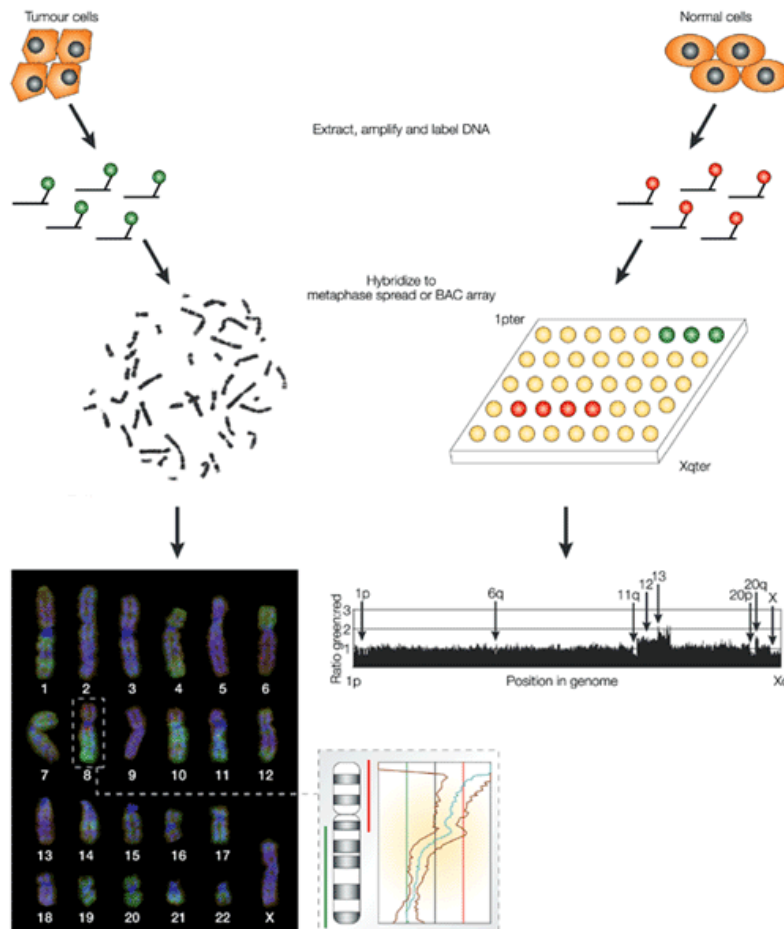


Copy number measurements  
(known as CGH)





# Array CGH (comparative genomic hybridization)

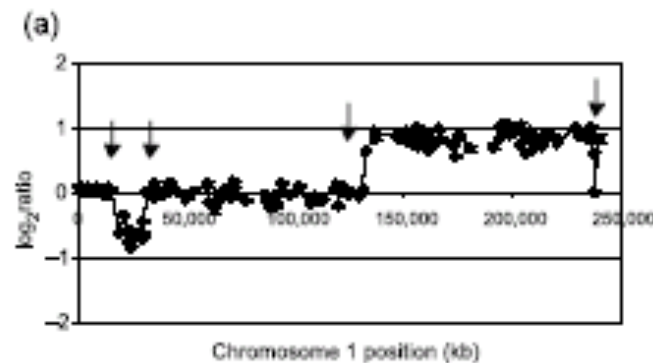


- The basic assumption of a CGH experiment is that the ratio of the binding of test and control DNA is proportional to the ratio of the copy numbers of sequences in the two samples.
- But various kinds of noises make the true observations less easy to interpret ...

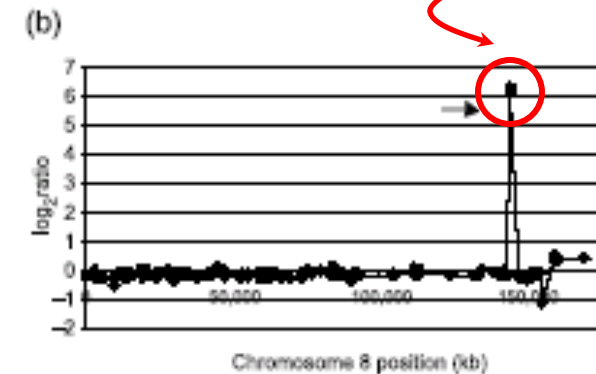


# DNA Copy number aberration types in breast cancer

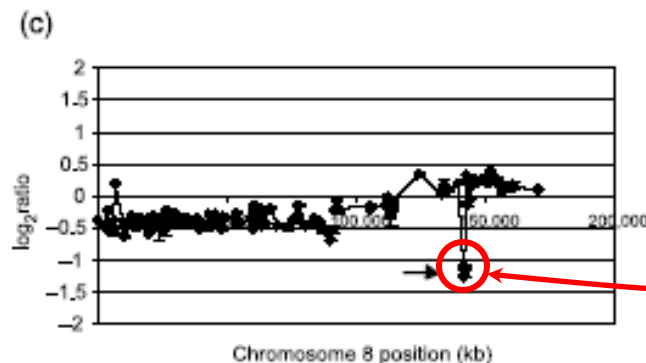
60-70 fold amplification of CMYC region



Copy number profile for chromosome 1 from 600 MPE cell line



Copy number profile for chromosome 8 from COLO320 cell line



Copy number profile for chromosome 8 in MDA-MB-231 cell line

deletion



# Suppose you were told about the following story before heading to Vegas...

## The Dishonest Casino !!!

A casino has two dice:

- Fair die

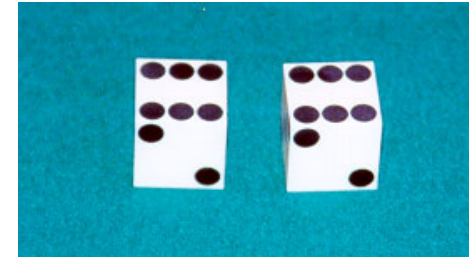
$$P(1) = P(2) = P(3) = P(5) = P(6) = 1/6$$

- Loaded die

$$P(1) = P(2) = P(3) = P(5) = 1/10$$

$$P(6) = 1/2$$

Casino player switches back-&-forth  
between fair and loaded die once  
every 20 turns



# Puzzles Regarding the Dishonest Casino

**GIVEN:** A sequence of rolls by the casino player

1245526462146146136136661664661636616366163616515615115146123562344

## QUESTION

- How likely is this sequence, given our model of how the casino works?
  - This is the **EVALUATION** problem
- What portion of the sequence was generated with the fair die, and what portion with the loaded die?
  - This is the **DECODING** question
- How “loaded” is the loaded die? How “fair” is the fair die? How often does the casino player change from fair to loaded, and back?
  - This is the **LEARNING** question



# Three Main Questions on HMMs

## 1. Evaluation

GIVEN an HMM  $M$ , and a sequence  $x$ ,  
FIND Prob ( $x$  |  $M$ )  
ALGO. Forward

## 2. Decoding

GIVEN an HMM  $M$ , and a sequence  $x$ ,  
FIND the sequence  $y$  of states that maximizes, e.g.,  $P(y | x, M)$ , or the most probable subsequence of states  
ALGO. Viterbi, Forward-backward

## 3. Learning

GIVEN an HMM  $M$ , with unspecified transition/emission probs., and a sequence  $x$ ,  
FIND parameters  $\theta = (\pi_i, a_{ij}, \eta_{ik})$  that maximize  $P(x | \theta)$   
ALGO. Baum-Welch (EM)



# Definition (of HMM)

- Observation space

Alphabetic set:

Euclidean space:

$$C = \{c_1, c_2, \dots, c_K\}$$
$$\mathbb{R}^d$$

- Index set of hidden states

$$I = \{1, 2, \dots, M\}$$

- Transition probabilities between any two states

$$p(y_t^j = 1 | y_{t-1}^i = 1) = a_{i,j},$$

or

$$p(y_t | y_{t-1}^i = 1) \sim \text{Multinomial}(a_{i,1}, a_{i,2}, \dots, a_{i,M}), \forall i \in I.$$

- Start probabilities

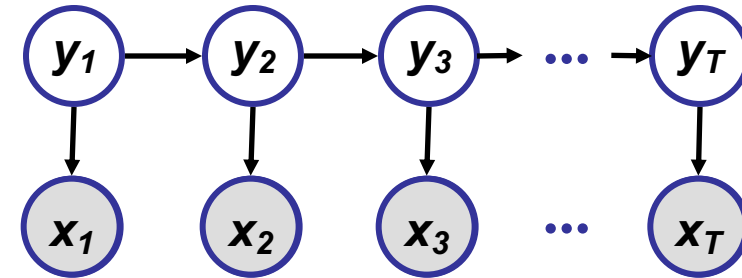
$$p(y_1) \sim \text{Multinomial}(\pi_1, \pi_2, \dots, \pi_M).$$

- Emission probabilities associated with each state

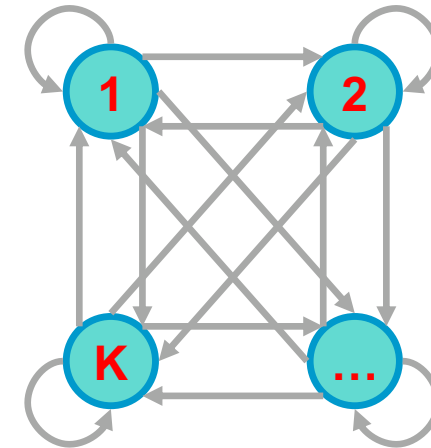
$$p(x_t | y_t^i = 1) \sim \text{Multinomial}(b_{i,1}, b_{i,2}, \dots, b_{i,K}), \forall i \in I.$$

or in general:

$$p(x_t | y_t^i = 1) \sim f(\cdot | \theta_i), \forall i \in I.$$



Graphical model



State automata



# Learning HMM: two scenarios

- Supervised learning: estimation when the “right answer” is known
  - Examples:
    - GIVEN: a genomic region  $x = x_1 \dots x_{1,000,000}$  where we have good (experimental) annotations of the CpG islands
    - GIVEN: the casino player allows us to observe him one evening, as he changes dice and produces 10,000 rolls
- Unsupervised learning: estimation when the “right answer” is unknown
  - Examples:
    - GIVEN: the porcupine genome; we don't know how frequent are the CpG islands there, neither do we know their composition
    - GIVEN: 10,000 rolls of the casino player, but we don't see when he changes dice
- **QUESTION**: Update the parameters  $\theta$  of the model to maximize  $P(x|\theta)$  --- Maximal likelihood (ML) estimation



# Supervised ML estimation

- Given  $\mathbf{x} = x_1 \dots x_N$  for which the true state path  $\mathbf{y} = y_1 \dots y_N$  is known,

- Define:

$A_{ij}$  = # times state transition  $i \rightarrow j$  occurs in  $\mathbf{y}$

$B_{ik}$  = # times state  $i$  in  $\mathbf{y}$  emits  $k$  in  $\mathbf{x}$

- We can show that the **maximum likelihood** parameters  $\theta$  are:

$$a_{ij}^{ML} = \frac{\#(i \rightarrow j)}{\#(i \rightarrow \bullet)} = \frac{\sum_n \sum_{t=2}^T y_{n,t-1}^i y_{n,t}^j}{\sum_n \sum_{t=2}^T y_{n,t-1}^i} = \frac{A_{ij}}{\sum_{j'} A_{ij'}}$$

$$b_{ik}^{ML} = \frac{\#(i \rightarrow k)}{\#(i \rightarrow \bullet)} = \frac{\sum_n \sum_{t=1}^T y_{n,t}^i x_{n,t}^k}{\sum_n \sum_{t=1}^T y_{n,t}^i} = \frac{B_{ik}}{\sum_{k'} B_{ik'}}$$

- What if  $\mathbf{x}$  is continuous? We can treat  $\{(x_{n,t}, y_{n,t}) : t = 1:T, n = 1:N\}$  as  $N$   $T$  observations of, e.g., a Gaussian, and apply learning rules for Gaussian ...





# Unsupervised ML estimation

- Given  $\mathbf{x} = \mathbf{x}_1 \dots \mathbf{x}_N$  for which the true state path  $\mathbf{y} = \mathbf{y}_1 \dots \mathbf{y}_N$  is unknown,

- EXPECTATION MAXIMIZATION

0. Starting with our best guess of a model  $\mathcal{M}$ , parameters  $\theta$ .
1. Estimate  $\mathbf{A}_{ij}$ ,  $\mathbf{B}_{ik}$  in the training data
  - How?
  - Update  $\theta$  according to  $\mathbf{A}_{ij}$ ,  $\mathbf{B}_{ik}$
  - Now a "supervised learning" problem

$$\mathbf{A}_{ij} = \sum_{n,t} \langle \mathbf{y}_{n,t-1}^i \mathbf{y}_{n,t}^j \rangle \quad \mathbf{B}_{ik} = \sum_{n,t} \langle \mathbf{y}_{n,t}^i \mathbf{x}_{n,t}^k \rangle$$

2. Repeat 1 & 2, until convergence

This is called the Baum-Welch Algorithm

We can get to a provably more (or equally) likely parameter set  $\theta$  each iteration



# The Baum Welch algorithm – an EM algorithm

- The complete log likelihood

$$\ell_c(\theta; \mathbf{x}, \mathbf{y}) = \log p(\mathbf{x}, \mathbf{y}) = \log \prod_n \left( p(y_{n,1}) \prod_{t=2}^T p(y_{n,t} | y_{n,t-1}) \prod_{t=1}^T p(x_{n,t} | x_{n,t}) \right)$$

- The expected complete log likelihood

$$\langle \ell_c(\theta; \mathbf{x}, \mathbf{y}) \rangle = \sum_n \left( \langle y_{n,1}^i \rangle_{p(y_{n,1} | \mathbf{x}_n)} \log \pi_i \right) + \sum_n \sum_{t=2}^T \left( \langle y_{n,t-1}^i y_{n,t}^j \rangle_{p(y_{n,t-1}, y_{n,t} | \mathbf{x}_n)} \log a_{i,j} \right) + \sum_n \sum_{t=1}^T \left( x_{n,t}^k \langle y_{n,t}^i \rangle_{p(y_{n,t} | \mathbf{x}_n)} \log b_{i,k} \right)$$

- EM

- The E step

$$\gamma_{n,t}^i = \langle y_{n,t}^i \rangle = p(y_{n,t}^i = 1 | \mathbf{x}_n)$$

$$\xi_{n,t}^{i,j} = \langle y_{n,t-1}^i y_{n,t}^j \rangle = p(y_{n,t-1}^i = 1, y_{n,t}^j = 1 | \mathbf{x}_n)$$

- The M step ("symbolically" identical to MLE)

$$\pi_i^{ML} = \frac{\sum_n \gamma_{n,1}^i}{N}$$

$$a_{ij}^{ML} = \frac{\sum_n \sum_{t=2}^T \xi_{n,t}^{i,j}}{\sum_n \sum_{t=1}^{T-1} \gamma_{n,t}^i}$$

$$b_{ik}^{ML} = \frac{\sum_n \sum_{t=1}^T \gamma_{n,t}^i x_{n,t}^k}{\sum_n \sum_{t=1}^{T-1} \gamma_{n,t}^i}$$

