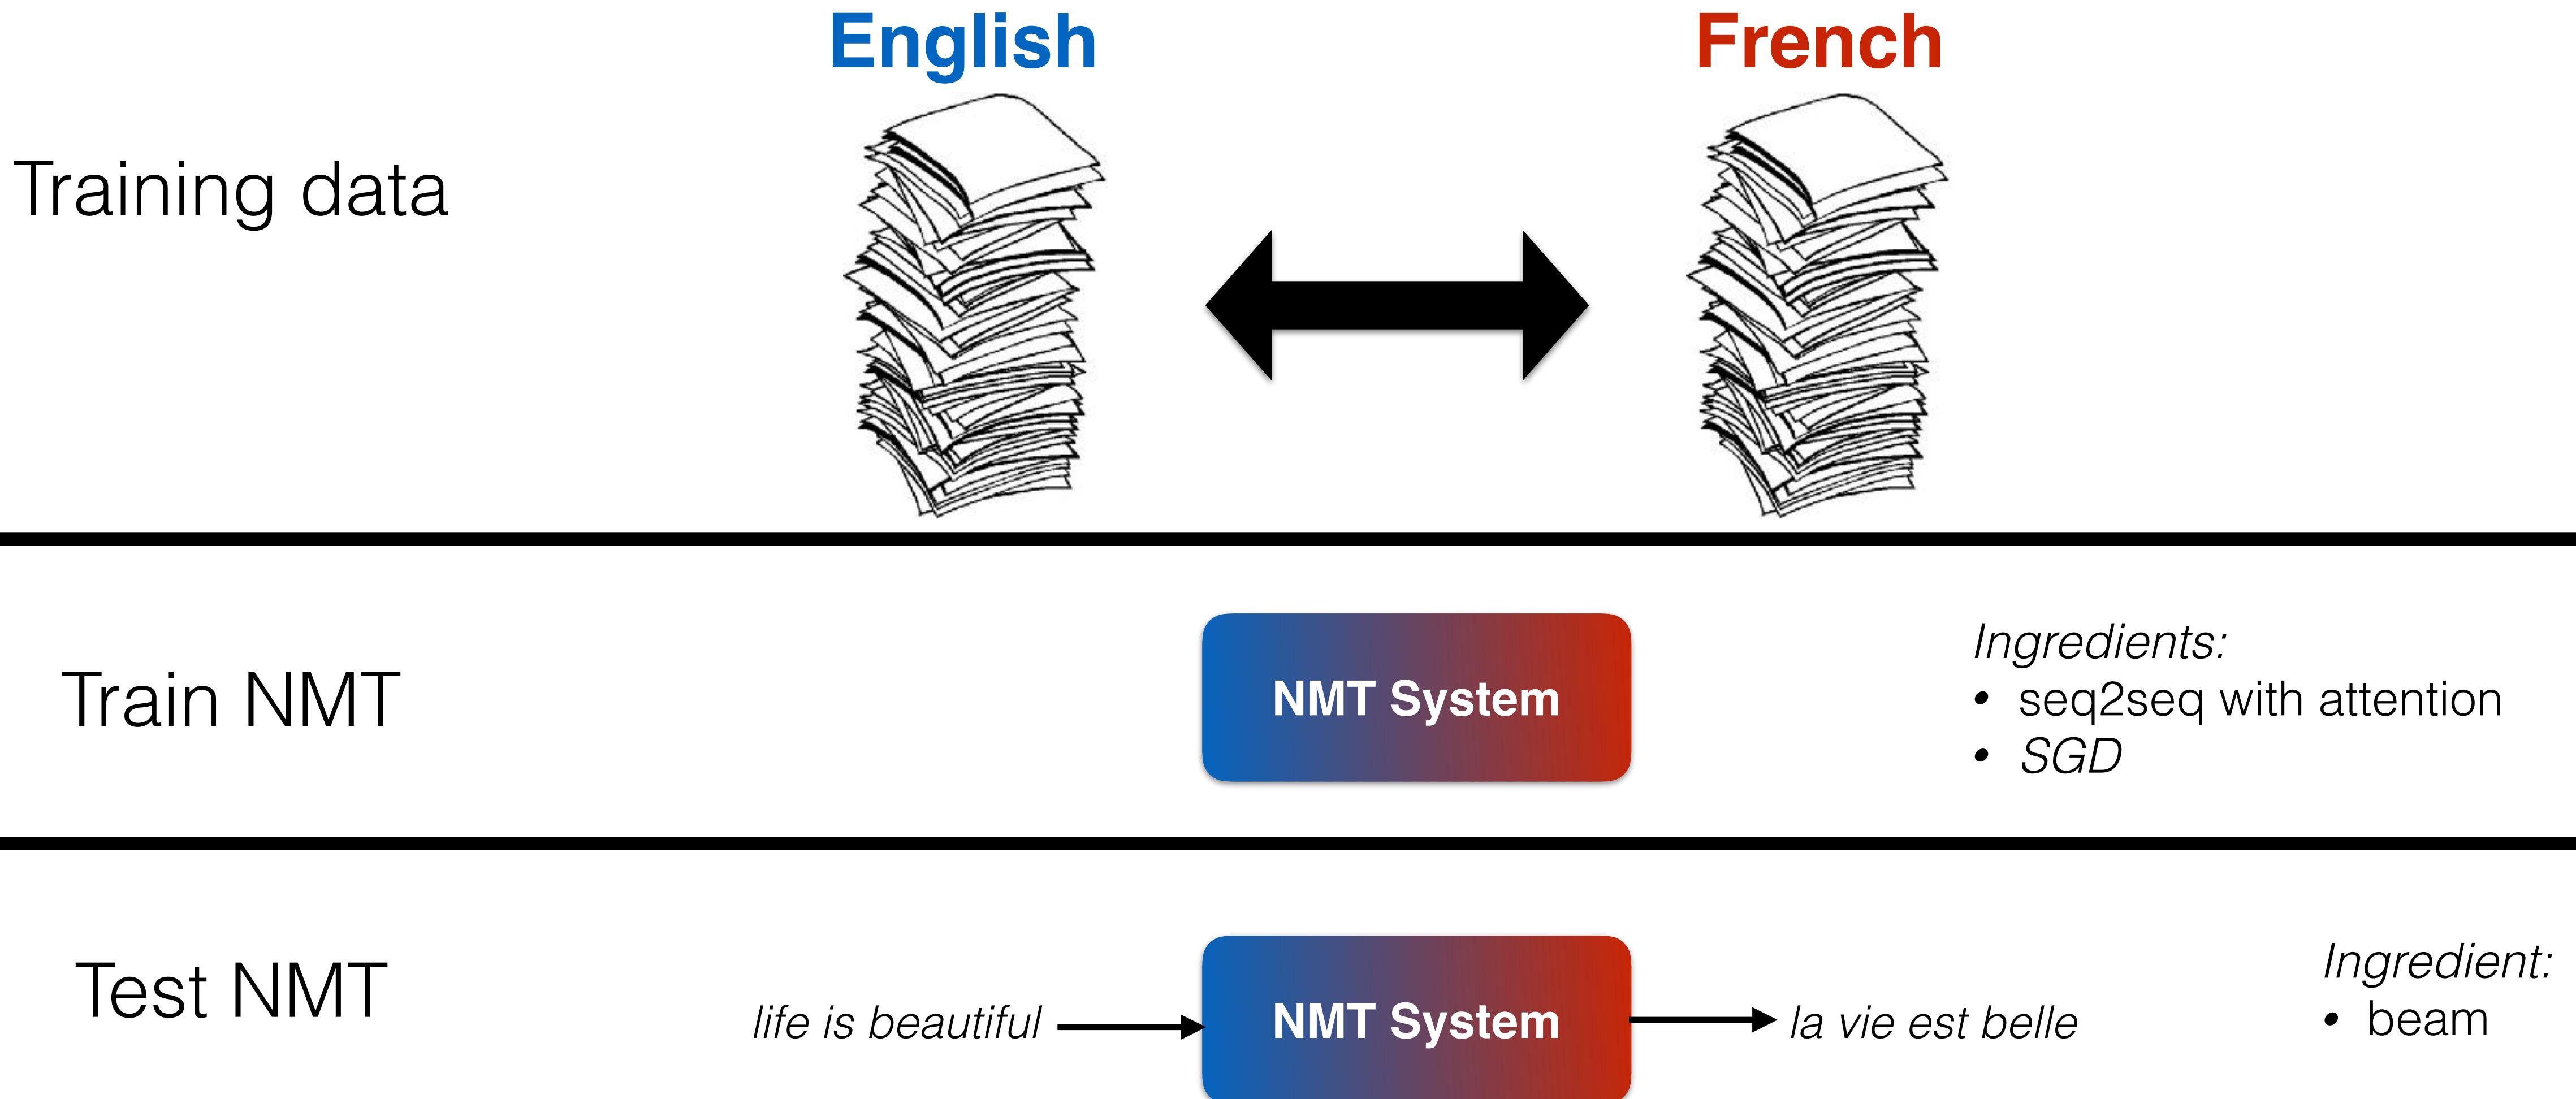


Low Resource Machine Translation

Marc'Aurelio Ranzato
Facebook AI Research - NYC
ranzato@fb.com

Machine Translation

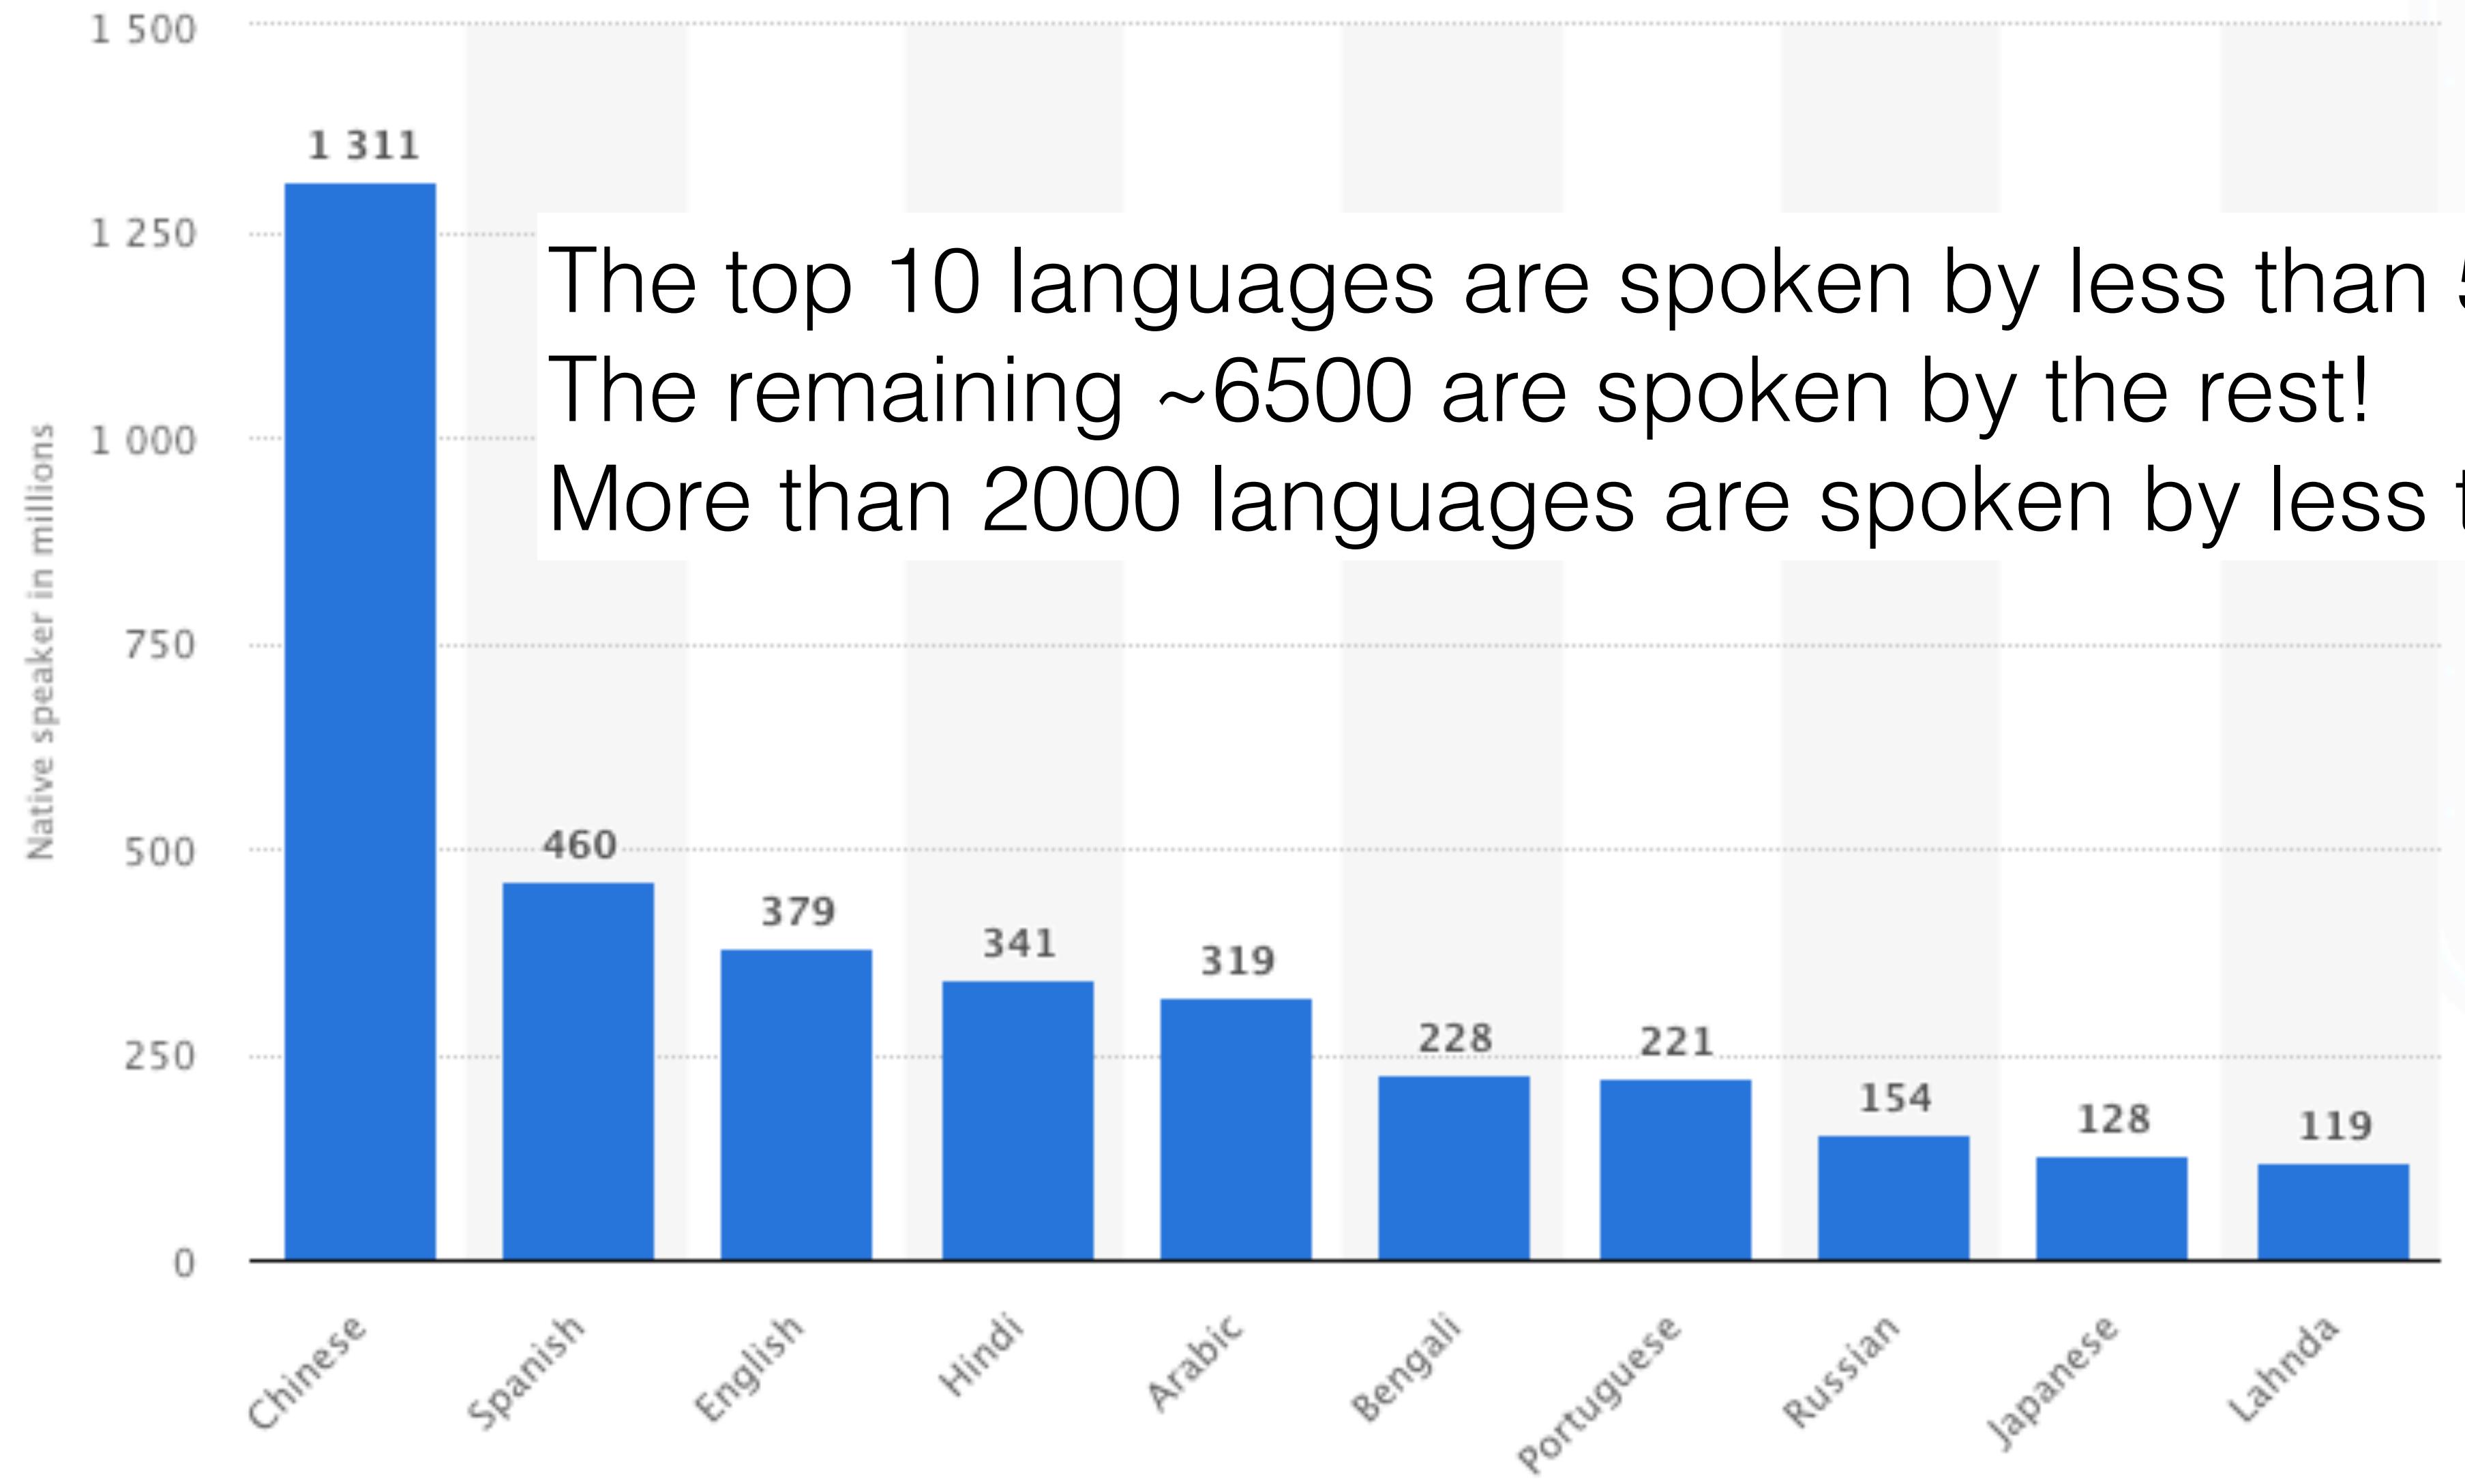


Some Stats

- 6000+ languages in the world
- 80% of the world population does not speak English
- Less than 5% of the people in the world are native English speakers.

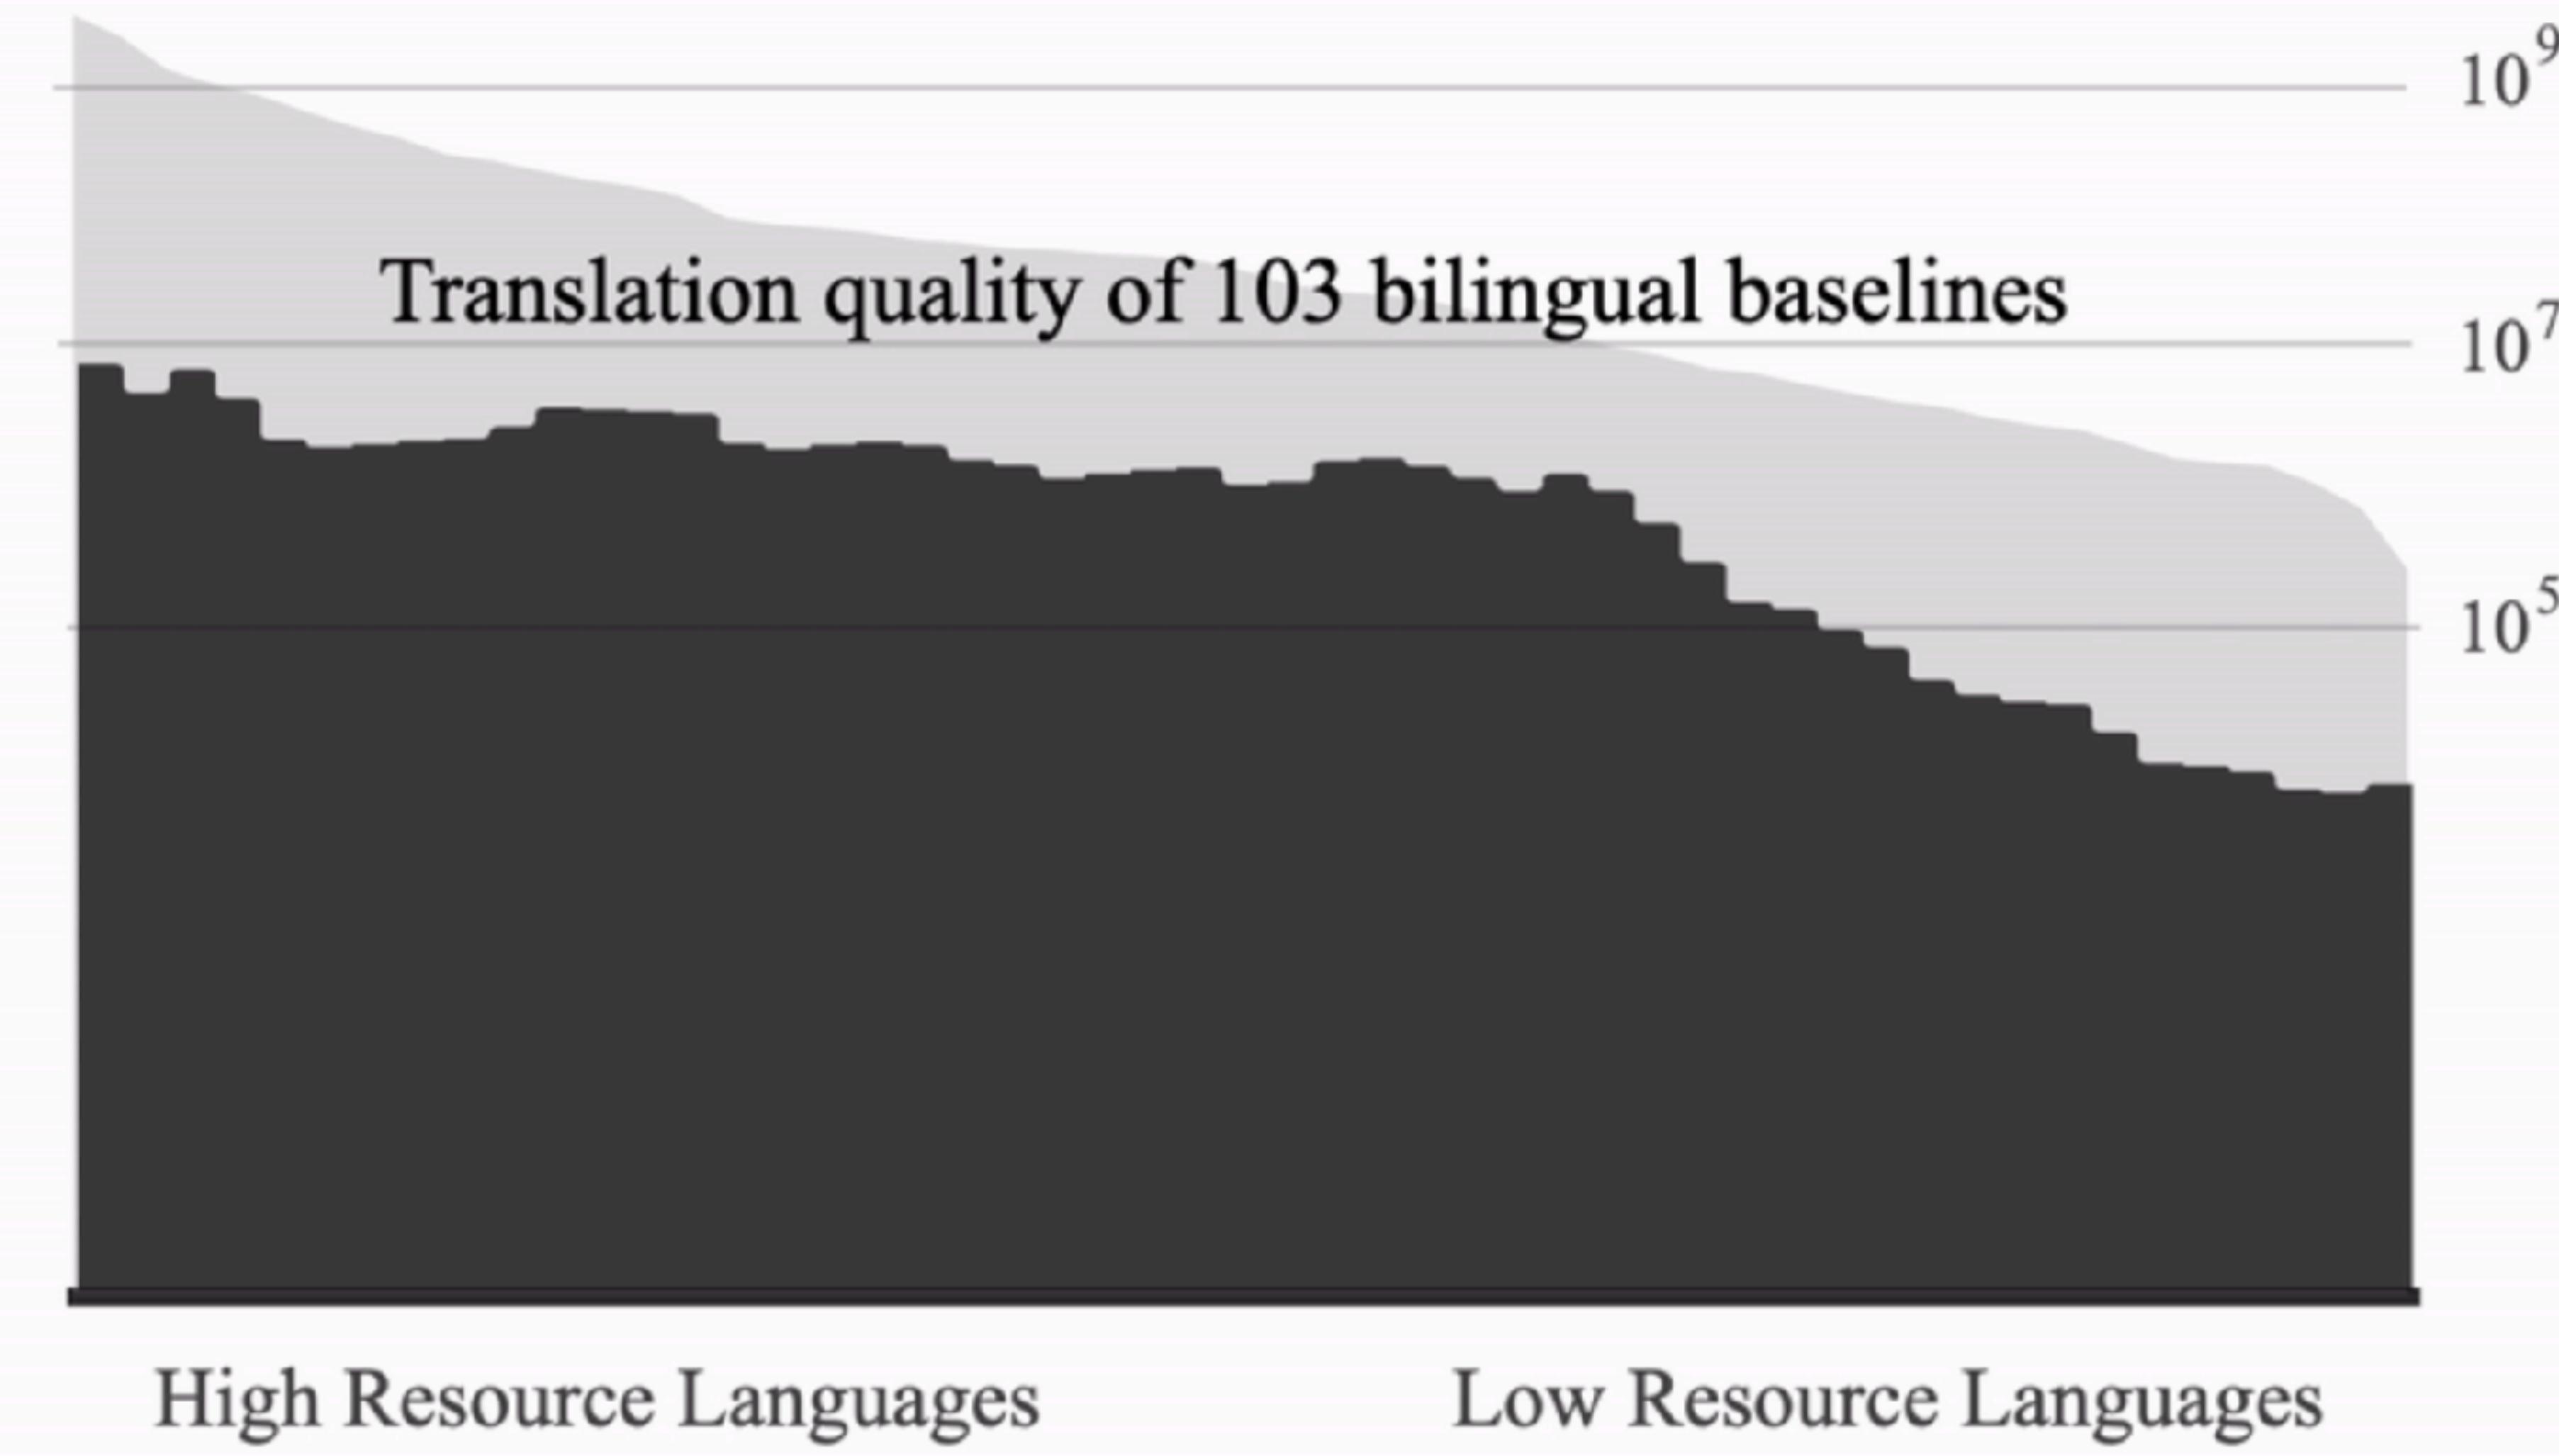


The Long Tail of Languages



The top 10 languages are spoken by less than 50% of the people.
The remaining ~6500 are spoken by the rest!
More than 2000 languages are spoken by less than 1000 people.

Data distribution over language pairs (X to English)



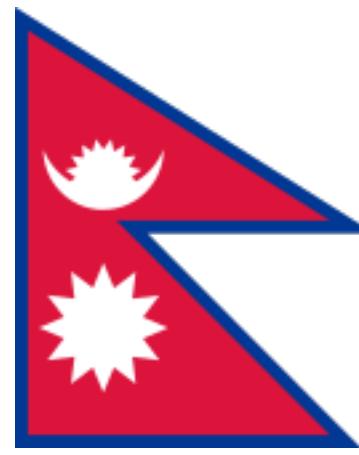
Machine Translation in Practice

Training data

English

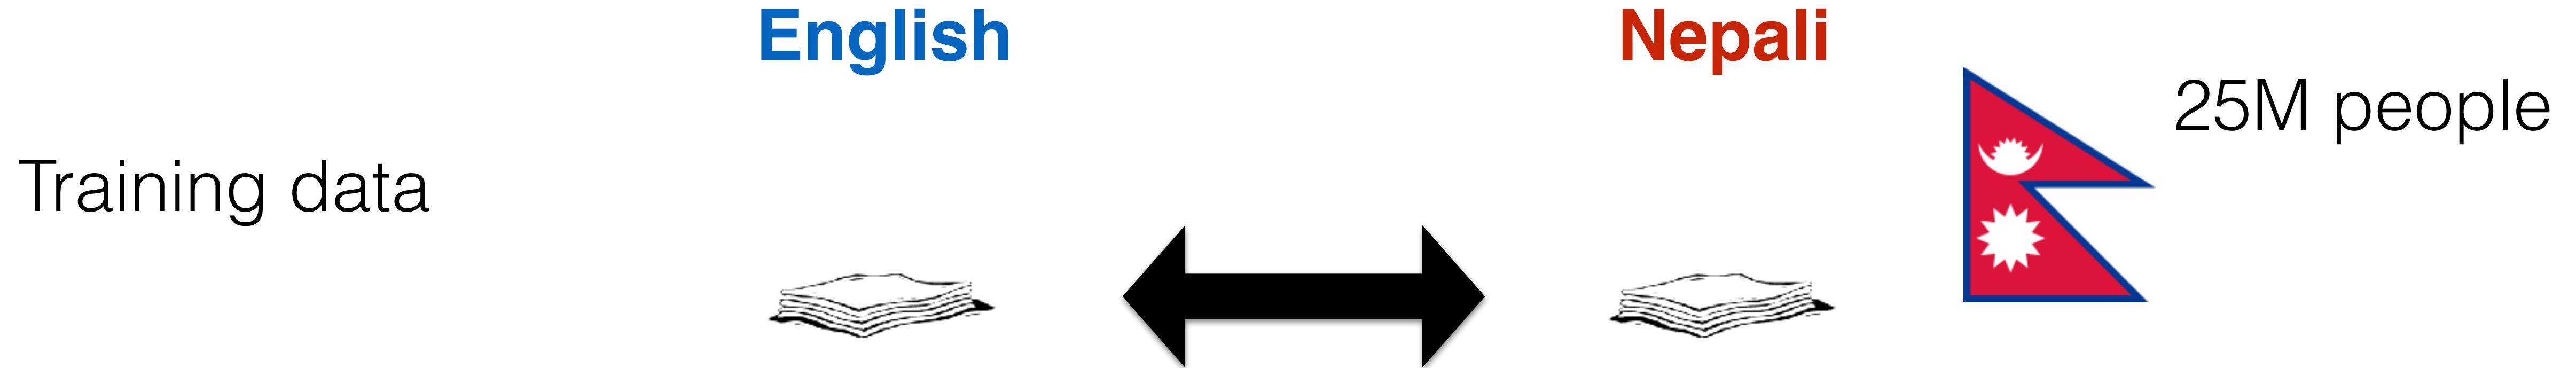


Nepali



25M people

Machine Translation in Practice



Parallel training data (collection of sentences with corresponding translation) is small!

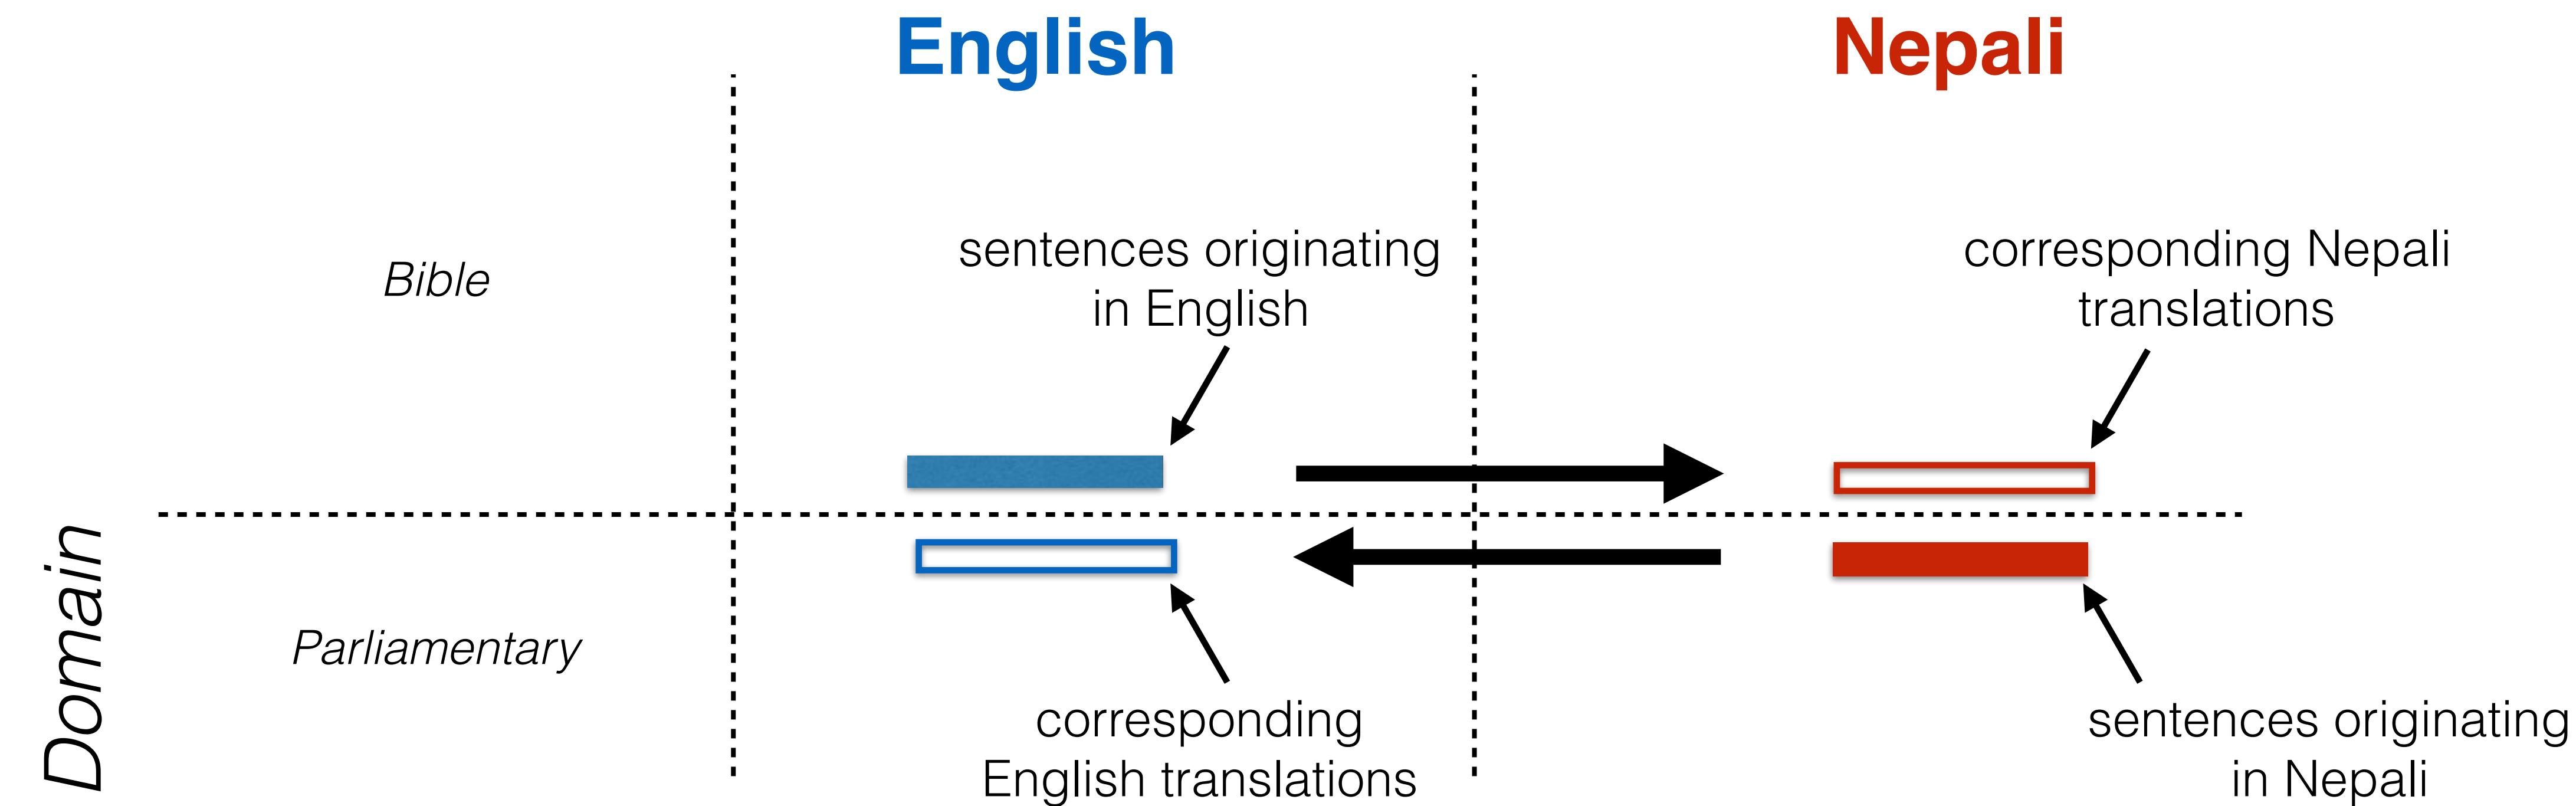
Machine Translation in Practice



Training data

Let's represent data with rectangles. The color indicates the language.

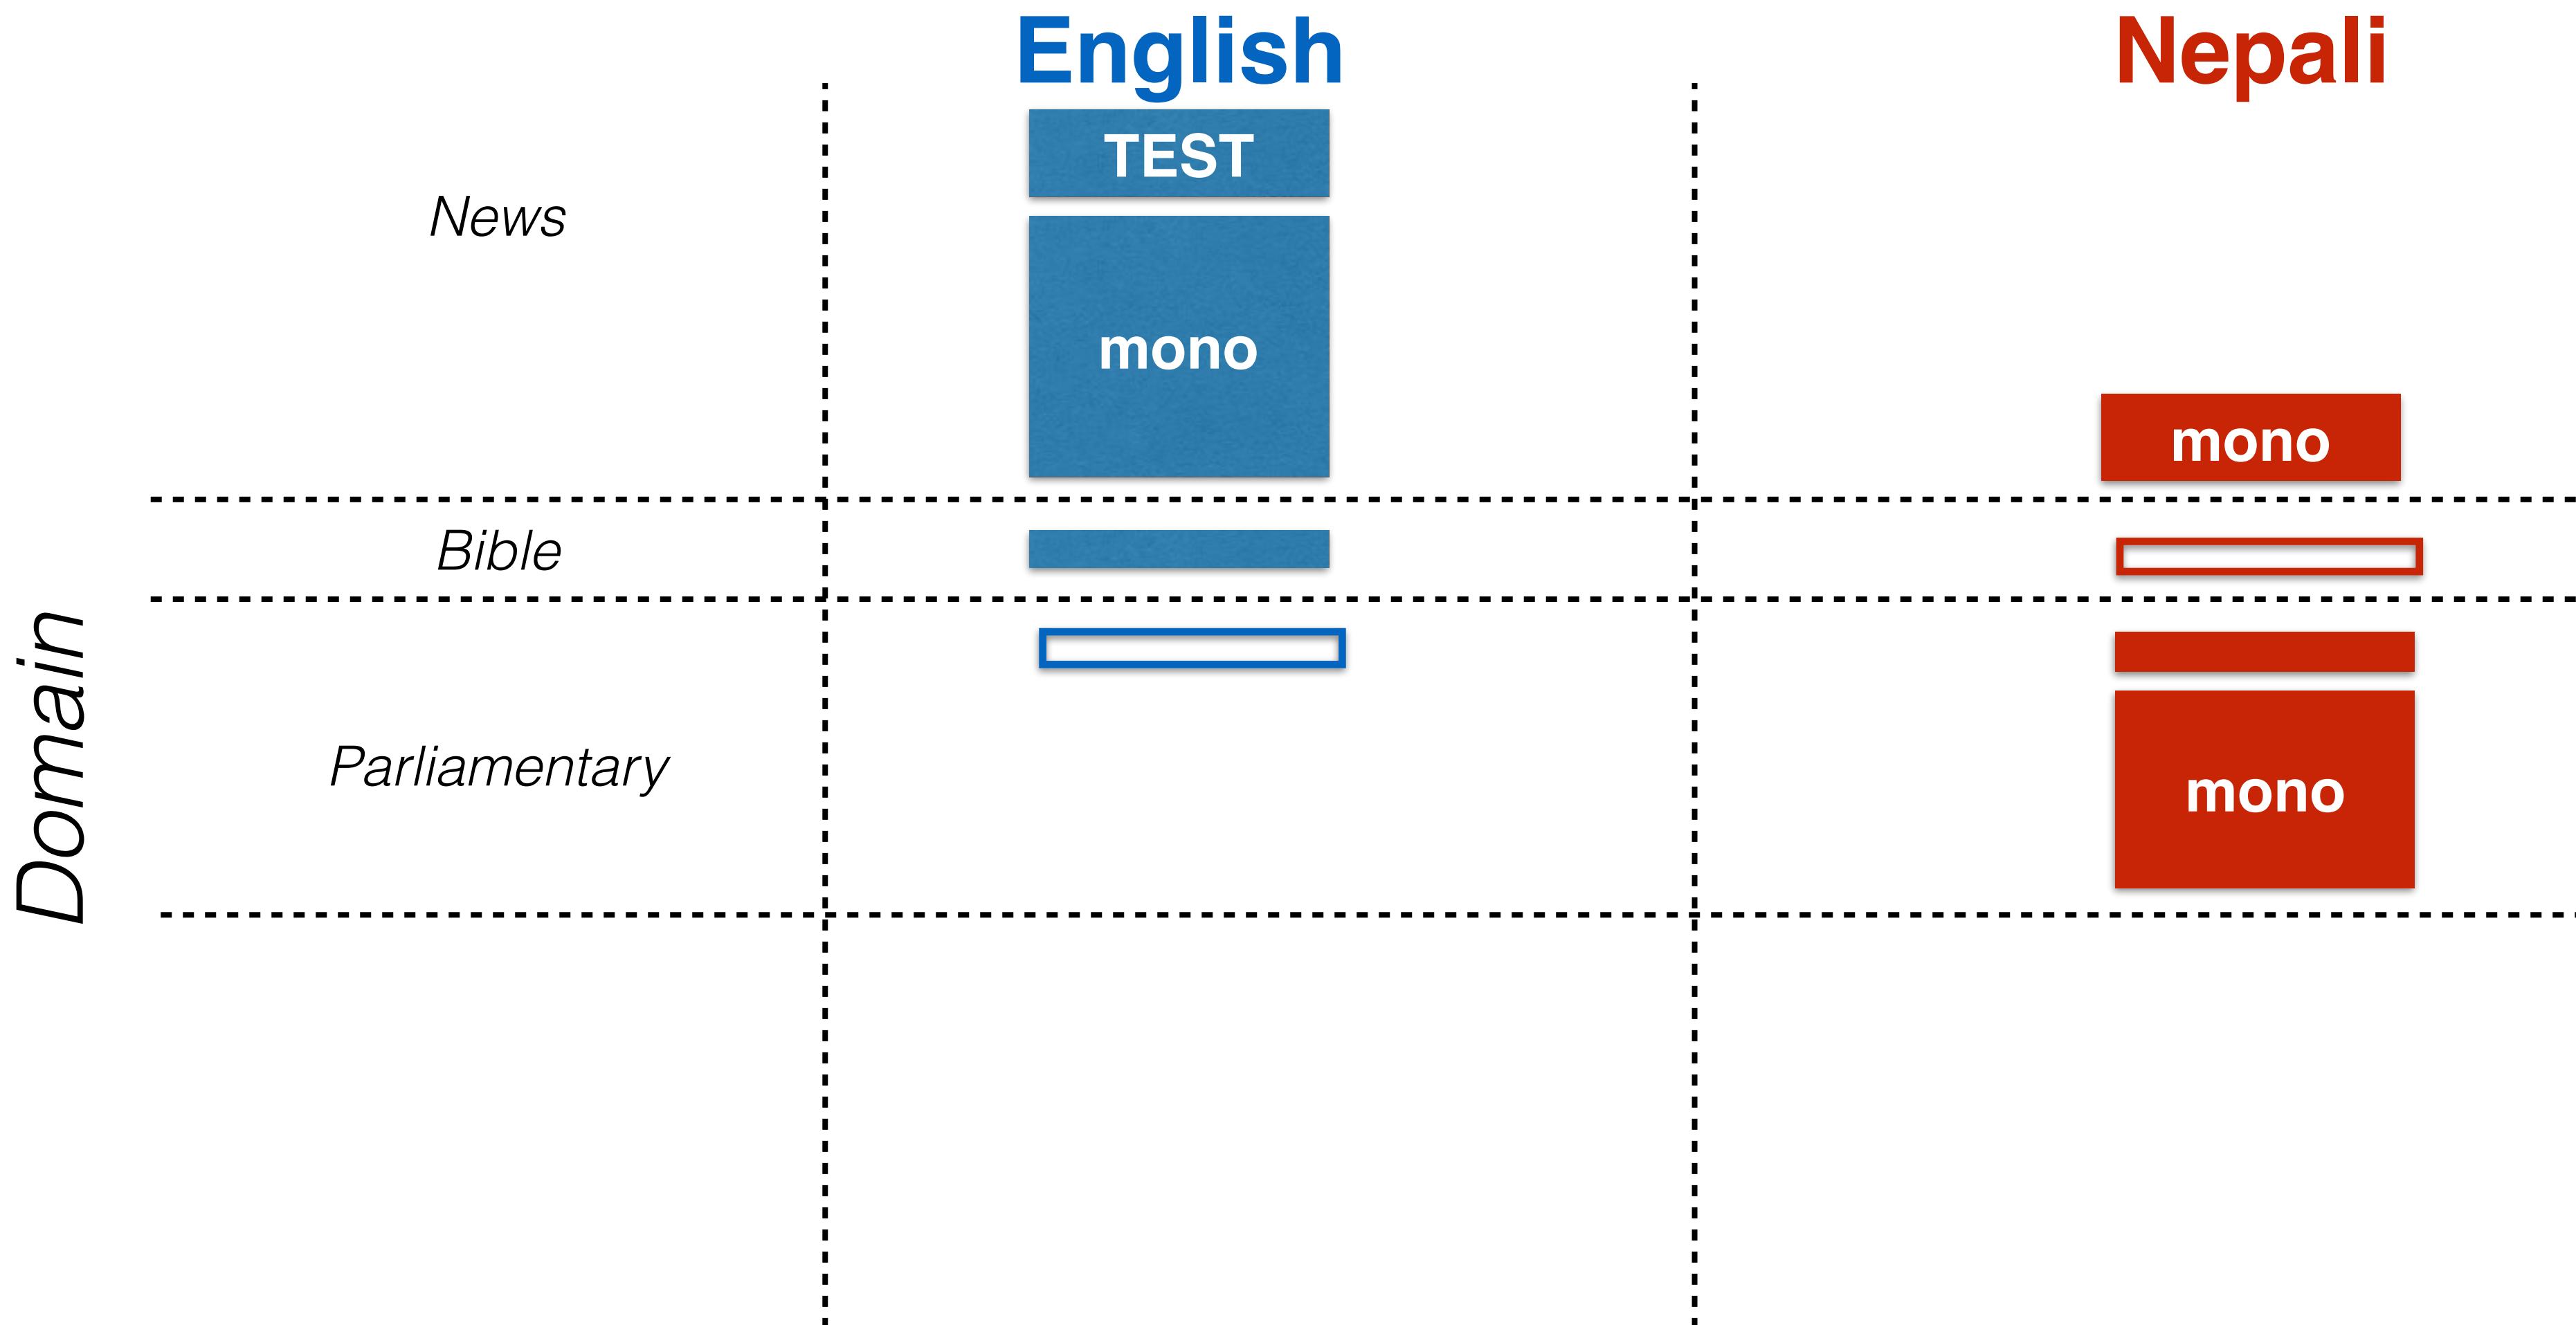
Machine Translation in Practice



Let's represent (human) translations with empty rectangles.

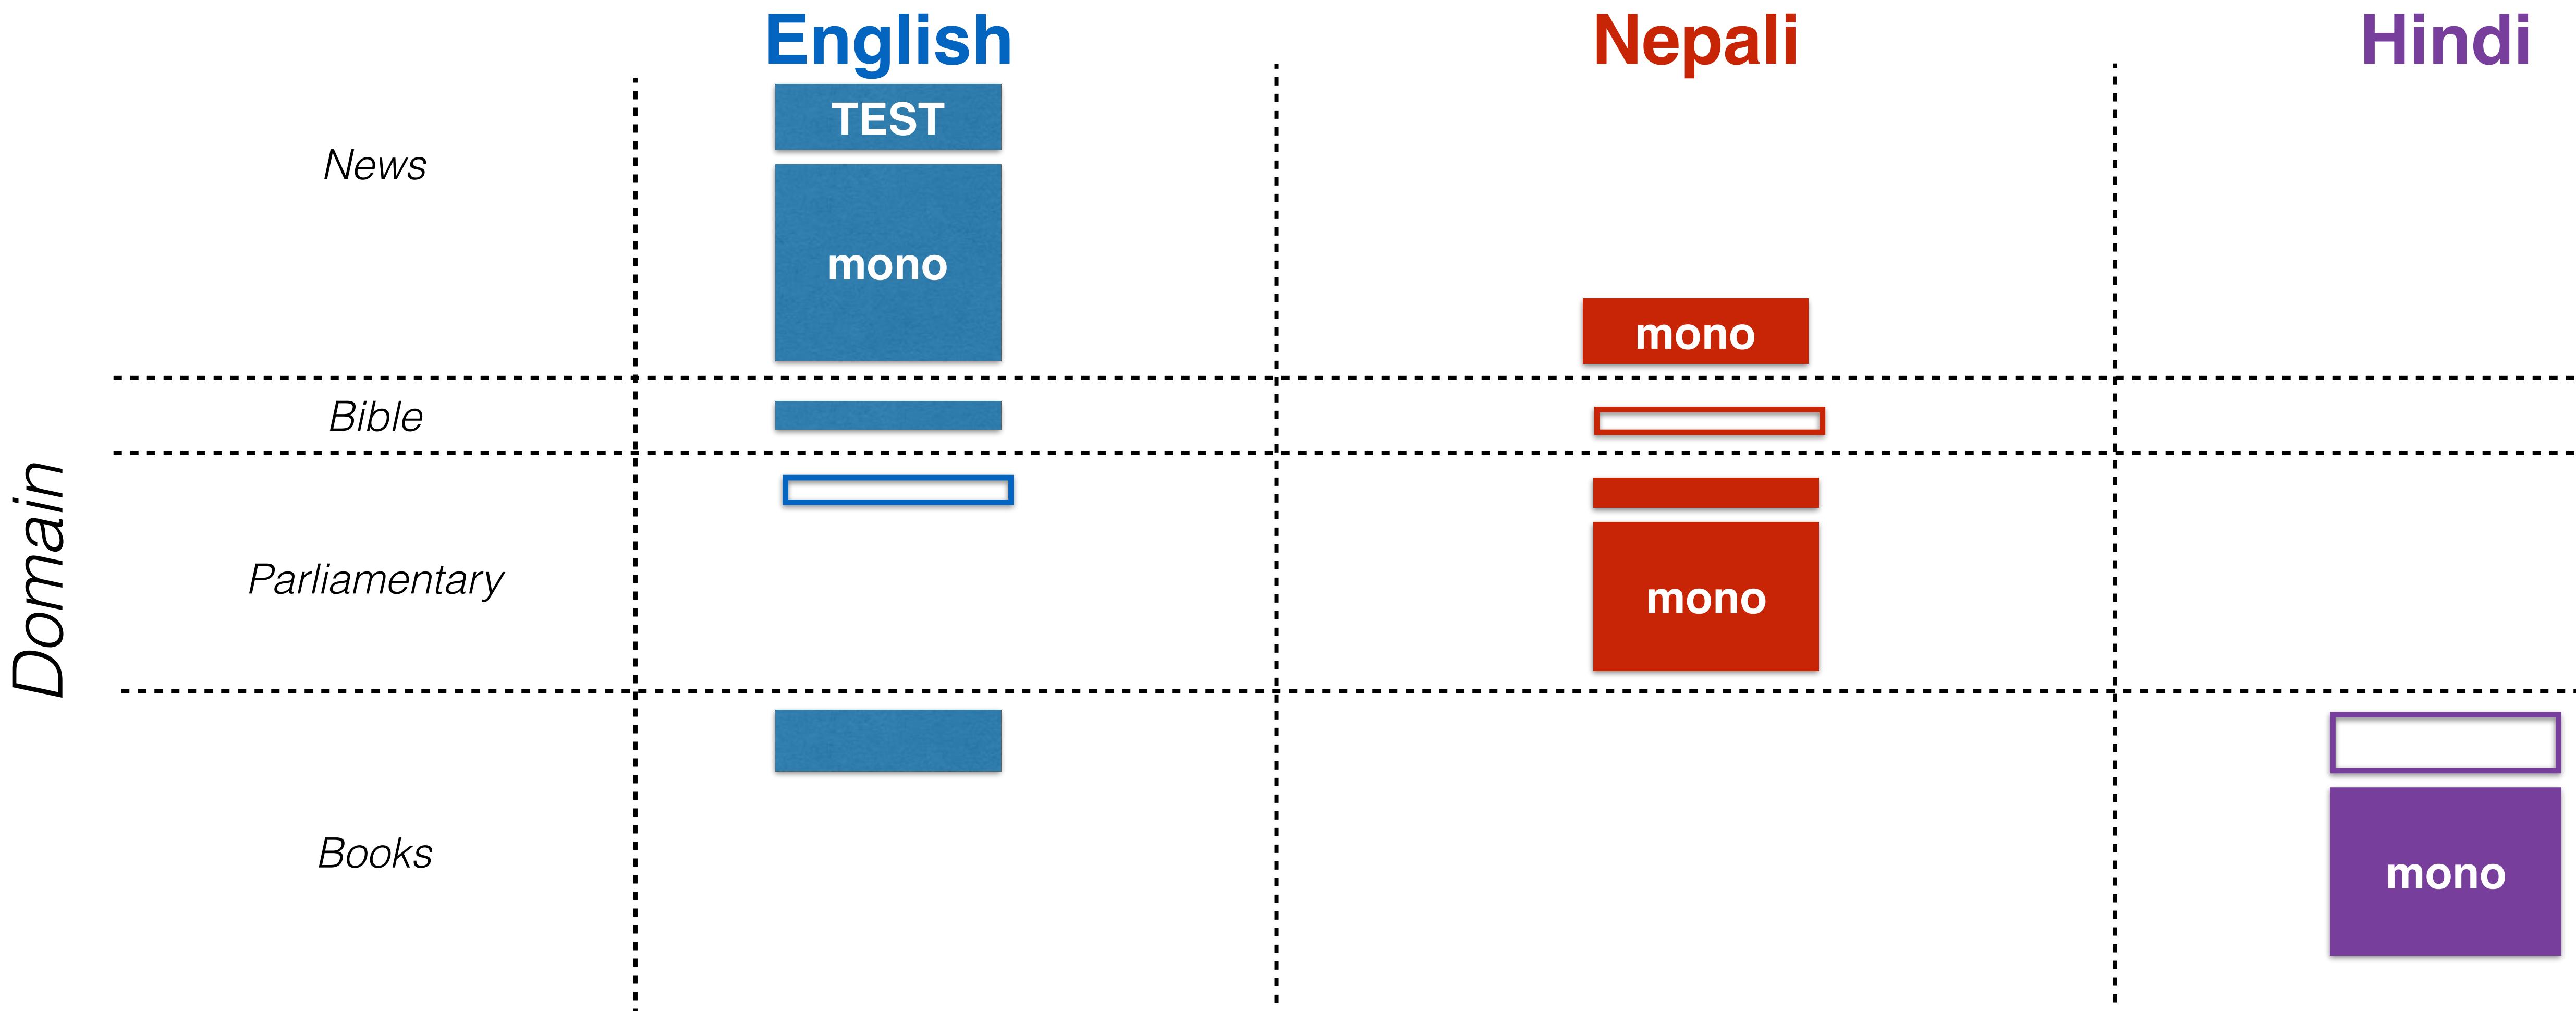
- Some parallel data originates in the source, some in the target language.
- Source and target domains may not match.

Machine Translation in Practice



- Test data might be in another domain.
- There might exist source side in-domain monolingual data.

Machine Translation in Practice



- There might be parallel and monolingual data with a high resource language close to the low resource language of interest. This data may belong to a different domain.

English

Nepali

Hindi

Sinhala

Bengali

Spanish

Tamil

Gujarati

TEST

Domain

the *Mondrian* like learning setting! ...



Low Resource Machine Translation

Loose definition: A language pair can be considered **low resource** when the number of parallel sentences is in the order of 10,000 or less.

Note: modern NMT systems have several hundred million parameters nowadays!

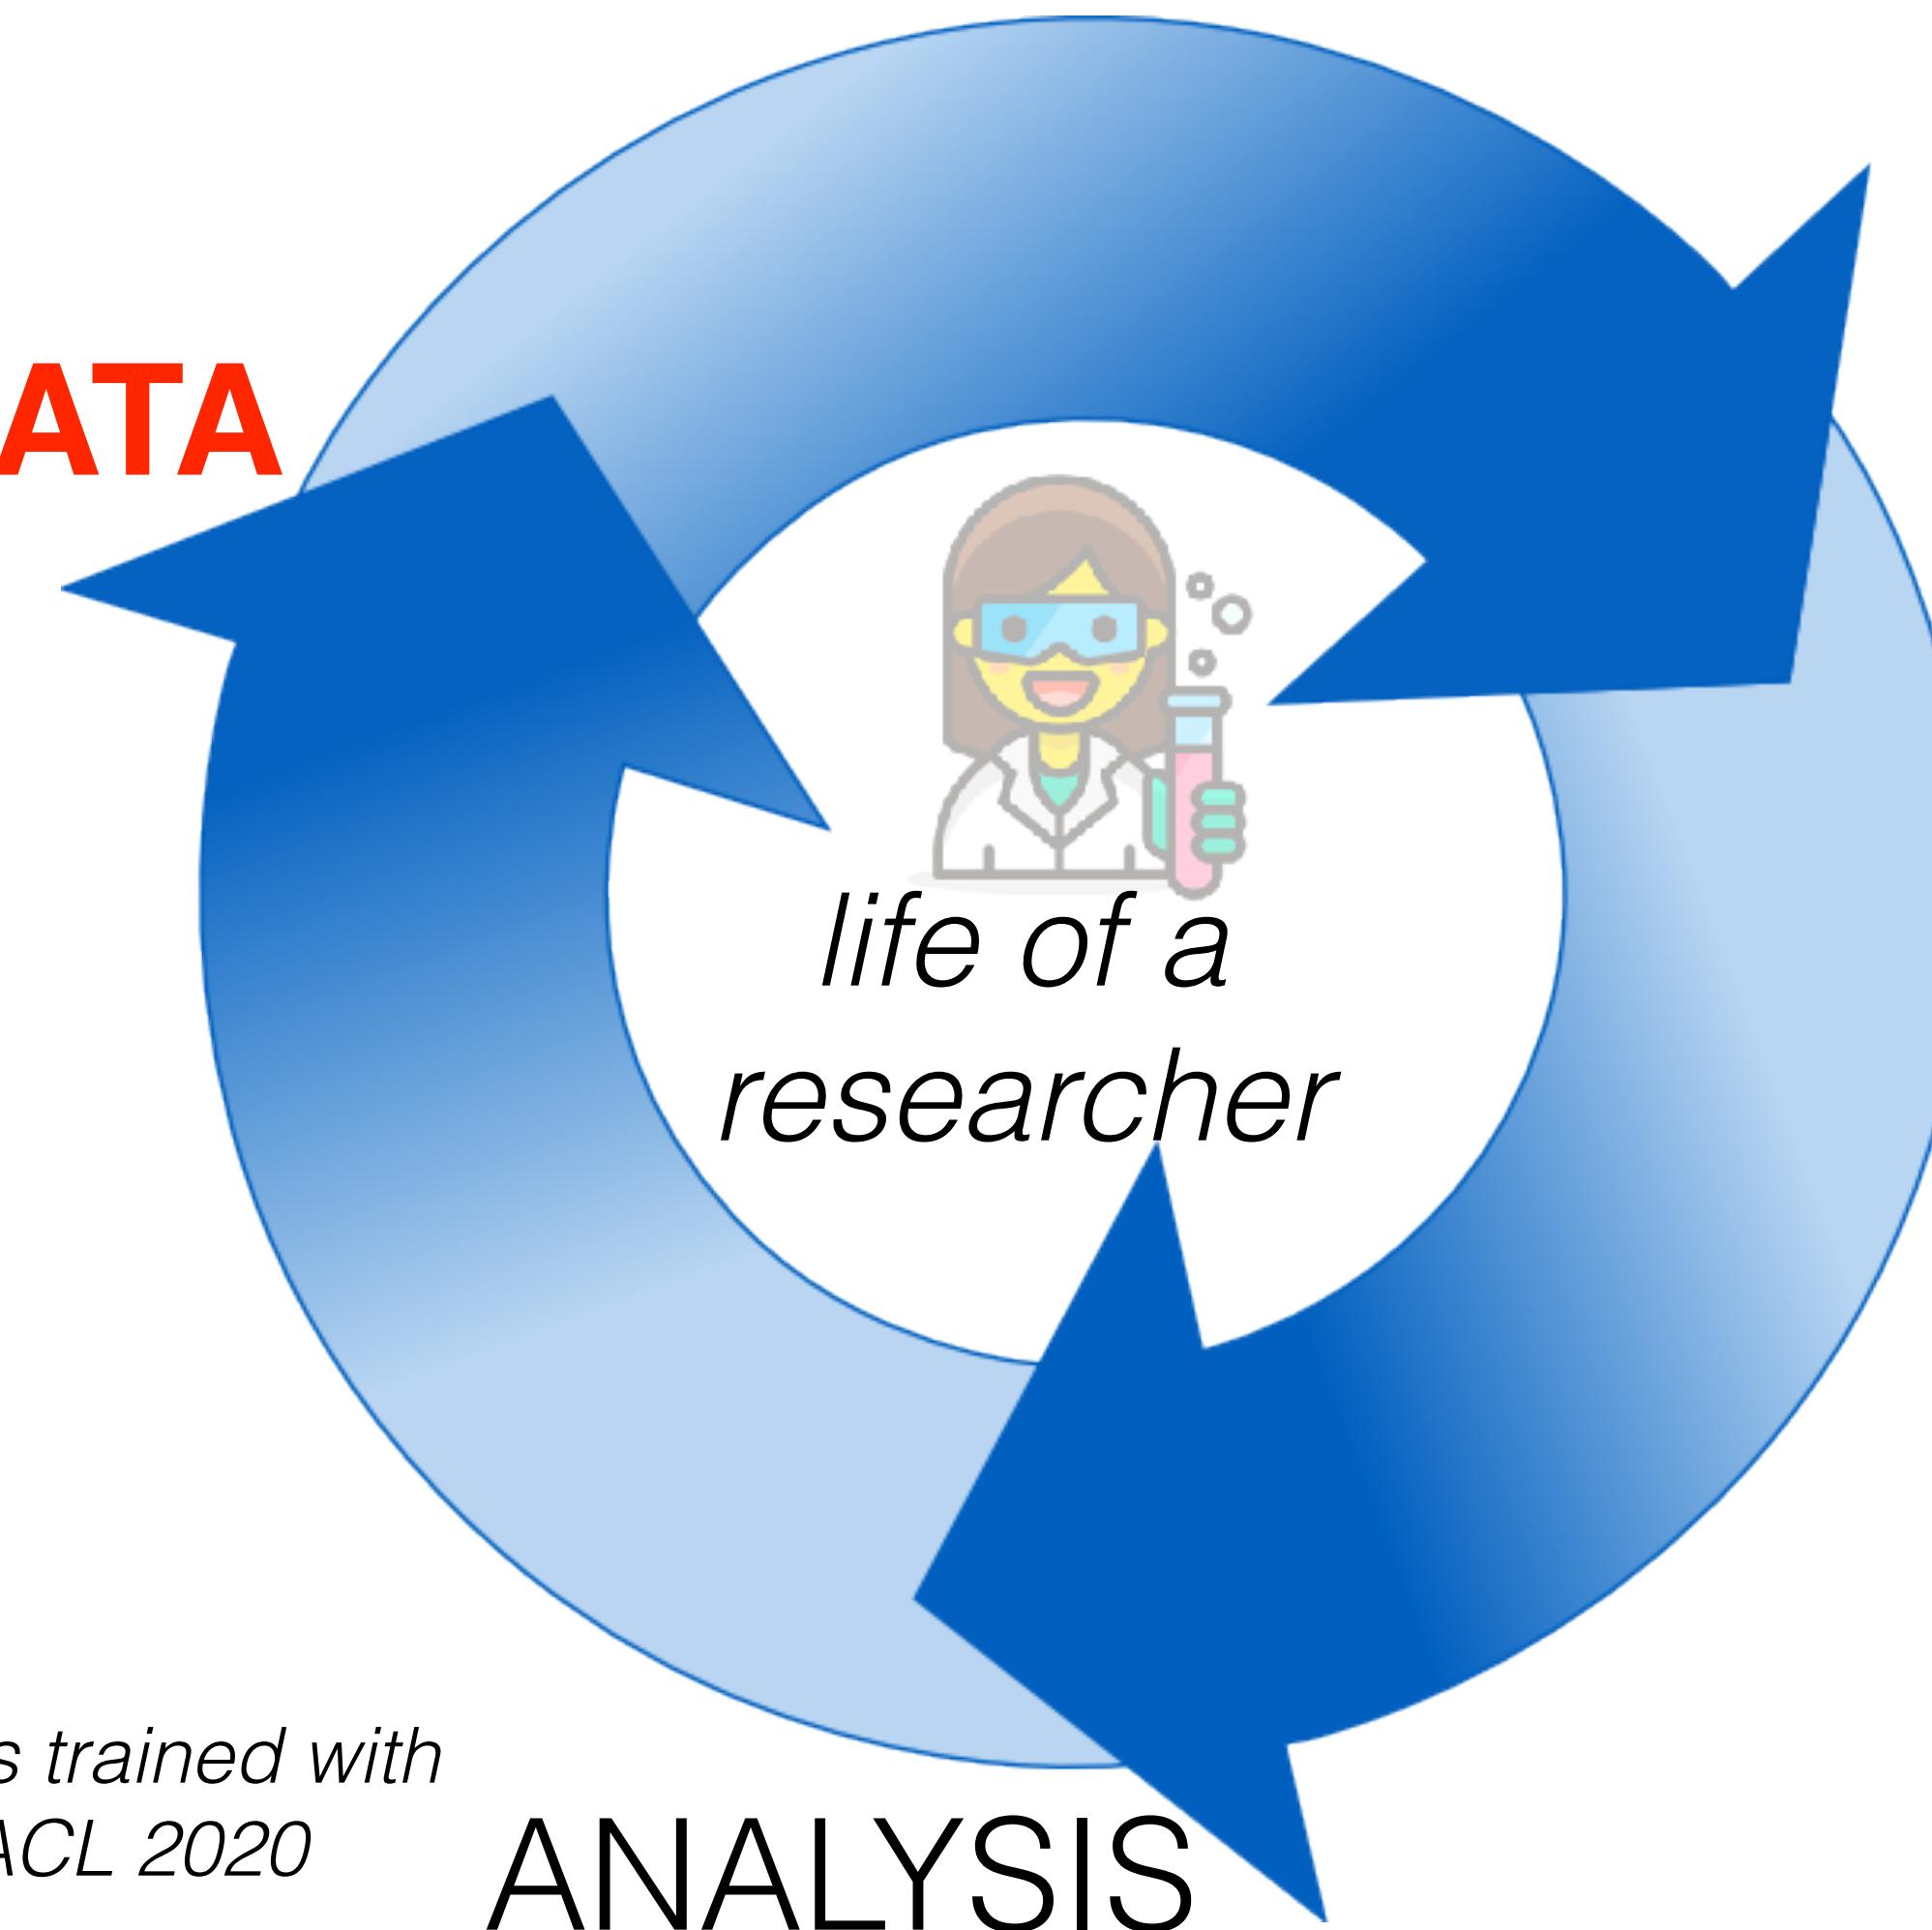
Challenges:

- data
 - sourcing data to train on
 - evaluation datasets
- modeling
 - unclear learning paradigm
 - domain adaptation
 - generalization

Why Low Resource MT Is Interesting?

- It is about learning with less labeled data.
- It is about modeling structured outputs and compositional learning.
- It is a real problem to solve.

Outline



"The FLoRes evaluation for low resource MT:..." Guzmán, Chen et al. 'EMNLP 2019'

"Analyzing uncertainty in NMT"
Ott et al. ICML 2018

"On the evaluation of MT systems trained with back-translation" Edunov et al. ACL 2020

"The source-target domain mismatch problem in MT" Shen et al. arXiv 1909.13151 2019

A Big “Small-Data” Challenge



... the open parallel corpus

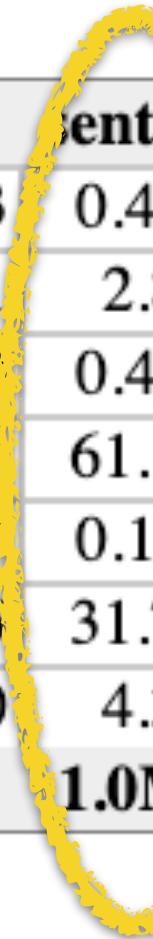
OPUS is a growing collection of translated texts from the web. In the OPUS project we try to convert and align free online data, to add linguistic annotation, and provide the community with a publicly available parallel corpus. OPUS is based on open source products and the corpus is also delivered as an open content package. We used several tools to compile the current collection. All pre-processing is done automatically. No manual corrections have been carried out.

The OPUS collection is growing! Check this page from time to time to see new data arriving ...
Contributions are very welcome! Please contact <jorg.tiedemann@helsinki.fi>

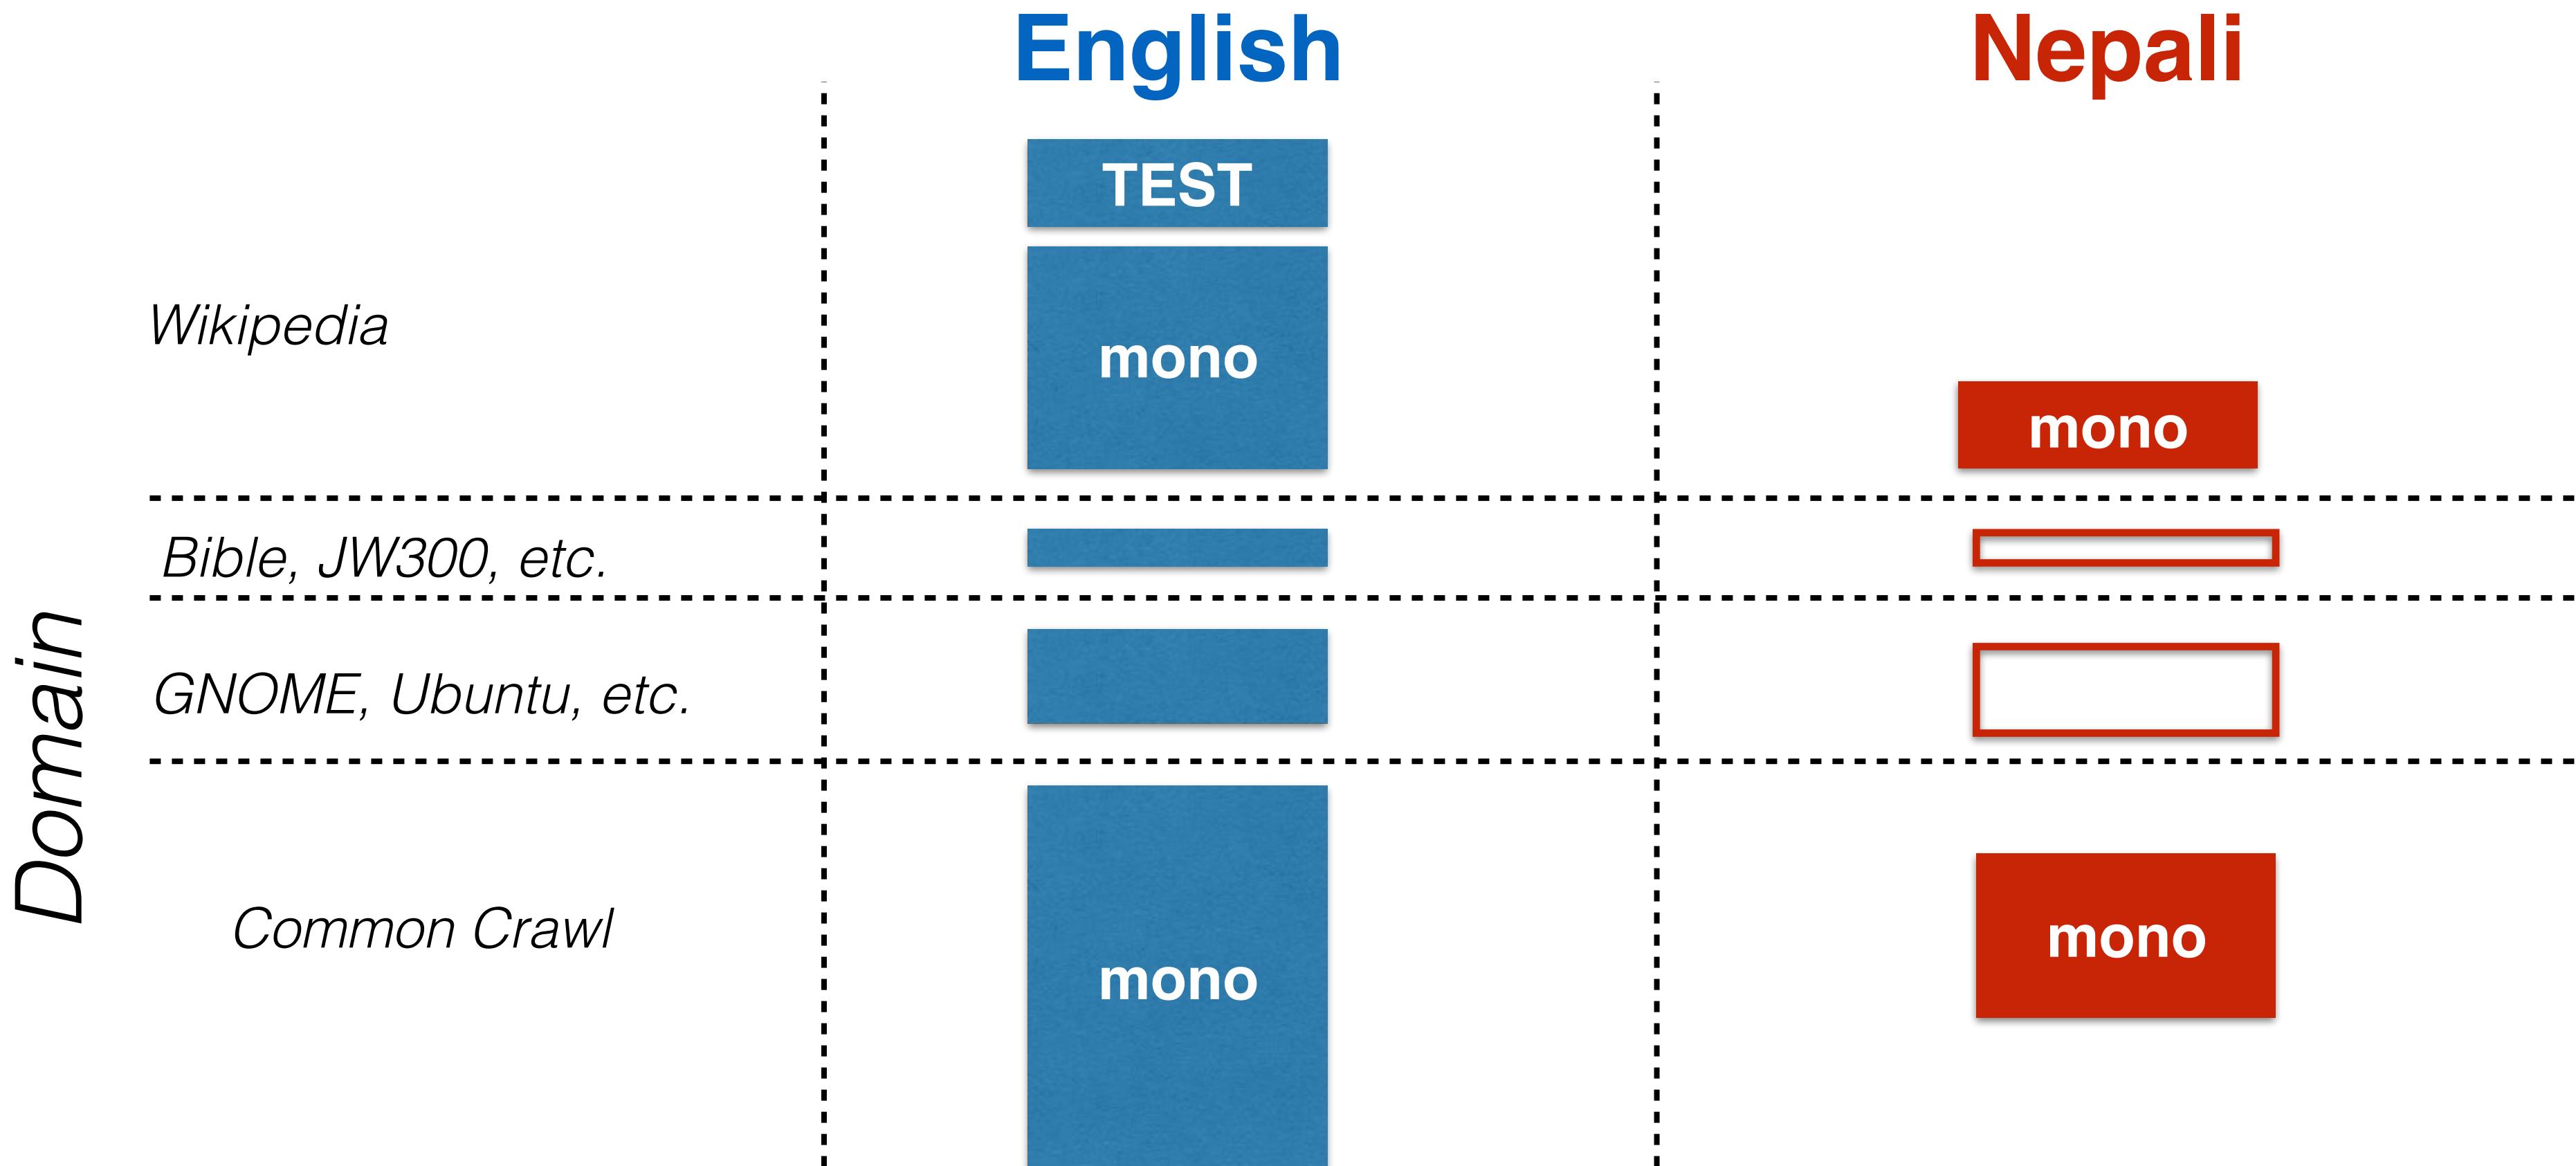
Search & download resources: en (English) ▾ ne (Nepali) ▾ all ▾ show all versions

Language resources: click on [tmx | moses | xces | lang-id] to download the data! (raw = untokenized, ud = parsed with universal dependencies, alg = word alignments and phrase tables)

corpus	doc's	sent's	en tokens	ne tokens	XCES/XML	raw	TMX	Moses	mono	raw	ud	alg	dic	freq	other files
JW300 v1	4663	0.4M	6.5M	5.5M	xces en ne	en ne			en ne	en ne			en ne		sample
wikimedia v20190628	1	2.8k	7.7M	1.2M	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt	dic	en ne		sample
GNOME v1	830	0.4M	1.8M	4.7M	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt		en ne		sample
bible-uedin v1	2	61.1k	1.8M	1.6M	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt	dic	en ne		sample
KDE4 v2	435	0.1M	0.6M	0.5M	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt	dic	en ne	query	sample
Ubuntu v14.10	155	31.7k	0.3M	0.2M	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt	dic	en ne		sample
QED v2.0a	60	4.3k	69.8k	41.7k	xces en ne	en ne	tmx	moses	en ne	en ne	alg smt	dic	en ne		sample
<i>total</i>	6146	1.0M	18.8M	13.7M	1.0M		0.6M	0.6M							



Case Study: En-Ne



In-domain data: no parallel, little monolingual.

Out-of-domain: little parallel, quite a bit monolingual

No translation originating from Nepali.

A Case Study: En-Ne

- Parallel Training data: versions of bible and ubuntu handbook (<1M sentences).
- Nepali Monolingual data: wikipedia (90K), common crawl (few millions).
- English Monolingual data: unlimited almost.
- Test data: ???

FLoRes Evaluation Benchmark

- Validation, test and hidden test set, each with 3000 sentences in English-Nepali and English-Sinhala.
- Sentences taken from *Wikipedia* documents.

Data Collection Process:

- Very expensive and slow.
- Very hard to produce high-quality translations:
 - automatic checks (language model filtering, transliteration filtering, length filtering, language id filtering, etc),
 - human assessment.

Examples

Si-En

අධි යාපනයෙන් පසු හෝ පවුලේ යුතුකම් ඉටු කරන්නට හෝ රෝග තත්ත්වයන් තිසා සිංහල උපසම්පූර්ණ නිතරම ඉවත් වෙති.

After education priests leave ordination in order to fulfill duties to the family or due to sickness.

තරතත , ගාරීරක හිංසනය , දේපල භාතිය , පහර දීම සහ මරාදැමීම මෙම දියුවම්ය .

Threatening, physical violence, property damage, assault and execution are these punishments.

En-Si

In Serious meets, the absolute score is somewhat meaningless.

සැබෑ තරග වලදී ලක්ණු සැසදීම තේරුමක් තැනි ක්රියාවකි .

Iphone users can and do access the internet frequently, and in a variety of places.

අයිලෝත් භාවිත කරන්නන්ට නිතරම සහ විවිධ ස්ථානවලදී අත්තරජාලයට පිවිසිය හැකිය .



Wikipedia originating in Si has different topics than Wikipedia originating in En

Examples

Ne-En

पुरानो समयमा राजालाई सल्लाह दिने समा 'संसद' कहलाउँथ्यो ।

In the past, the assembly that advised the king were called 'parliament'.

कार्यकर्ताका रूपमा अफ्रिकन नेशनल कंग्रेसमा आबद्ध भए ।

As a worker African Mandela joined the Congress party.

En-Ne

The academic research tended toward the improvement of basic technologies, rather than their specific applications.

शैक्षिक अनुसन्धानले उनीहरूको विशिष्ट अनुप्रयोगहरूको सट्टा आधारभूत प्रविधिको सुधारको पक्षमा जोड दिए ।

It has automatic spell checking and correction, predictive word capabilities, and a dynamic dictionary that learns new words.

यसमा स्वचालित हिज्जे जाँच र सुधार छ , भविष्यवाणी शब्द क्षमताहरू , र गतिशील शब्दकोश हुन्छ जसले नयाँ शब्दहरू सिक्छ ।



- Useful to evaluate truly low resource language pairs.
 - WMT 2019 and WMT 2020 shared filtering task.
 - Several publications.
- Sustained effort, more to come...

FLoRes Low Resource MT Benchmark

This repository contains data and baselines from the paper:

[The FLoRes Evaluation Datasets for Low-Resource Machine Translation: Nepali-English and Sinhala-English](#).

The data can be downloaded directly at:

https://github.com/facebookresearch/flores/raw/master/data/wikipedia_en_ne_si_test_sets.tgz

Baselines

The following instructions will can be used to reproduce the baseline results from the paper.

Requirements

The baseline uses the [Indic NLP Library](#) and [sentencepiece](#) for preprocessing; [fairseq](#) for model training; and [sacrebleu](#) for scoring.

Dependencies can be installed via pip:

```
$ pip install fairseq sacrebleu sentencepiece
```

The Indic NLP Library will be cloned automatically by the `prepare-{ne,si}en.sh` scripts.

Download and preprocess data

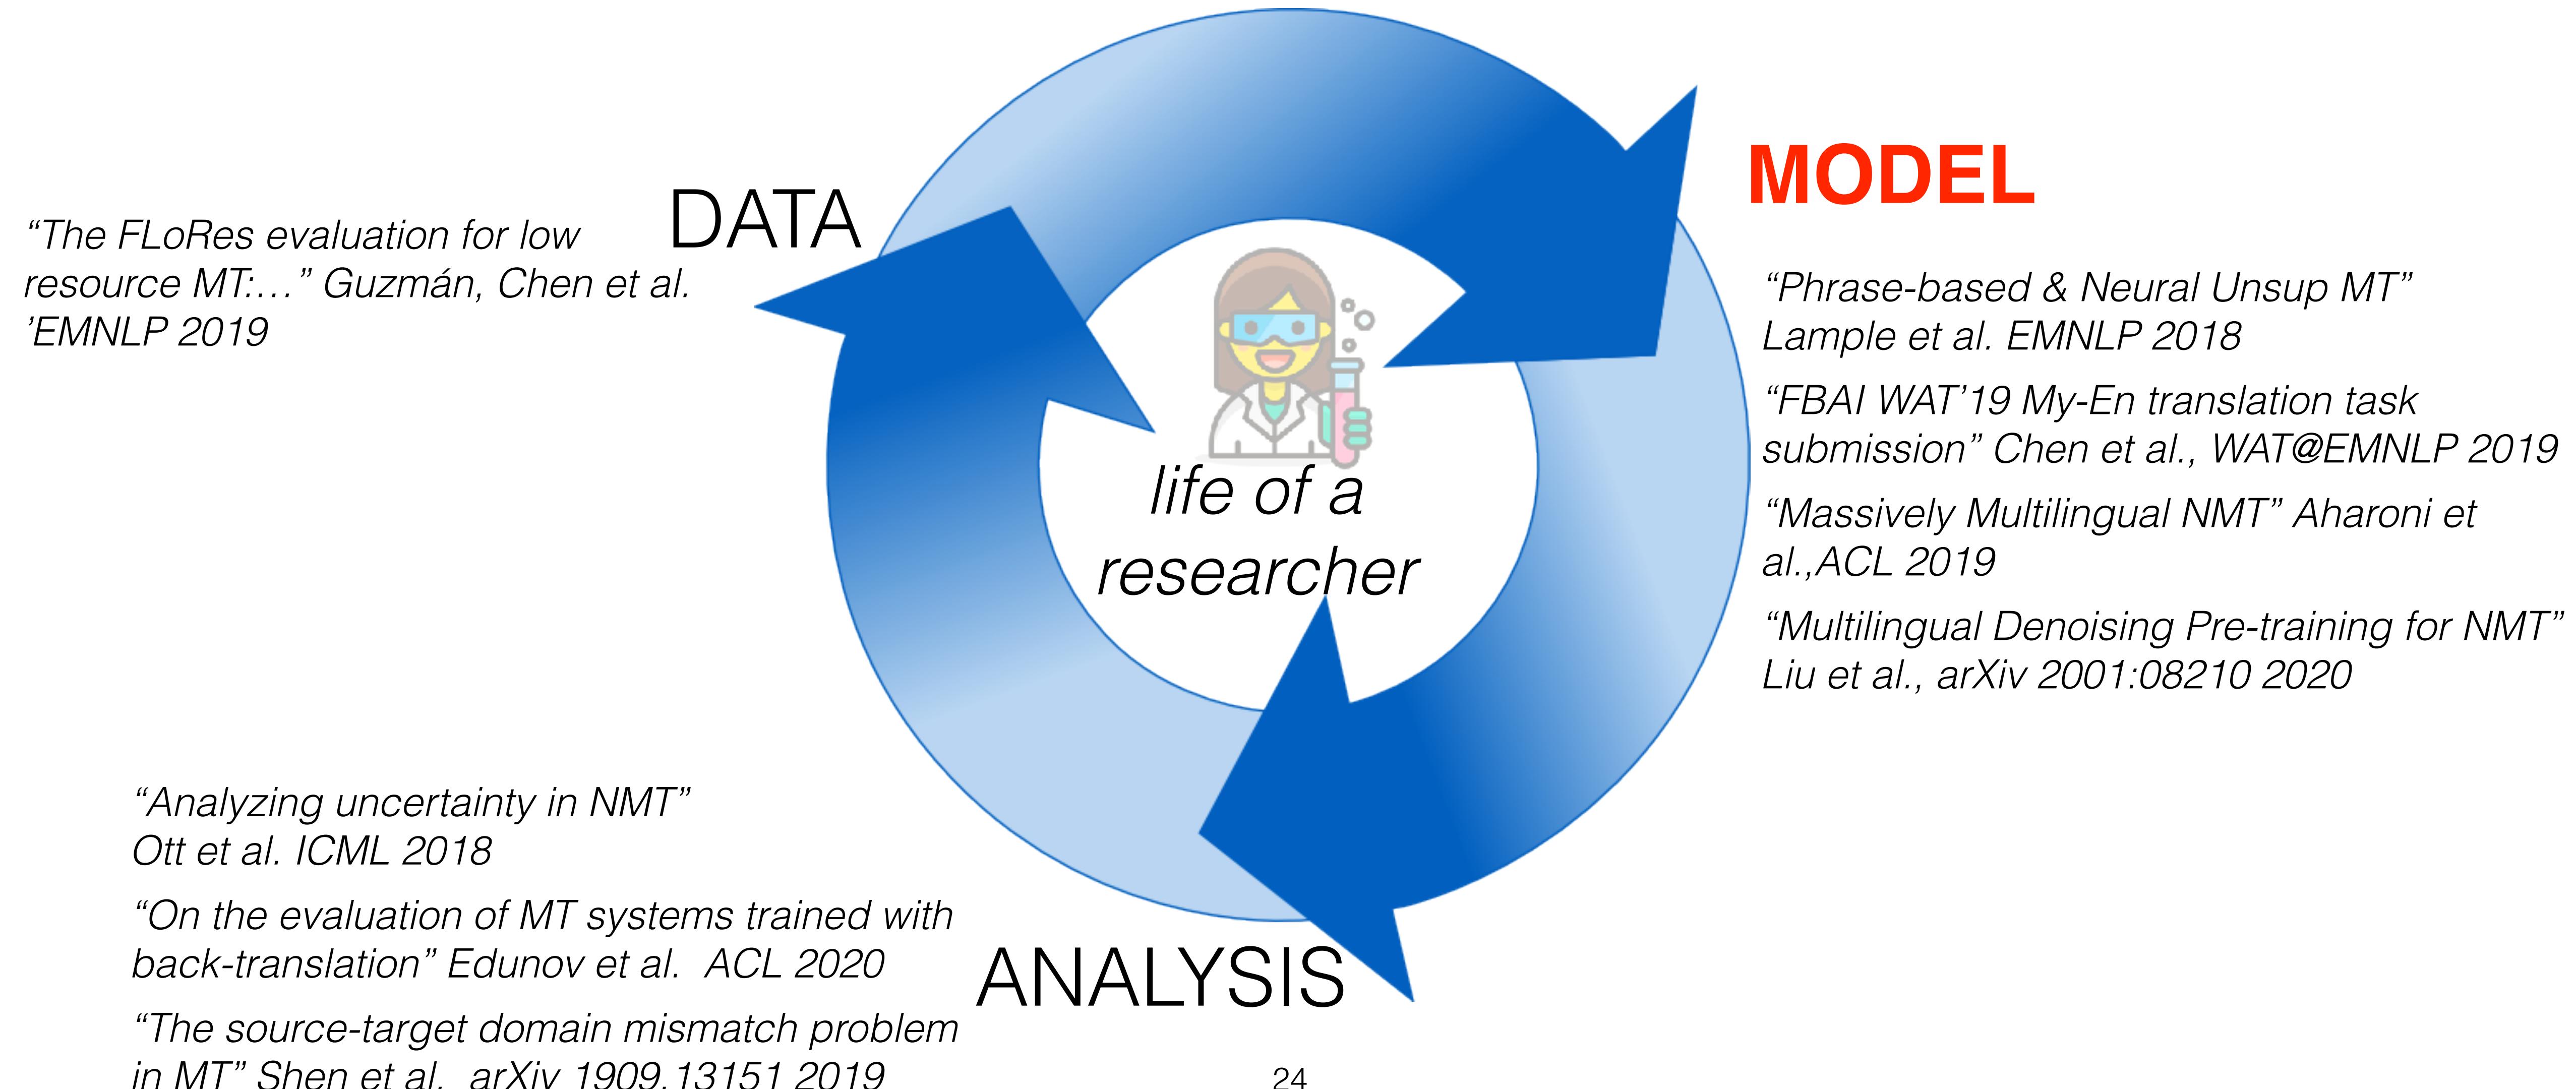
<https://github.com/facebookresearch/flores>

data & baseline models

What Did We Learn?

- Data is often as or more important than designing a model.
- Collecting data is not trivial.
- Look at the data!!

Outline



English

Nepali

Hindi

Sinhala

Bengali

Spanish

Tamil

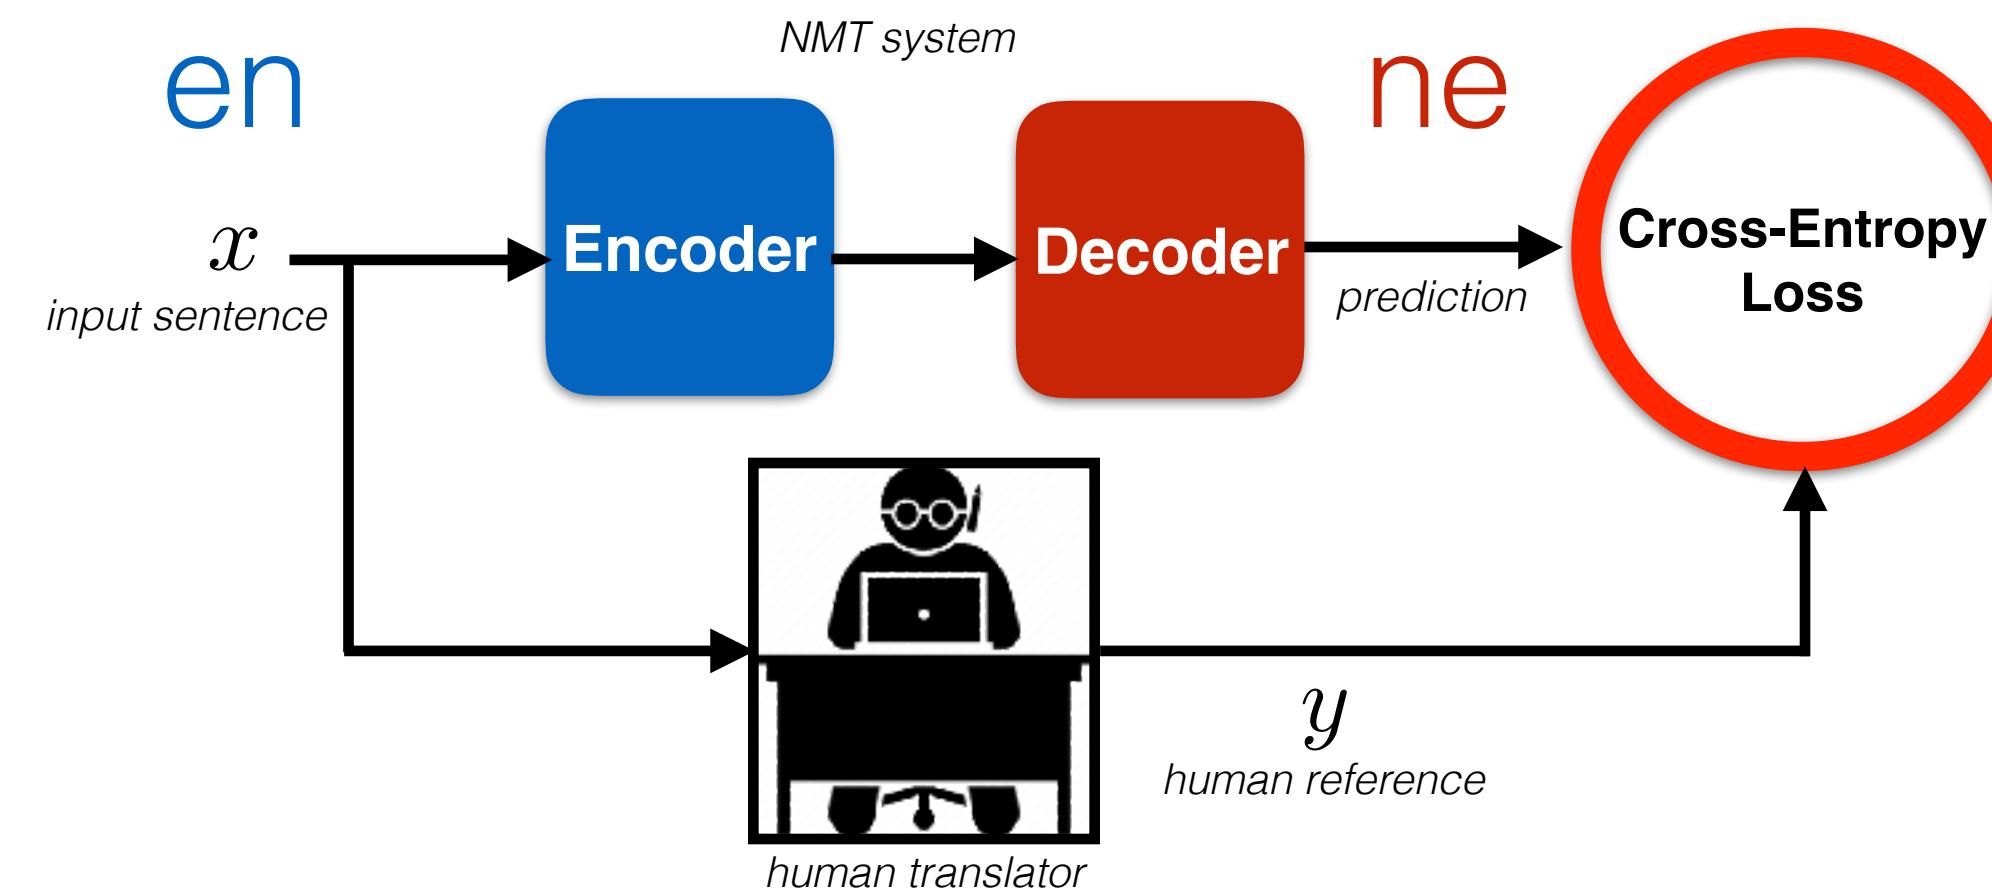
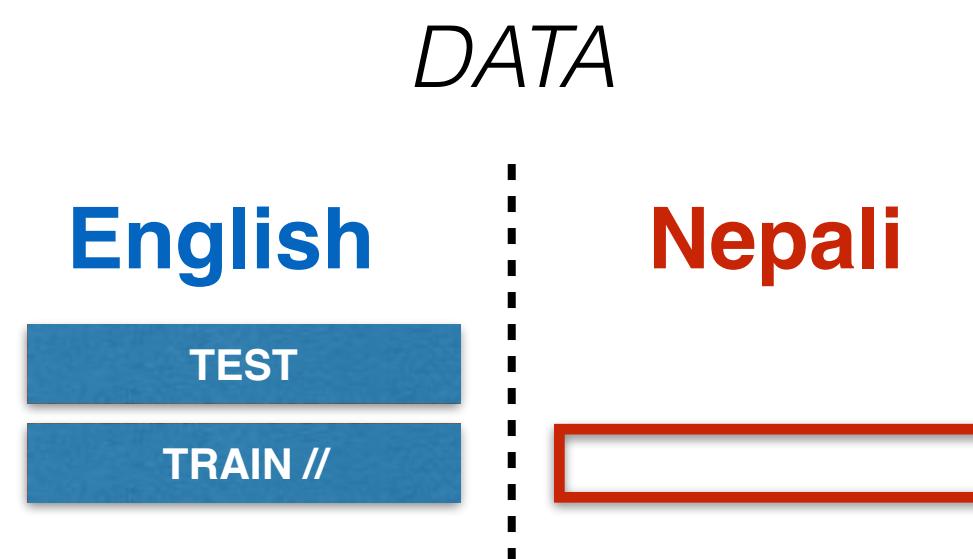
Gujarati

Domain

TEST



ML Perspective: Supervised Learning



Training Dataset

$$\mathcal{D} = \{(x, y)_i\}_{i=1,\dots,N}$$

If N is small, how can we further regularize the model?

- dropout [1]
- label smoothing [2]

Learning Framework: Supervised Learning.

Per-sample loss:

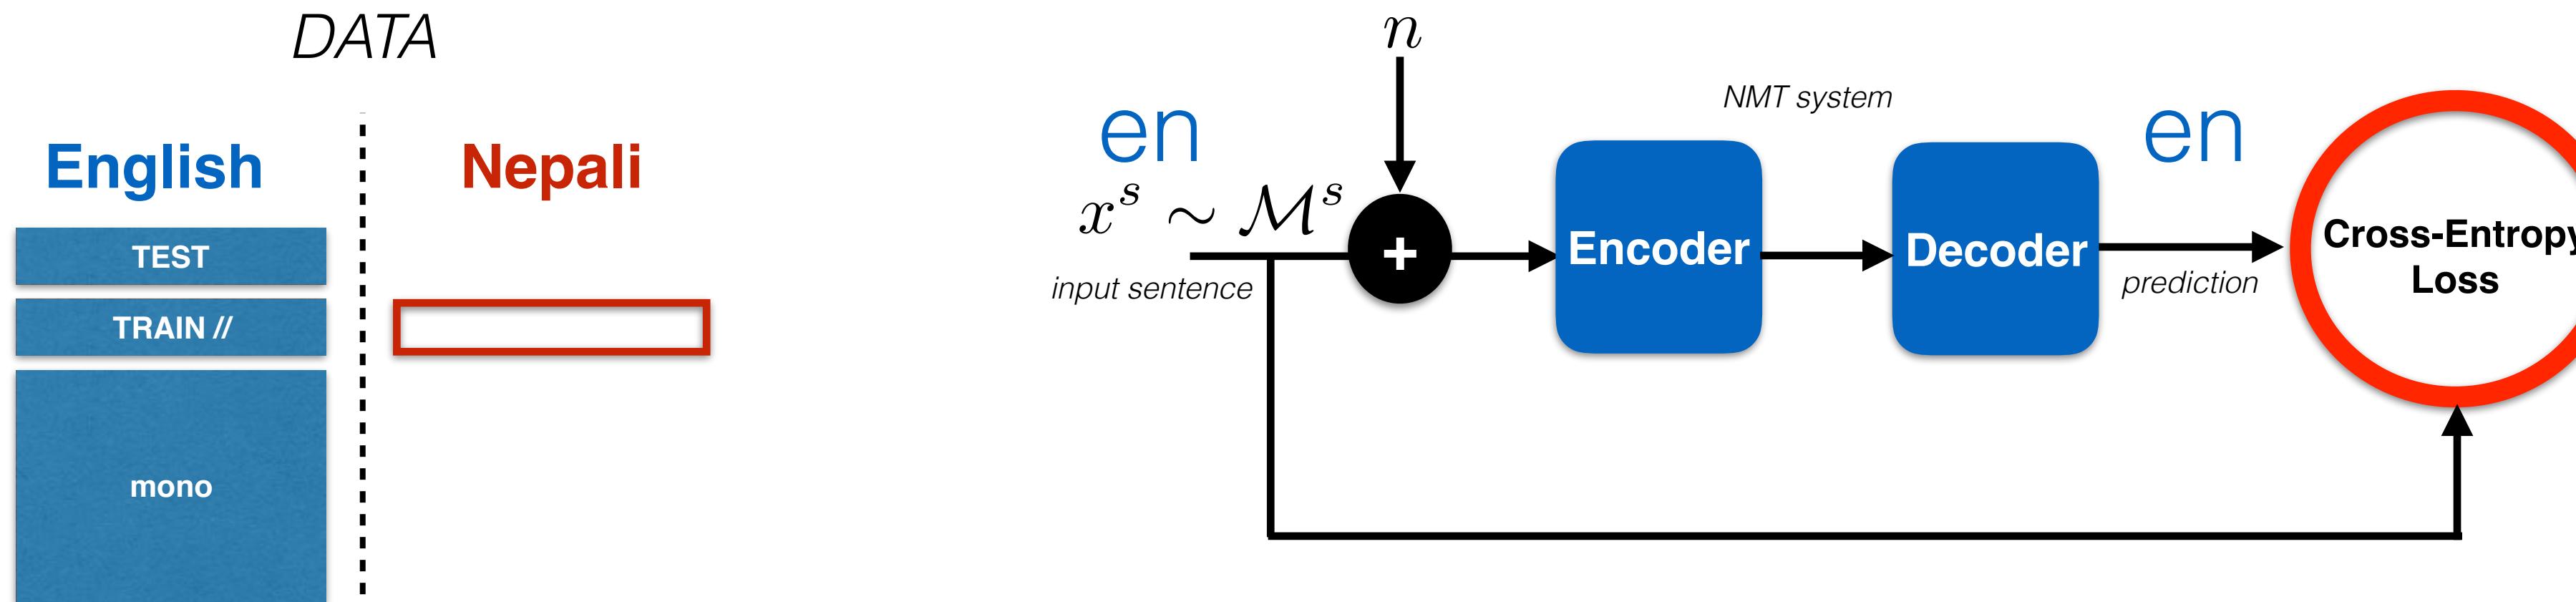
$$\mathcal{L}(\theta) = -\log p(y|x)$$

usual attention-based transformer

[1] Srivastava et al. "Dropout: a simple way to prevent neural networks from overfitting" JMLR 2014

[2] Szegedy et al. "Rethinking the inception architecture for computer vision" CVPR 2016

ML Perspective: Semi-Supervised Learning



Training Dataset

$$\mathcal{D} = \{(x, y)_i\}_{i=1,\dots,N}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,\dots,M_s}$$

Adding source-side monolingual data.

Idea: model $p(x)$.

Noise: word drop, swap, etc.

E.g.: *The cat the on sat mat.*

The cat sat on the.

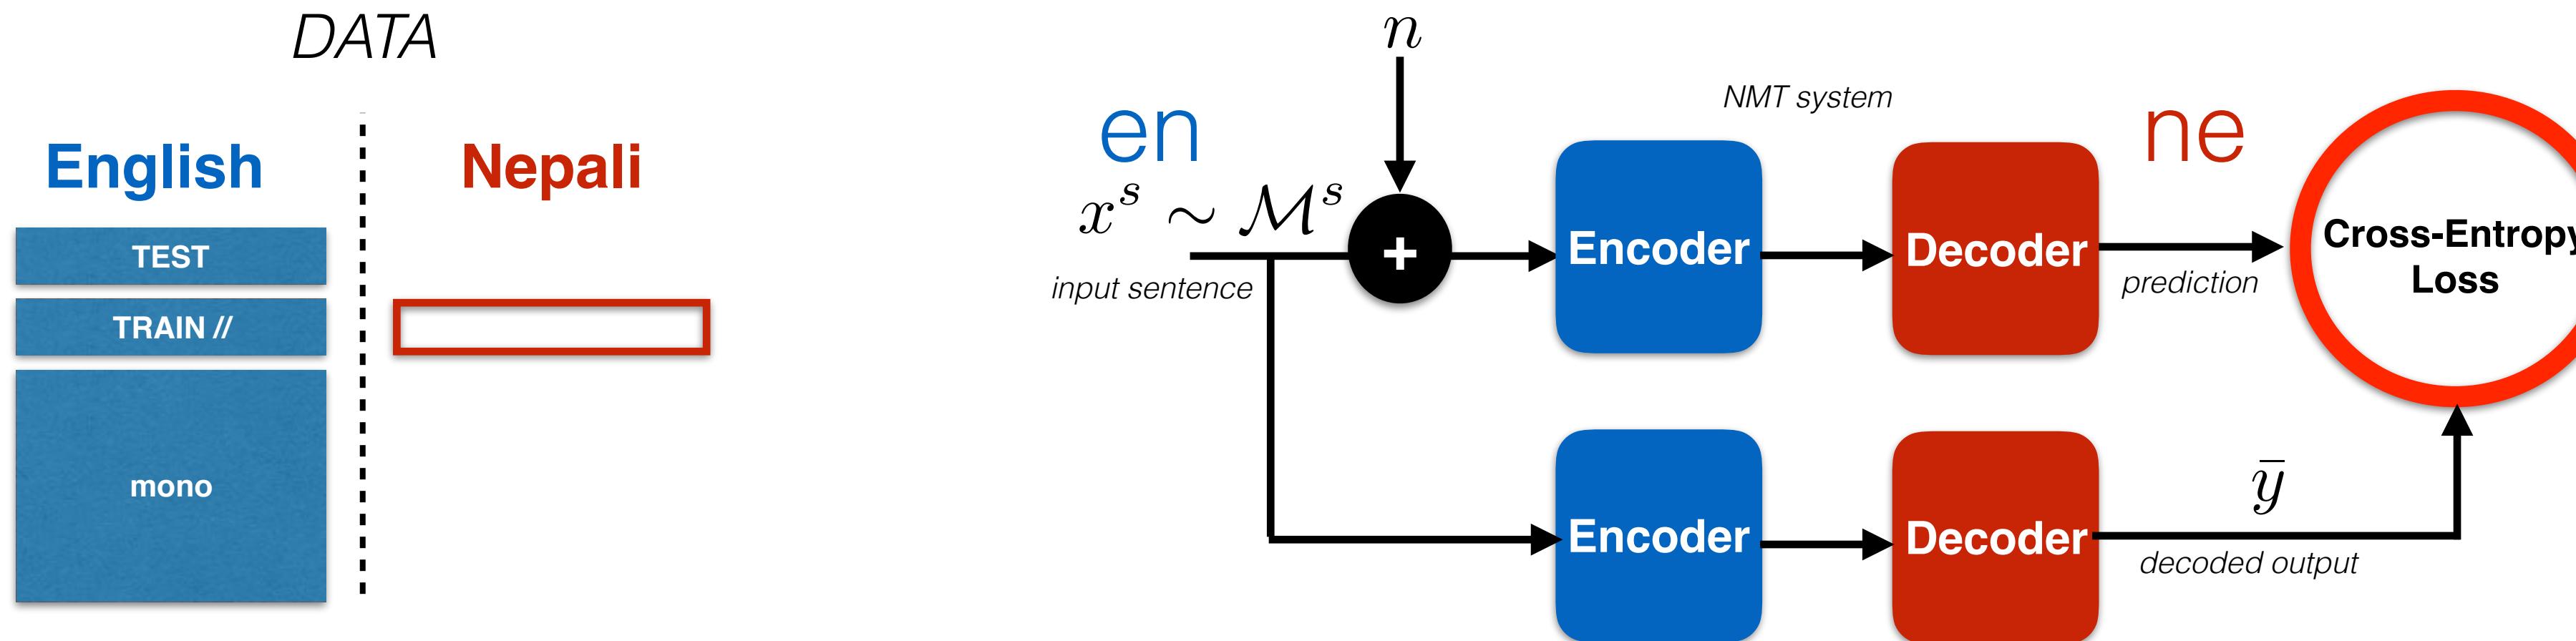


Learning Framework: DAE
Either pre-train or add a DAE loss to the supervised cross-entropy term.

$$\mathcal{L}^{DAE}(\theta) = -\log p(x|x + n)$$

Vincent et al. "Stacked denoising auto-encoders..." JMLR 2010
Liu et al. "Multilingual denoising pretraining for NMT" arXiv:2001.08210 2020

ML Perspective: Semi-Supervised Learning



Training Dataset

$$\mathcal{D} = \{(x, y)_i\}_{i=1,\dots,N}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,\dots,M_s}$$

Adding source-side monolingual data.

An alternative approach to DAE.

$$\mathcal{L}^{ST}(\theta) = -\log p(\bar{y}|x + n)$$

$$\mathcal{L}(\theta) = \mathcal{L}^{\text{sup}}(\theta) + \lambda \mathcal{L}^{ST}(\theta)$$

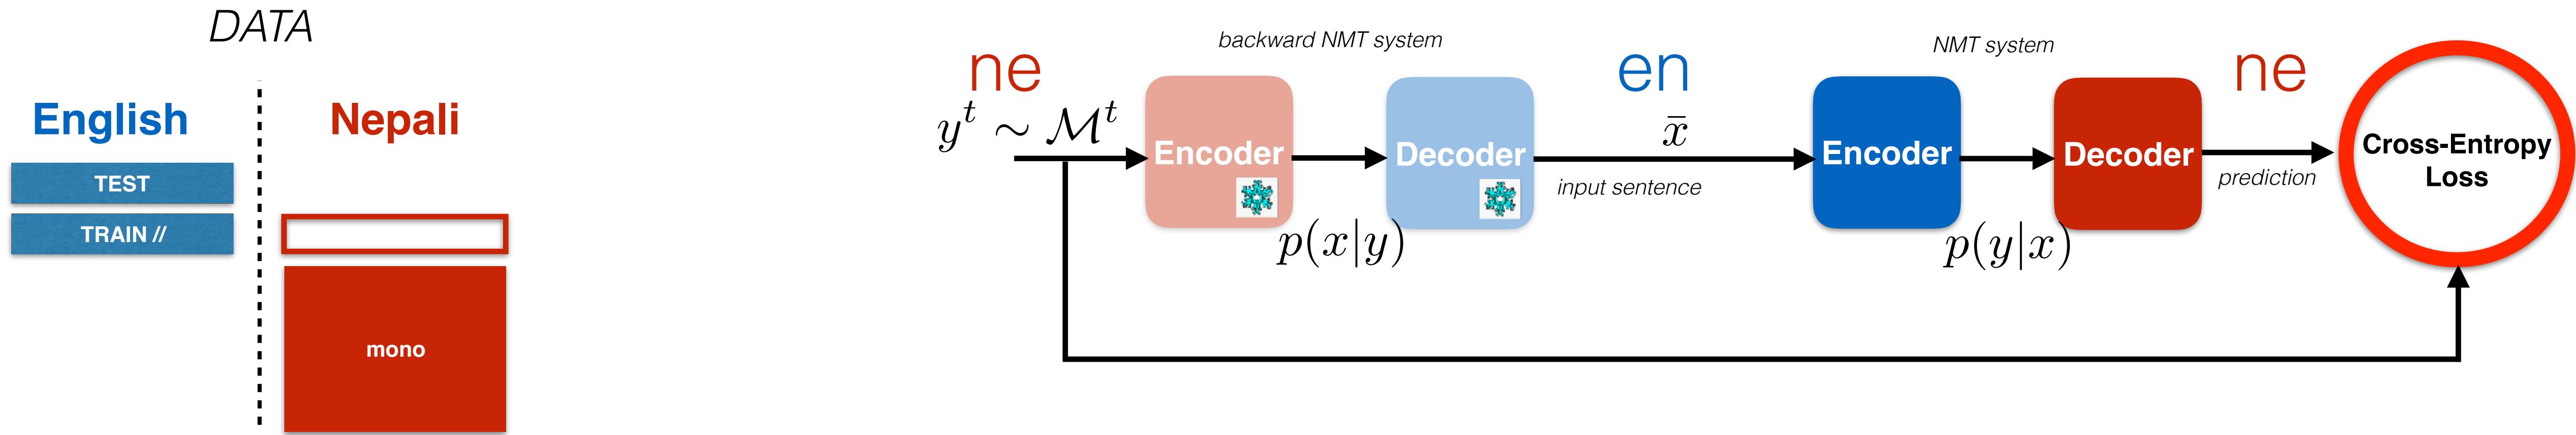
Learning Framework: Self-Training (ST).

ALGORITHM

- train model $p(y|x)$ on \mathcal{D}
- repeat
 - decode $x^s \sim \mathcal{M}^s$ to \bar{y} and create additional dataset $\mathcal{A}^s = \{(x_j^s, \bar{y}_j)\}_{j=1,\dots,M_s}$
 - retrain model on: $\mathcal{D} \cup \mathcal{A}^s$

Key elements: decoding and training noise.

ML Perspective: Semi-Supervised Learning



Training Dataset

$$\mathcal{D} = \{(x, y)_i\}_{i=1,\dots,N}$$

$$\mathcal{M}^t = \{y_k^t\}_{k=1,\dots,M_t}$$

Adding target-side monolingual data.

Two benefits:

- Decoder learns a good language model.
- Better generalization via data augmentation.
- Unlike ST, target is correct but input is not.

$$\mathcal{L}^{BT}(\theta) = -\log p(y|\bar{x})$$

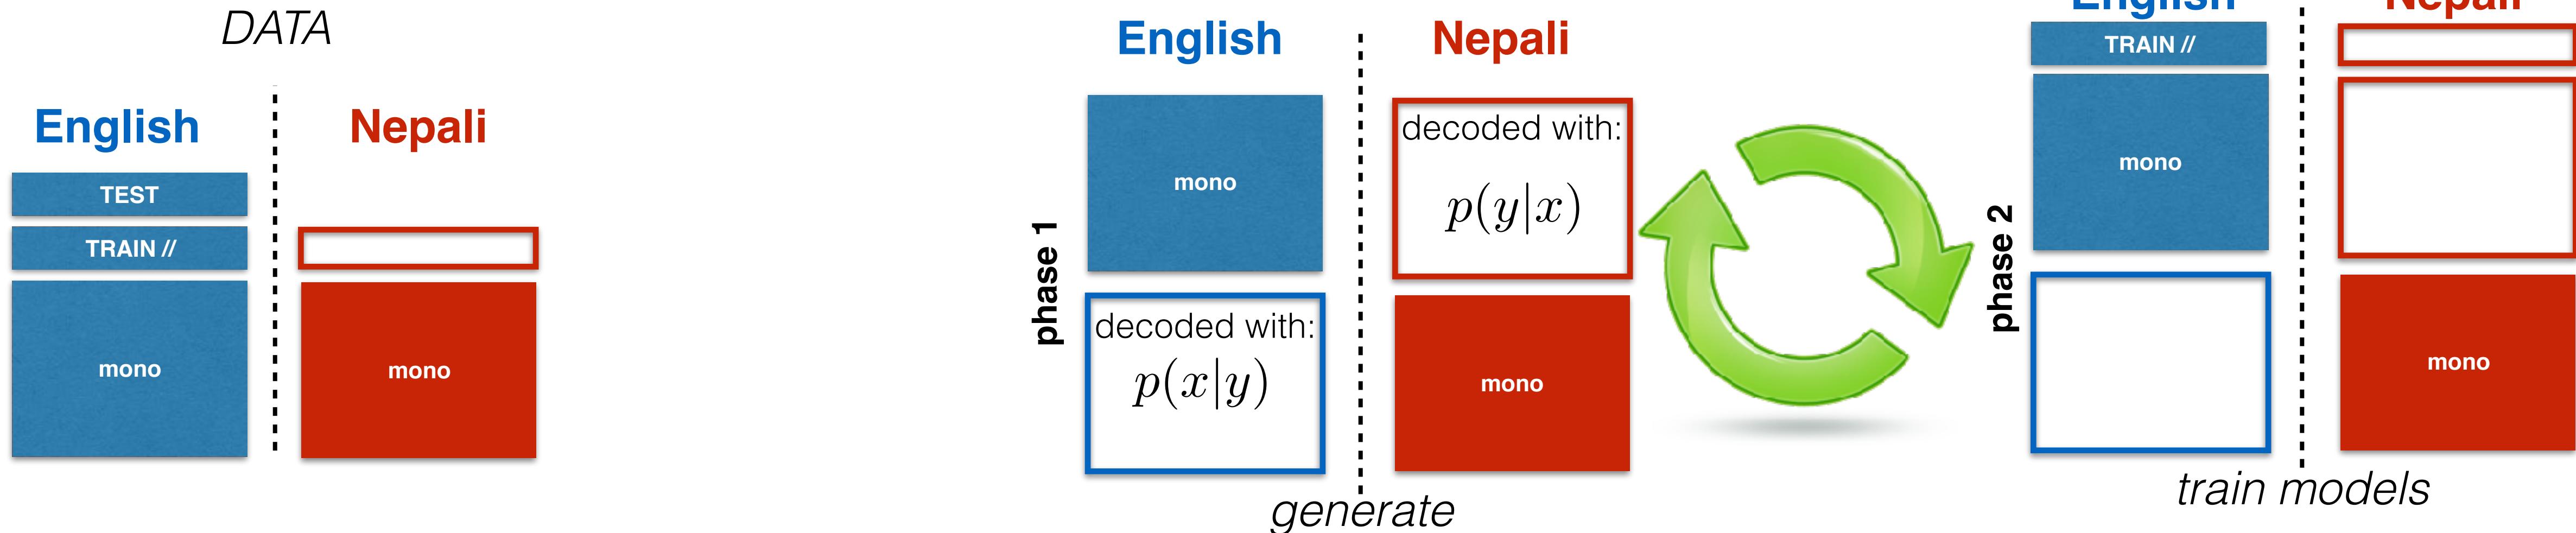
$$\mathcal{L}(\theta) = \mathcal{L}^{sup}(\theta) + \lambda \mathcal{L}^{BT}(\theta)$$

Learning Framework: Back-Translation (BT).

ALGORITHM

- train model $p(x|y)$ and $p(y|x)$ on \mathcal{D}
- decode $y^t \sim \mathcal{M}^t$ to \bar{x} with $p(x|y)$, create additional dataset $\mathcal{A}^t = \{(\bar{x}_k, y_k^t)\}_{k=1,\dots,M_t}$
- retrain model $p(y|x)$ on: $\mathcal{D} \cup \mathcal{A}^t$

ML Perspective: Semi-Supervised Learning



Training Dataset

$$\mathcal{D} = \{(x, y)_i\}_{i=1,\dots,N}$$

$$\mathcal{M}^t = \{y_k^t\}_{k=1,\dots,M_t}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,\dots,M_s}$$

Adding both source & target-side monolingual data.

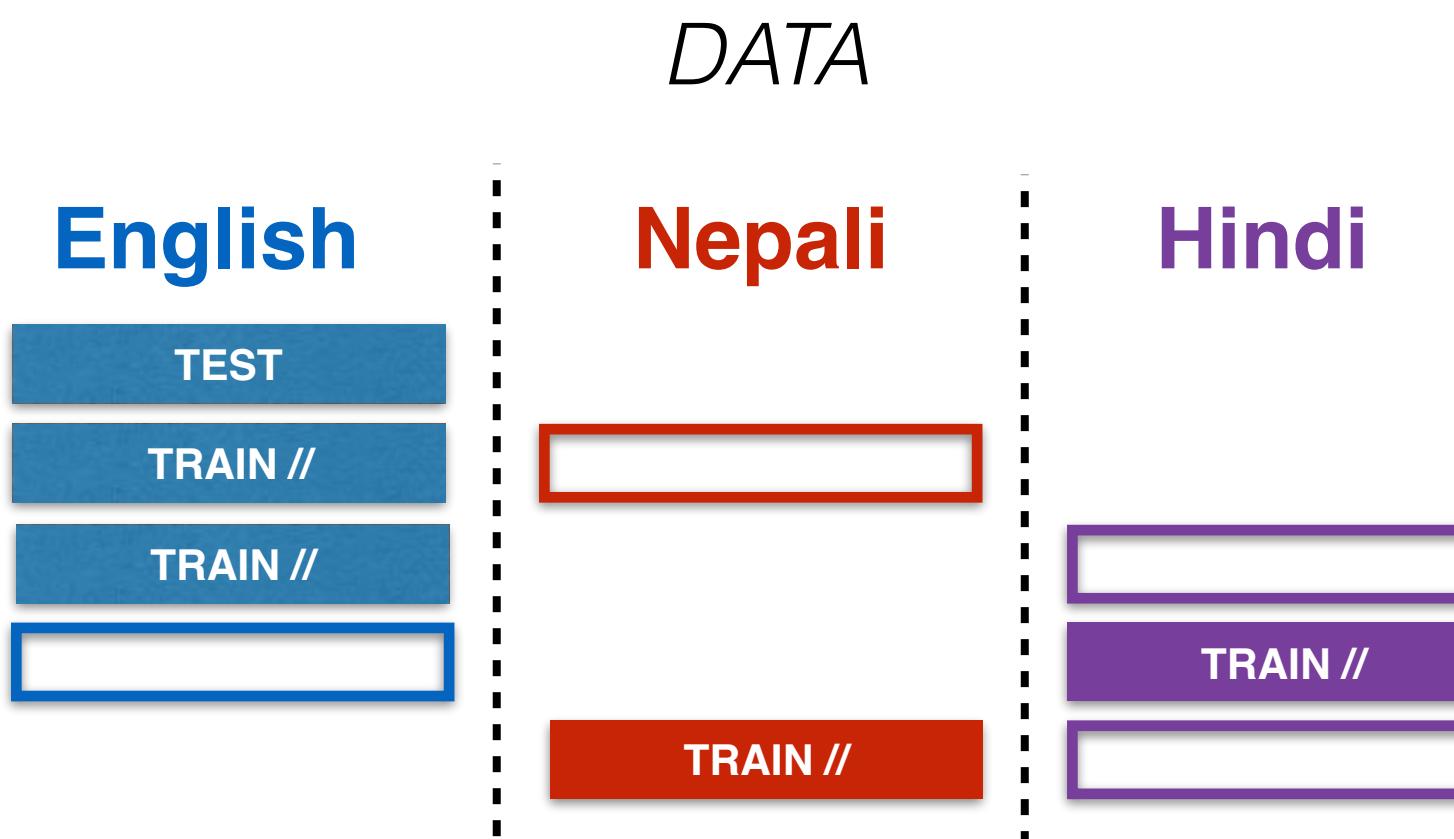
$$\mathcal{L}^{\text{total}}(\theta) = -\log p(y|x) - \lambda_1 \log p(y^t|\bar{x}^t) - \lambda_2 \log p(\bar{y}^s|x^s)$$

Learning Framework: Iterative ST & BT.

ALGORITHM

- train model $p(x|y)$ and $p(y|x)$ on \mathcal{D}
- repeat
 - decode $y^t \sim \mathcal{M}^t$ to \bar{x} with $p(x|y)$, create additional dataset $\mathcal{A}^t = \{(\bar{x}_k, y_k^t)\}_{k=1,\dots,M_t}$
 - decode $x^s \sim \mathcal{M}^s$ to \bar{y} with $p(y|x)$, create additional dataset $\mathcal{A}^s = \{(x_j^s, \bar{y}_j)\}_{j=1,\dots,M_s}$
 - retrain both $p(y|x)$ and $p(x|y)$ on: $\mathcal{D} \cup \mathcal{A}^t \cup \mathcal{A}^s$

ML Perspective: Multi-Task/Multi-Modal

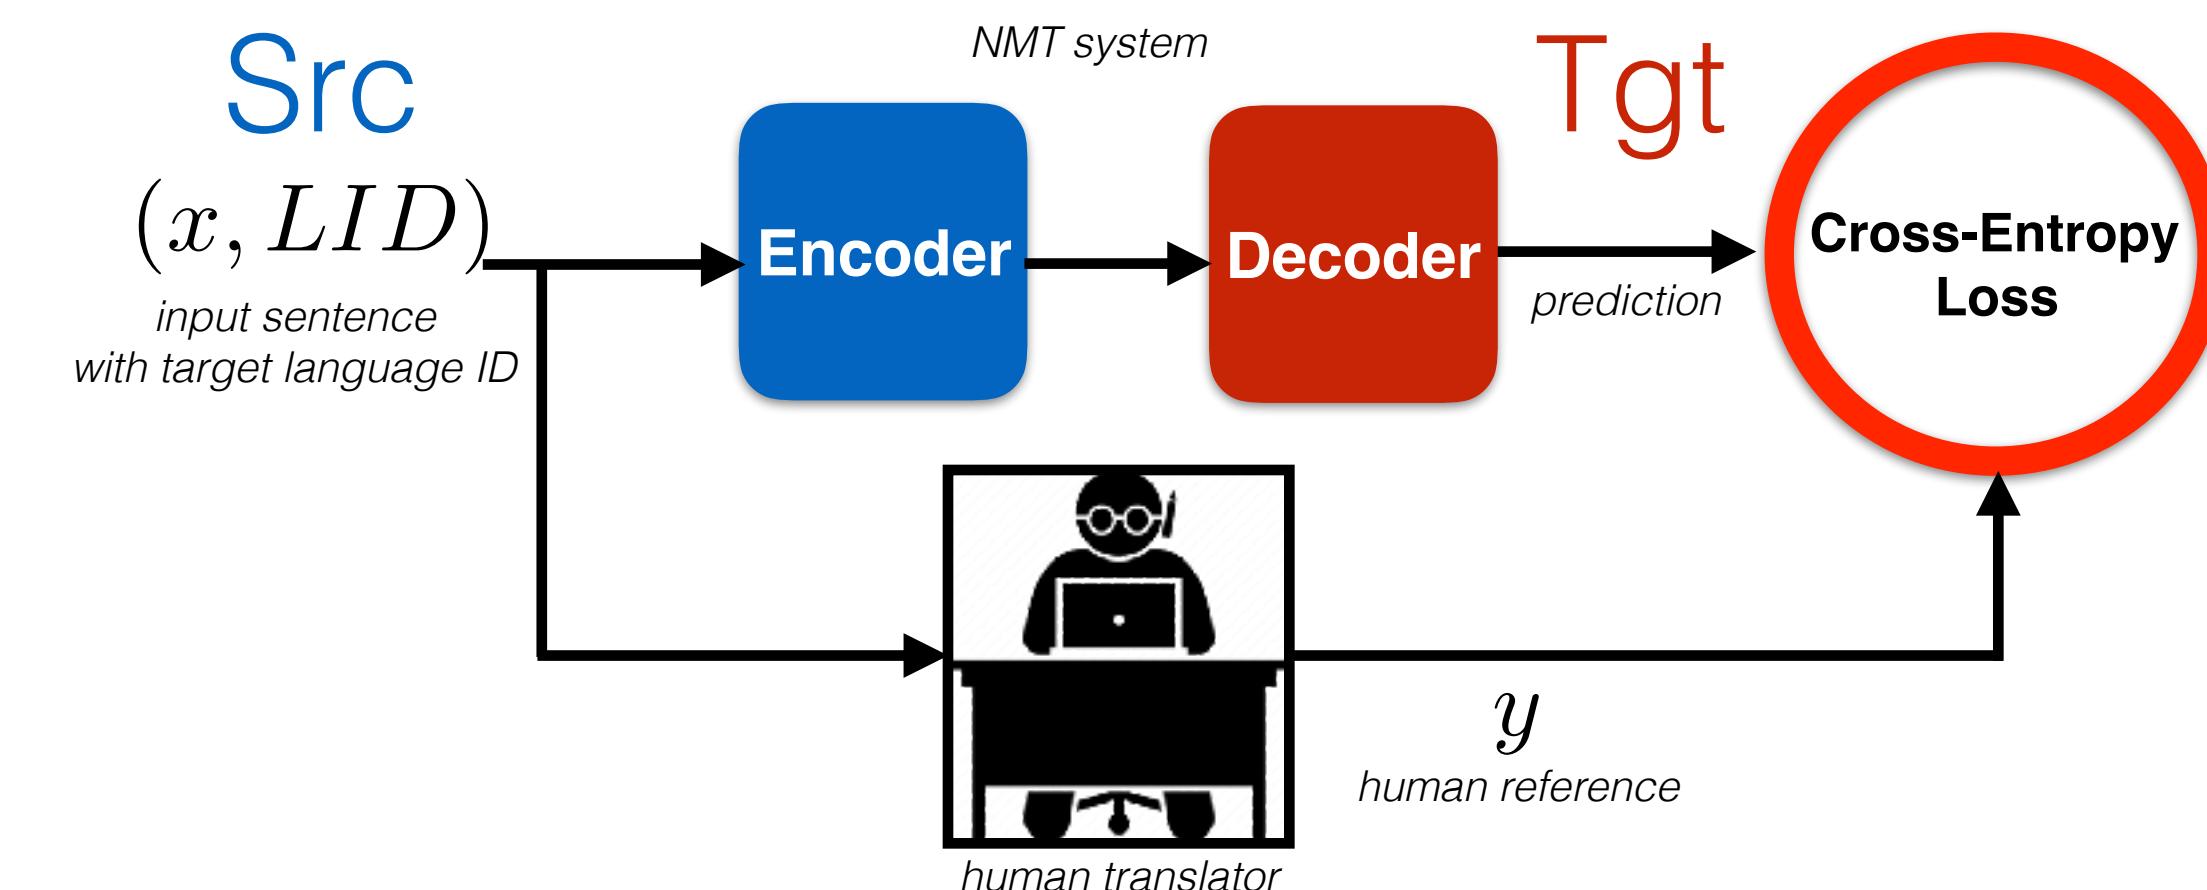


Training Dataset

$$\mathcal{D}_{\text{en,ne}} \cup \mathcal{D}_{\text{en,hi}} \cup \mathcal{D}_{\text{hi,en}} \cup \mathcal{D}_{\text{ne,hi}}$$

Adding parallel data in other languages.

$$\mathcal{L}(\theta) = - \sum_{s,t} \mathbb{E}_{(x,y) \sim \mathcal{D}_{s,t}} [\log p(y|x; \theta)]$$



Learning Framework: Multilingual Training

Share the same encoder and the same decoder with all the language pairs.

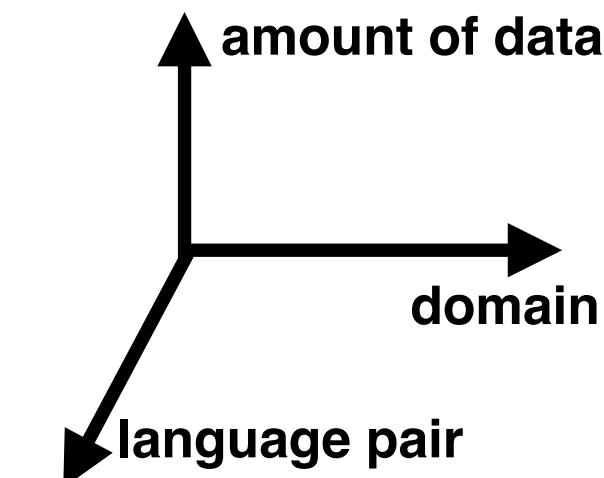
Prepend a target language identifier to the source sentence to inform decoder of desired language.

Concatenate all the datasets together.

Train using standard cross-entropy loss.

Conclusion so far...

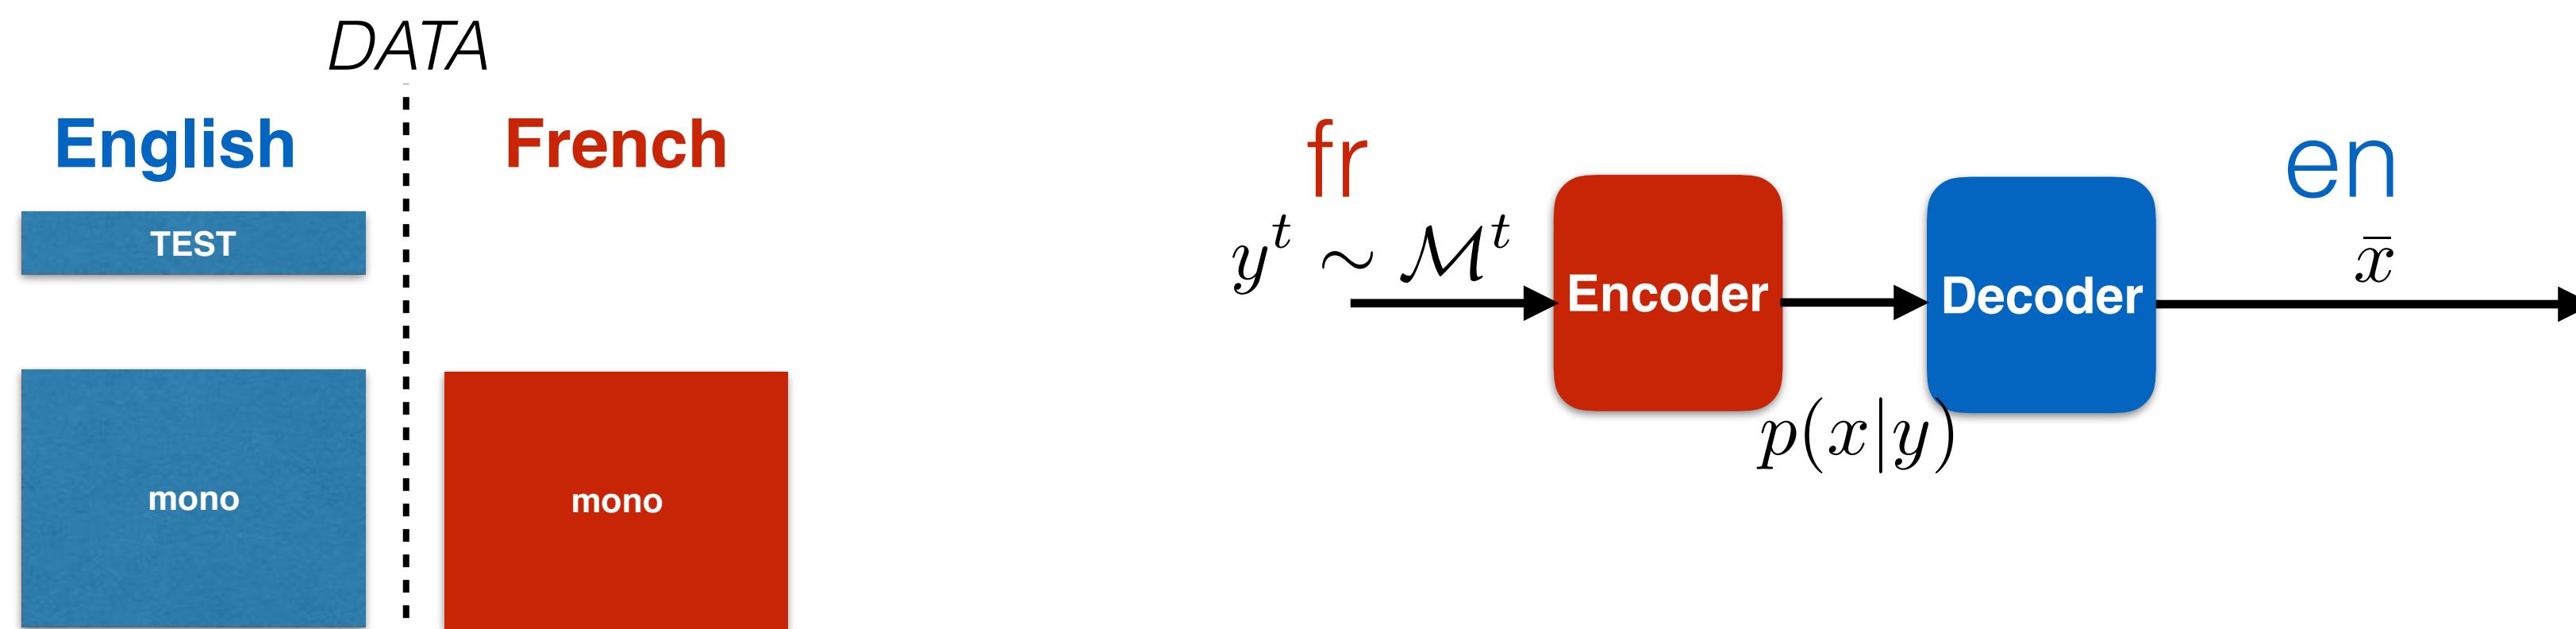
- Assuming no domain effect, there are lots of training paradigms depending on the available data.
- It is hard to predict what works best.
- In general, DAE pretraining, (iterative) BT and multi-lingual training perform strongly on low resource languages.
- All these methods can be combined together, but it requires some level of craftsmanship...
- Final touch: ensembling, fine-tuning, distillation, etc.



Open Challenges

- Diversity of domains and varying translation quality.
- Wildly varying dataset sizes.
- Diversity of language pairs.
- Training large models to handle large datasets, as soon as we put together lots of language pairs and monolingual data.

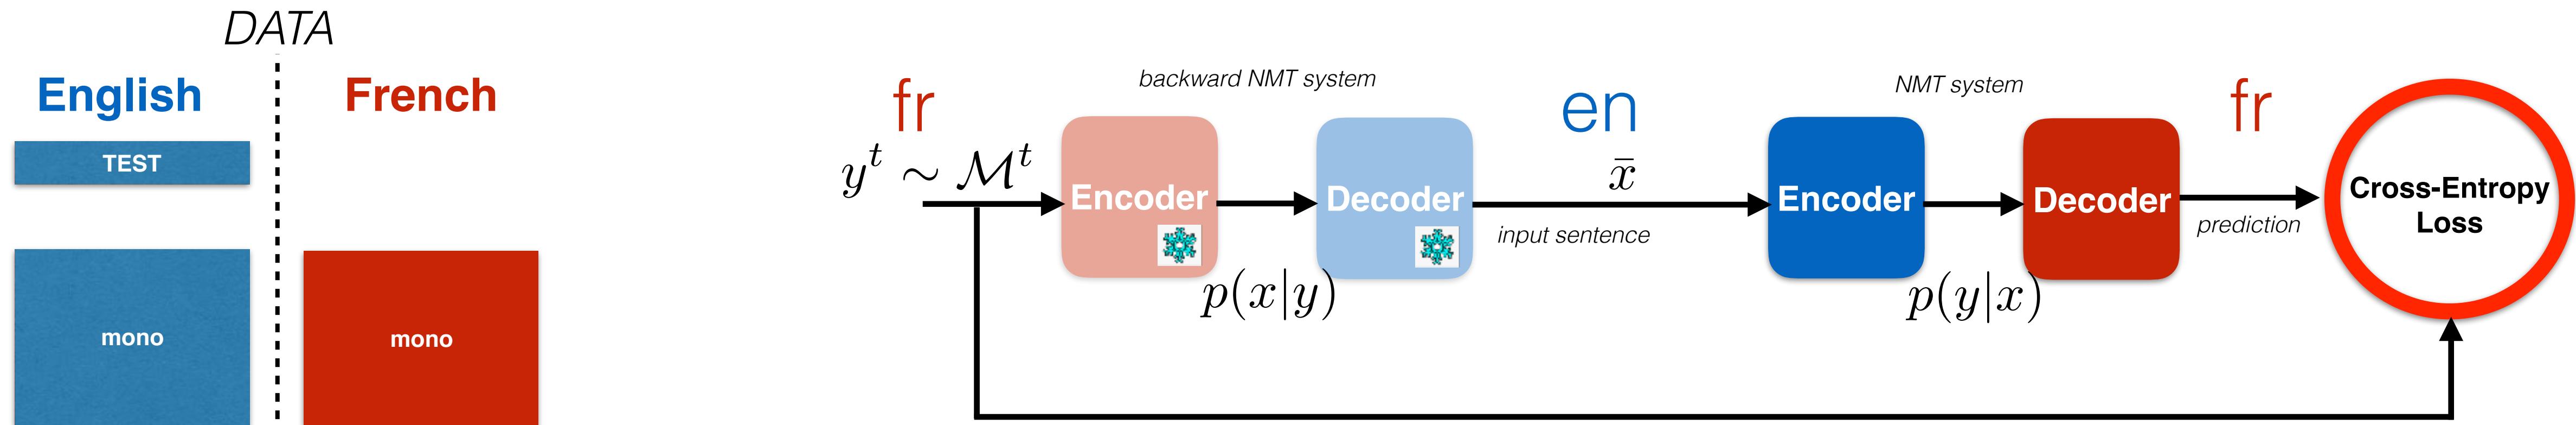
Case Study #1: Unsupervised MT



$$\mathcal{M}^t = \{y_k^t\}_{k=1,..,M_t}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,..,M_s}$$

Case Study #1: Unsupervised MT



$$\mathcal{M}^t = \{y_k^t\}_{k=1,..,M_t}$$
$$\mathcal{M}^s = \{x_j^s\}_{j=1,..,M_s}$$

...and vice versa starting from English.

This is an example of auto-encoding or cycle consistency.

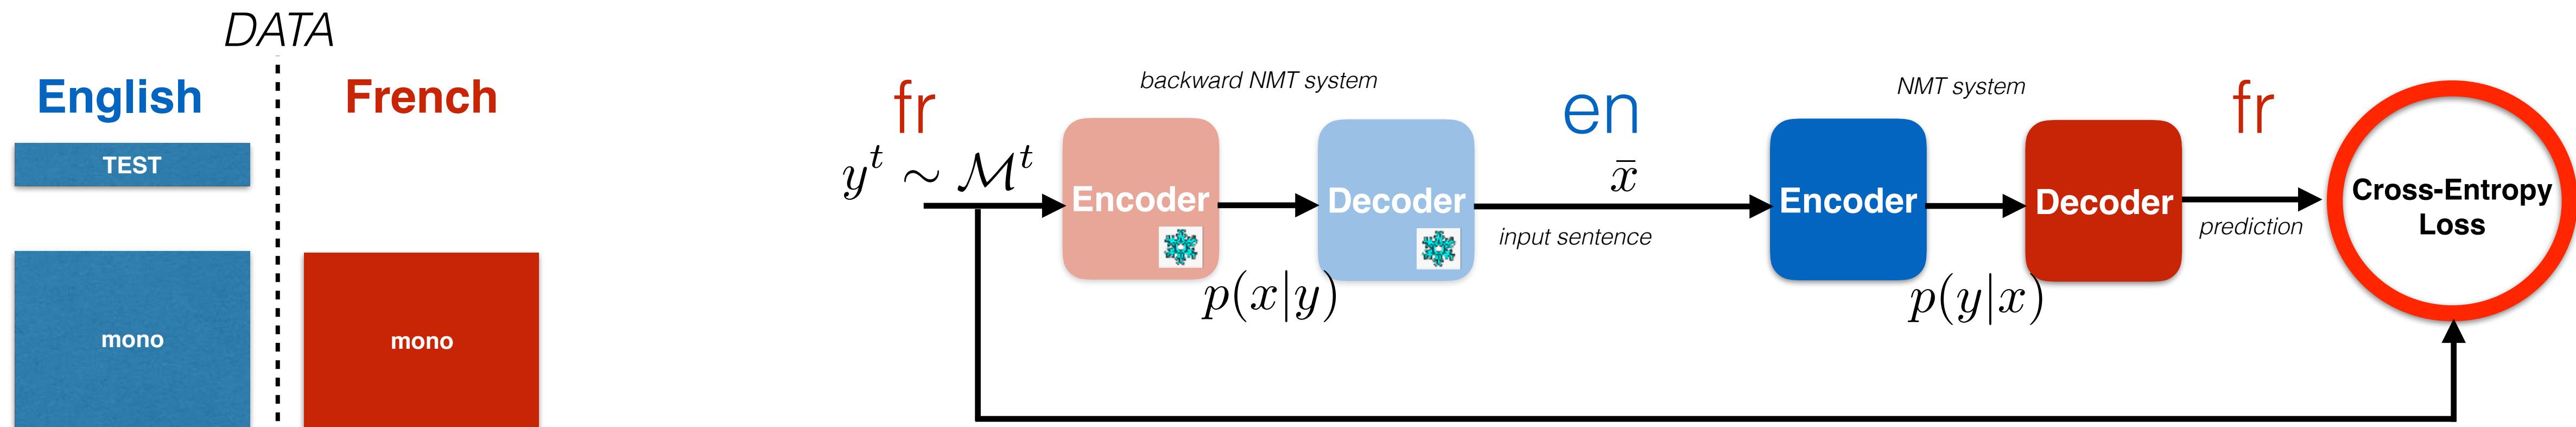
Unpaired Image-to-Image Translation
using Cycle-Consistent Adversarial Networks

Jun-Yan Zhu* Taesung Park* Phillip Isola Alexei A. Efros
Berkeley AI Research (BAIR) laboratory, UC Berkeley



Problem: lack of constraint on \bar{x}

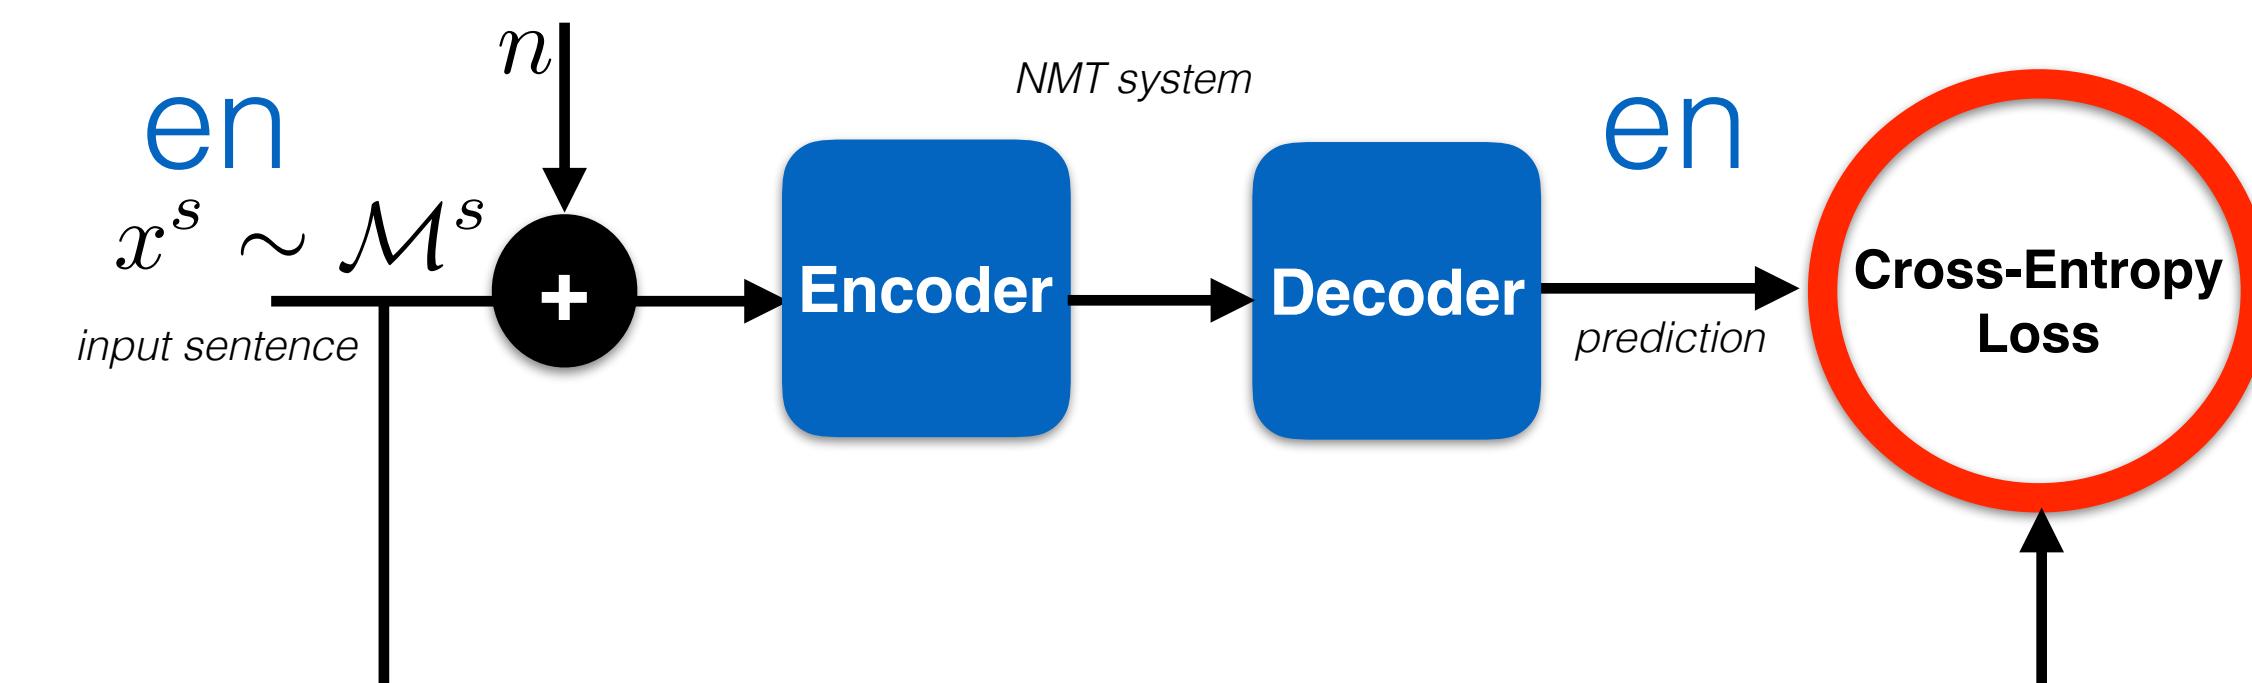
Case Study #1: Unsupervised MT



$$\mathcal{M}^t = \{y_k^t\}_{k=1,..,M_t}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,..,M_s}$$

DAE makes sure decoder outputs fluently in the desired language.



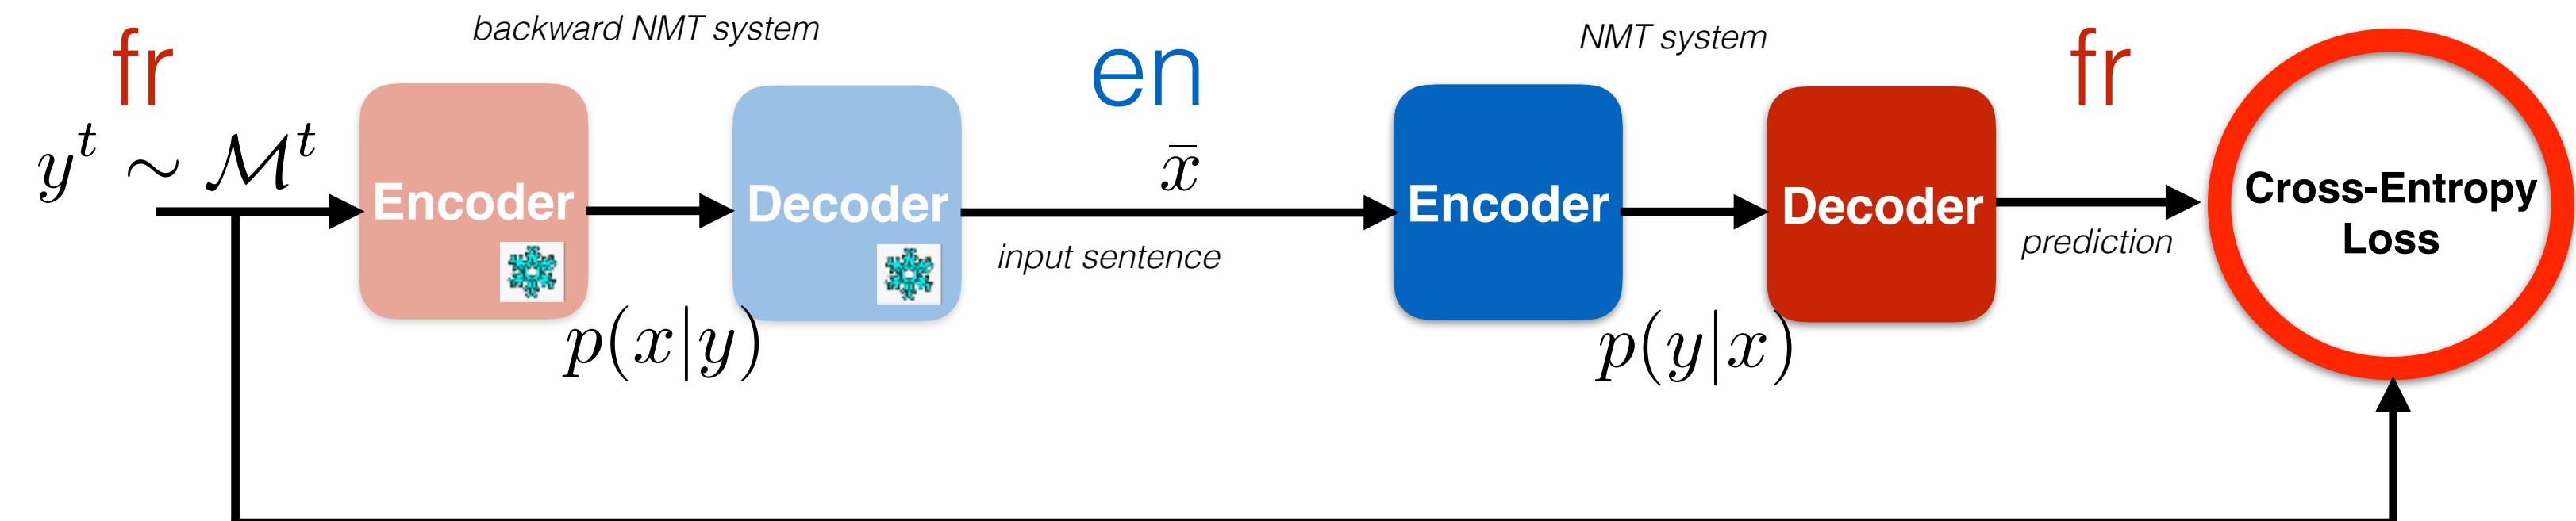
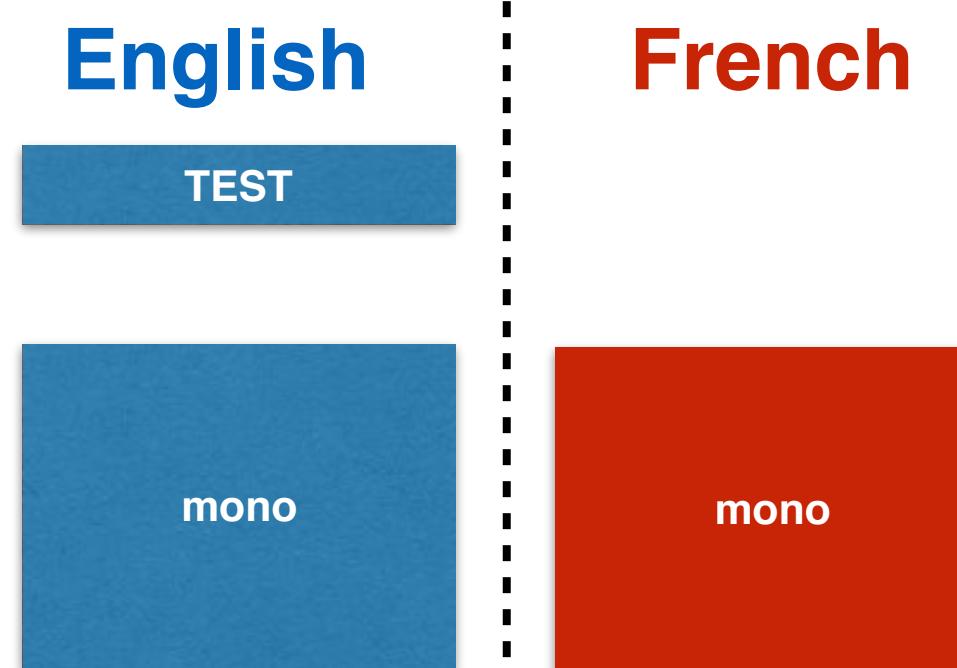
Problem: lack of modularity.

Decoder may behave differently when fed with representations from French encoder VS English encoder.

Lample et al. "Phrase-based and neural unsupervised MT" EMNLP 2018
Artetxe et al. "An effective approach to unsupervised MT" ACL 2019

Case Study #1: Unsupervised MT

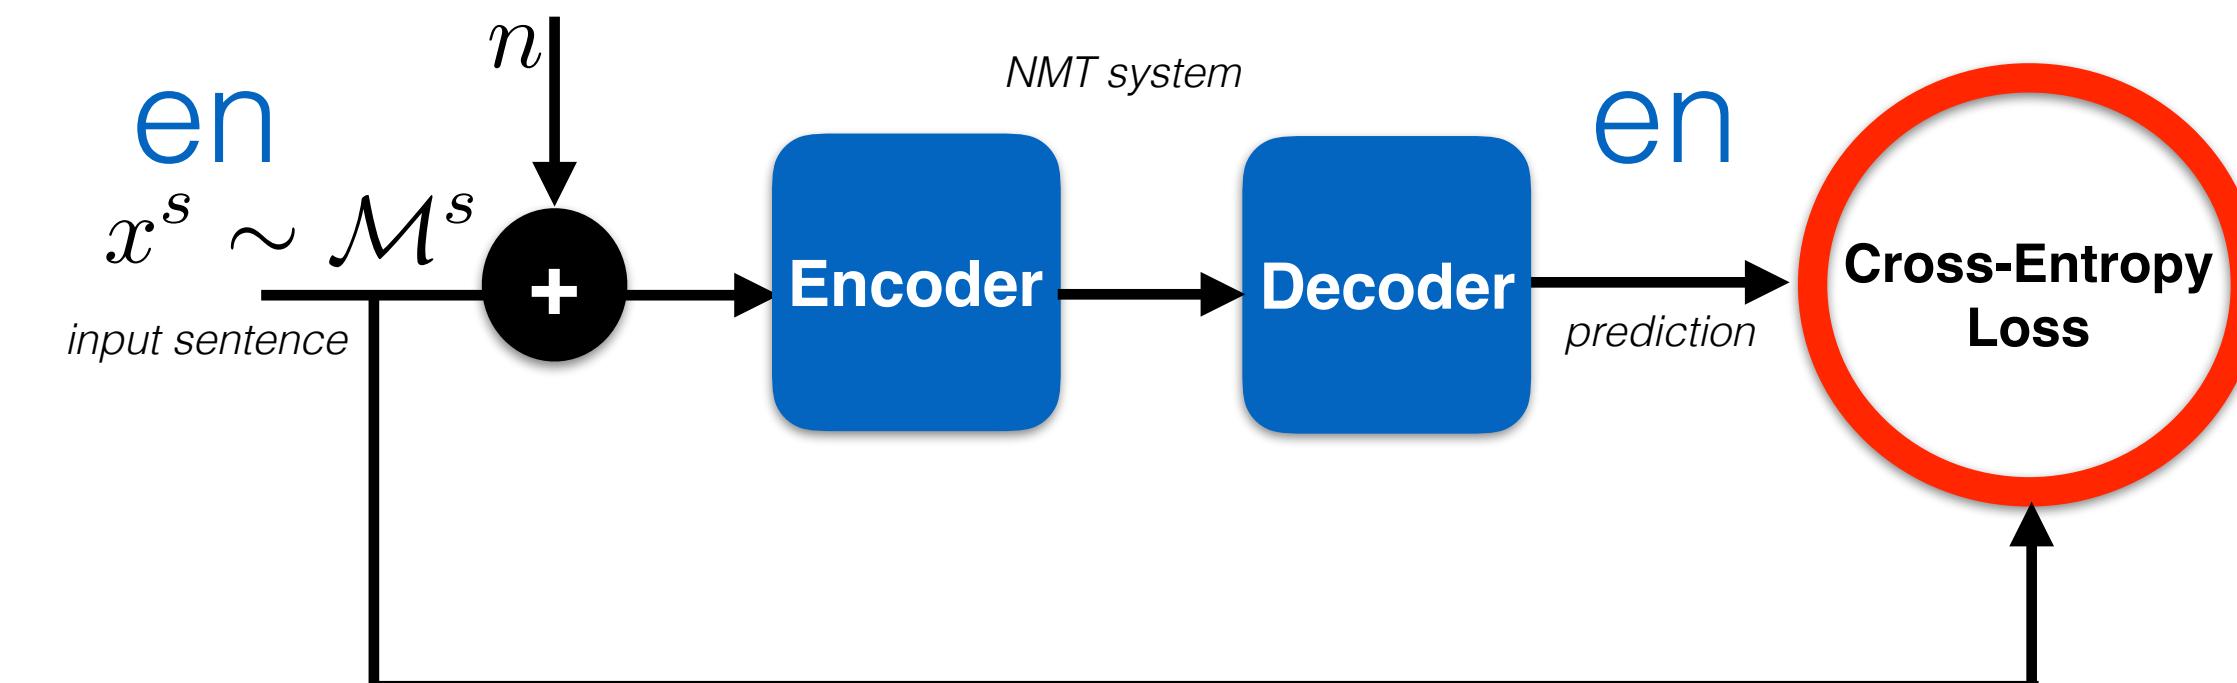
DATA



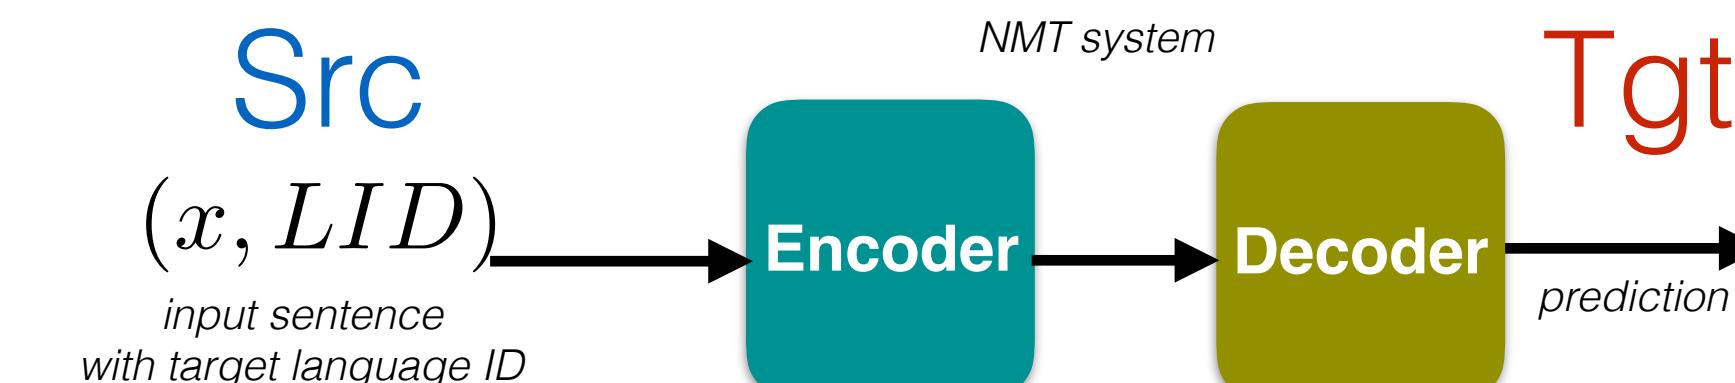
$$\mathcal{M}^t = \{y_k^t\}_{k=1,..,M_t}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,..,M_s}$$

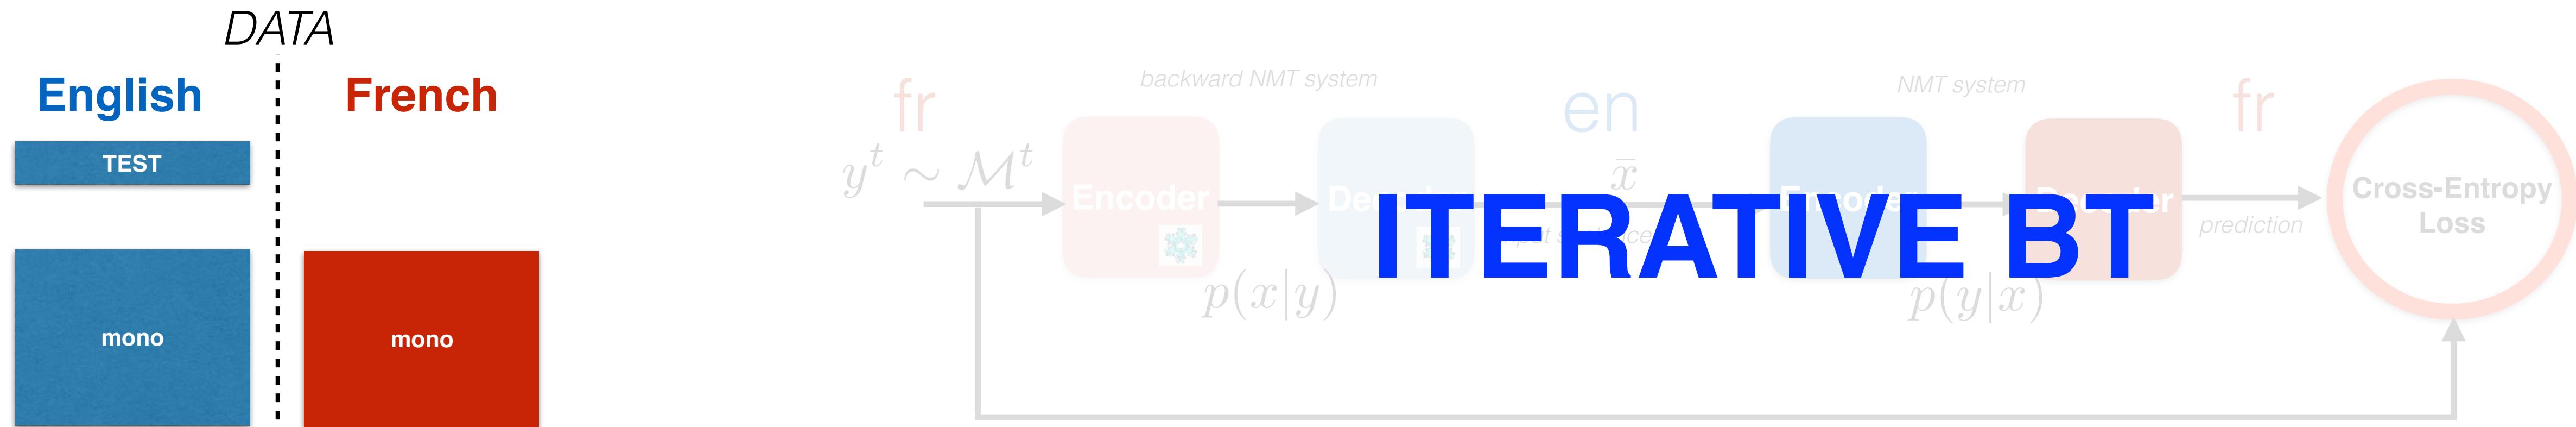
DAE makes sure decoder outputs fluently in the desired language.



Like in multilingual NMT, share encoder and decoder parameters. Encoder is encouraged to produce shared representations (particularly if pre-trained).



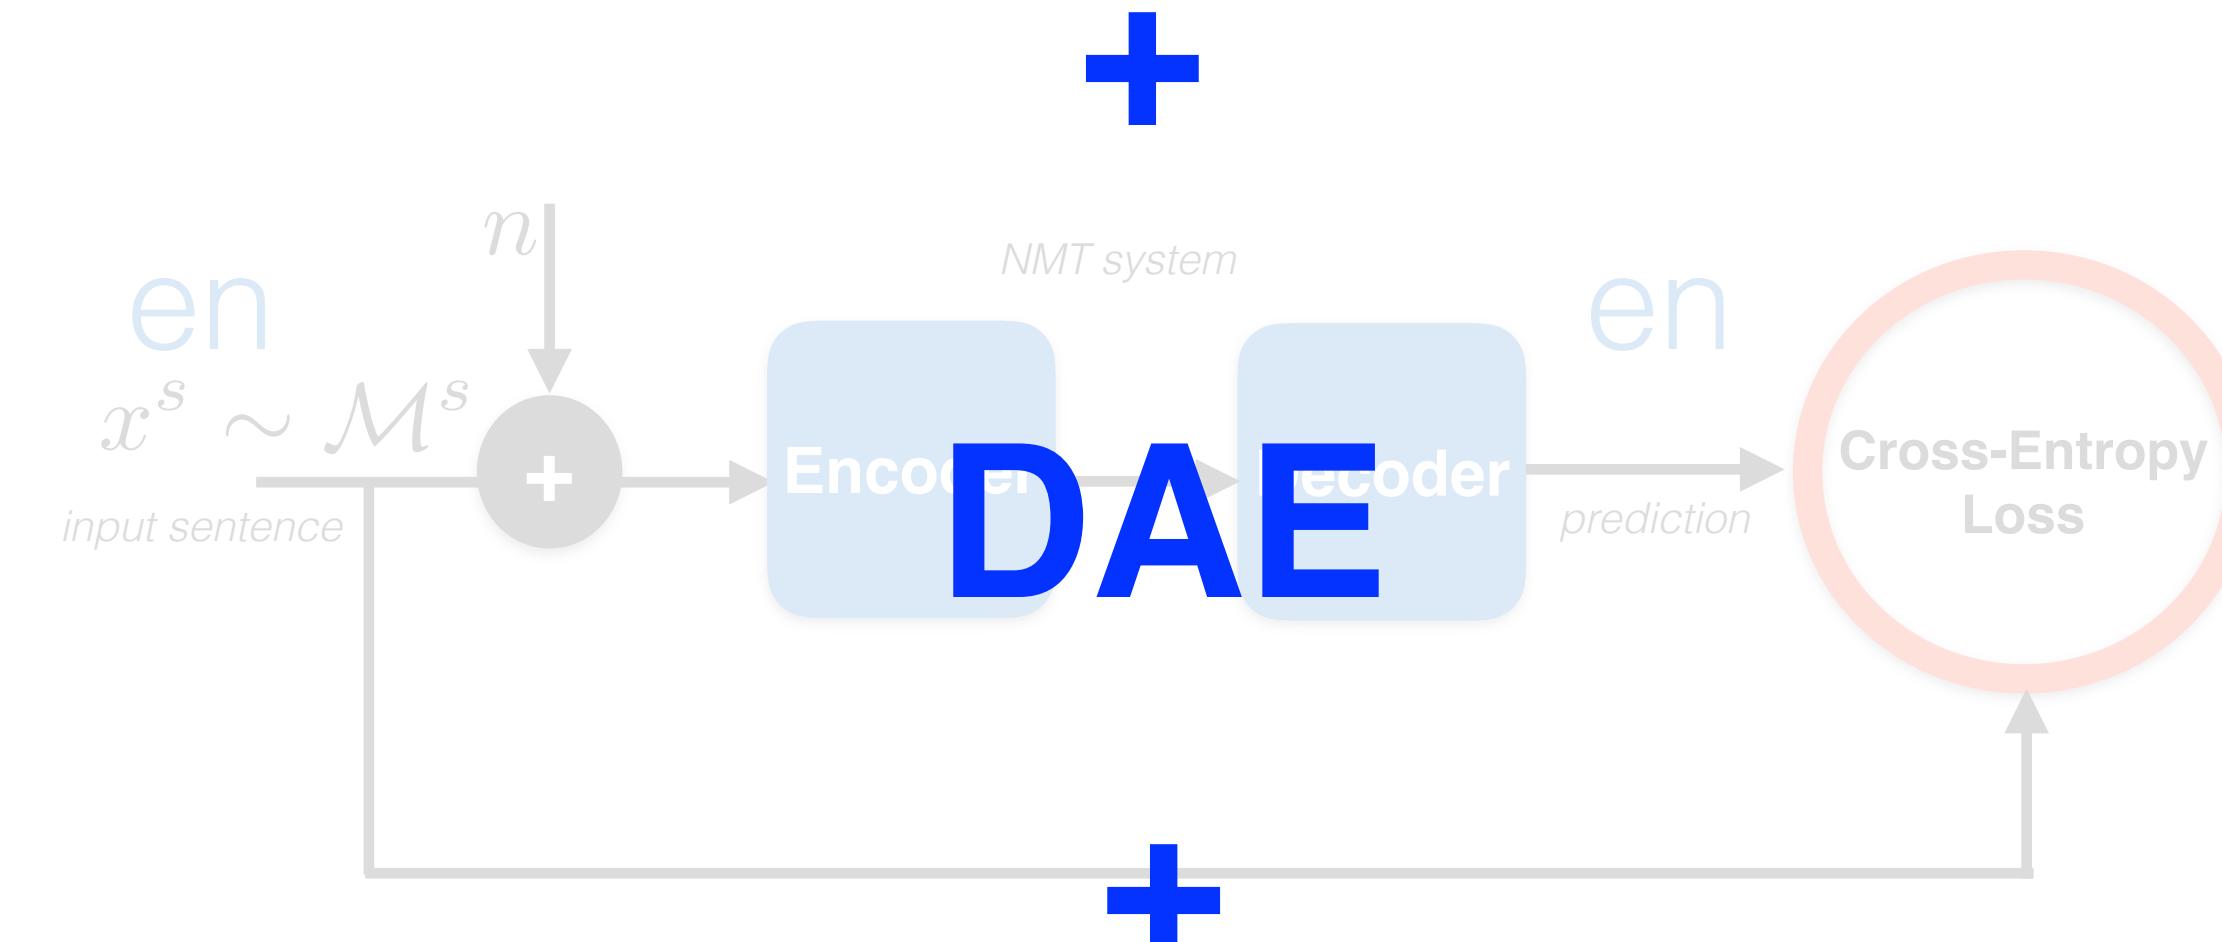
Case Study #1: Unsupervised MT



$$\mathcal{M}^t = \{y_k^t\}_{k=1,..,M_t}$$

$$\mathcal{M}^s = \{x_j^s\}_{j=1,..,M_s}$$

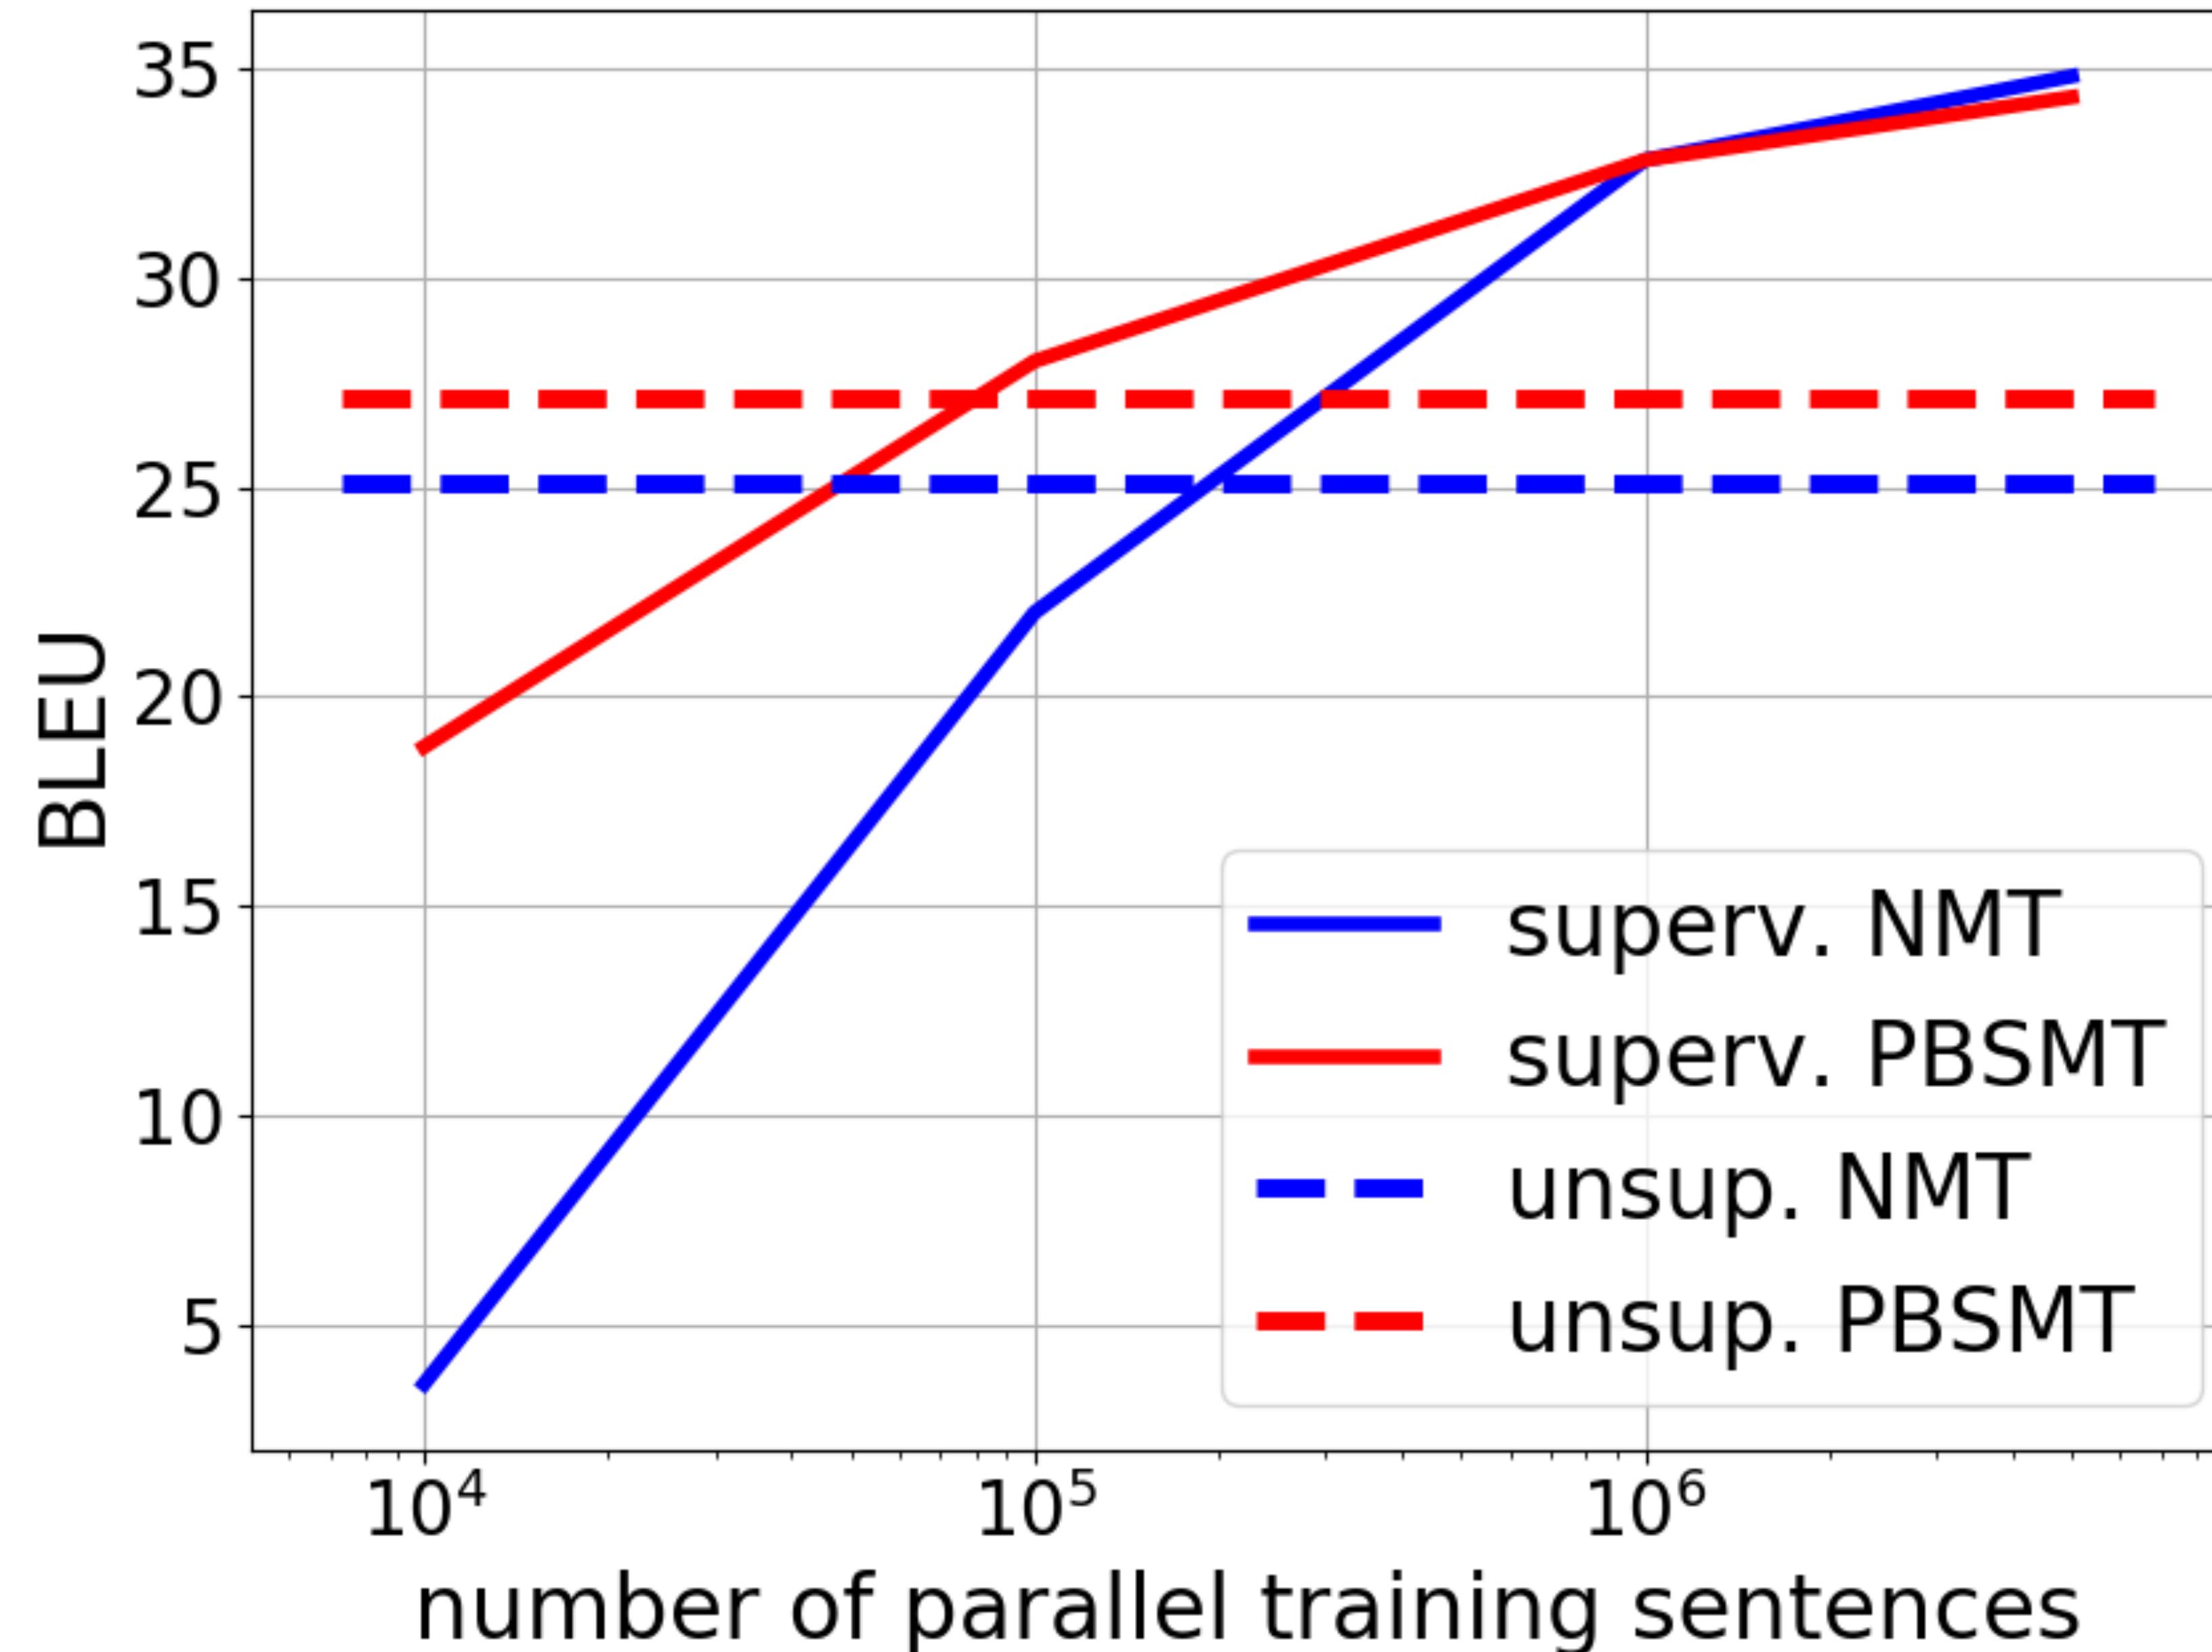
DAE makes sure decoder outputs fluently in the desired language.



Like in multilingual NMT, share encoder and decoder parameters. Encoder is encouraged to produce shared representations (particularly if pre-trained).



WMT'14 En-Fr



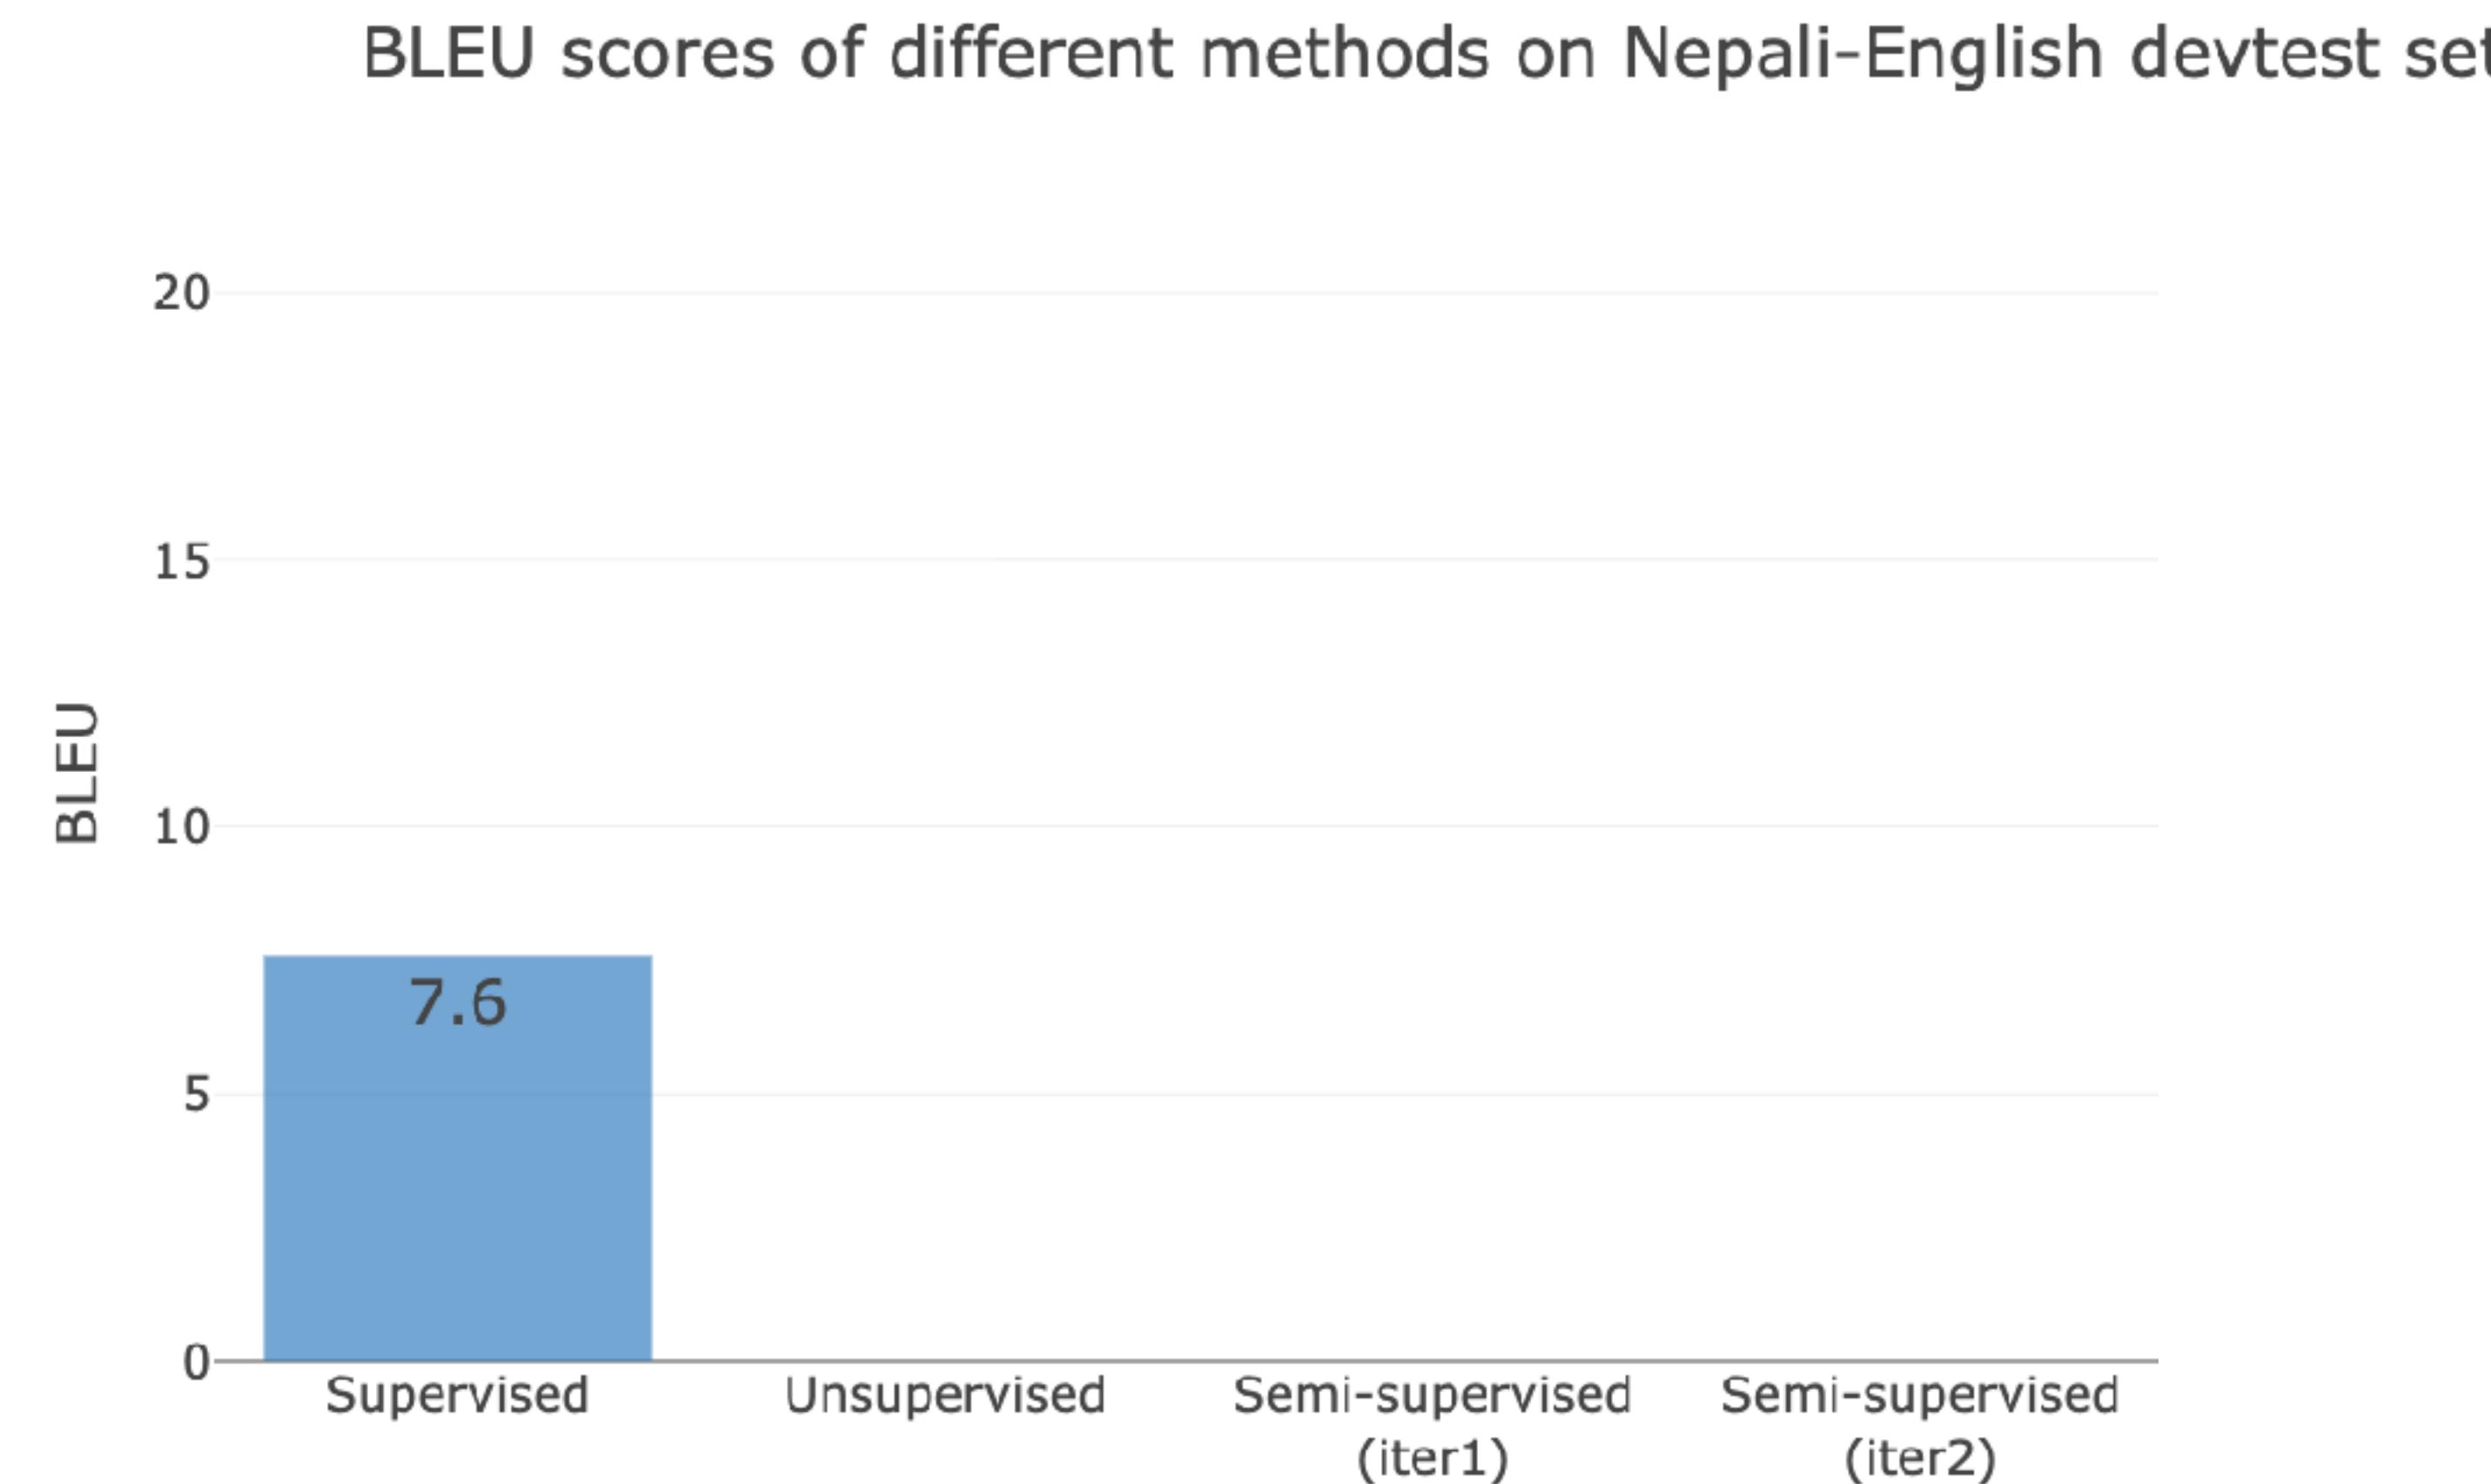
Same ideas can be applied to phrase-based statistical MT systems (PBSMT). NMT and PBSMT can be combined for even better results.

Since unsupMT was trained on about 10M sentences, each parallel sentence is worth 100 monolingual sentences (for this dataset and language pair).

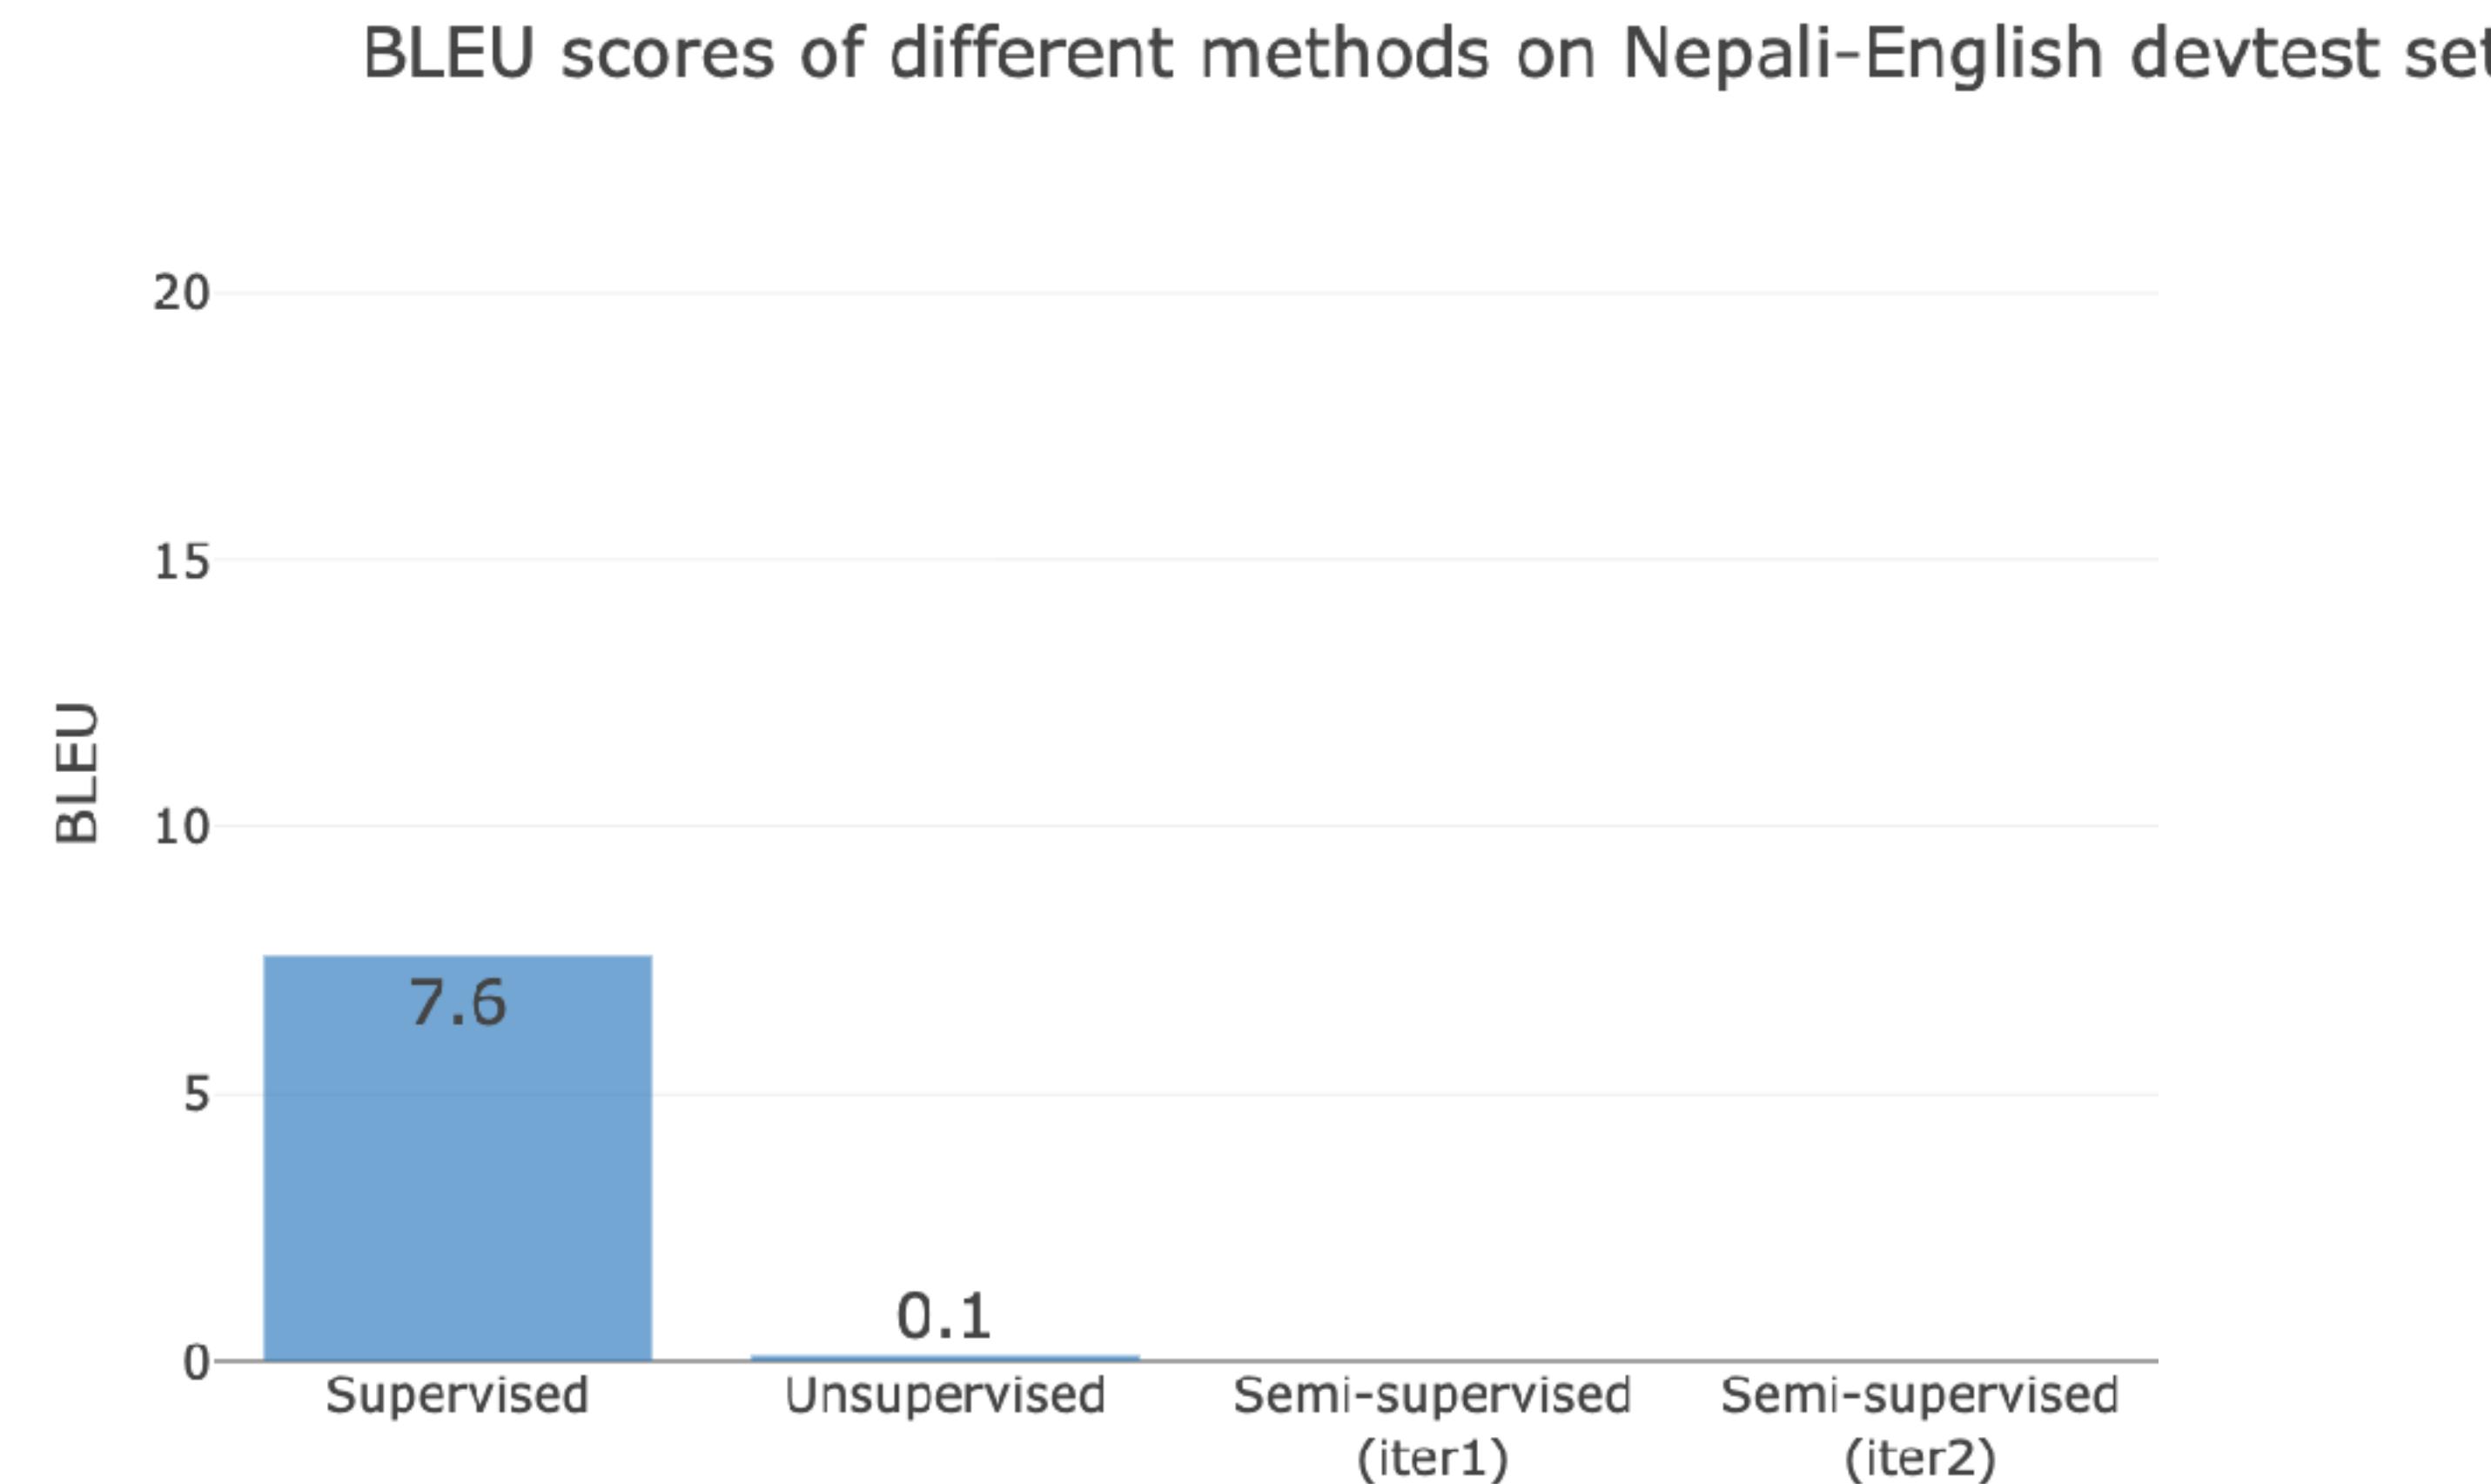
Case Study #2: FLoRes Ne-En

	In-domain (Wikipedia)	Out-of-domain
Parallel	None	500K sentences (Bible, GNOME/Ubuntu, OpenSubtitle, ...) *Hindi: 1.5M
Monolingual	100K sentences	~5M sentences (CommonCrawl) *Hindi: 45M

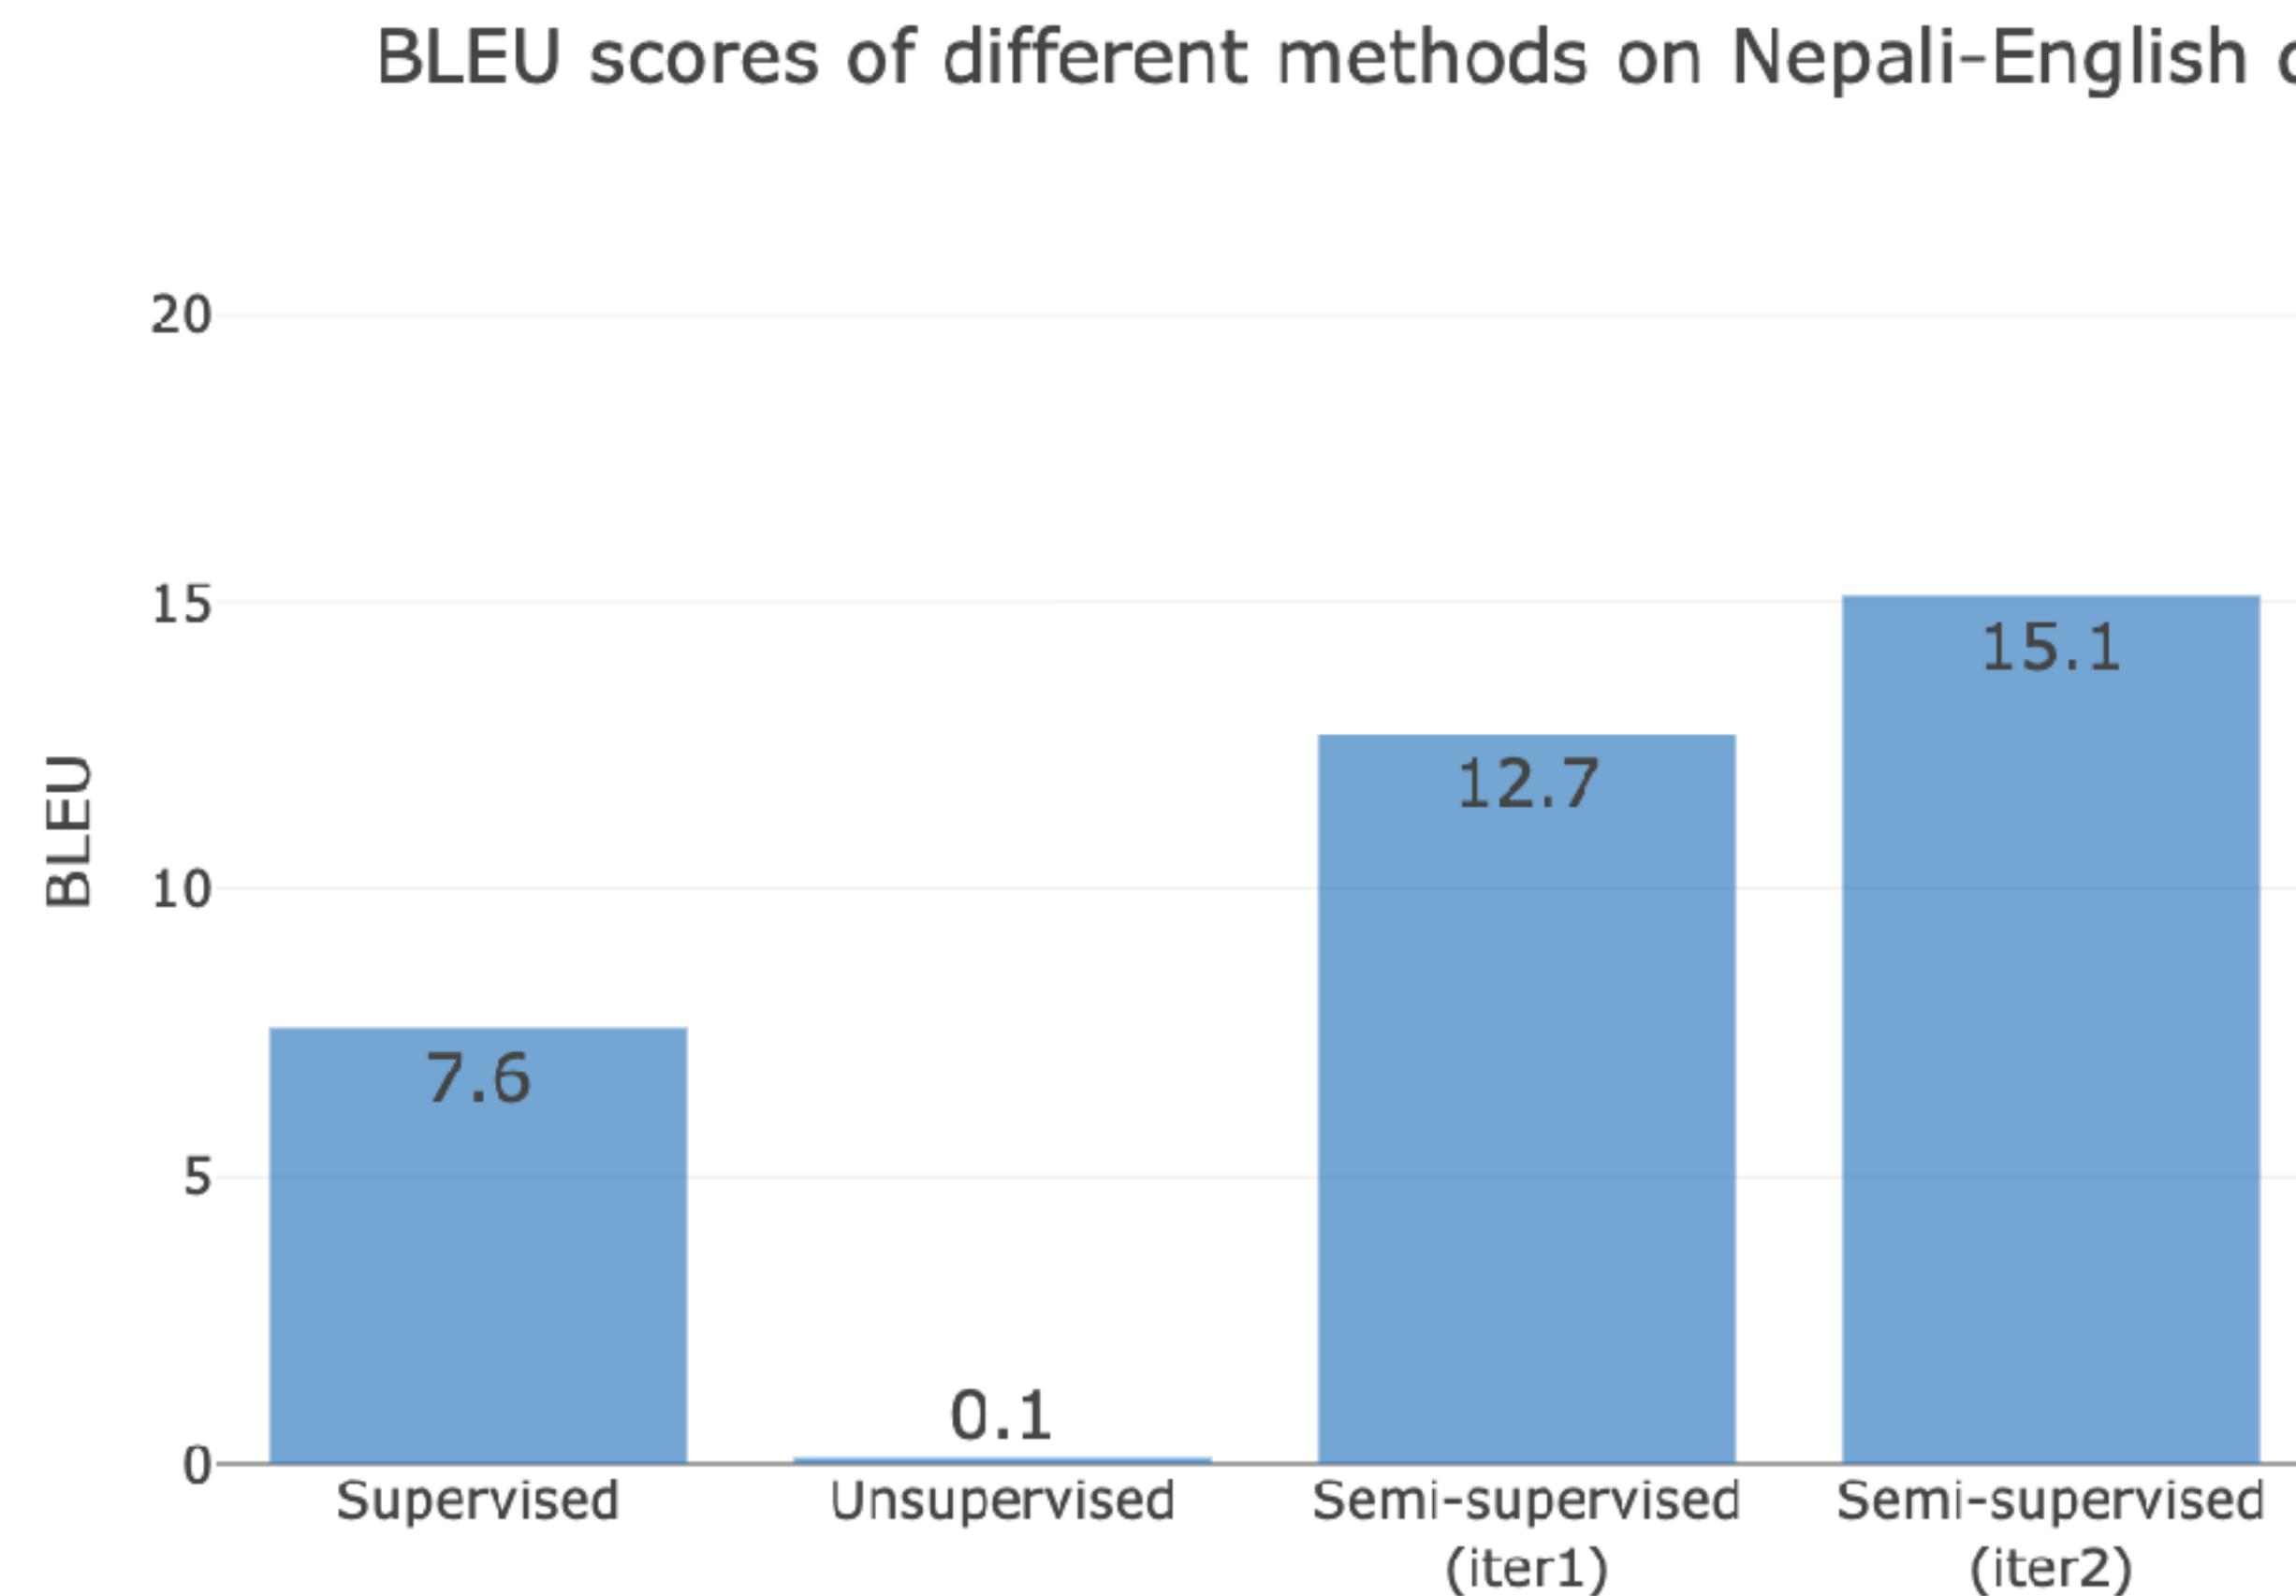
Results on FLoRes: Ne-En



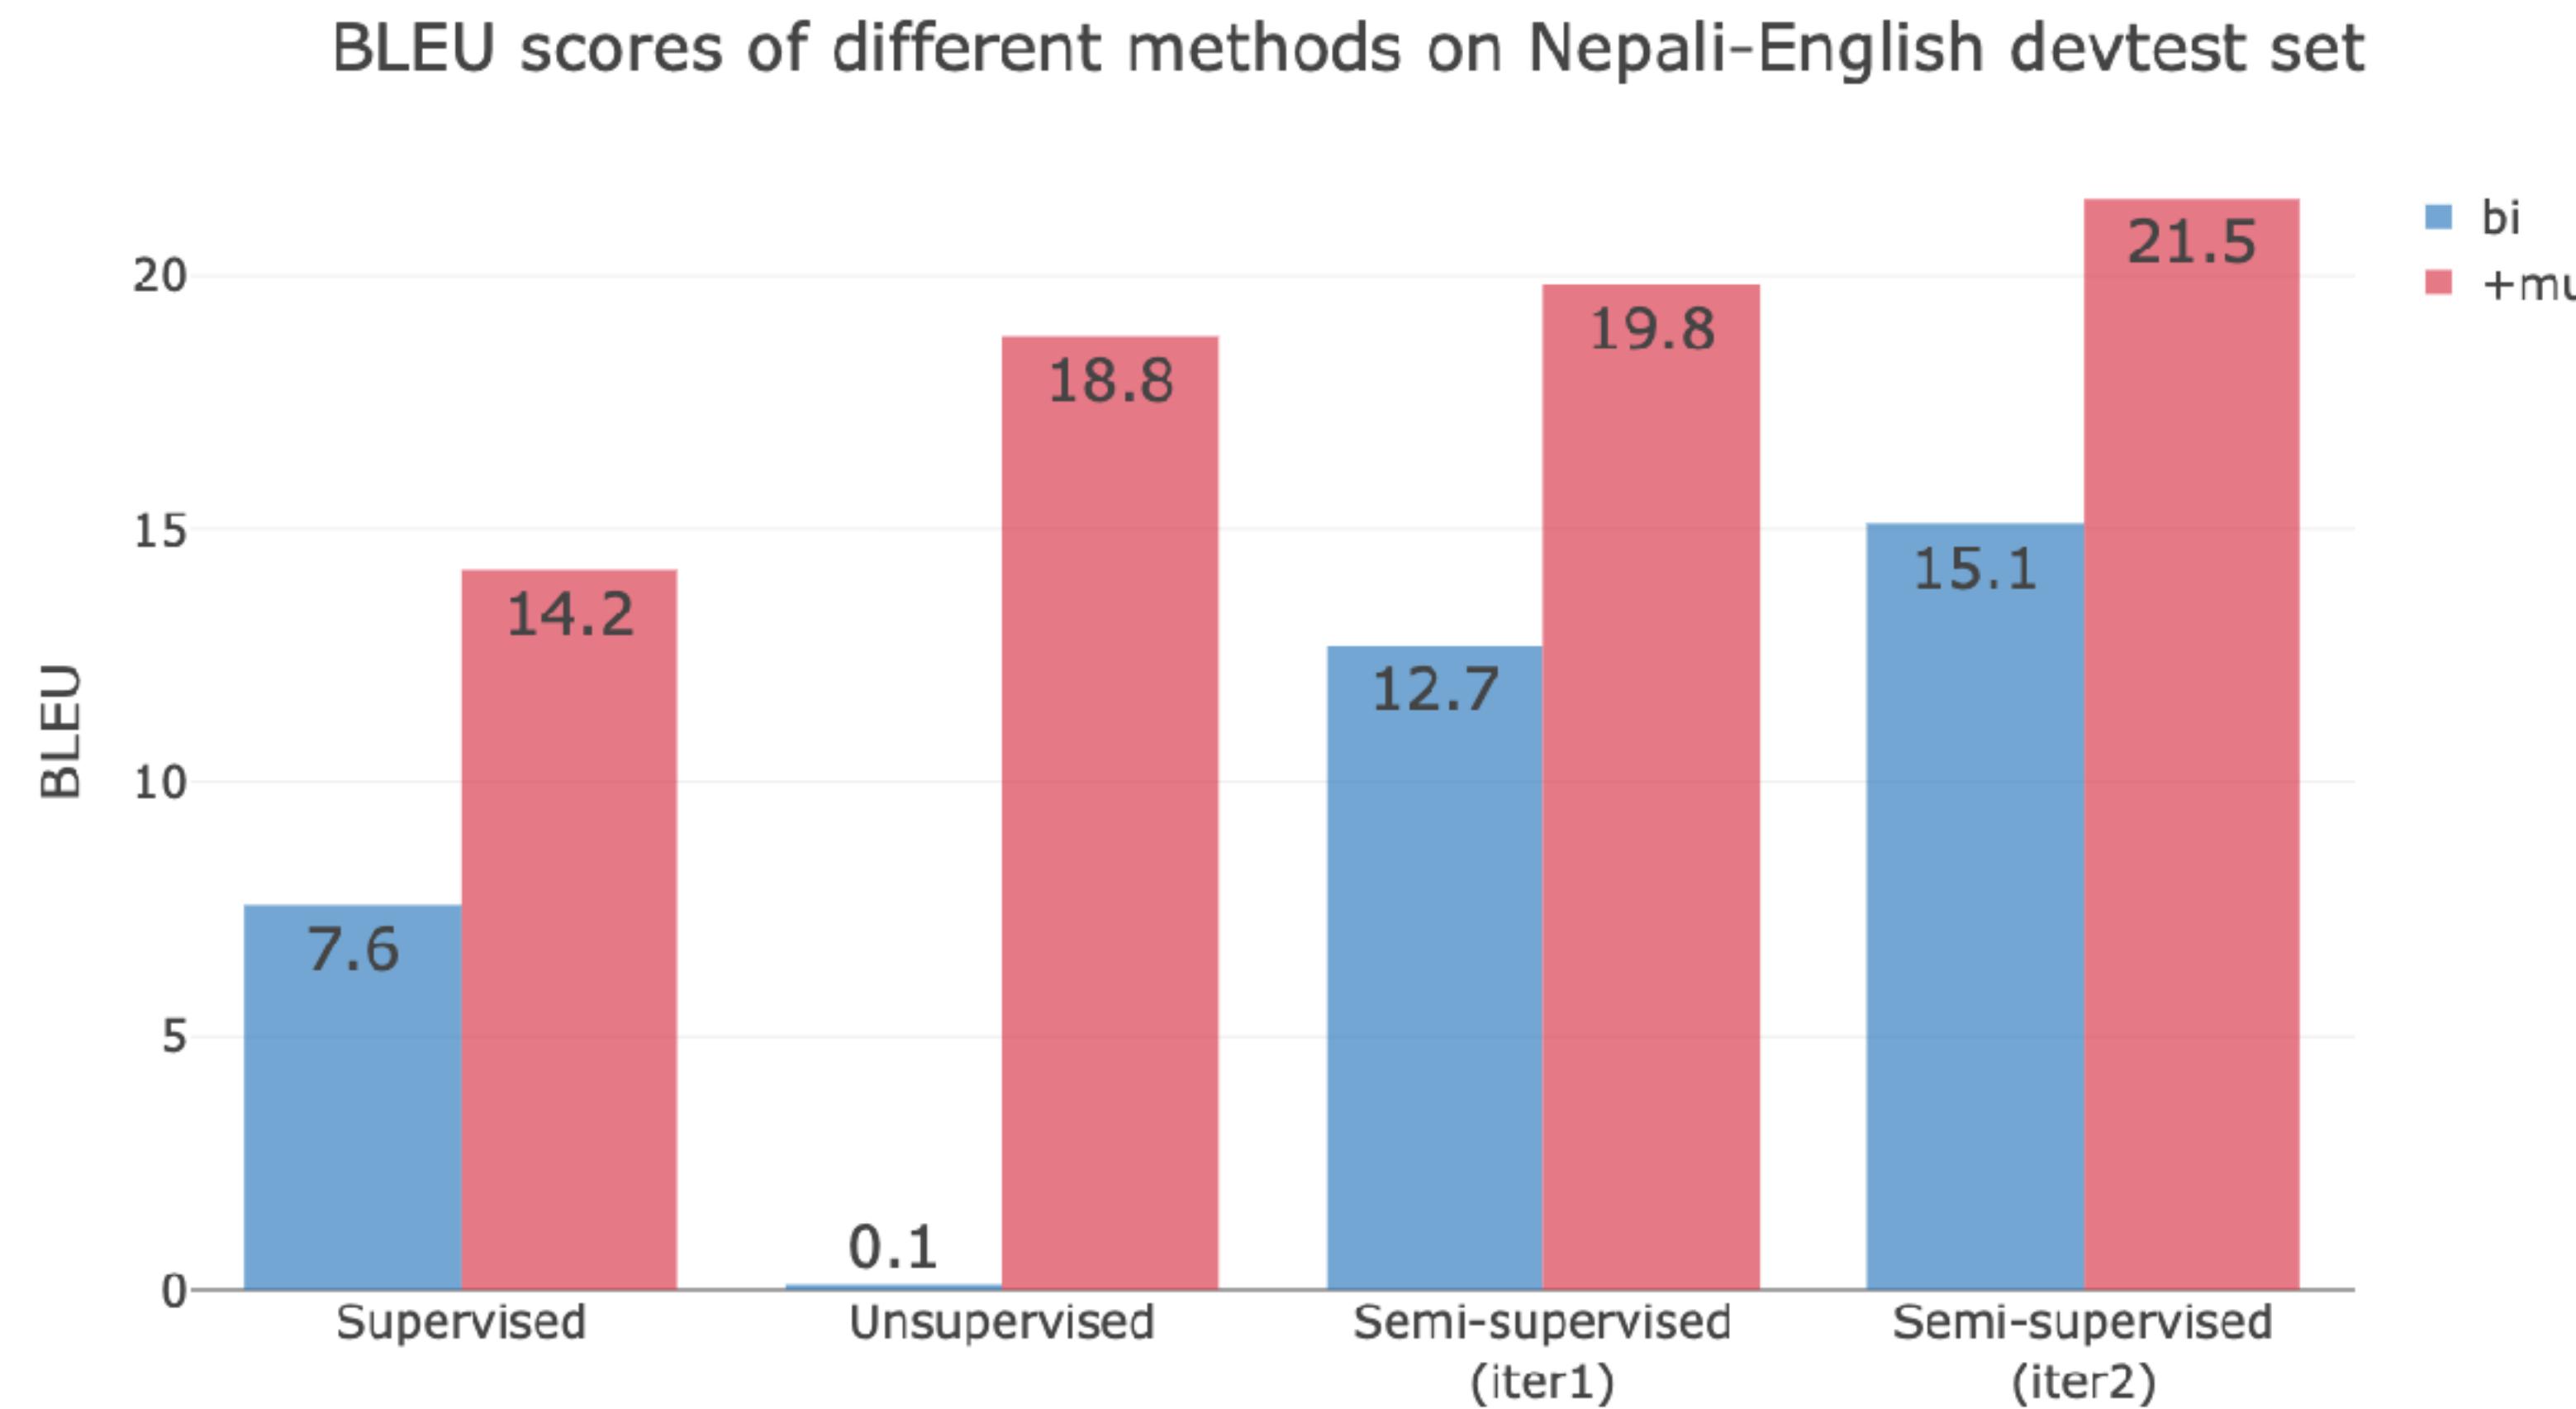
Results on FLoRes: Ne-En



Results on FLoRes: Ne-En



Results on FLoRes: Ne-En



Case-Study #3: English-Burmese

ပဟိုစာမျက်နှာ

ဆွေးနွေးချက်

ဖတ်ရန်

ရင်းမြစ်ကို ကြည့်ရန်

ရာဇ်ဝင်ကြည့်ရန်

ဝိကီပီးဒီးယား တွင် ရှာဖွေရန်



ပဟိုစာမျက်နှာ

ဝိကီပီးဒီးယားမှ ကြိုဆိုပါသည်။

မည်သူမဆုံး ကြည့်ရပ်ဆင်နိုင်သော အခမဲ့လွှတ်လပ်စွယ်စုံကျမ်း ဖြစ်ပါသည်။
အကြောင်းအရာပေါင်း ၄၄၈၁၄ ခုကို မြန်မာဘာသာဖြင့် ဖတ်ရှုနိုင်ပါသည်။



အထူးအကြောင်းအရာ

သိန့်နယ် သည် ယခင် ရှမ်းပအော်ရာအိပ်ပြည်နယ်များတွင် ပါဝင်ခဲ့သော ပြည်နယ်တစ်ခု ဖြစ်ပါသည်။ သိန့်နယ်ကို ခရစ် ၁၈၈၈ ခုနှစ် အားလုံးတို့ ဝင်ရောက်သီမ်းပိုက်ပြီးနောက်မှ မြောက်သိန့်နယ် (သိန့်နယ်)နှင့် တောင်သိန့်နယ်(မိုင်းရယ်နယ်)ဟု ခွဲခြားအုပ်ချုပ်ခဲ့သည်။ ရေးအဓိကသမယက သိန့်နယ်ကြီးသည် အစိတ်စိတ်ကွဲပြားခြင်းမရှိဘဲ ရှမ်းပြည်နယ်တရာ်များတွင် အကျယ်ပြန့်စုံမှုးစုံးသော နယ်ကြီးဖြစ်ခဲ့သည်။ သို့သော် မြန်မာဘုရင်များ ဝင်ရောက်တိုက်ခိုက် သီမ်းပိုက်ပြီးသည့်နောက် အုပ်ချုပ်ရေးဝါဒအရ ရာထူးလှသည့်နယ်ရှင် တော်ဘွားတို့ကြောင့် သိန့်နယ်ကြီးသည် ပါးနယ်အထိ အစိတ်စိတ်ကွဲပြားခဲ့လေသည်။ ထိုအတွင်း တော်ဘွားအချင်းချင်း စိတ်ဝမ်းကွဲကာ တစ်ဦးနှင့်တစ်ဦးတိုက်ခိုက်၏၍ ဆိုင်ရာ နယ်ပယ်များကို အုပ်ချုပ်က သည်။ နောက်စုံးခရစ် ၁၈၈၈ ခုနှစ်၊ အားလုံးတို့ဝင်ရောက်လာမှ အထက်ပါ အတိုင်းနှစ်နယ်ခွဲ၏၍ အုပ်ချုပ်ခဲ့သည်။ နှစ်နယ်ခွဲ၏၍ အုပ်ချုပ်စက ခွန်ဆိုင်တံ့ဟမ်းအား မြောက်သိန့်နယ်အတွက် တော်ဘွားအဖြစ်လည်းကောင်း၊ ဆိုင်နော်ဖော်သား နော်မိုင်းအား တောင်သိန့်နယ် တော်ဘွားအဖြစ်လည်းကောင်း၊ ခန့်အပ်ခဲ့လေသည်။ ၁၉၂၅ ခုနှစ်တွင် မြောက်သိန့်နယ်ကို စပ်ဟုံဖက် တော်ဘွားအဖြစ် ဆောင်ရွက်ခဲ့လေသည်။

- အနုပညာ
- အစွဲပွဲစွဲ
- ပထဝိဝင်

- သမိုင်း
- သချာ
- သိပ္ပါယ်

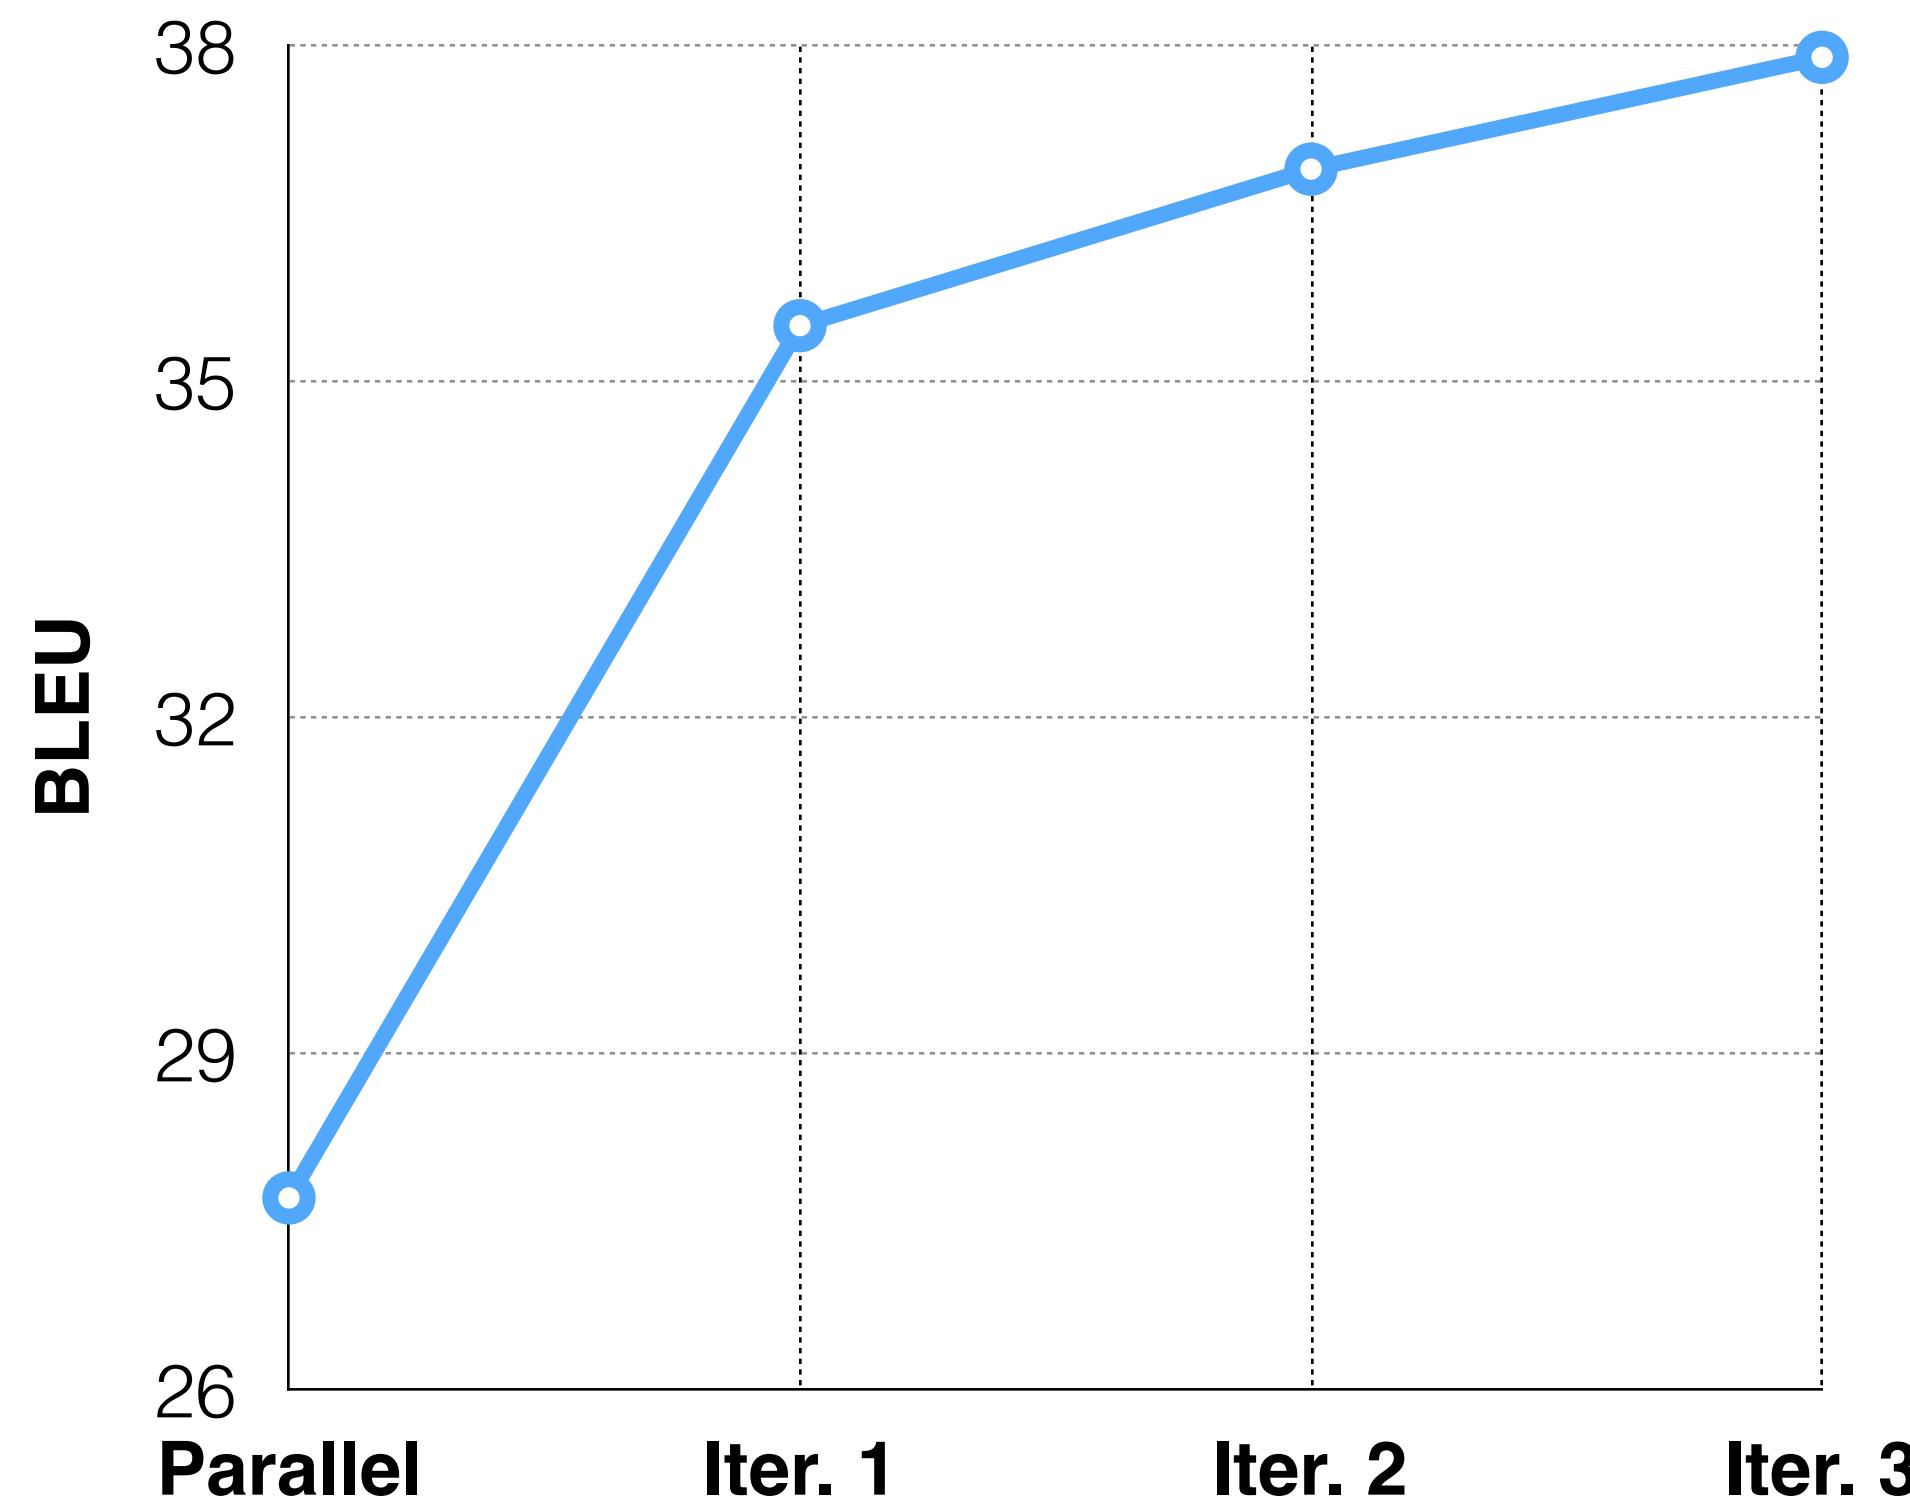
- လူမှုရေး
- နည်းပညာ
- မှုပိုးအားလုံး

Workshop on Asian Translation 2019: English-Burmese

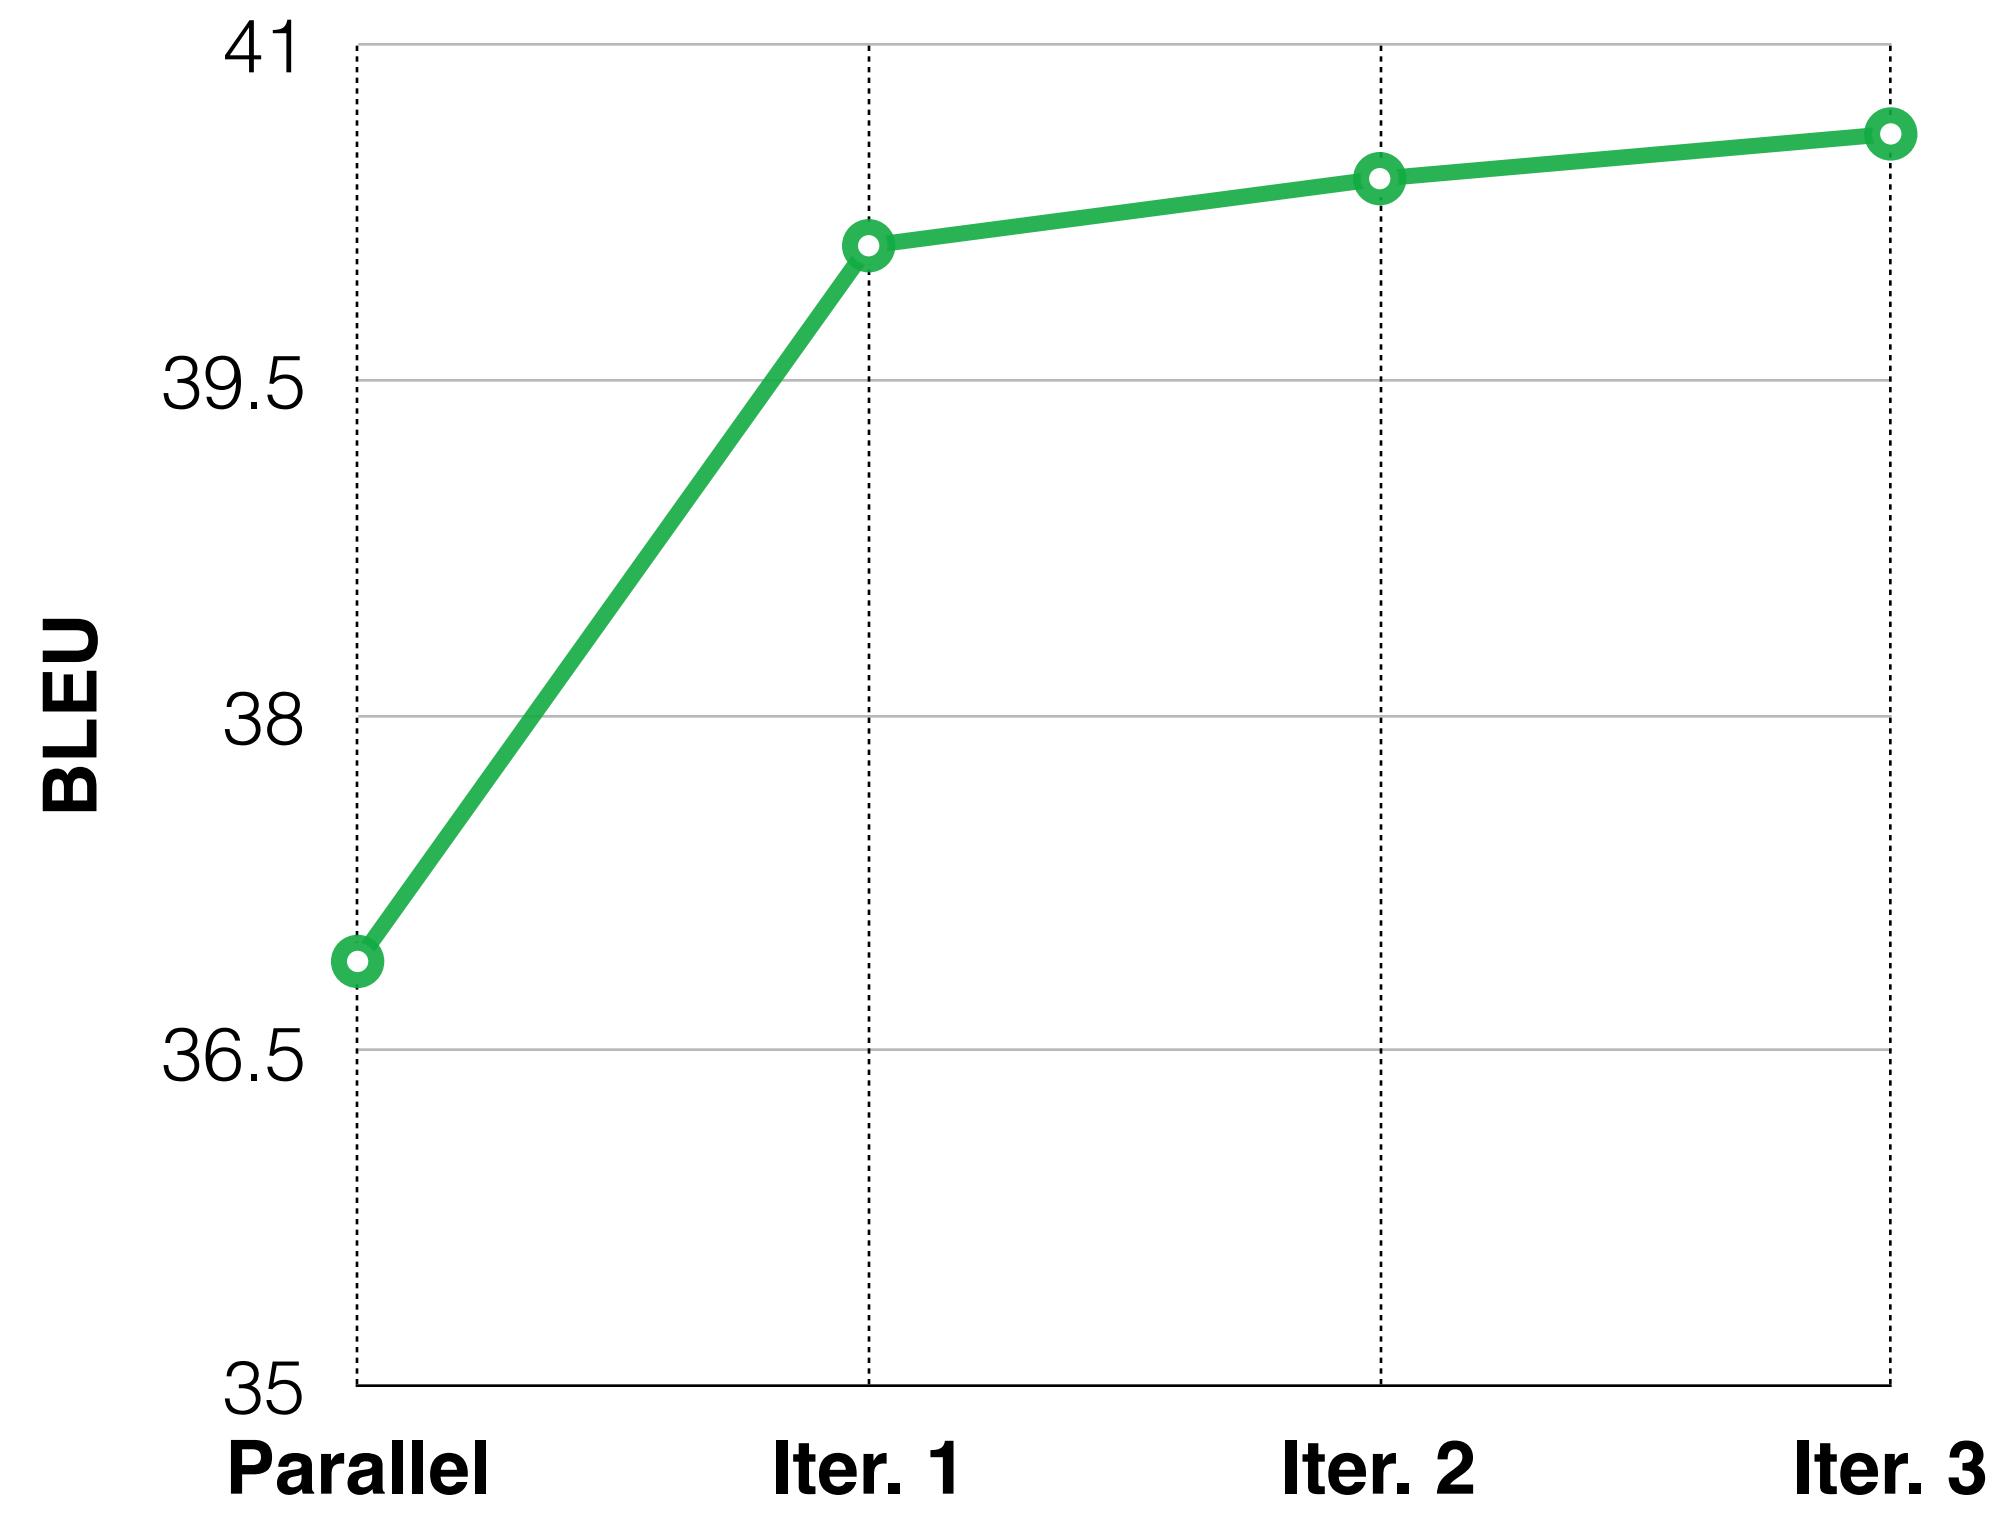
	In-domain (News)	Out-of-domain
Parallel	20K sentences	200K sentences
Monolingual	~79M sentences (En only)	~23M sentences (My only)

Results: Iterative ST+BT

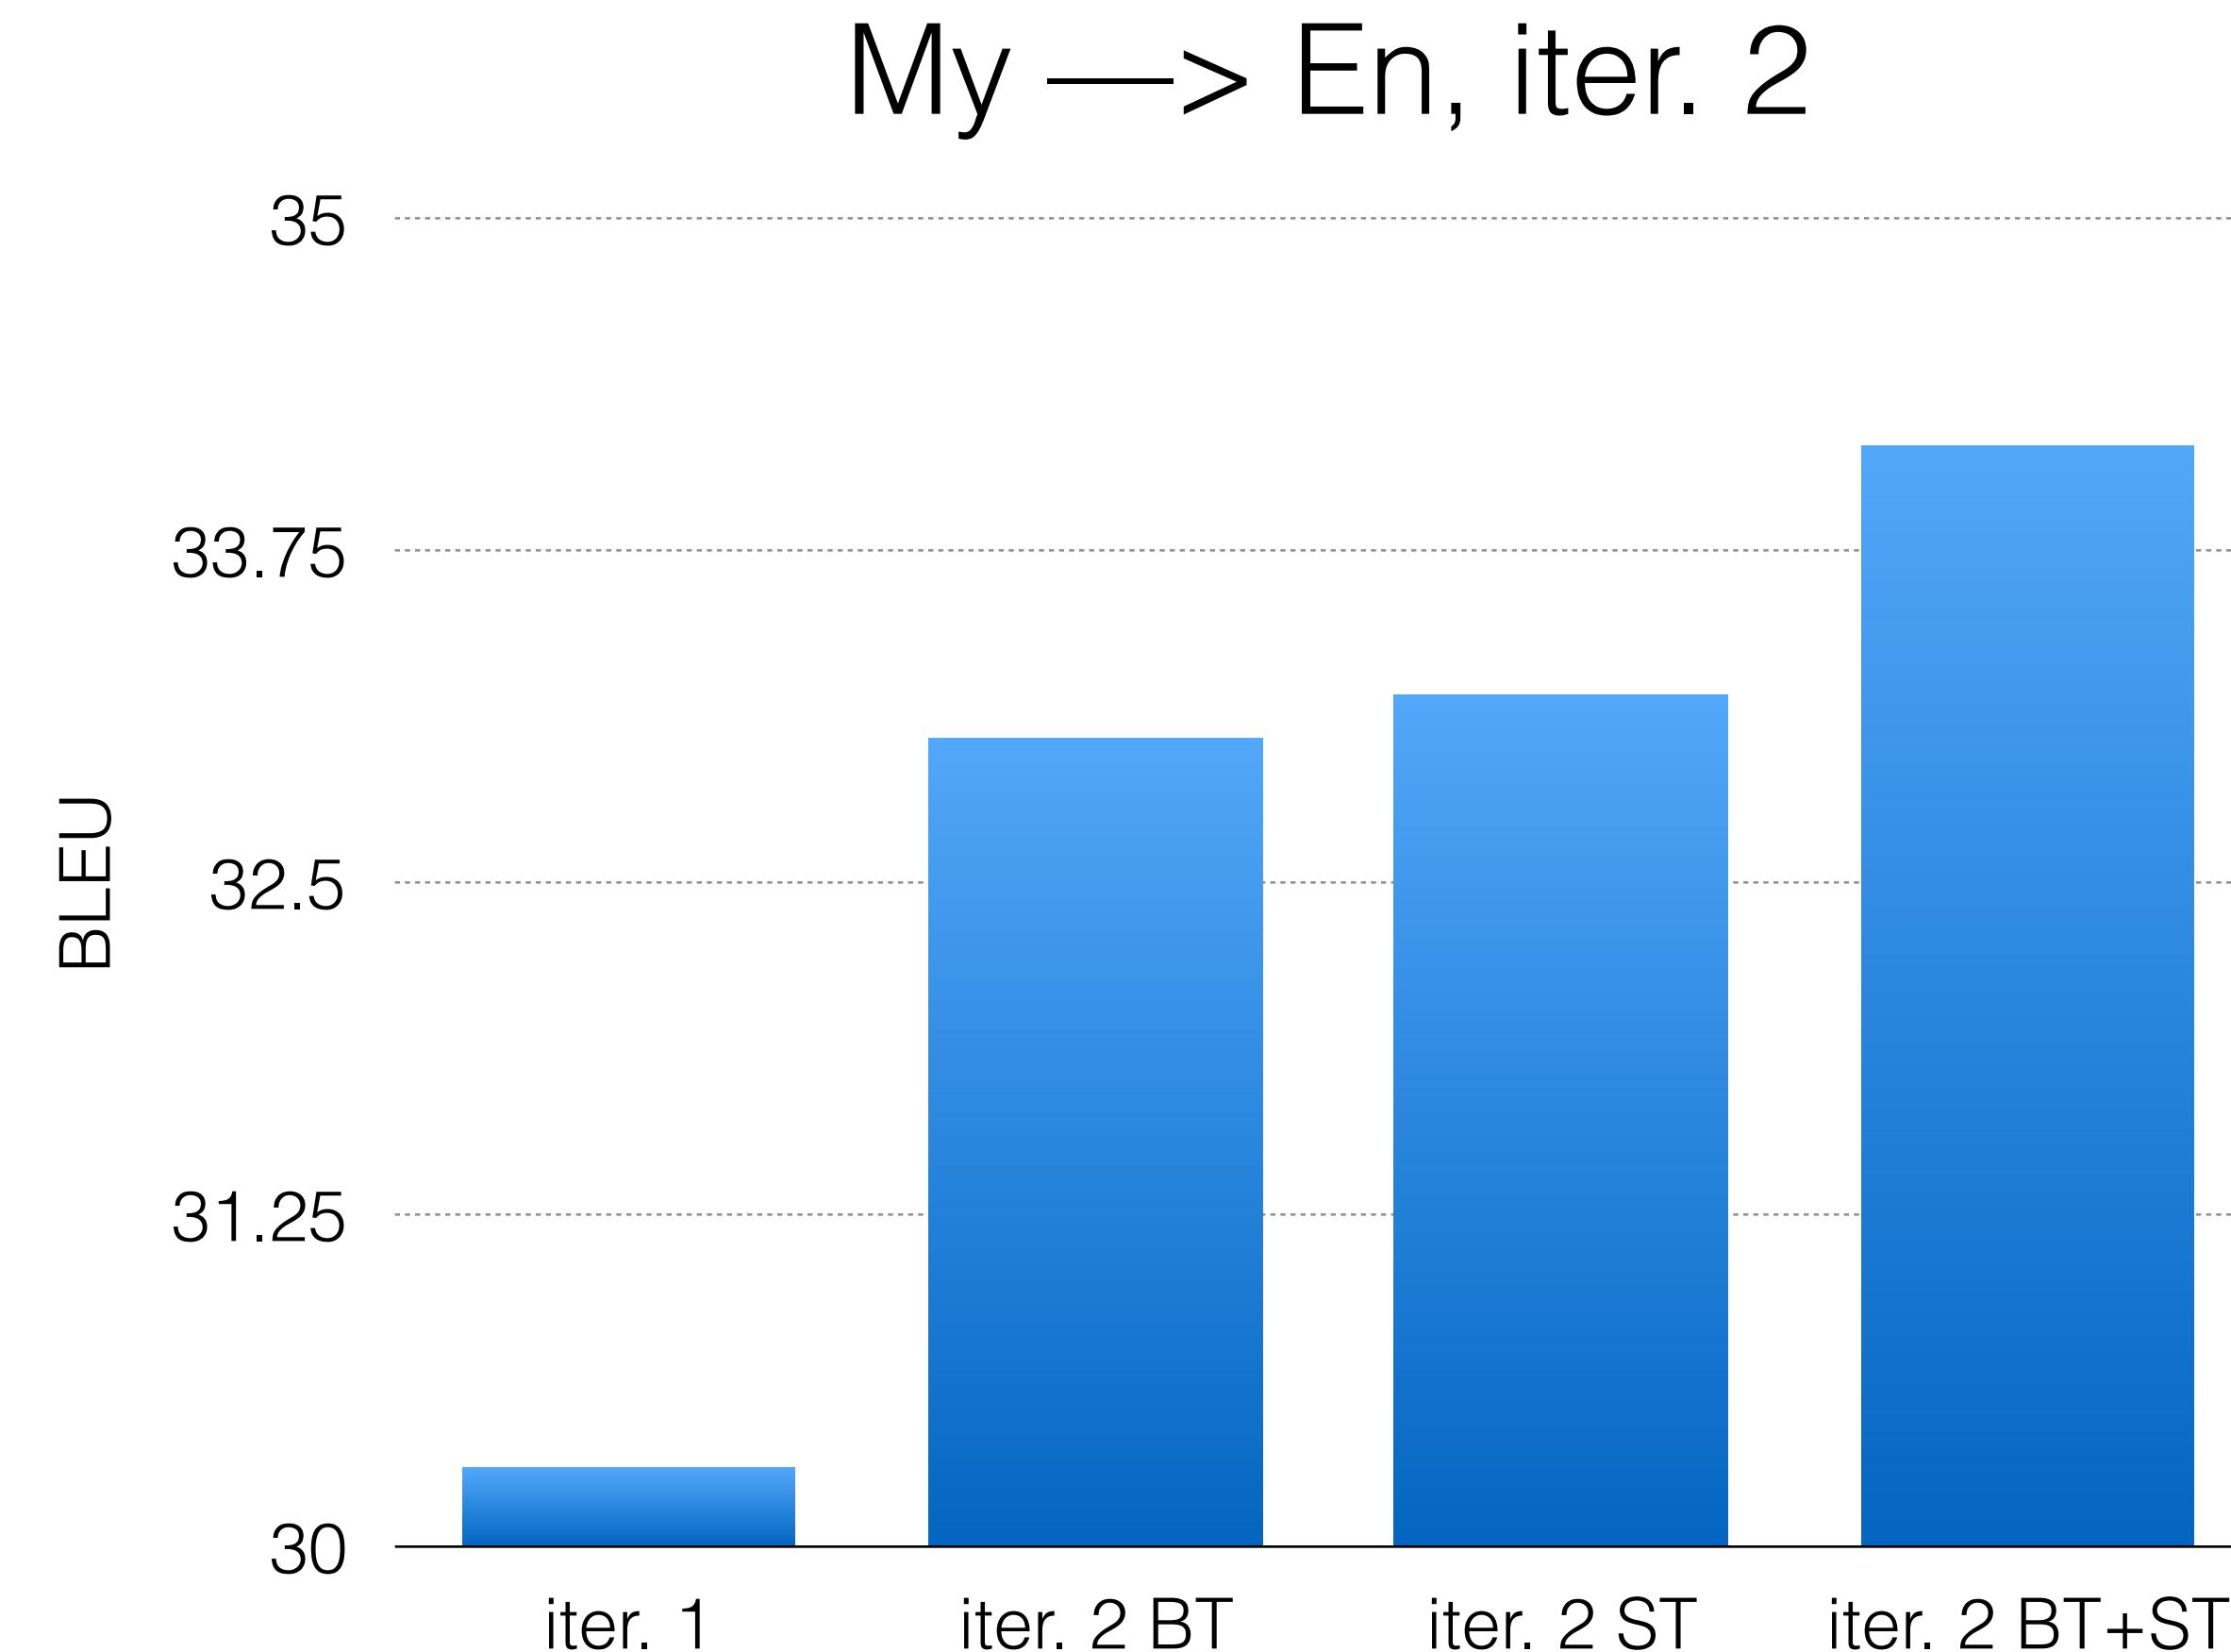
My → En



En → My

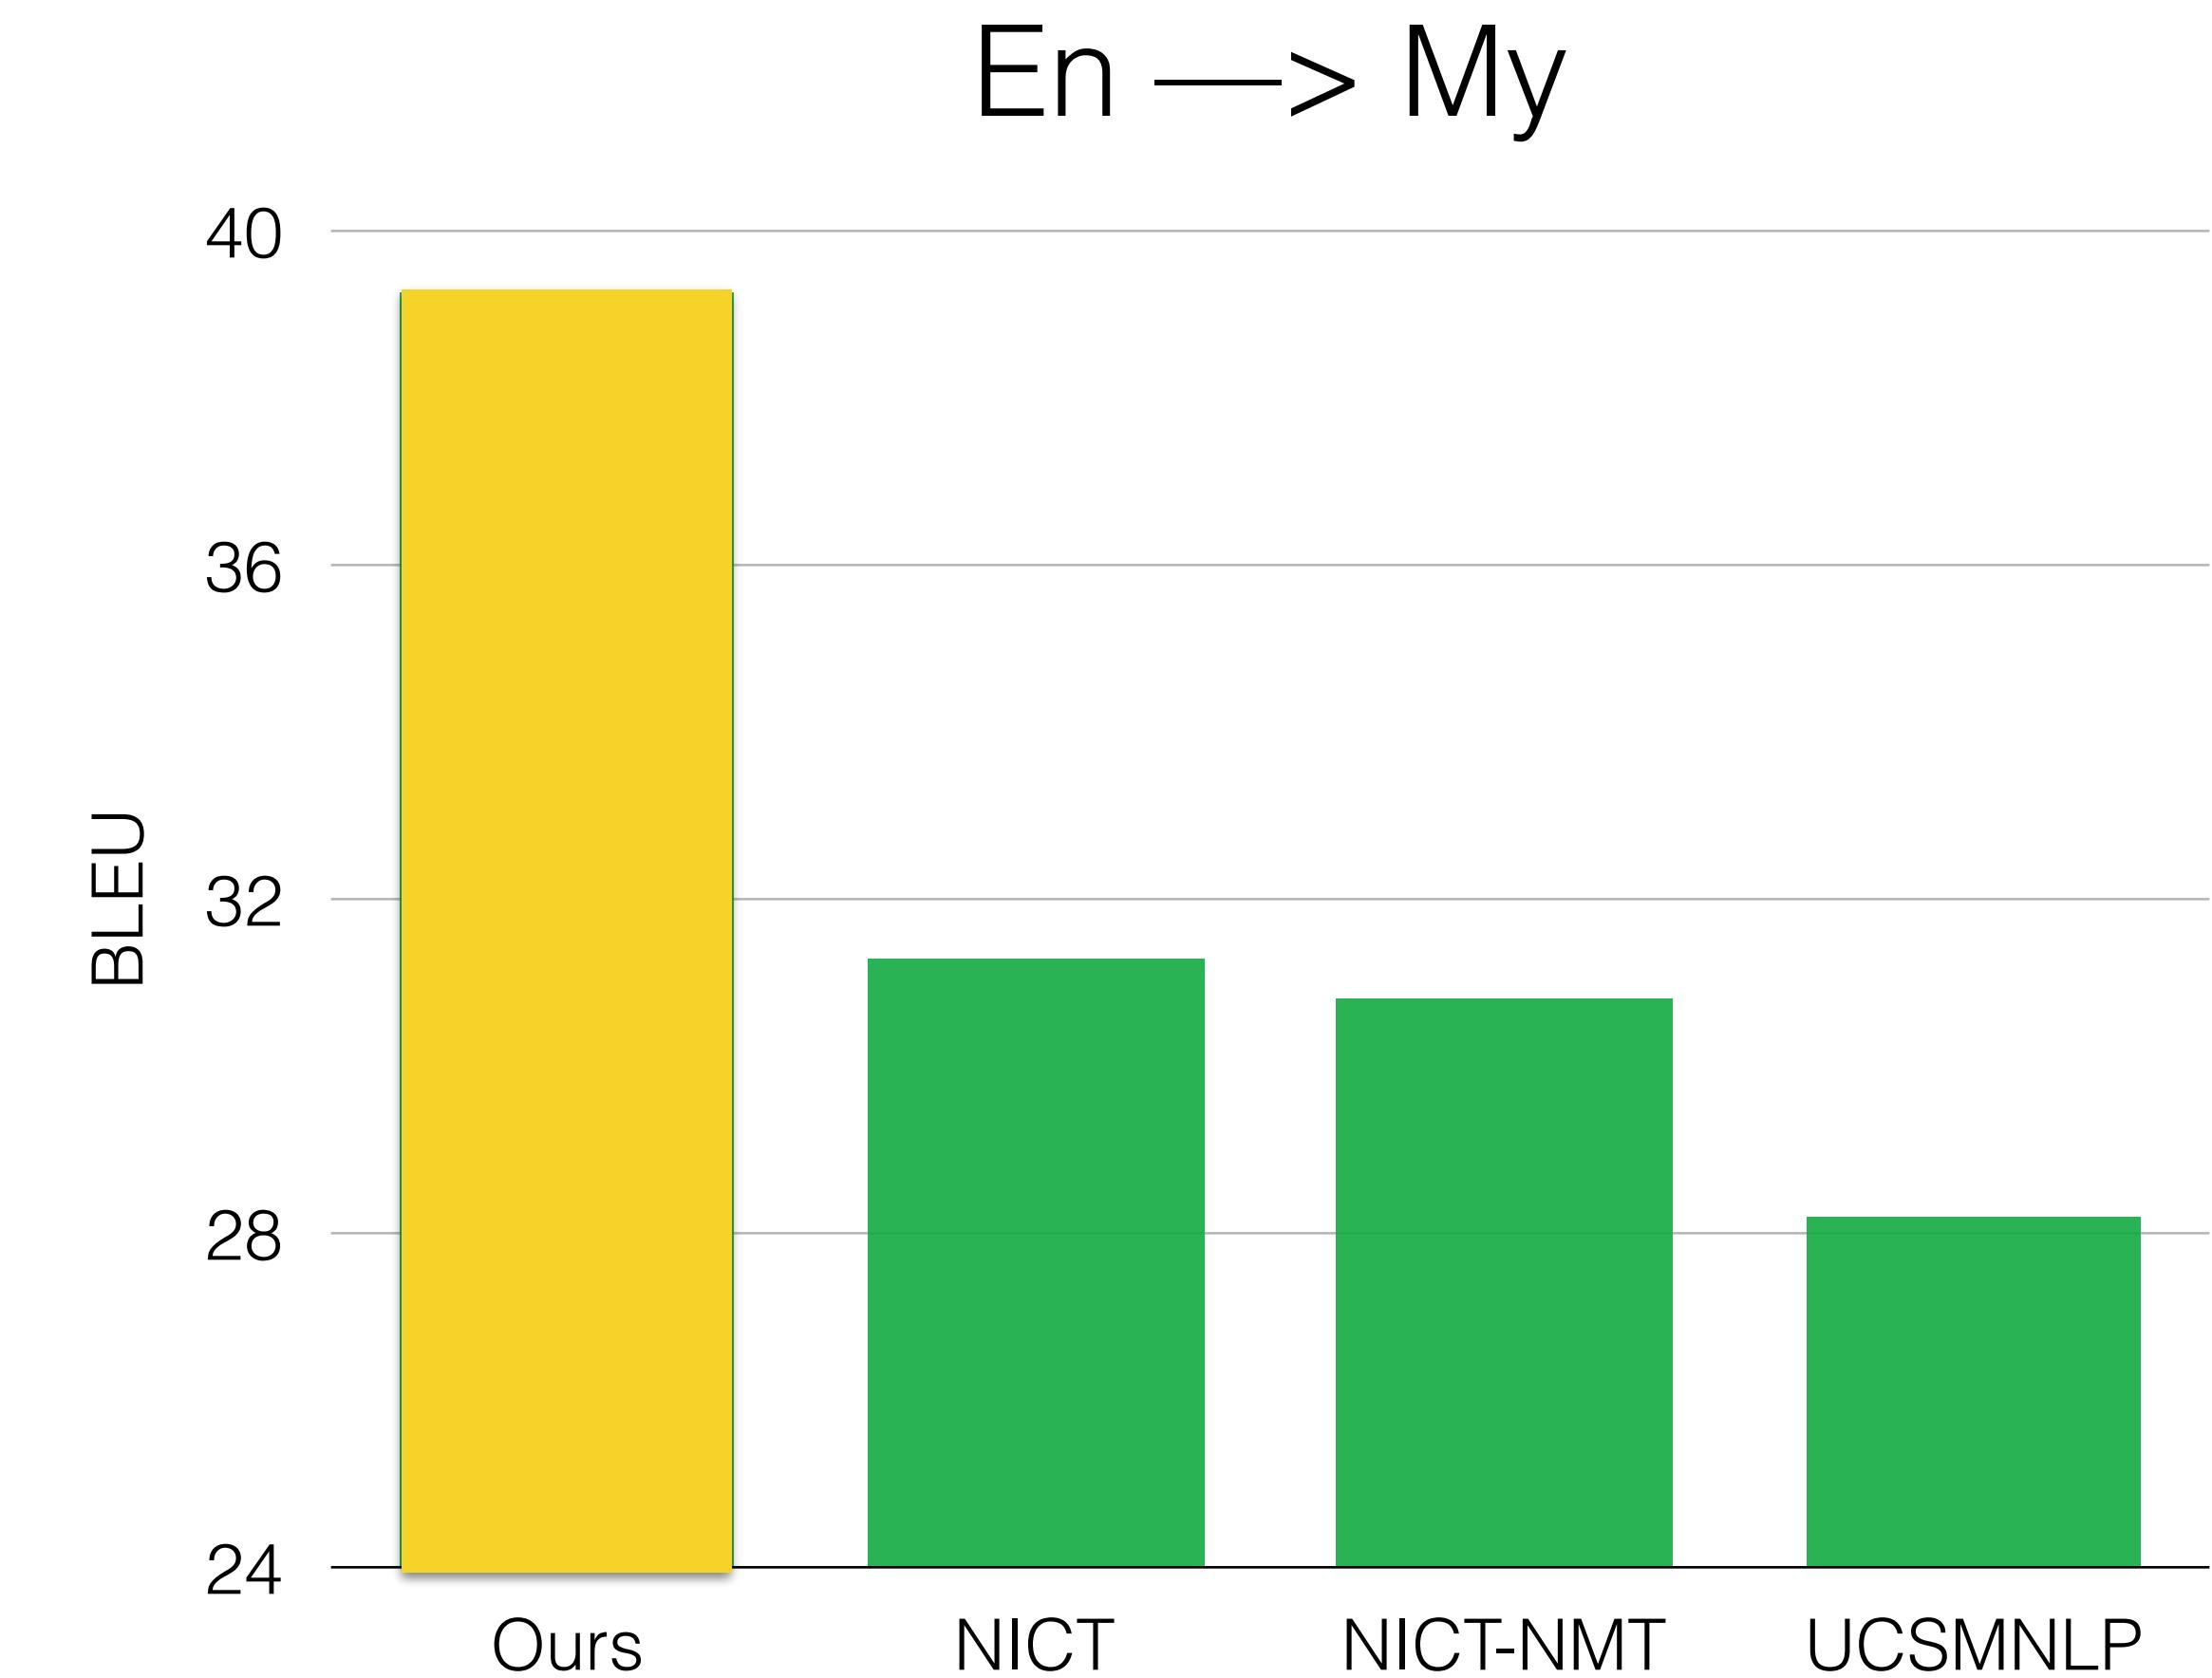
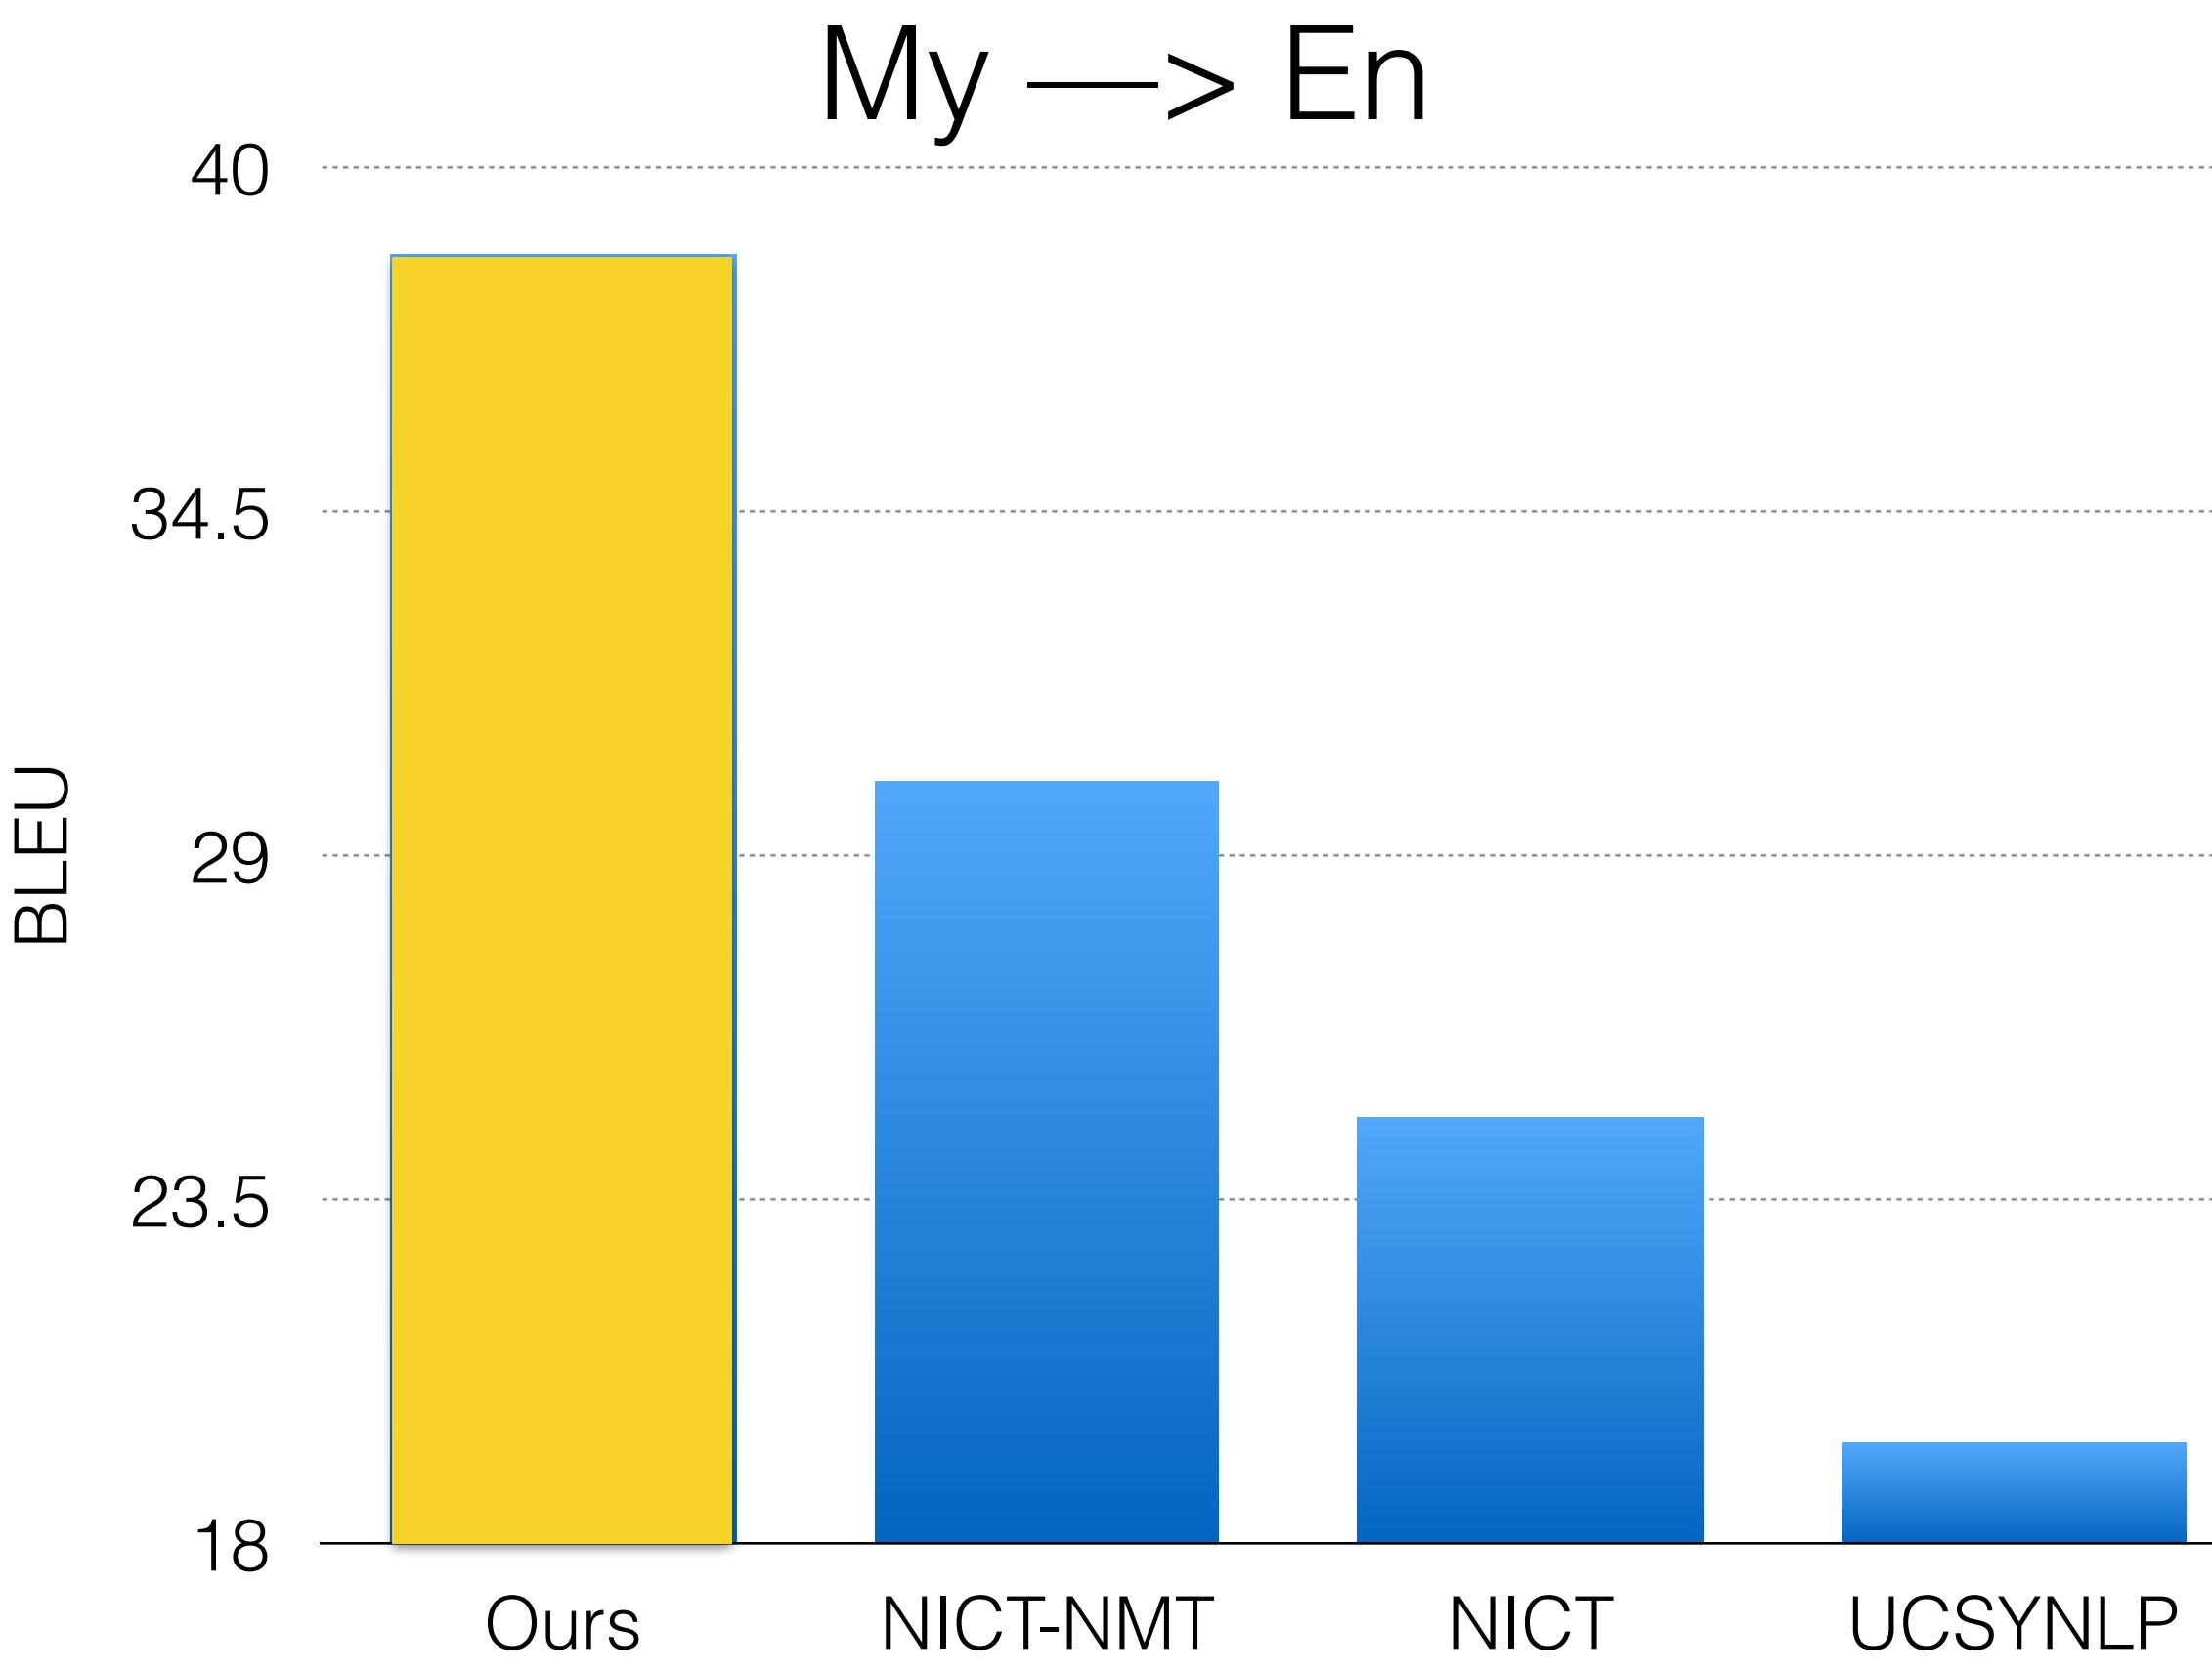


Results: BT vs ST vs BT+ST



Final Results of 2019 Competition

+8 BLEU compared to second best



Demo (Burmese → English)



ဗြိတ်သူ နိုင်ငံရေး အကျပ်အတည်း

① 4 စက်တင်ဘာ 2019

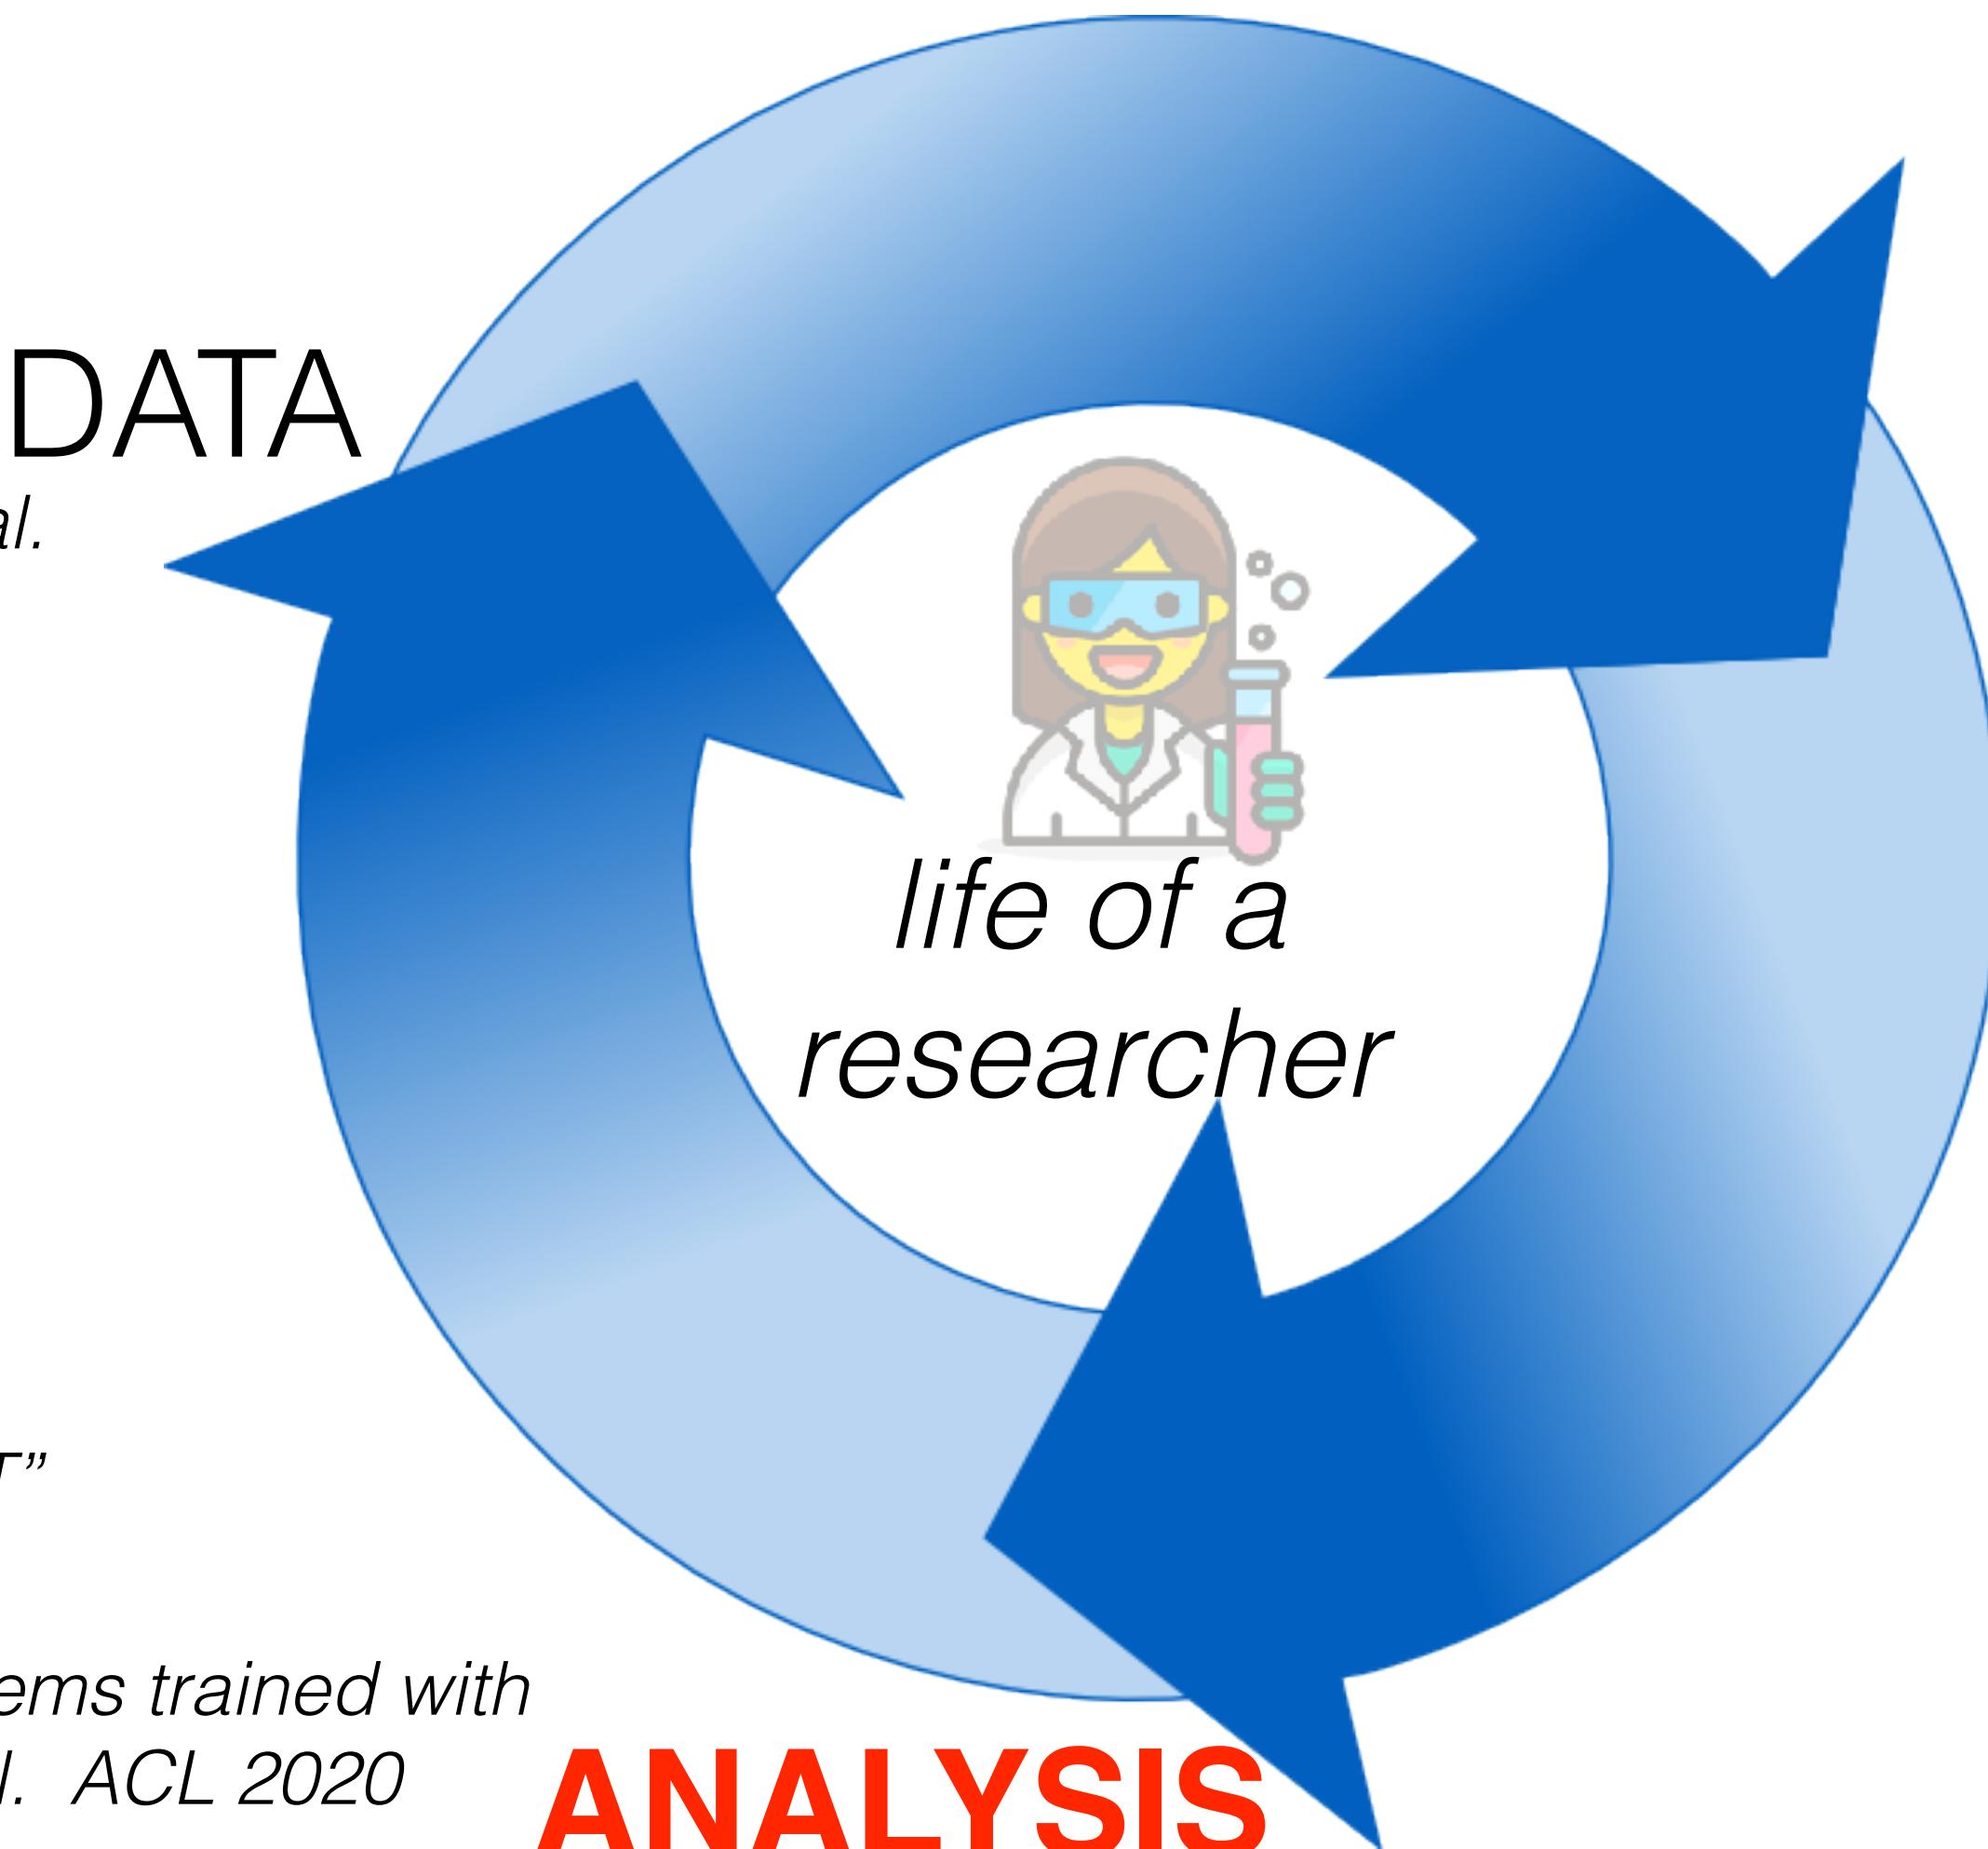
f ၁၀၂၂၂

ဗြိတ်သူ ဝန်ကြီးချုပ် ဘောဂစ်ဂျွန်ဆင်ဟာ ရွှေးကောက်ပွဲ ကျင်းပဖို့ အတွက် ပါလီမန်ကို ဒီဇန်၊
တောင်းဆို ဖို့ များနေတယ်လိုအနုံမျိုးထားကြပါတယ်။

Conclusion so far...

- Iterative back-translation, multi-lingual training work remarkably well.
- By feeding more data (BT, ST, pre-training, multi-lingual training) we can afford training bigger models. Bigger models train on more data generalize better.
- Low-resource MT requires big compute! Remember that about 100 monolingual sentences give the same training signal as a single pair of parallel sentence.

Outline



"The FLoRes evaluation for low resource MT:..." Guzmán, Chen et al. 'EMNLP 2019'

"Analyzing uncertainty in NMT"
Ott et al. ICML 2018

"On the evaluation of MT systems trained with back-translation" Edunov et al. ACL 2020

"The source-target domain mismatch problem in MT" Shen et al. arXiv 1909.13151 2019

MODEL

"Phrase-based & Neural Unsup MT"
Lample et al. EMNLP 2018

"FBAI WAT'19 My-En translation task submission" Chen et al., WAT@EMNLP 2019

"Massively Multilingual NMT" Aharoni et al., ACL 2019

"Multilingual Denoising Pre-training for NMT"
Liu et al., arXiv 2001:08210 2020

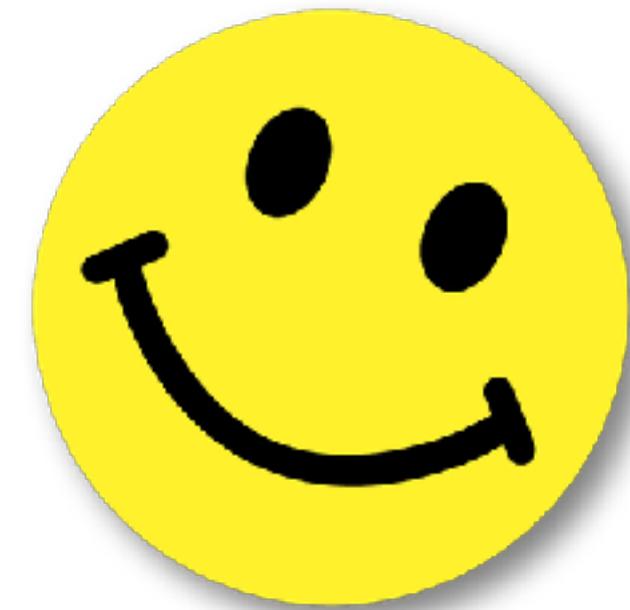
ANALYSIS

Simulating Low-Resource MT

Simulating low-resource MT with a high resource language:
using *EuroParl* data with **20K** parallel sentences and **100K** monolingual target sentences.

EuroParl Fr → En	
only parallel data	30.4 BLEU
parallel data + BT	33.8 BLEU

+3.4 BLEU!



A Worrisome Finding

BT sometimes yields very mild improvements.

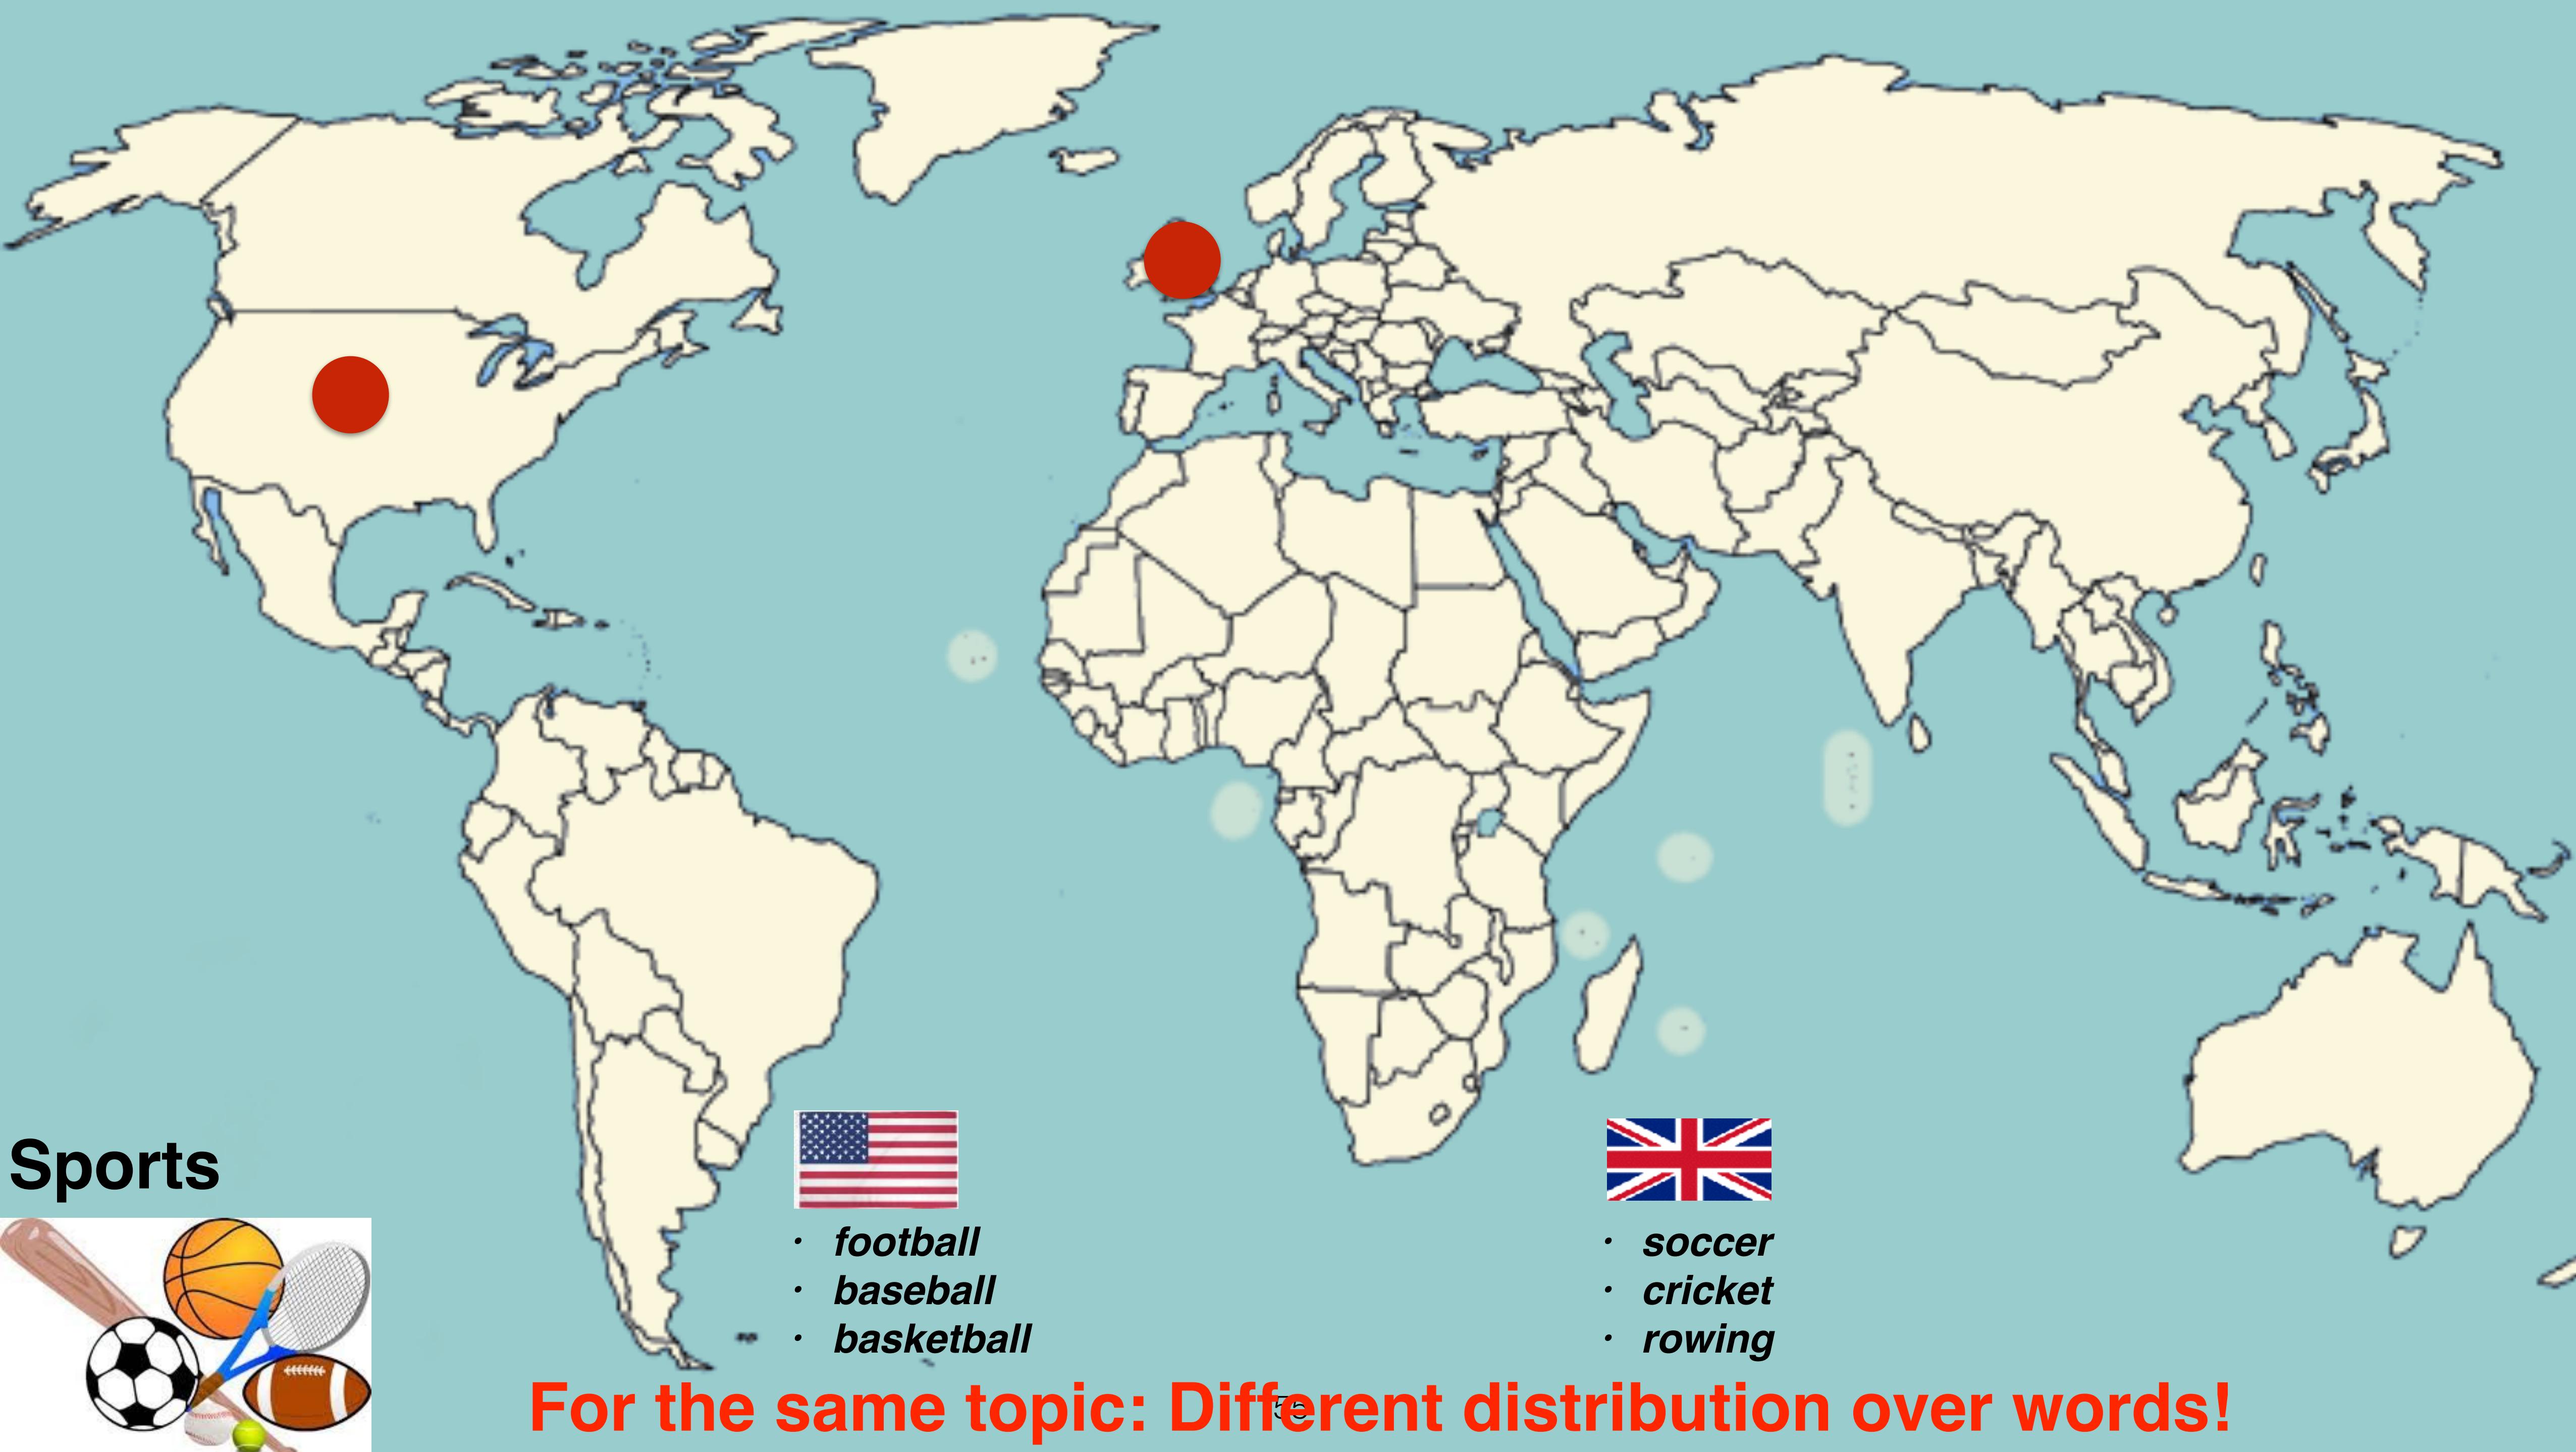
Example

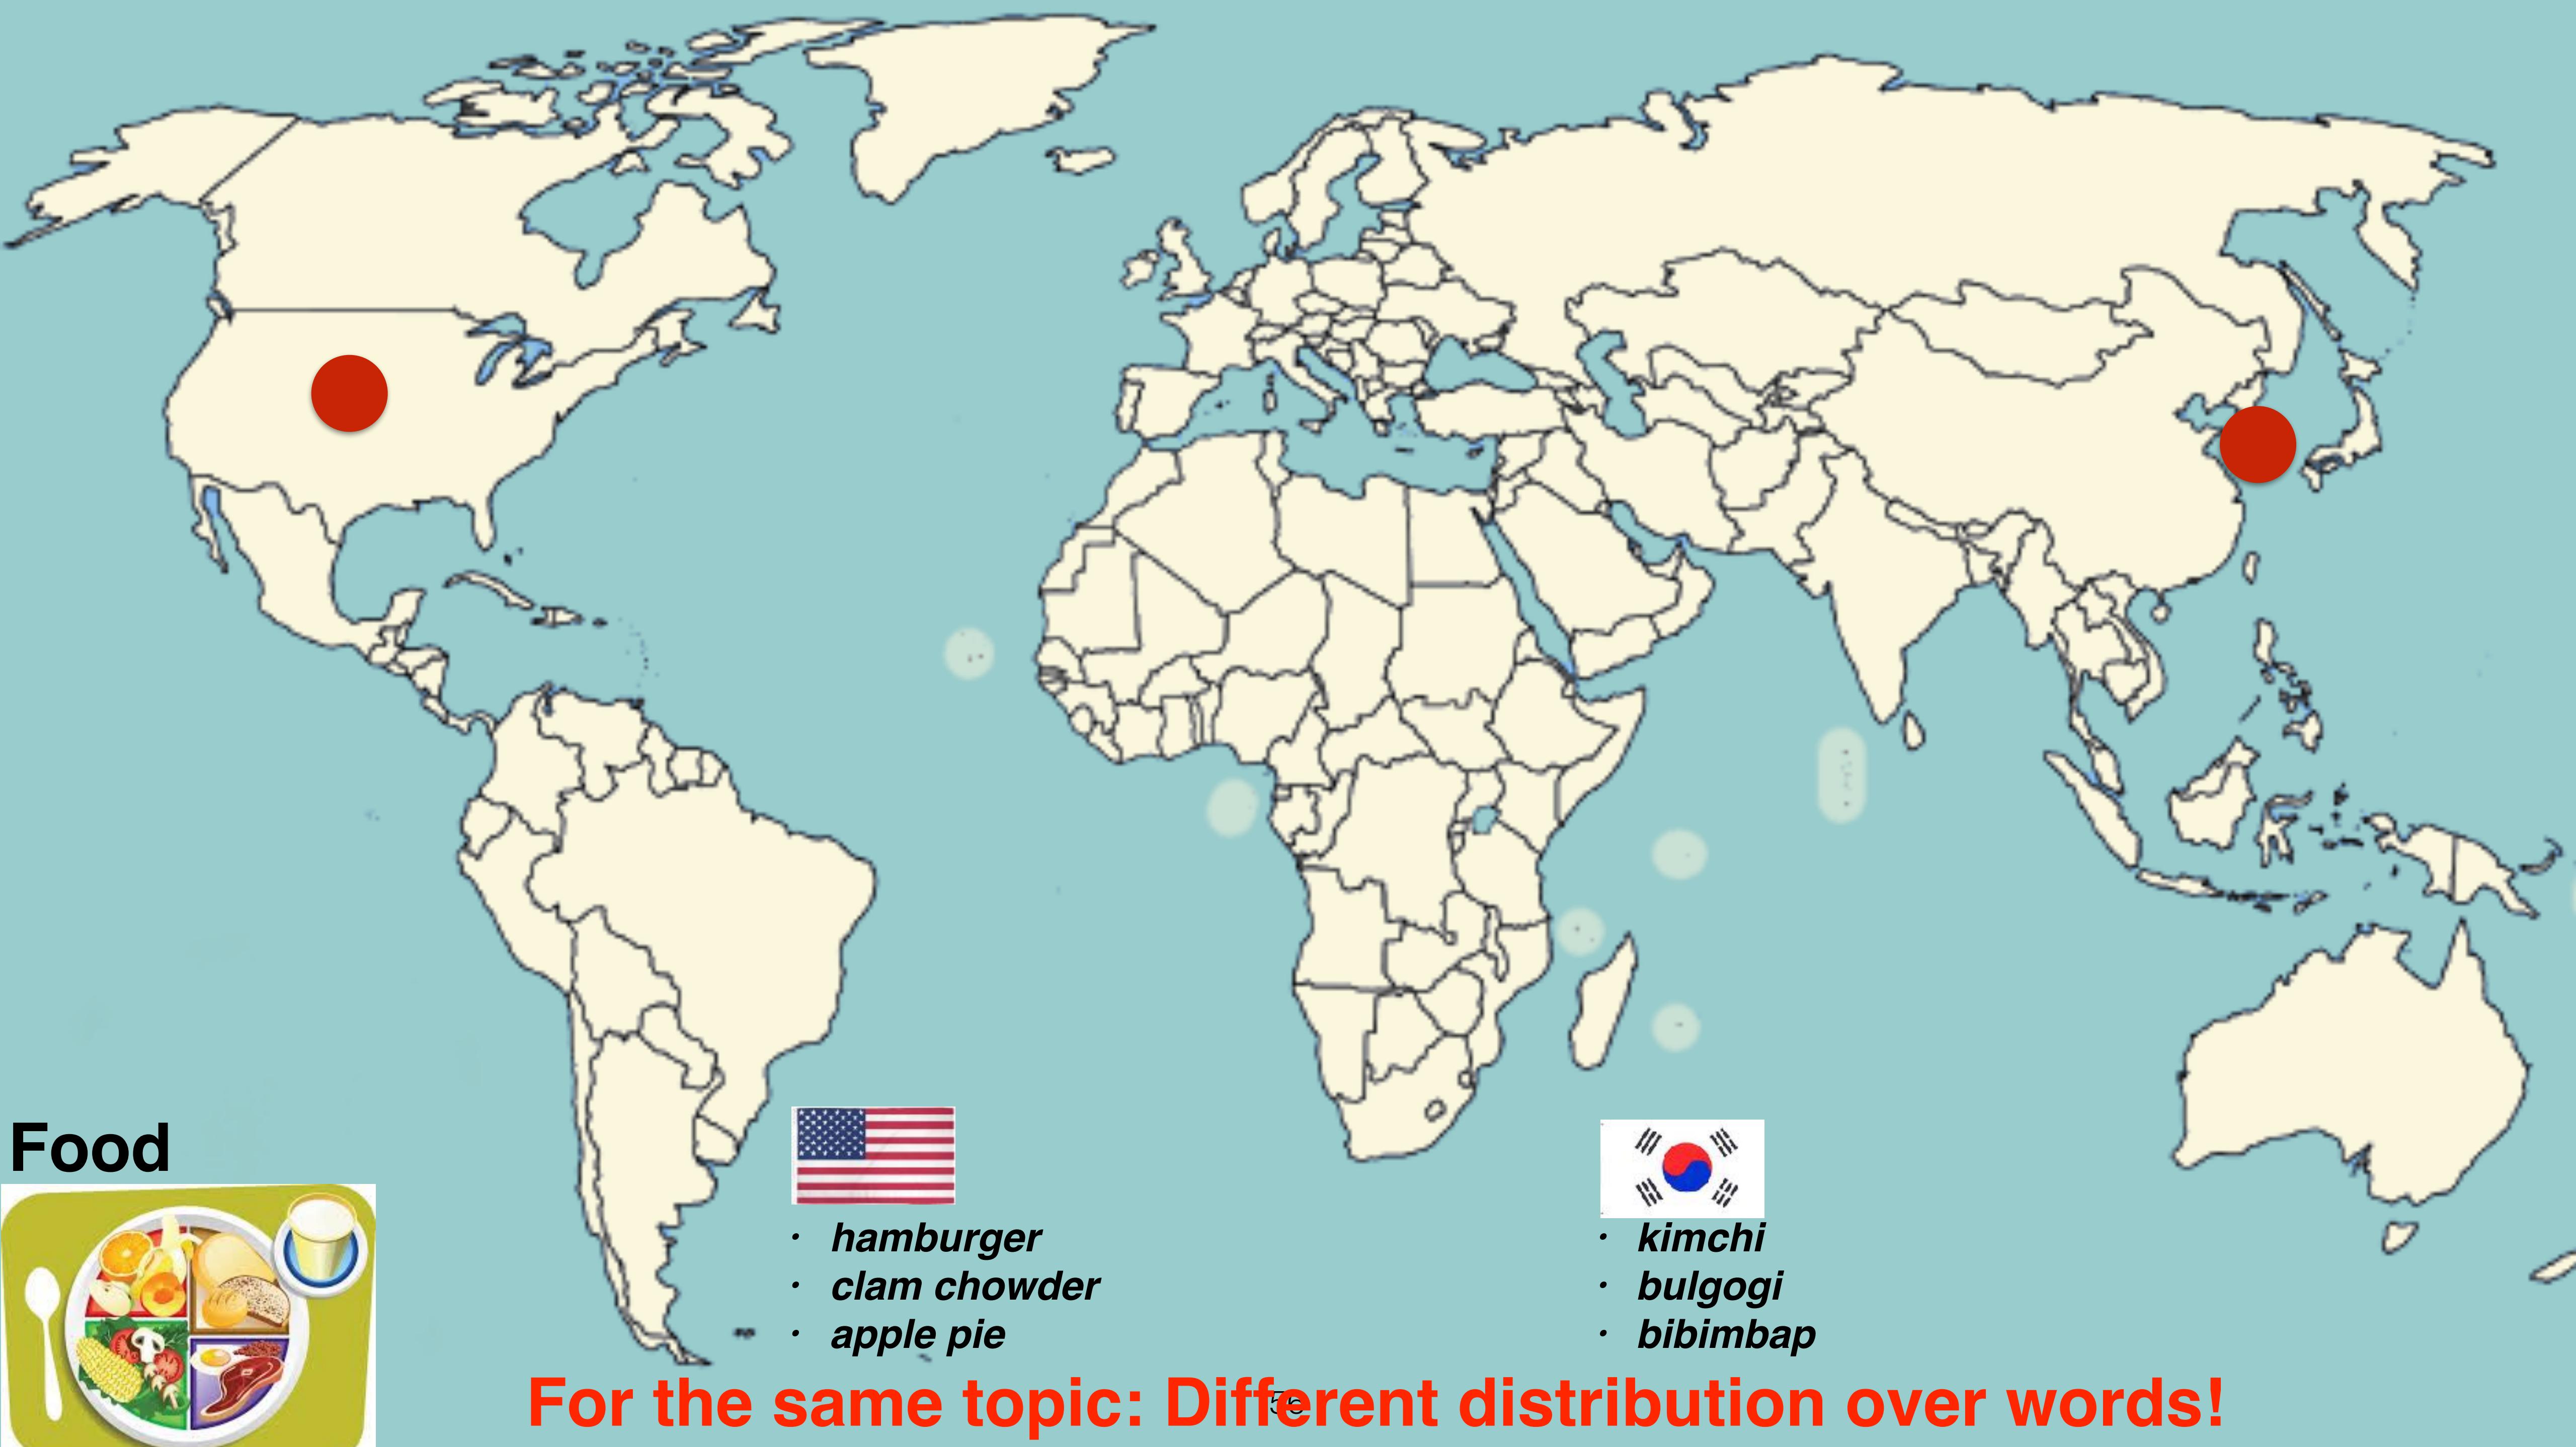
<i>FB public posts En—>My</i>	
only parallel data	15.2 BLEU
parallel data + BT	15.3 BLEU

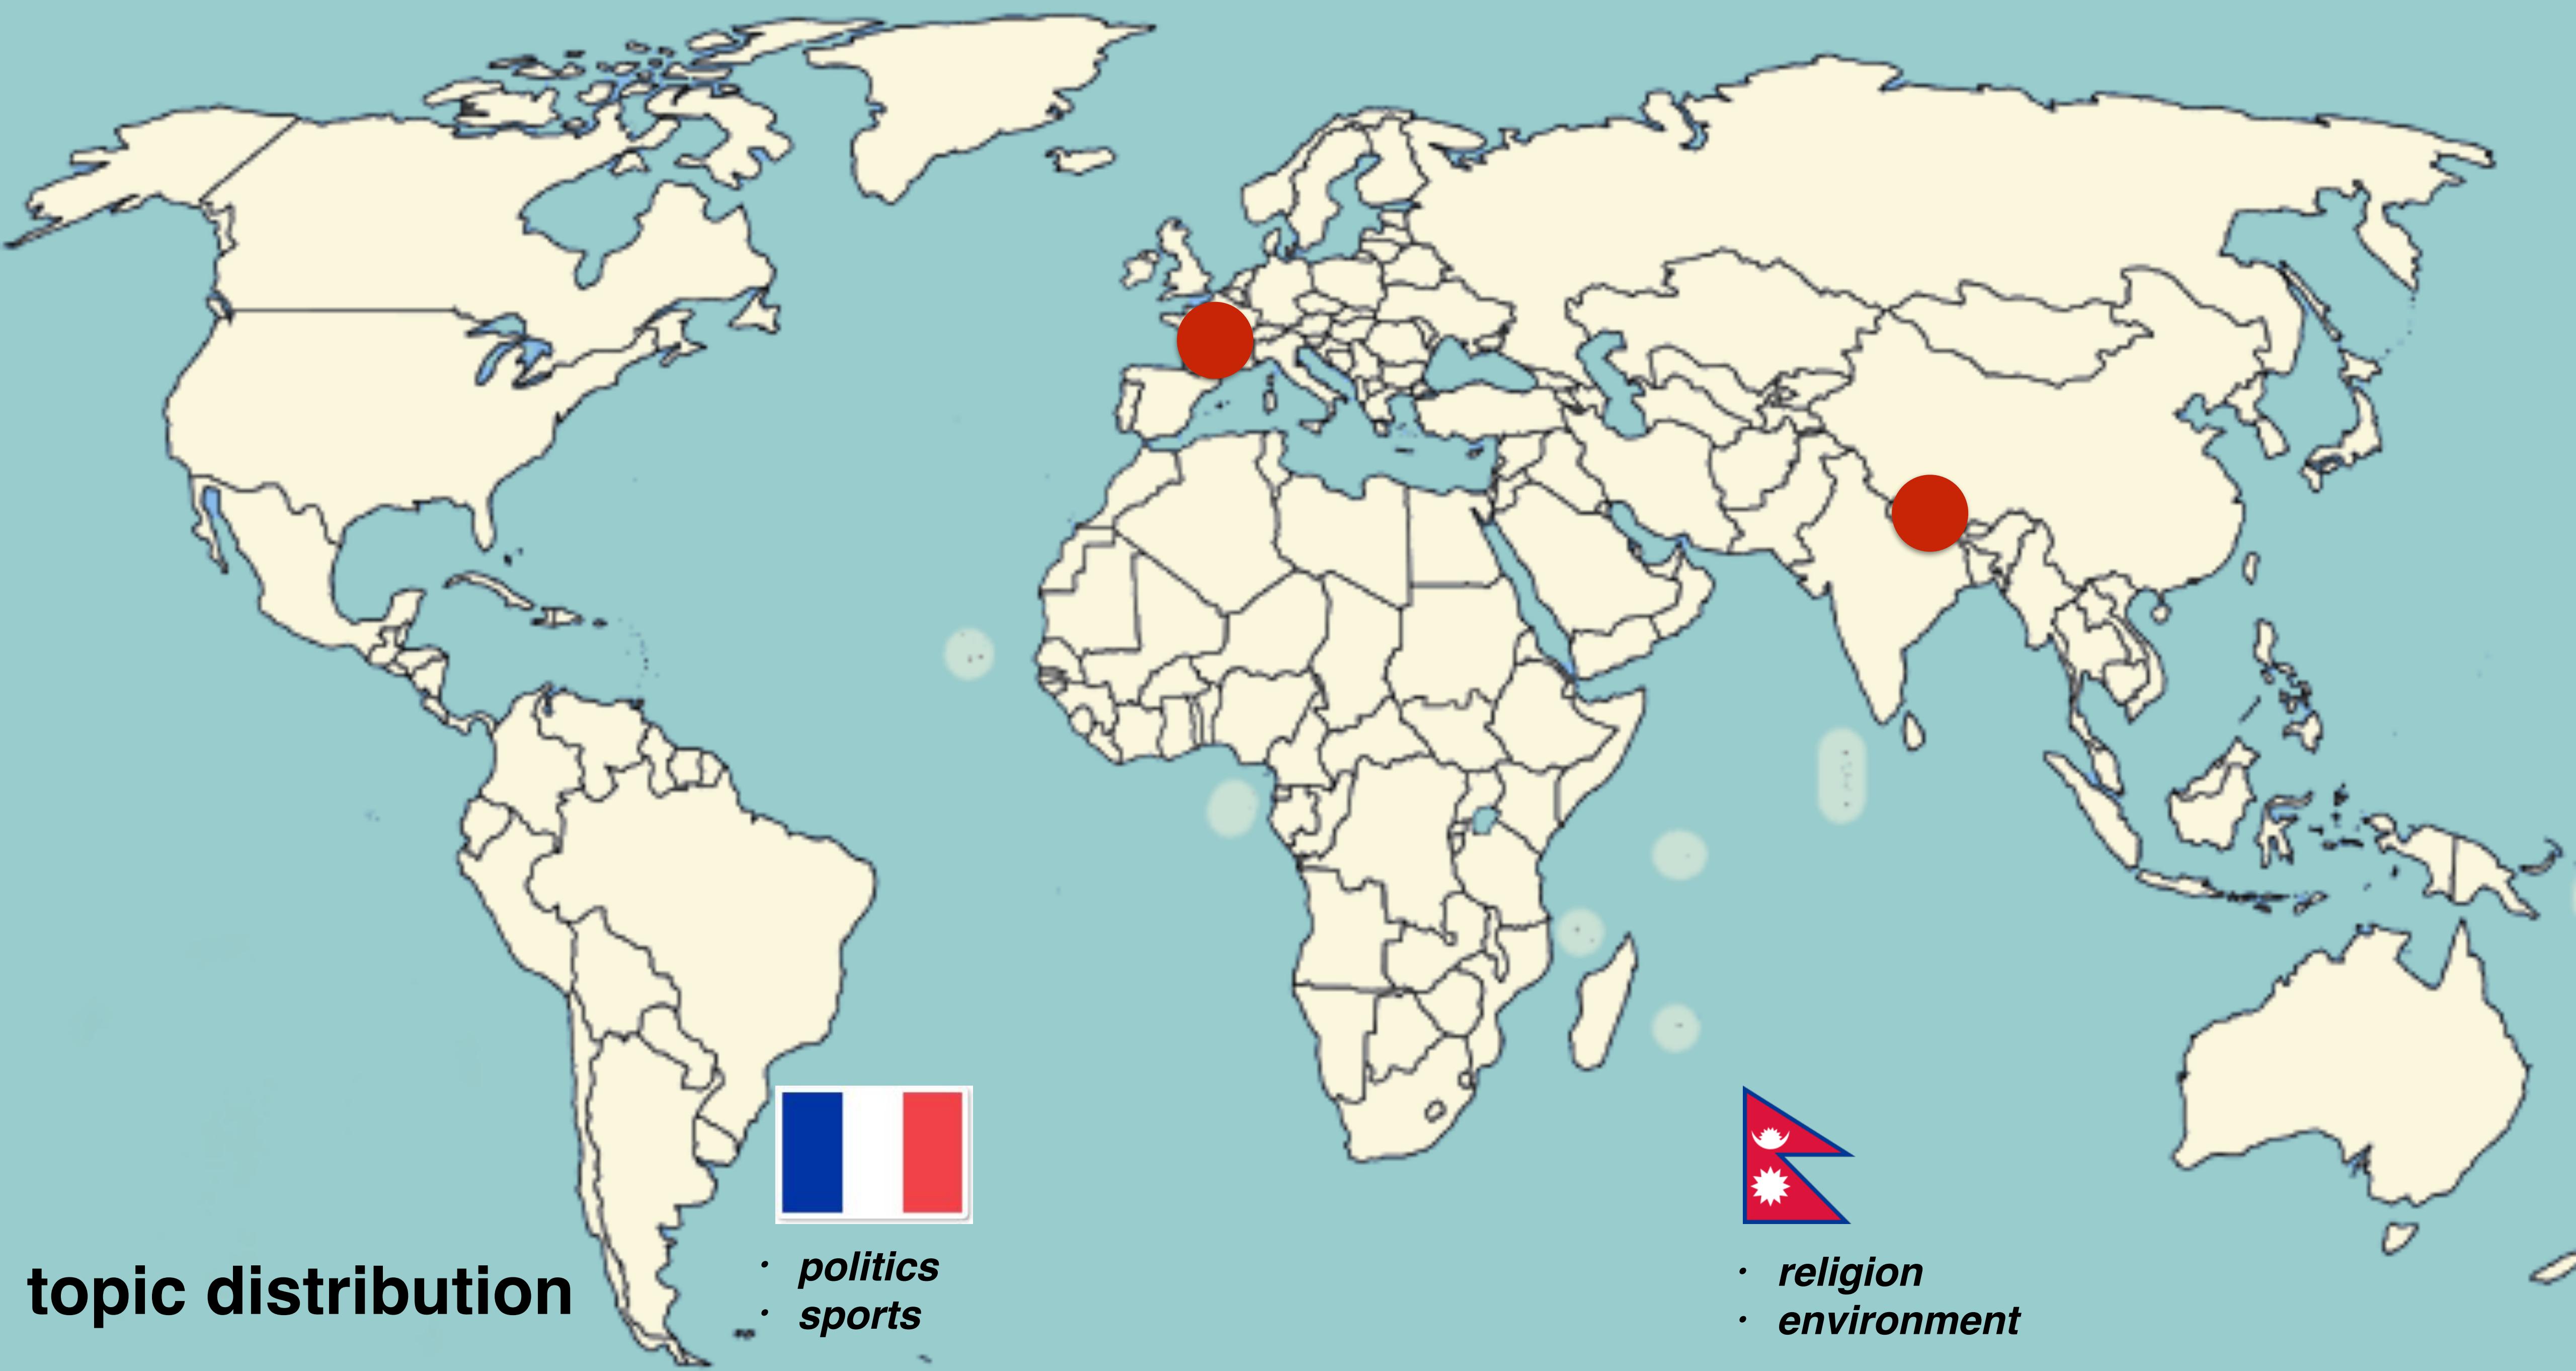
+0.1 BLEU!



Why is BT not working as well?







Domains differ in the topic distribution

Examples from FLoRes

Si-En

අධි යාපනයෙන් පසු හෝ පවුලේ යුතුකම් ඉටු කරන්නට හෝ රෝග තත්ත්වයන් තිසා සිංහල උපසම්පෑළවෙන් නිතරම ඉවත් වෙති.

After education priests leave ordination in order to fulfill duties to the family or due to sickness.

තරතත , ගාරීරක හිංසනය , දේපල භාතිය , පහර දීම සහ මරාදැමීම මෙම දියුවම්ය .

Threatening, physical violence, property damage, assault and execution are these punishments.

Wikipedia in Sinhala has different topic distribution.

En-Si

In Serious meets, the absolute score is somewhat meaningless.

සැබැං තරග වලදී ලකුණු සැසදීම තේරුමක් තැනි ක්රියාවකි .

Iphone users can and do access the internet frequently, and in a variety of places.

අයිලෝත් භාවිත කරන්නන්ට නිතරම සහ විවිධ ස්ථානවලදී අත්තරජාලයට පිවිසිය හැකිය .

original

translation

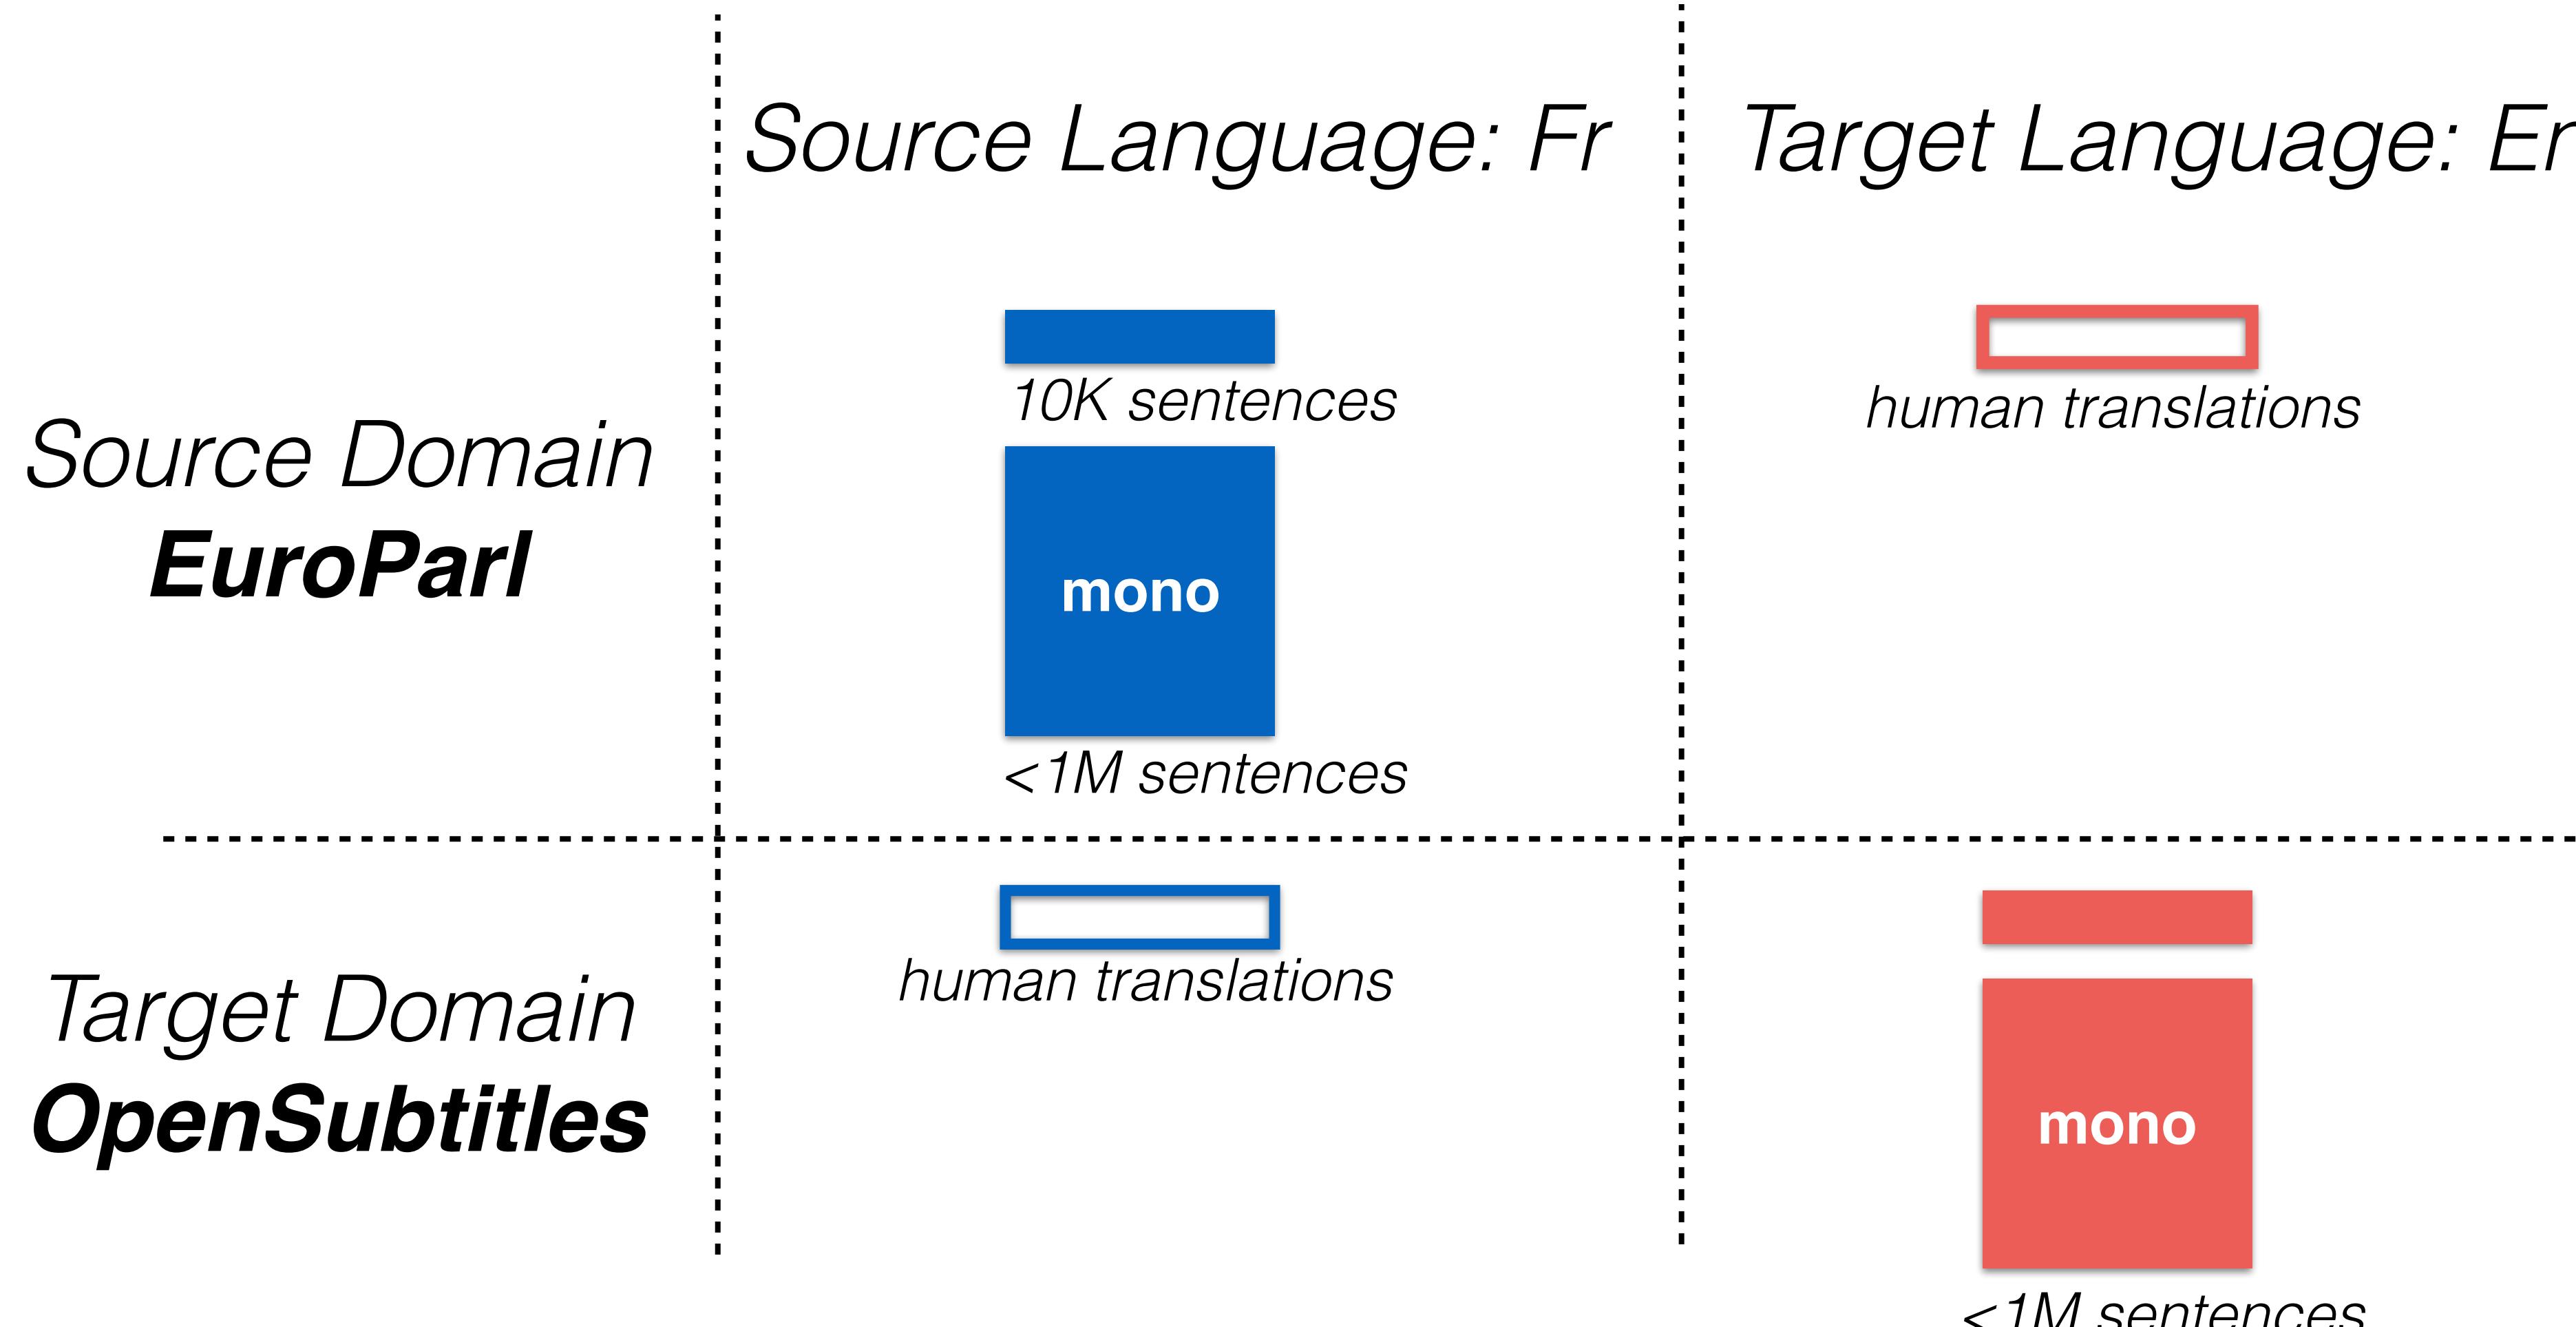
Source-Target Domain Mismatch (STDM)

- Def.: Content produced in blogs, social networks, news outlets, etc. varies with the geographic location.
- STDM is even more pronounced in low resource MT, where source & target geographic locations are typically farther apart and cultures have more distinct traits.
- STDM manifests itself in terms of: Different distribution of topics, and for the same topic different distribution over words.
- STDM is hard to measure because of the unknown effect of translationese.
- STDM makes the MT problem even harder: source originating data and target originating data are not comparable in general. BT won't work as well. UnsupMT won't work as well.

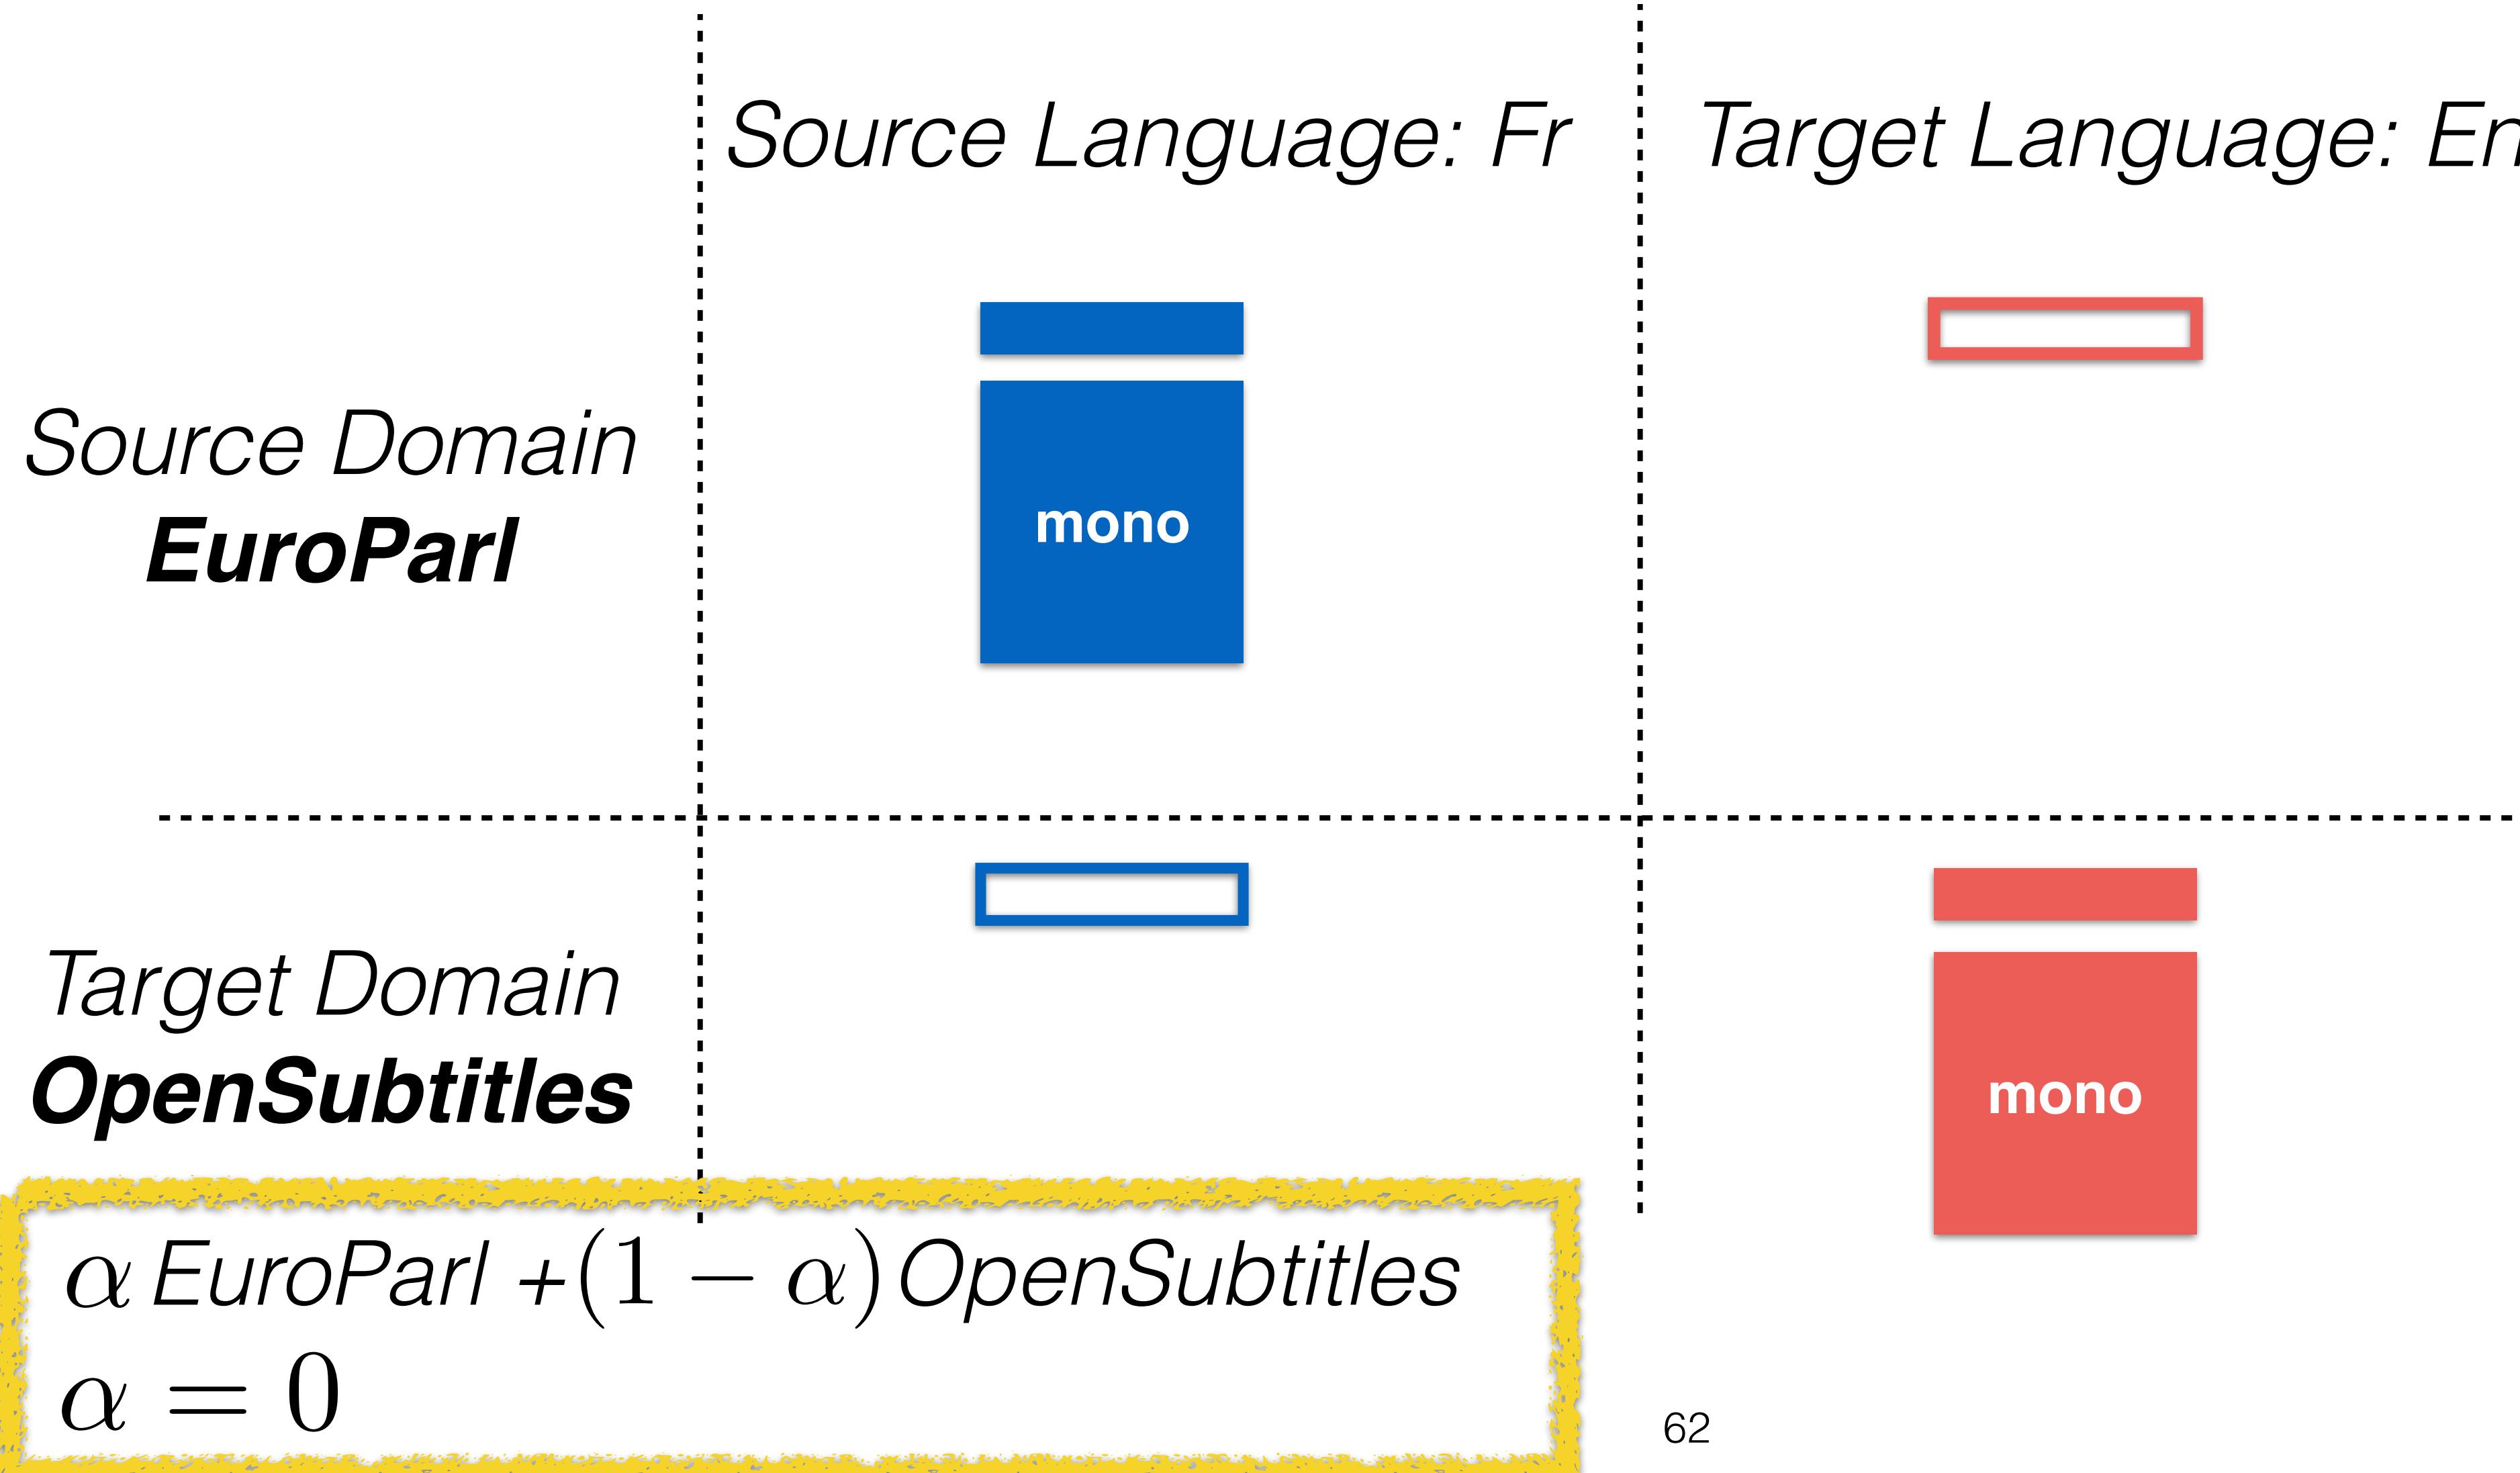
Questions

- Is it true that BT is less effective when there is STDM?
- What baselines shall we consider when there is STDM?
- What are general best practices when there is STDM?
- How to study STDM in a controlled setting?

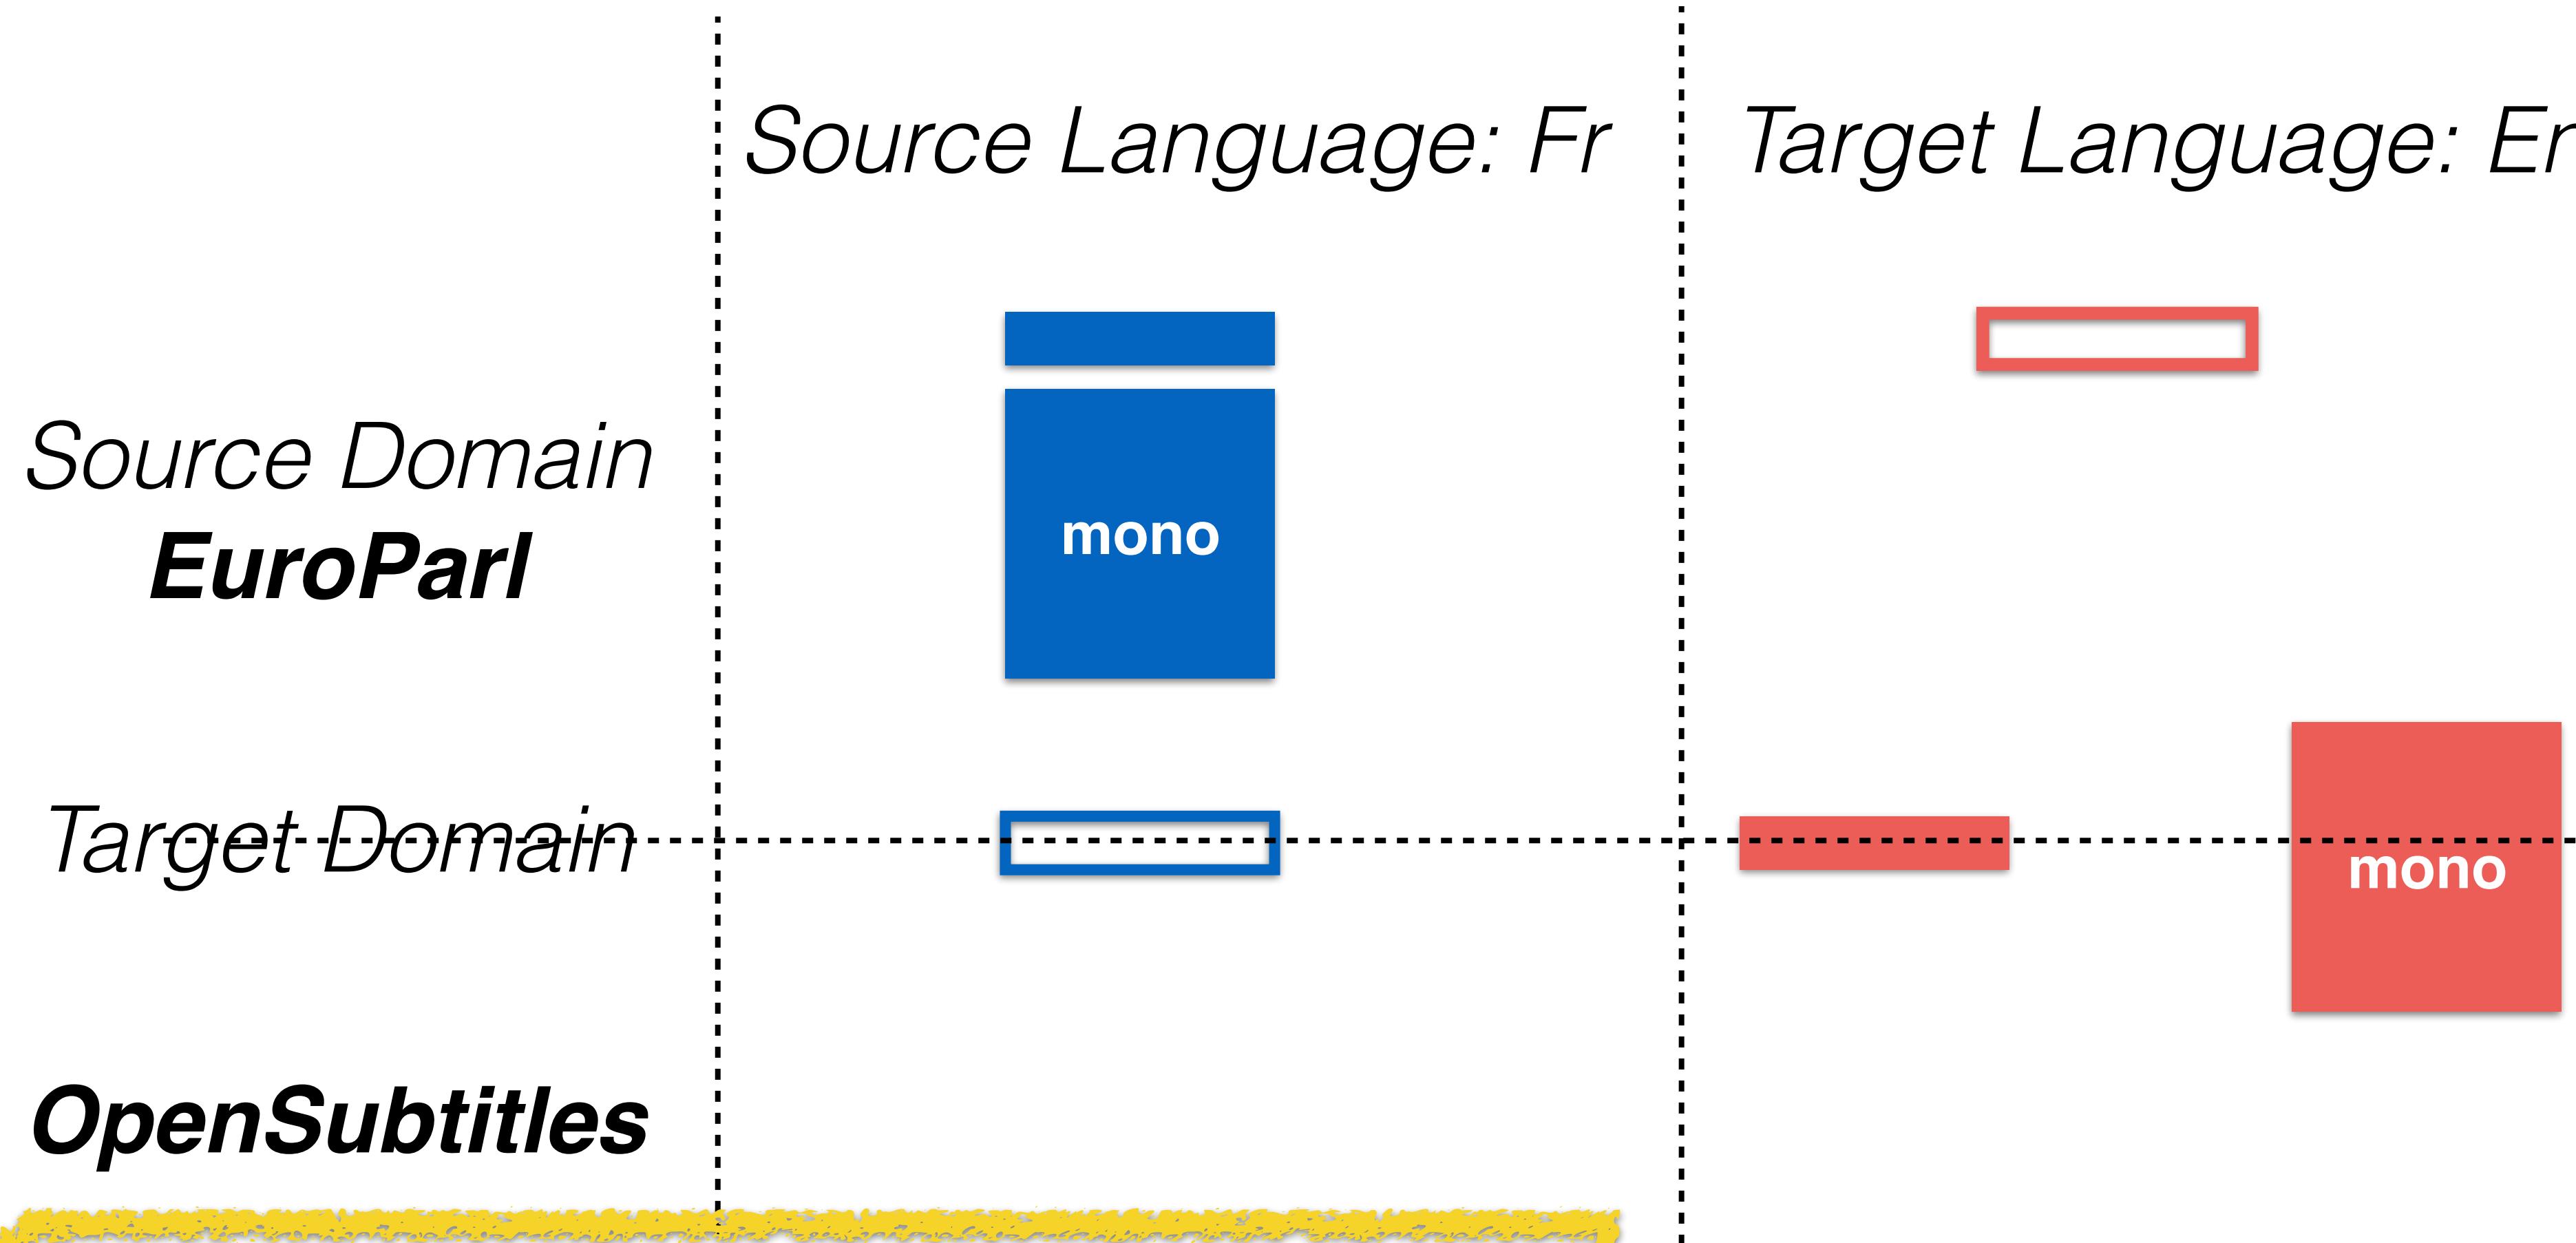
Controlled Setting



Controlled Setting



Controlled Setting



Controlled Setting

Source Domain &
Target Domain
EuroParl

Source Language: Fr



Target Language: En



$\alpha \text{EuroParl} + (1 - \alpha) \text{OpenSubtitles}$
 $\alpha = 1$

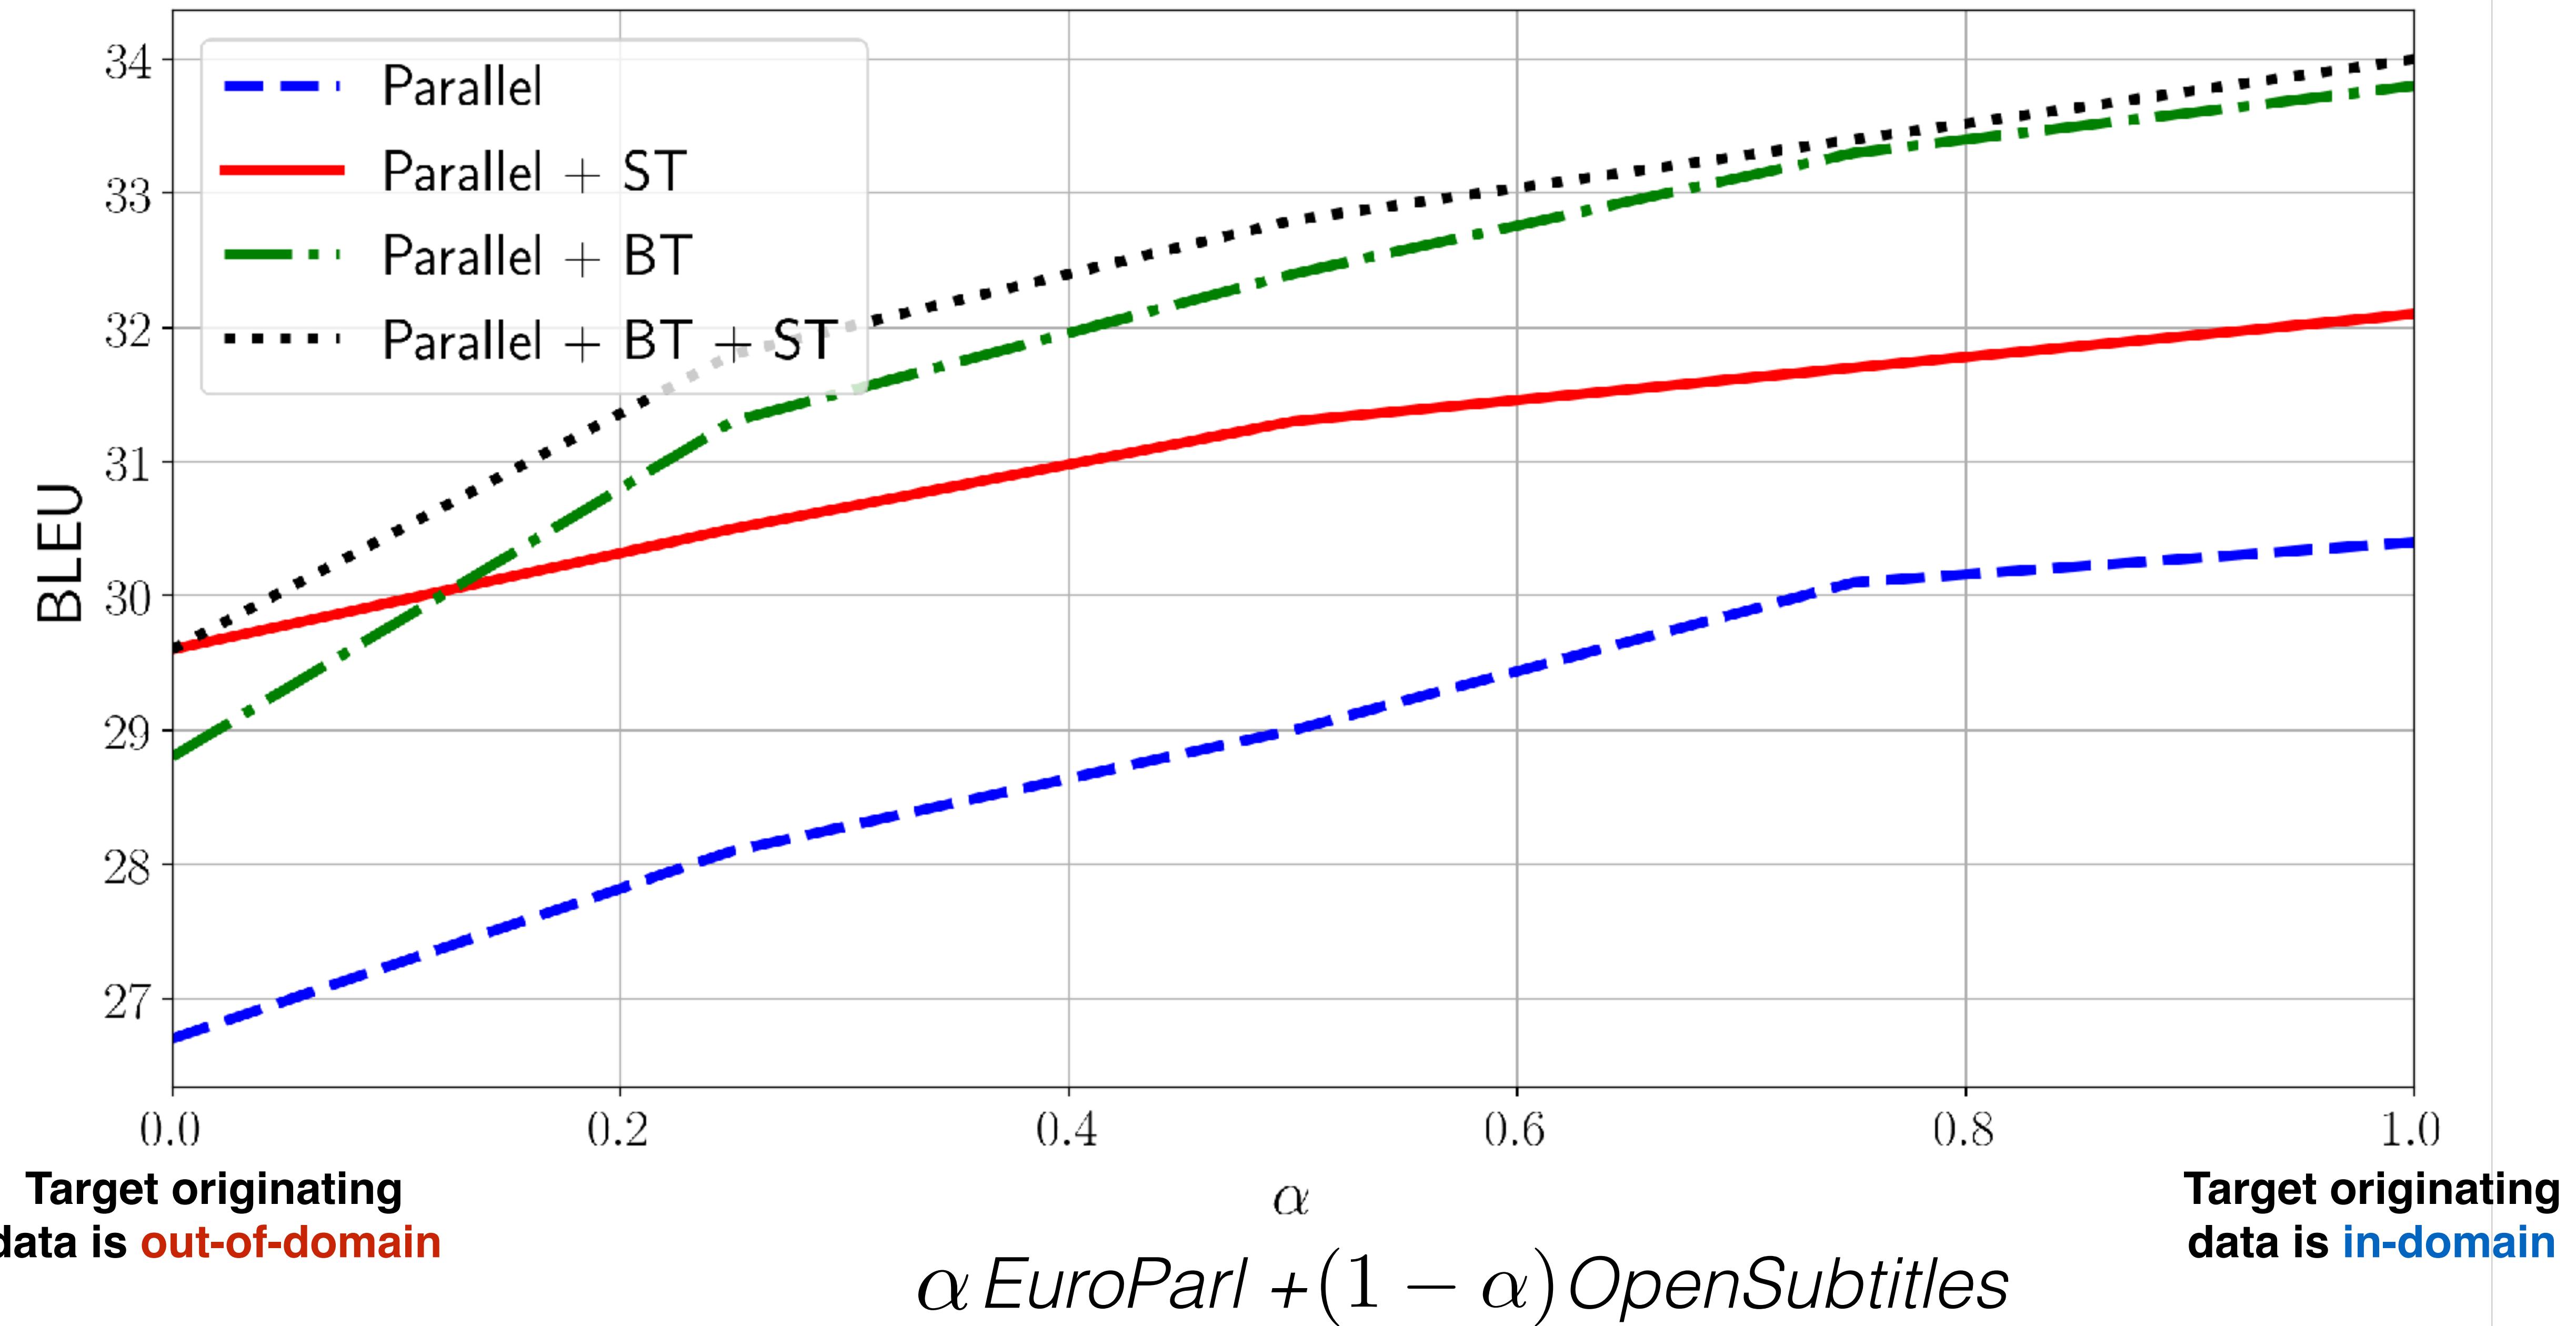
	<i>target mono data</i>	<i>source mono data</i>	<i>in-domain mono data</i>
• Back-Translation:	✓	✗	✗
• Self-Training:	✗	✓	✓

Q.: Is it better to have clean targets but out-of-domain data, or noisy targets but in-domain data?

Q.: What's the effect of amount of parallel/monolingual data?

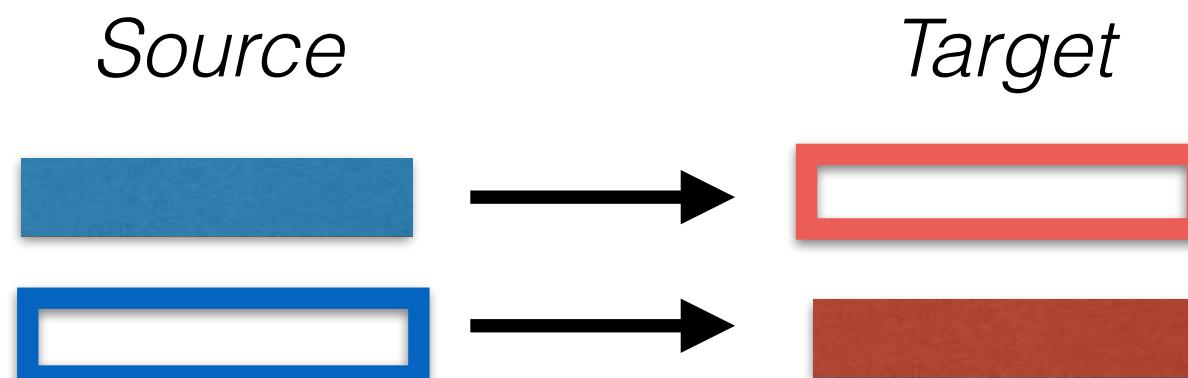
Q.: What's the effect of the quality of the model forward model when training with ST?

Varying Domain of Target Originating Data

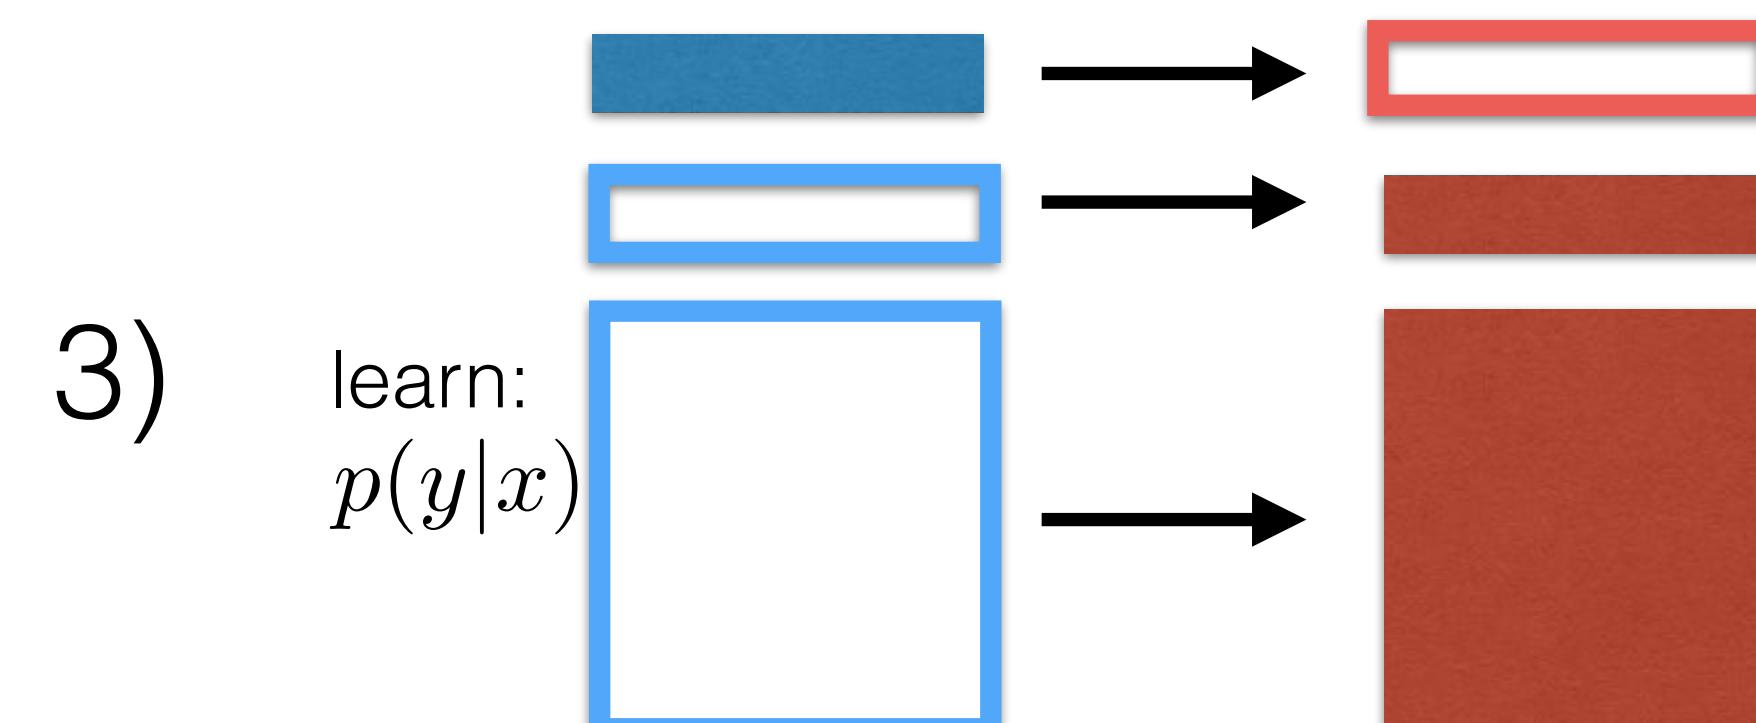
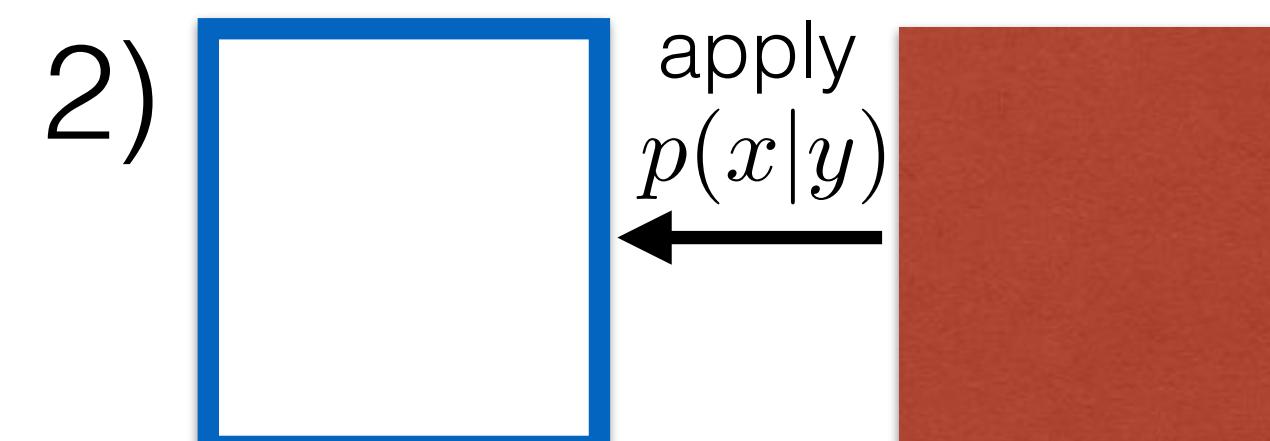
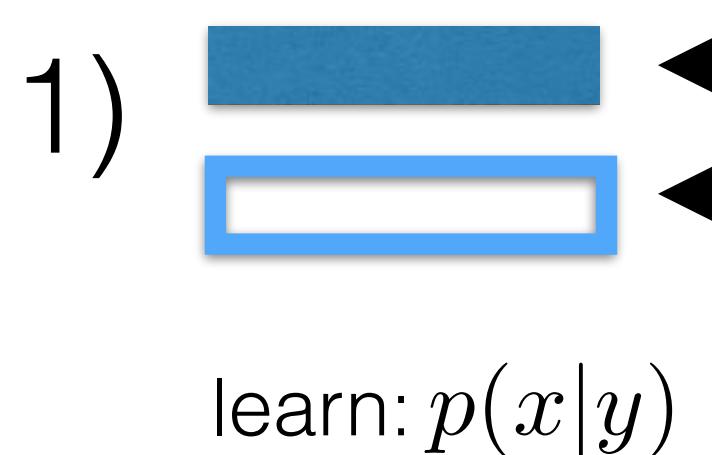


Baseline Approaches

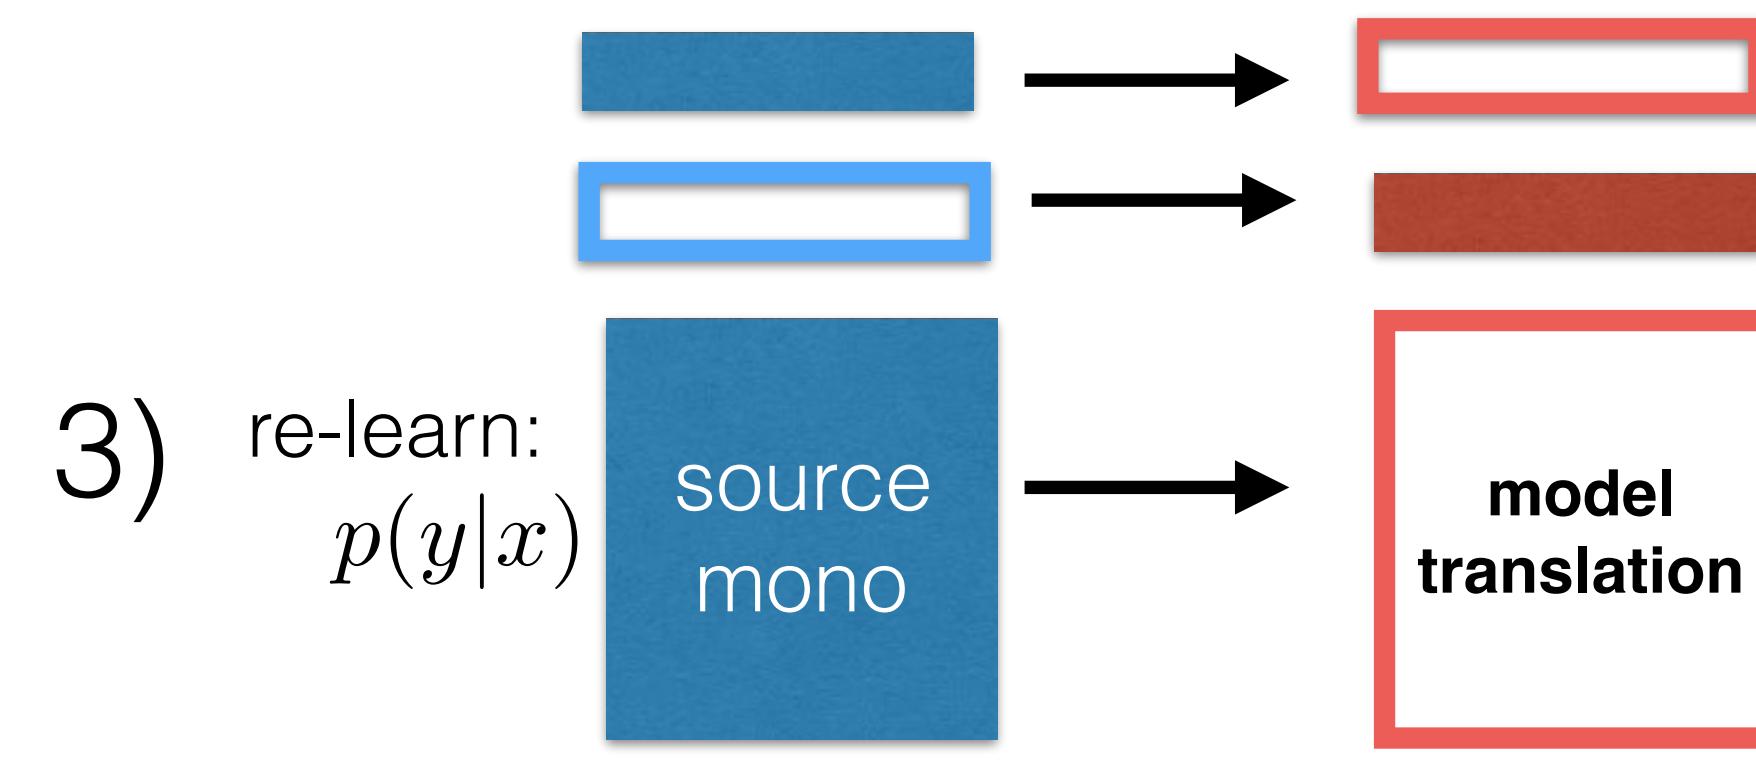
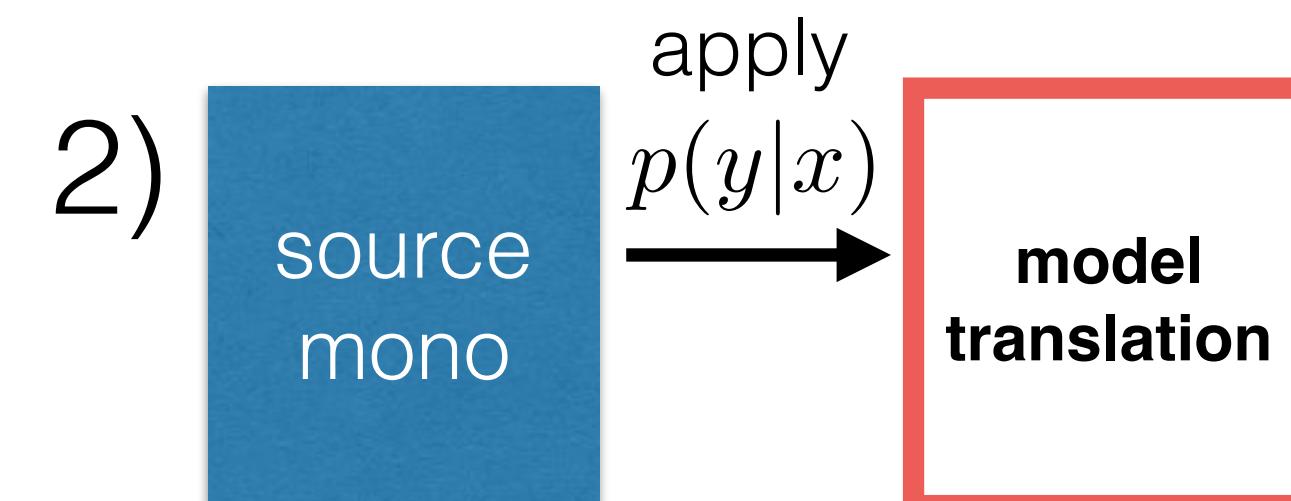
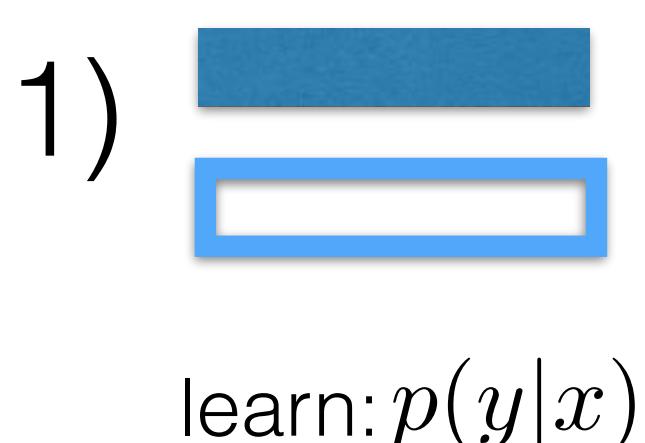
- Bitext only:



- Back-Translation:

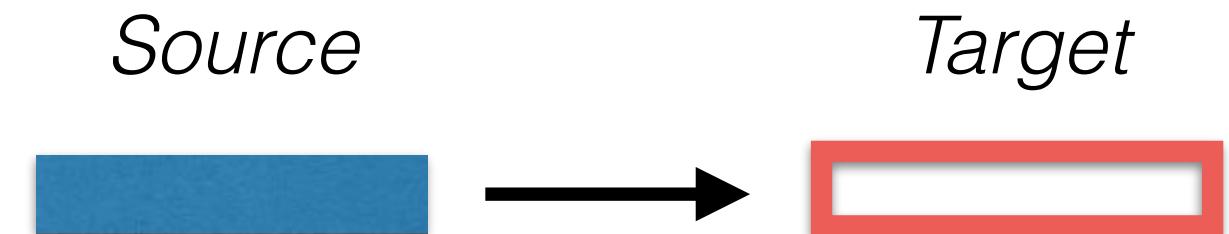


- Self-Training:

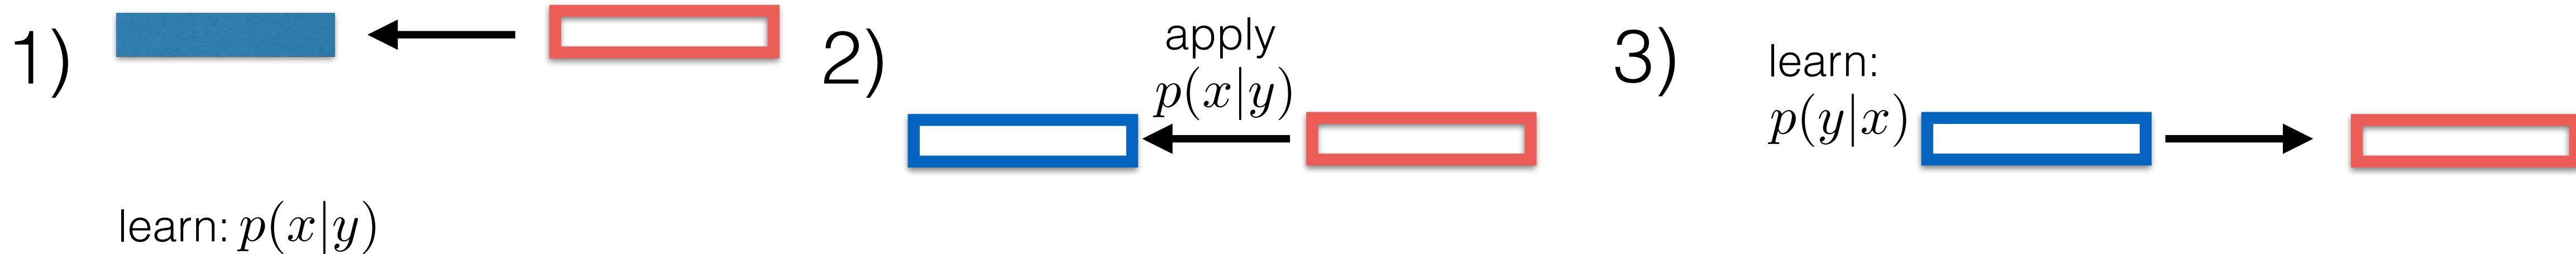


Baseline Approaches: Only In-Domain Data

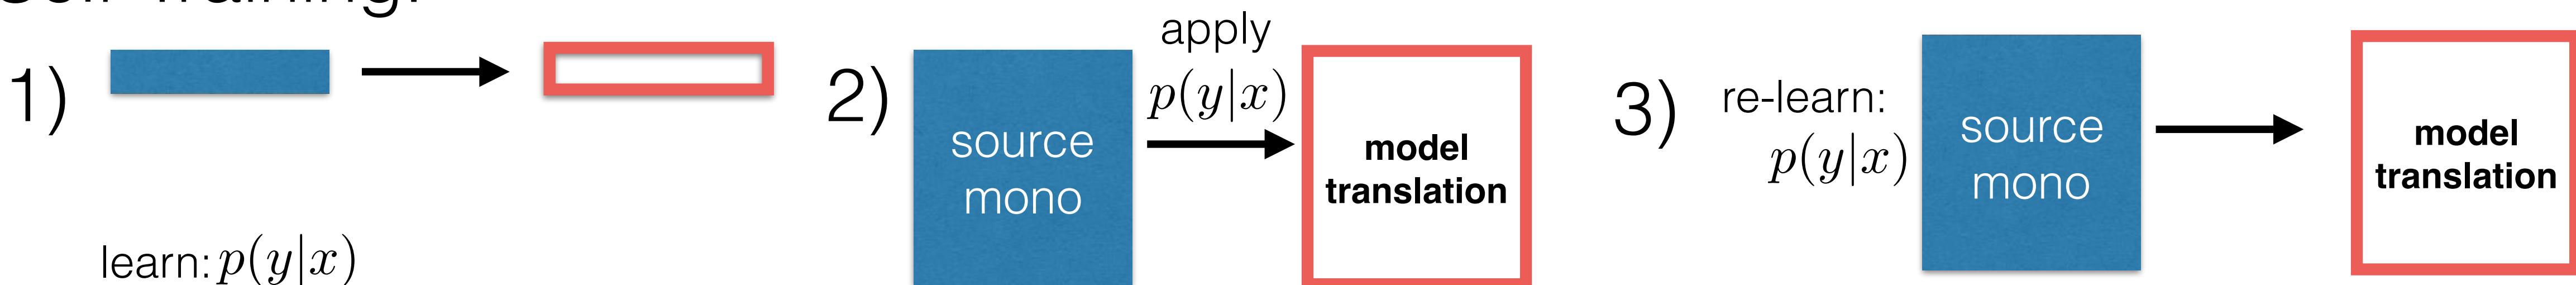
- Bitext only:



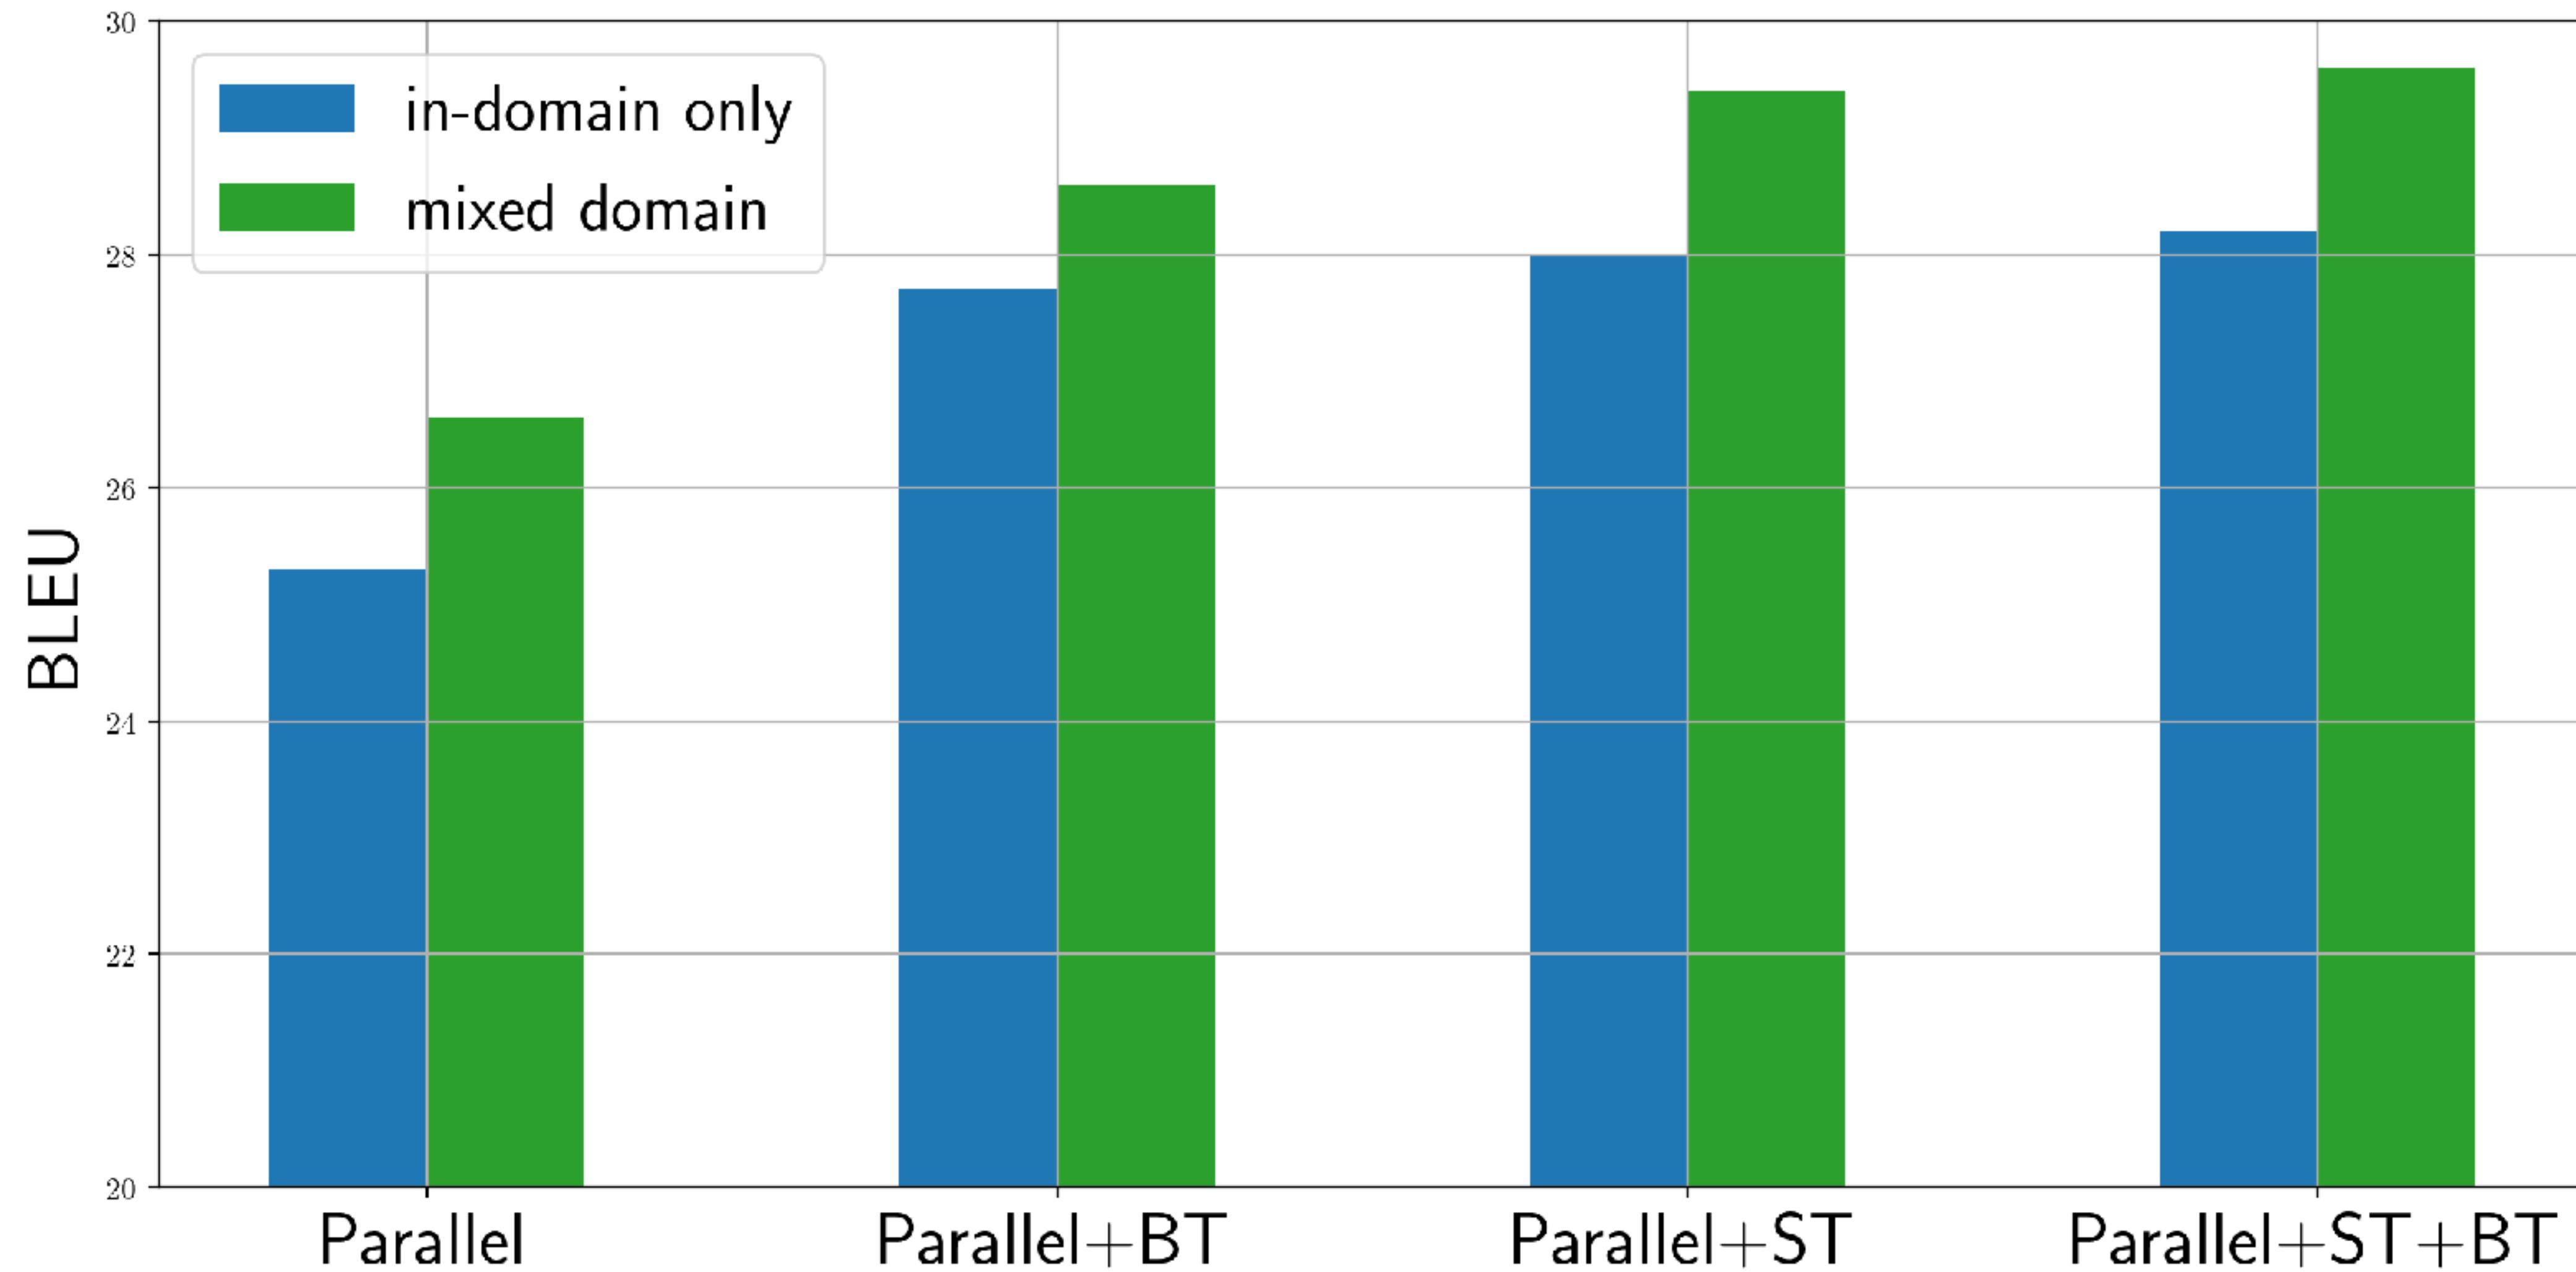
- Back-Translation:



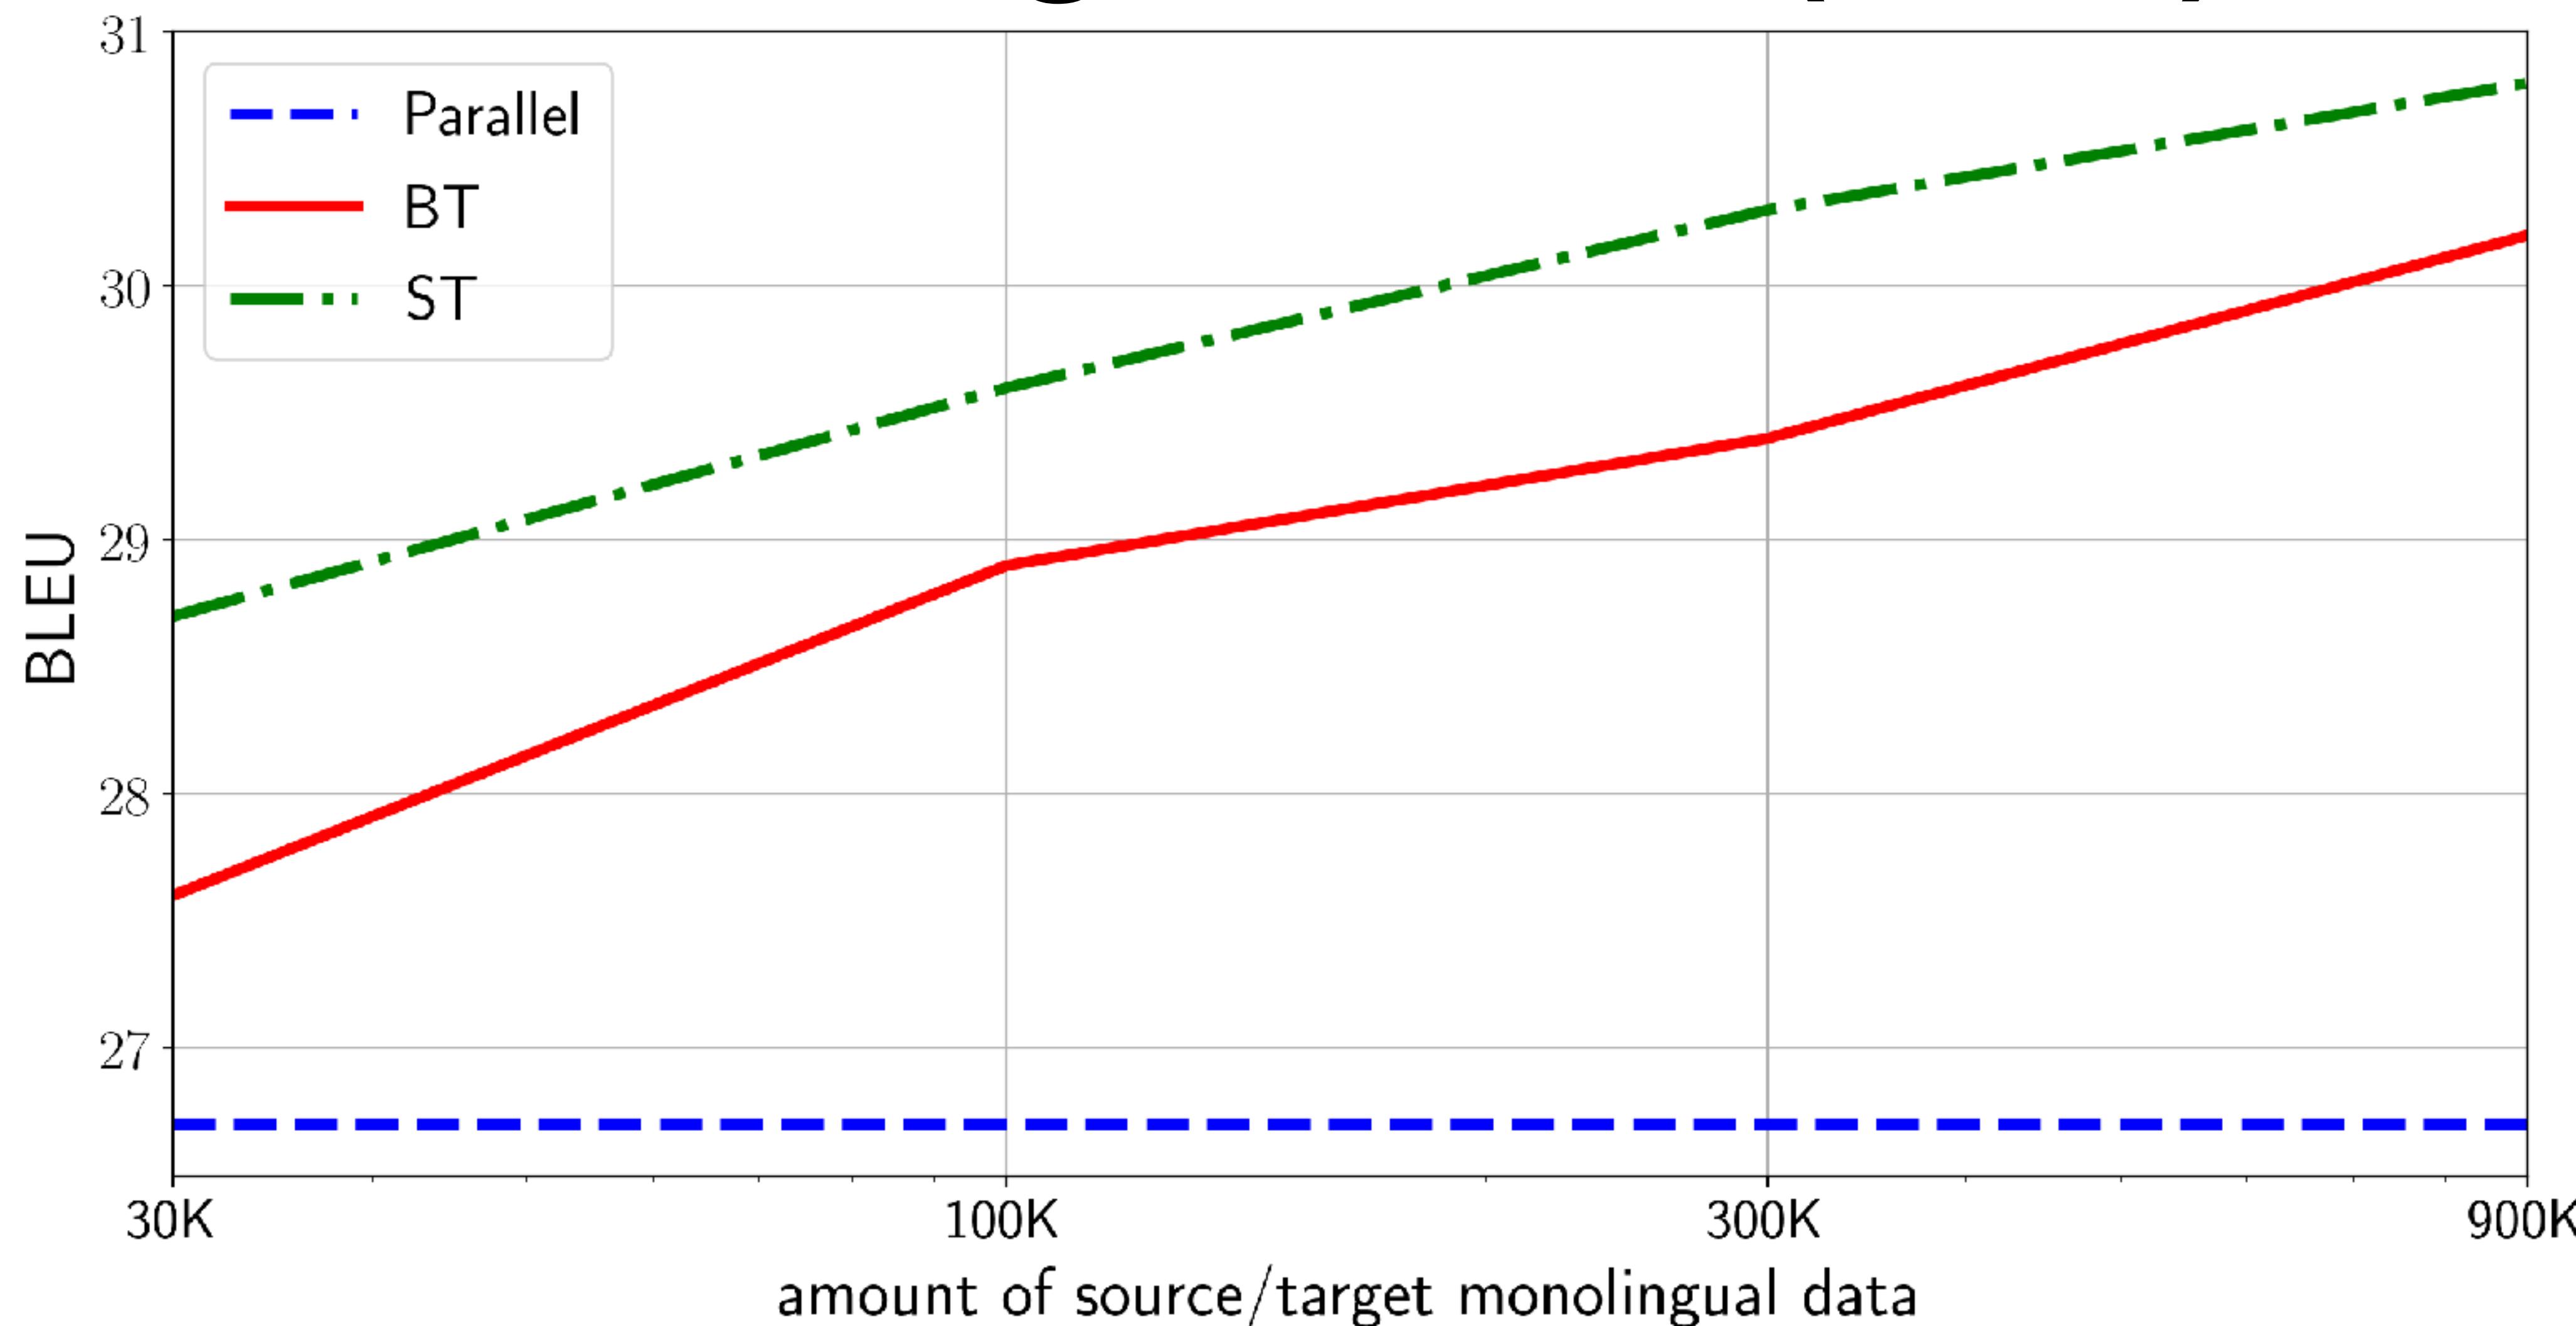
- Self-Training:



In-Domain Only VS. Mixed Domain



Varying Amount of Monolingual Data ($\alpha = 0$)

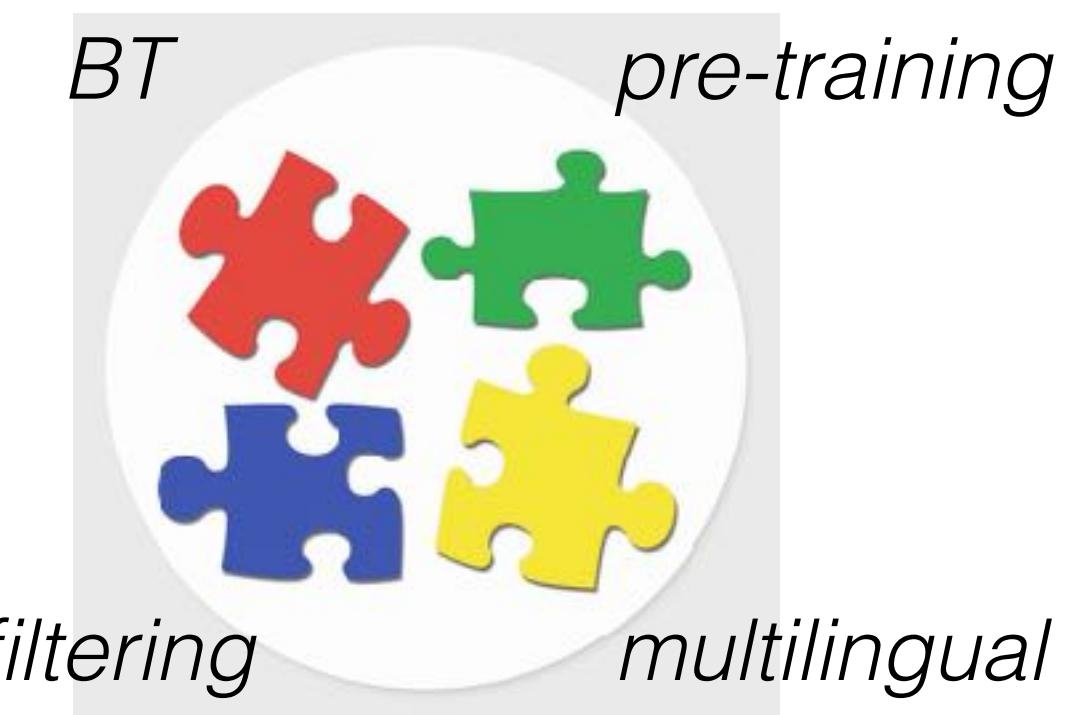


Conclusion

- STDM is particularly significant in low resource language pairs.
- Controlled setting helps studying STDM in isolation.
- ST is more robust than BT to STDM. We have already seen that combining ST & BT worked best in En-My.
- In practice, the influence of STDM depends on several factors, such as the amount of parallel and monolingual data, the domains, etc. In particular, if domains are not too distinct, STDM may even help regularizing!

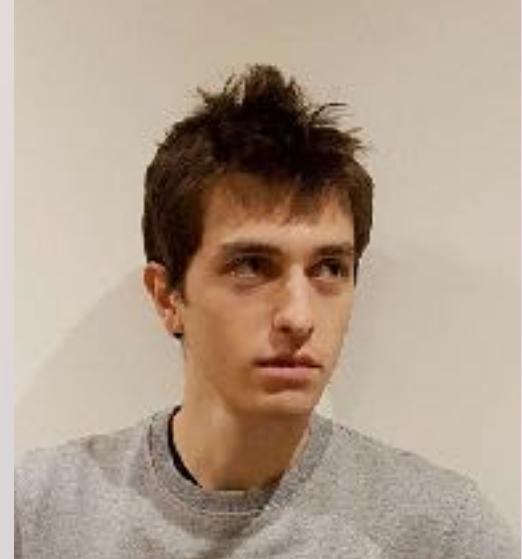
What I did not talk about: Filtering

- Data: extract clean version of common crawl.
- Learn a joint embedding space.
- Use approximate nearest neighbor methods to find closest matching sentence in embedding space.
- Translation quality on several languages is even higher than using actual existing bilingual + mono data!



Final Remarks

- Low resource MT is a good use case applications of several long standing ML problems: aligning domains, learning with less supervision, leveraging compositionality, etc.
- The importance and difficulty of data collection should not be under-estimated.
- A healthy cycle of research: data, modeling, analysis.
- Low resource MT key idea: use as many auxiliary tasks and data.
- Low resource MT requires lots of data and compute.
- Lots of open challenges.
 - Specific to low resource: dealing with all sorts of domain mismatch, learning from little data, quality of evaluation sets...
 - General of text generation: better use of context, common sense, striking a good trade-off between accuracy and speed, controllability, safety, biases...



Guillaume Lample



Ludovic Denoyer



Myle Ott



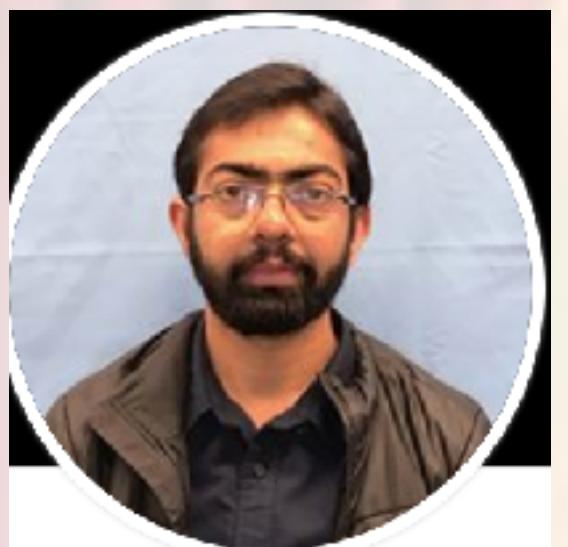
Peng-Jen Chen



Paco Guzmán



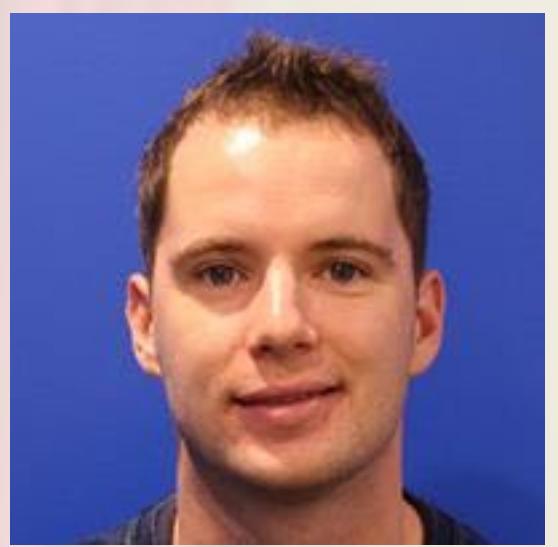
Jiajun Shen



Naman Goyal



Jiatao Gu



Alexis Conneau



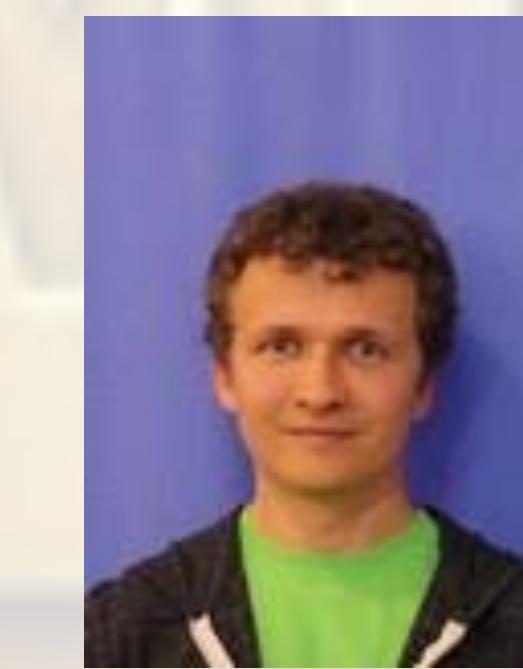
Philipp Kohen



Michael Auli



Junxian He



Sergey Edunov



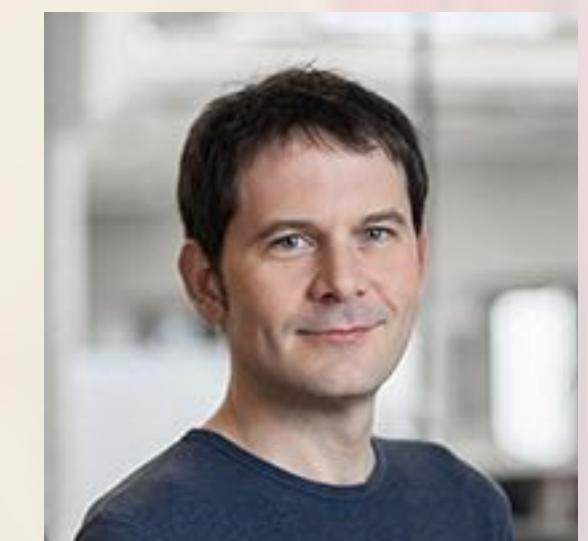
Xian Li



Juan-Miguel Pino



Vishrav Chaudhary



Hervé Jegou

Questions?

Вопросы?

¿Preguntas?

Domande?